

# 从 C++ 到 AI 计算：第三册实践与代码卷

本地量化推理引擎、完整源码、KV Cache、后端优化与验收

CPU Performance Study

本卷是《从 C++ 到 AI 计算：第三册》的配套实践与代码卷，围绕 `linux_cpu_inference` 贯穿项目展开：模型资产、tokenizer、张量布局、int8 kernel、reference decoder、KV cache、trace、7B 账本、服务调度、CPU/GPU 后端边界和工程验收都要落到可构建代码、完整源码、测试和报告。

# 前言

第三册原理卷把 AI 推理系统拆成模型资产、typed graph、张量布局、量化、KV cache、后端、服务运行时和验收体系。本实践与代码卷只做一件事：把这些概念压到同一个可运行项目里，并把完整源码阅读路径并入同一目录，让读者知道每个概念对应哪段代码、哪条命令、哪种输入、哪份报告和哪条失败路径。

贯穿项目是 [books/cpu-volume-3/labs/linux\\_cpu\\_inference](https://github.com/QuintusLiu/books/cpu-volume-3/labs/linux_cpu_inference)。它不是玩具目录，而是第三册已经持续维护的本地 CPU 量化推理切片：包含教学模型格式、tokenizer、safetensors manifest、gate、int8 linear、packed/AVX2 路径、tiny reference decoder、GPT-2 reference path、KV cache trace、7B compute ledger、serving SLO probe、自研 GPT-2 smoke、真实 small 开源模型 smoke、CLI、benchmark 和 CTest。本卷不维护第二份实现；源码章节通过 [lstinputlisting](#) 直接读取同一份工程文件，所有练习、讲解和完整代码都回到同一份源码。

本卷的学习顺序是从小证据到大系统：先让一个 tiny MLP 算对，再解释张量地址流和 int8 累加；再让 reference decoder、sampler 和 KV cache trace 固定语义；然后接模型导入、shape verifier、memory planner 和 IR lowering；最后用 CPU/GPU 后端边界、工业专项和回归门禁收束。代码走读不再被放到最后当附录，而是插入到相同主题部分中：讲完项目入口就读构建文件和调用链，讲完量化就读公共合同、tiny MLP 和 int8 kernel，讲到真实模型时读 tokenizer、safetensors 和 GPT-2，收束验收时读工具、benchmark、smoke 和测试。每章都要求读者交付代码运行结果或报告字段，而不是只抄概念定义。

## 核心思想

第三册实践的目标不是“写一个看起来像大模型的 demo”，而是建立一条能被复查的推理证据链：资产能校验，shape 能验证，kernel 能对照，trace 能定位，性能能降级，服务指标能解释，回归门禁能长期维护。

# 常见缩写和术语先导

| 缩写       | 全称                               | 本卷中的含义                          |
|----------|----------------------------------|---------------------------------|
| AI       | Artificial Intelligence          | 这里特指本地推理工程，不讨论泛泛应用口号。           |
| LLM      | Large Language Model             | decoder-only 文本生成模型是本卷主线。       |
| KV Cache | Key/Value Cache                  | attention 历史 K/V 状态，属于请求生命周期。   |
| GEMM     | General Matrix Multiply          | 矩阵乘，是预填充和训练常见核心算子。              |
| GEMV     | General Matrix Vector Multiply   | 矩阵向量乘，是单 token decode 的核心形态。    |
| SIMD     | Single Instruction Multiple Data | 一条指令处理多个 lane 的 CPU 执行方式。       |
| AVX2     | Advanced Vector Extensions 2     | x86-64 上的 256-bit SIMD 指令集之一。   |
| FMA      | Fused Multiply-Add               | 乘加融合指令，常出现在浮点 kernel 中。         |
| SLO      | Service Level Objective          | 服务目标，例如 TTFT、TPOT、p95/p99。      |
| TTFT     | Time To First Token              | 从请求进入到首 token 输出的时间。            |
| TPOT     | Time Per Output Token            | decode 阶段每个输出 token 的平均或尾部时间。   |
| IR       | Intermediate Representation      | 编译器内部表示，用于 pass、lowering 和计划生成。 |
| RAG      | Retrieval-Augmented Generation   | 检索增强生成，需要把外部文档和 prompt 编排起来。    |
| MoE      | Mixture of Experts               | 多专家模型，会改变路由、负载和内存账本。            |
| LoRA     | Low-Rank Adaptation              | 低秩适配权重，常用于微调和多租户加载。             |

这些缩写不要死记。每个缩写在正文里都要落到一个代码对象或报告字段：KV cache 对应 trace 里的 slot 和 position，SLO 对应 serving probe 的 CSV，IR 对应 executable plan，SIMD/AVX2 对应 kernel 反汇编和 benchmark。

## 本卷结构

本卷按第三册原书的四个大块组织，但目录不再采用“前面做实践、后面贴源码”的割裂方式。每个大块都把实践任务、代码走读和验收证据放在同一条主线里：[linux\\_cpu\\_inference](#)。

第一部分从模型资产到量化 kernel。你会运行 tiny MLP，先读构建入口、项目边界和调用链总图，再检查 tokenizer 和 safetensors manifest，手算张量地址，比较 scalar、packed 和 AVX2 路径，并在公共合同、tiny MLP 与 int8 kernel 源码里解释 scale、accumulator 和 requantize 的边界。

第二部分进入 Transformer、KV cache、sampler 和服务运行时。你会跑 reference decoder trace，解释每个 checkpoint 的 shape、stride、dtype 和 layout，构造 KV cache 状态表，分析 admission、queue wait、TTFT、TPOT、取消和拒绝。

第三部分处理模型导入、真实模型加载闭环、memory planner 和 IR lowering。你会把 manifest、shape verifier、tensor role、state edge、workspace、liveness 和 executable plan 写成可检查报告，并在同一部分中读 reference decoder、tokenizer、safetensors 和 GPT-2 reference path 的完整实现，显式运行 LCQI 自研 GPT-2 safetensors smoke 和一个 small 级别开源 GGUF 模型 smoke，而不是只停留在 tiny fixture。

第四部分收束到后端优化和验收。你会比较 CPU kernel 证据、GPU 后端边界、工业专项能力和回归门禁，并直接读工具、smoke、benchmark、ledger、serving probe 和测试源码。最终交付不是“跑通一次”，而是可复查的工程档案：命令、输入、输出、trace、benchmark、不可验证边界和回归规则。

# 目录

|                                                          |     |
|----------------------------------------------------------|-----|
| 前言                                                       | ii  |
| 常见缩写和术语先导                                                | iii |
| 本卷结构                                                     | iv  |
| <b>第一部分 推理链路、张量账本与量化</b>                                 |     |
| 第 1 章实践：端到端推理链路和项目总入口                                    | 2   |
| 一、对应原书问题：概念必须落到对象和命令                                     | 2   |
| 二、从特例推到规律：先抓一个能手算的切片                                     | 2   |
| 三、实践任务：把观察固定成报告字段                                        | 2   |
| 四、验收和递进大作业                                               | 2   |
| 代码走读：构建入口、项目边界与调用链总图                                     | 3   |
| 源码阅读覆盖范围                                                 | 3   |
| 根构建入口                                                    | 3   |
| LCQI 目标定义                                                | 6   |
| 项目说明与模型资产                                                | 9   |
| 为什么要先讲调用链                                                | 15  |
| 一、全工程入口：哪些文件决定能不能复现                                      | 16  |
| 二、Tiny MLP 调用链：最短闭环怎样跑通                                  | 16  |
| 三、Int8 Kernel 调用链：reference、packed、AVX2 和 benchmark 怎样对照 | 17  |
| 四、Tokenizer、Safetensors 与 Reference Decoder：真实模型前的三道门    | 17  |
| 五、GPT-2 调用链：从 HuggingFace 资产到生成 token                    | 18  |
| 六、报告、服务探针、外部对标和测试                                        | 19  |
| 第 2 章实践：张量布局、地址流与第一个 MatVec                              | 20  |
| 一、对应原书问题：概念必须落到对象和命令                                     | 20  |
| 二、从特例推到规律：先抓一个能手算的切片                                     | 20  |
| 三、实践任务：把观察固定成报告字段                                        | 20  |
| 四、验收和递进大作业                                               | 20  |
| 第 3 章实践：GEMM、packing 与 SIMD 证据                           | 21  |
| 一、对应原书问题：概念必须落到对象和命令                                     | 21  |
| 二、从特例推到规律：先抓一个能手算的切片                                     | 21  |
| 三、实践任务：把观察固定成报告字段                                        | 21  |
| 四、验收和递进大作业                                               | 21  |
| 第 4 章实践：int8 量化、累加器与重新量化                                 | 22  |
| 一、对应原书问题：概念必须落到对象和命令                                     | 22  |
| 二、从特例推到规律：先抓一个能手算的切片                                     | 22  |
| 三、实践任务：把观察固定成报告字段                                        | 22  |
| 四、验收和递进大作业                                               | 22  |
| 代码走读：公共合同、Tiny MLP 与 Int8 Kernel                         | 23  |
| 为什么先读头文件                                                 | 23  |
| Tiny MLP 与推理入口                                           | 23  |
| Int8 Kernel、Tokenizer 与 Safetensors                      | 24  |
| Reference Decoder 与 GPT-2 合同                             | 27  |
| 最短闭环：从模型文件到 logits                                       | 33  |
| 模型加载与基础推理                                                | 33  |
| Int8 Kernel 基础实现                                         | 40  |
| <b>第二部分 Transformer、KV Cache 与服务运行时</b>                  |     |
| 第 5 章实践：Transformer 切片与 KV Cache 状态                      | 52  |
| 一、对应原书问题：概念必须落到对象和命令                                     | 52  |
| 二、从特例推到规律：先抓一个能手算的切片                                     | 52  |
| 三、实践任务：把观察固定成报告字段                                        | 52  |
| 四、验收和递进大作业                                               | 52  |
| 第 6 章实践：Reference Decoder、logits 与采样                     | 53  |
| 一、对应原书问题：概念必须落到对象和命令                                     | 53  |
| 二、从特例推到规律：先抓一个能手算的切片                                     | 53  |
| 三、实践任务：把观察固定成报告字段                                        | 53  |

|                                                      |     |
|------------------------------------------------------|-----|
| 四、验收和递进大作业                                           | 53  |
| 第 7 章实践：服务调度、SLO 与资源预算                               | 54  |
| 一、对应原书问题：概念必须落到对象和命令                                 | 54  |
| 二、从特例推到规律：先抓一个能手算的切片                                 | 54  |
| 三、实践任务：把观察固定成报告字段                                    | 54  |
| 四、验收和递进大作业                                           | 54  |
| 第 8 章实践：RAG、Agent 与状态图边界                             | 55  |
| 一、对应原书问题：概念必须落到对象和命令                                 | 55  |
| 二、从特例推到规律：先抓一个能手算的切片                                 | 55  |
| 三、实践任务：把观察固定成报告字段                                    | 55  |
| 四、验收和递进大作业                                           | 55  |
| <b>第三部分 模型导入、AI 编译器与内存计划</b>                         |     |
| 第 9 章实践：模型导入、manifest 与 Shape Verifier               | 57  |
| 一、对应原书问题：概念必须落到对象和命令                                 | 57  |
| 二、从特例推到规律：先抓一个能手算的切片                                 | 57  |
| 三、实践任务：把观察固定成报告字段                                    | 57  |
| 四、验收和递进大作业                                           | 57  |
| 第 10 章实践：真实模型加载闭环与首 Token Trace                      | 58  |
| 一、对应原书问题：概念必须落到对象和命令                                 | 58  |
| 二、从特例推到规律：先抓一个能手算的切片                                 | 58  |
| 三、自研 GPT-2 reference smoke：先把标准 safetensors 跑通       | 58  |
| 四、真实 small 模型 smoke：不用 tiny，先让开源模型实际跑起来              | 59  |
| 五、实践任务：把观察固定成报告字段                                    | 60  |
| 六、验收和递进大作业                                           | 60  |
| 代码走读：Reference Decoder、Tokenizer、Safetensors 与 GPT-2 | 61  |
| 从 tiny MLP 走向 decoder-only 模型                        | 61  |
| Tiny Reference Decoder                               | 61  |
| Tokenizer 与 Safetensors 实现                           | 78  |
| GPT-2 Reference Path                                 | 94  |
| 第 11 章实践：运行时内存计划、liveness 与 Workspace                | 139 |
| 一、对应原书问题：概念必须落到对象和命令                                 | 139 |
| 二、从特例推到规律：先抓一个能手算的切片                                 | 139 |
| 三、实践任务：把观察固定成报告字段                                    | 139 |
| 四、验收和递进大作业                                           | 139 |
| 第 12 章实践：IR、Pass、Lowering 与 Executable Plan          | 140 |
| 一、对应原书问题：概念必须落到对象和命令                                 | 140 |
| 二、从特例推到规律：先抓一个能手算的切片                                 | 140 |
| 三、实践任务：把观察固定成报告字段                                    | 140 |
| 四、验收和递进大作业                                           | 140 |
| <b>第四部分 后端优化、工业专项与验收</b>                             |     |
| 第 13 章实践：CPU 后端、微内核与性能证据                             | 142 |
| 一、对应原书问题：概念必须落到对象和命令                                 | 142 |
| 二、从特例推到规律：先抓一个能手算的切片                                 | 142 |
| 三、实践任务：把观察固定成报告字段                                    | 142 |
| 四、验收和递进大作业                                           | 142 |
| 第 14 章实践：GPU 后端边界、copy 与同步成本                         | 143 |
| 一、对应原书问题：概念必须落到对象和命令                                 | 143 |
| 二、从特例推到规律：先抓一个能手算的切片                                 | 143 |
| 三、实践任务：把观察固定成报告字段                                    | 143 |
| 四、验收和递进大作业                                           | 143 |
| 第 15 章实践：工业 LLM 专项能力对照                               | 144 |
| 一、对应原书问题：概念必须落到对象和命令                                 | 144 |
| 二、从特例推到规律：先抓一个能手算的切片                                 | 144 |
| 三、实践任务：把观察固定成报告字段                                    | 144 |
| 四、验收和递进大作业                                           | 144 |
| 第 16 章实践：工程验收、回归门禁与最终项目                              | 145 |
| 一、对应原书问题：概念必须落到对象和命令                                 | 145 |
| 二、从特例推到规律：先抓一个能手算的切片                                 | 145 |
| 三、实践任务：把观察固定成报告字段                                    | 145 |
| 四、验收和递进大作业                                           | 145 |

|                                                  |     |
|--------------------------------------------------|-----|
| 代码走读：工具、Smoke、Benchmark 与回归测试 . . . . .          | 146 |
| 工具不是附属品 . . . . .                                | 146 |
| Benchmark、Trace、Ledger 与 Serving Probe . . . . . | 146 |
| 外部模型 Smoke 与对比脚本 . . . . .                       | 161 |
| 完整测试 . . . . .                                   | 184 |
| 实践与代码总索引 . . . . .                               | 202 |
| 一、实践主题到代码走读的合并位置 . . . . .                       | 202 |
| 二、最小复现命令 . . . . .                               | 202 |
| 术语表 . . . . .                                    | 203 |





第零部分

# 推理链路、张量账本与量化

## 第 1 章实践：端到端推理链路和项目总入口

本章对应原书中的第三册“端到端链路：从模型文件到下一个 token”。不要把它当作概念复述，本章要把模型目录、manifest、tokenizer、typed graph、backend plan、request session、KV cache、oracle 和 profiler 放到同一个可运行项目里检查。贯穿代码仍然是 `books/cpu-volume-3/labs/linux_cpu_inference`；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `tiny_mlp_i8.txt`，核心命令或测试入口是 `lcqi_cli`。

Listing 1.1 第 1 章实践：端到端推理链路和项目总入口的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  ↪ DCMMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_cli / tiny_mlp_i8.txt
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：只跑一个 tiny MLP 推理请求。读者要先手写预期结果，再运行项目观察输出。动作是：把 `readfiles->tokenize->run_inference->logits->predicted_class` 串起来。如果输出里没有 `predicted_class` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：模型文件存在不等于推理链路可信；只有 CLI 输出、测试断言、输入合同和失败路径同时成立，才算完成入口闭环这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 `books/cpu-volume-3/labs/linux_cpu_inference/README.md`、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：最终大作业的总目录：每章报告都放入 `reports/volume3-practice/`，第一章提交项目地图、命令记录和一次成功/失败输入对照。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。

# 代码走读：构建入口、项目边界与调用链总图

## 源码阅读覆盖范围

LCQI 的源码边界由三个层次组成。第一层是仓库构建入口，它决定 CMake 怎样把第三册的实验代码纳入构建和测试。第二层是 `labs/linux_cpu_inference`，它包含真正的库、CLI、benchmark、trace、ledger、serving probe 和测试目标。第三层是工具脚本，它们下载或复用外部模型资产，生成 smoke 和 benchmark 报告。

### 源码阅读任务

先读本章的 CMake 文件，再读 README。读者要能回答三个问题：一是 `lcqi` 静态库包含哪些实现文件；二是哪些可执行文件只依赖自研代码，哪些依赖可选外部库；三是真实 GPT-2 和 SmolLM2 smoke 为什么不进入默认 CTest。

## 根构建入口

第三册根目录的 CMake 文件只做项目级选择：是否构建第三册 labs，以及如何把 `labs/linux_cpu_inference` 接入。它应该保持薄层，避免把具体目标定义分散到多个地方。

Listing 1.2 `cpu-volume-3/README.md`

```
1 # 从 C++ 到 AI 计算：第三册
2
3 第三册用于承接 AI、AI 编译器、算子开发和本地大模型推理引擎。第二册只讨论硬件原
  ↳ 理、多核、高性能计算、并发、锁、无锁数据结构、并行算法和分布式计算，不再
  ↳ 直接进入 AI。
4
5 本册贯穿项目是一台能推理开源 7B 左右 decoder-only 大模型的本地引擎，并把 AI 编
  ↳ 译器作为引擎内部主线：模型图导入、typed tensor IR、算子融合、内存规划、
  ↳ CPU/GPU lowering、kernel selection 和 runtime dispatch 都要能落回编译原理
  ↳ 与底层机器账本。工程路线从 CPU 上可解释的 reference 和量化切片开始，逐步
  ↳ 扩展到 tokenizer、模型加载、Transformer 层、KV cache、CPU SIMD/多线程后
  ↳ 端、GPU 后端、正确性报告和性能对标。
6
7 正式书稿工程位于：
8
9 ```text
10 source/latex/
11 ```
12
13 构建命令：
14
15 ```bash
16 make -C source/latex check
17 make -C source/latex pdf
18 make -C source/latex epub
19 ```
20
21 生成文件：
22
23 ```text
```

```

24 source/latex/main.pdf
25 source/latex/main.epub
26 ```
27
28 当前保留的第一阶段切片：
29
30 ```text
31 labs/linux_cpu_inference/
32 ```
33
34 该目录目前包含最小 MLP/int8 权重教学切片、tiny reference decoder、GPT-2 ↵
    ↳ reference path、int8 scalar baseline、packed layout、AVX2 SIMD 路径和 ↵
    ↳ shape sweep benchmark，用来验证张量、线性层、量化权重、tokenizer、↵
    ↳ safetensors、GPT-2 forward、KV cache、oracle 和 benchmark 的基本边界。当↵
    ↳ 前 x86-64 AVX2 结果见：
35
36 ```text
37 results/lcqi-x86-avx2-benchmark-2026-06-25.csv
38 results/lcqi-x86-avx2-decode4096-2026-06-25.csv
39 ```
40
41 自研 GPT-2 smoke 入口如下：
42
43 ```bash
44 make cpu3-gpt2-smoke
45 make cpu3-gpt2-benchmark-compare
46 ```
47
48 第一条命令下载 `openai-community/gpt2` 的标准 safetensors 资产到用户缓存目录，↵
    ↳ 用 LCQI C++ reference path 生成 1 个 token，并把报告写入：
49
50 ```text
51 results/lcqi-gpt2-smoke.txt
52 ```
53
54 第二条命令用同一个 GPT-2 F32 模型对比 LCQI full-prefix 和 cached-KV 生成路径，↵
    ↳ 并在本机存在 llama.cpp/SmolLM2 缓存时附带成熟引擎参考，报告写入：
55
56 ```text
57 results/lcqi-gpt2-benchmark-compare.txt
58 ```
59
60 它仍然不是生产级 AI Infra。生产级门禁见：
61
62 ```text
63 reports/ai-infra-maturity-audit.md
64 ```
65
66 第三册后续必须把 lab 推进到“小型 CPU 推理 runtime + 手写高性能算子 + 严格 ↵
    ↳ benchmark 报告 + 真实系统对标”，否则只能证明系统学习，不能证明生产级能↵
    ↳ 力。

```

Listing 1.3 cpu-volume-3/CMakeLists.txt

```

1 cmake_minimum_required(VERSION 3.31)
2 project(CPUVolume3 LANGUAGES CXX)
3
4 set(CMAKE_CXX_STANDARD 20)
5 set(CMAKE_CXX_STANDARD_REQUIRED ON)
6 set(CMAKE_CXX_EXTENSIONS OFF)
7 set(CMAKE_EXPORT_COMPILE_COMMANDS ON)
8
9 option(CPU_VOLUME_3_BUILD_LABS "Build CPU Volume 3 labs" ON)
10
11 if(CPU_VOLUME_3_BUILD_LABS)
12     enable_testing()
13     add_subdirectory(labs/linux_cpu_inference)
14 endif()

```

实践与代码卷也提供一个薄 CMake 入口，它把同一份 LCQI 源码作为可测试对象接入。这样本卷不会复制实现，只复用第三册主工程。

Listing 1.4 cpu-volume-3-practice/CMakeLists.txt

```

1 cmake_minimum_required(VERSION 3.31)
2 project(CPUVolume3Practice LANGUAGES CXX)
3
4 set(CMAKE_CXX_STANDARD 20)
5 set(CMAKE_CXX_STANDARD_REQUIRED ON)
6 set(CMAKE_CXX_EXTENSIONS OFF)
7 set(CMAKE_EXPORT_COMPILE_COMMANDS ON)
8
9 enable_testing()
10 add_subdirectory(..cpu-volume-3/labs/linux_cpu_inference linux_cpu_inference)

```

本卷自身的 Makefile 负责构建 PDF/EPUB，并通过 `test` 目标回到第三册真实 LCQI 工程做 CMake/CTest 验证。实践与代码卷如果只排版、不测试，就不能宣称“完整源码可运行”。

Listing 1.5 cpu-volume-3-practice/Makefile

```

1 LATEX_DIR := source/latex
2
3 .PHONY: all check pdf epub test clean
4
5 all: pdf epub
6
7 check:
8     $(MAKE) -C $(LATEX_DIR) check
9
10 pdf:
11     $(MAKE) -C $(LATEX_DIR) pdf
12
13 epub:
14     $(MAKE) -C $(LATEX_DIR) epub
15

```

```

16 test:
17     cmake -S . -B build/lcqi-debug -DCMAKE_BUILD_TYPE=Debug
18     cmake --build build/lcqi-debug
19     ctest --test-dir build/lcqi-debug --output-on-failure
20
21 clean:
22     $(MAKE) -C $(LATEX_DIR) clean

```

Listing 1.6 cpu-volume-3-practice/source/latex/Makefile

```

1 LATEXMK ?= latexmk
2 ENGINE ?= -xelatex
3 MAIN ?= main.tex
4 EPUB_BUILDER ?= ../../../../cpu-volume-1/source/latex/scripts/build_epub.py
5
6 .PHONY: all pdf epub clean check
7
8 all: pdf epub
9
10 pdf:
11     $(LATEXMK) $(ENGINE) -interaction=nonstopmode -halt-on-error $(MAIN)
12
13 epub:
14     @BOOK_TITLE="从 C++ 到 AI 计算：第三册实践与代码卷" \
15     BOOK_IMPORT_TITLE="CPU Volume 3 Practice Code" \
16     BOOK_ID_SEED="cpu-volume-3-practice-code-weread-20260629" \
17     EPUB_ASCII_TITLES=1 \
18     BOOK_SUBTITLE="本地量化推理引擎、完整源码、KV Cache、后端优化与验收" \
19     BOOK_AUTHOR="CPU Performance Study" \
20     python3 $(EPUB_BUILDER) --book-dir . --output main.epub --no-cover-page --↵
    ↵ forbid-images --linearize-tables --wechat-compatible --no-embed-fonts
21
22 clean:
23     $(LATEXMK) -C
24
25 check:
26     @python3 scripts/check_inputs.py
27     @command -v xelatex >/dev/null 2>&1 || { echo "missing xelatex"; exit 1; }
28     @command -v latexmk >/dev/null 2>&1 || { echo "missing latexmk"; exit 1; }

```

## LCQI 目标定义

`linux_cpu_inference/CMakeLists.txt` 是代码走读中的第一份核心文件。这里能看到 `lcqi` 库、`lcqi_cli`、`lcqi_bench`、`lcqi_trace`、`lcqi_gpt2`、`lcqi_ledger`、`lcqi_serving_probe`、可选 `lcqi_onednn_bench` 和 `lcqi_tests` 的完整关系。

Listing 1.7 linux\_cpu\_inference/CMakeLists.txt

```

1 add_library(lcqi STATIC
2     src/gpt2_reference.cpp
3     src/inference.cpp
4     src/int8_kernels.cpp

```

```

5     src/model.cpp
6     src/reference_decoder.cpp
7     src/safetensors.cpp
8     src/tokenizer.cpp
9 )
10
11 if(CMAKE_SYSTEM_PROCESSOR MATCHES "x86_64|AMD64|i[3-6]86")
12     if(CMAKE_CXX_COMPILER_ID MATCHES "Clang|GNU")
13         target_sources(lcqi PRIVATE src/int8_kernels_avx2.cpp)
14         set_source_files_properties(src/int8_kernels_avx2.cpp PROPERTIES ↵
↵ COMPILE_OPTIONS "-mavx2;-mfma")
15         target_compile_definitions(lcqi PRIVATE LCQI_ENABLE_AVX2=1)
16     endif()
17 endif()
18
19 if(NOT CMAKE_SYSTEM_NAME STREQUAL "Linux")
20     message(WARNING "LCQI is written for the book's local Linux CPU target; ↵
↵ current system is ${CMAKE_SYSTEM_NAME}.")
21 endif()
22
23 target_include_directories(lcqi
24     PUBLIC
25     ${CMAKE_CURRENT_SOURCE_DIR}/include
26 )
27
28 target_compile_features(lcqi PUBLIC cxx_std_20)
29
30 if(CMAKE_CXX_COMPILER_ID MATCHES "Clang|GNU")
31     target_compile_options(lcqi PRIVATE -Wall -Wextra -Wpedantic)
32 endif()
33
34 add_executable(lcqi_cli
35     src/main.cpp
36 )
37
38 target_link_libraries(lcqi_cli PRIVATE lcqi)
39
40 if(CMAKE_CXX_COMPILER_ID MATCHES "Clang|GNU")
41     target_compile_options(lcqi_cli PRIVATE -Wall -Wextra -Wpedantic)
42 endif()
43
44 add_executable(lcqi_bench
45     src/bench.cpp
46 )
47
48 target_link_libraries(lcqi_bench PRIVATE lcqi)
49
50 if(CMAKE_CXX_COMPILER_ID MATCHES "Clang|GNU")
51     target_compile_options(lcqi_bench PRIVATE -Wall -Wextra -Wpedantic)
52 endif()
53
54 add_executable(lcqi_trace

```

```

55     src/trace.cpp
56 )
57
58 target_link_libraries(lcqi_trace PRIVATE lcqi)
59
60 if(CMAKE_CXX_COMPILER_ID MATCHES "Clang|GNU")
61     target_compile_options(lcqi_trace PRIVATE -Wall -Wextra -Wpedantic)
62 endif()
63
64 add_executable(lcqi_gpt2
65     src/gpt2_cli.cpp
66 )
67
68 target_link_libraries(lcqi_gpt2 PRIVATE lcqi)
69
70 if(CMAKE_CXX_COMPILER_ID MATCHES "Clang|GNU")
71     target_compile_options(lcqi_gpt2 PRIVATE -Wall -Wextra -Wpedantic)
72 endif()
73
74 add_executable(lcqi_ledger
75     src/ledger.cpp
76 )
77
78 if(CMAKE_CXX_COMPILER_ID MATCHES "Clang|GNU")
79     target_compile_options(lcqi_ledger PRIVATE -Wall -Wextra -Wpedantic)
80 endif()
81
82 add_executable(lcqi_serving_probe
83     src/serving_probe.cpp
84 )
85
86 if(CMAKE_CXX_COMPILER_ID MATCHES "Clang|GNU")
87     target_compile_options(lcqi_serving_probe PRIVATE -Wall -Wextra -Wpedantic)
88 endif()
89
90 find_package(dnnl QUIET)
91 if(dnnl_FOUND)
92     add_executable(lcqi_onednn_bench
93         src/onednn_bench.cpp
94     )
95     target_link_libraries(lcqi_onednn_bench PRIVATE DNNL::dnnl)
96     target_compile_features(lcqi_onednn_bench PRIVATE cxx_std_20)
97     if(CMAKE_CXX_COMPILER_ID MATCHES "Clang|GNU")
98         target_compile_options(lcqi_onednn_bench PRIVATE -Wall -Wextra ↵
99         ↵ Wpedantic)
100     endif()
101 endif()
102
103 add_executable(lcqi_tests
104     tests/lcqi_tests.cpp
105 )

```



```

106 target_link_libraries(lcqi_tests PRIVATE lcqi)
107
108 if(CMAKE_CXX_COMPILER_ID MATCHES "Clang|GNU")
109     target_compile_options(lcqi_tests PRIVATE -Wall -Wextra -Wpedantic)
110 endif()
111
112 add_test(NAME lcqi_tests COMMAND ${<TARGET_FILE:lcqi_tests>})
113 add_test(NAME lcqi_gpt2_tiny COMMAND ${<TARGET_FILE:lcqi_gpt2>})
114 set_tests_properties(lcqi_gpt2_tiny PROPERTIES
115     PASS_REGULAR_EXPRESSION "generated_ids 1 2 3 3 3"
116 )
117 add_test(NAME lcqi_trace COMMAND ${<TARGET_FILE:lcqi_trace>})
118 set_tests_properties(lcqi_trace PROPERTIES
119     PASS_REGULAR_EXPRESSION "tokenizer,0,-1,5,1,i32,contiguous,10,4,0;1;2;3;4"
120 )
121 add_test(NAME lcqi_ledger COMMAND ${<TARGET_FILE:lcqi_ledger>})
122 set_tests_properties(lcqi_ledger PROPERTIES
123     PASS_REGULAR_EXPRESSION "bandwidth_limit,int4_at_100_gbps,23.602,tokens/s"
124 )
125 add_test(NAME lcqi_serving_probe COMMAND ${<TARGET_FILE:lcqi_serving_probe>})
126 set_tests_properties(lcqi_serving_probe PROPERTIES
127     PASS_REGULAR_EXPRESSION "request,req_too_long,0,memory_budget_exceeded"
128 )

```

## 项目说明与模型资产

README 是复现实验的入口。它写清楚当前能力、后续边界、构建命令、tiny 推理命令、benchmark 字段、reference trace 字段、GPT-2 smoke、7B ledger、serving SLO probe 和 SmolLM2 small smoke。读源码前先看 README，可以避免把 tiny fixture、GPT-2 reference path 和外部 llama.cpp smoke 混成同一个能力等级。

**Listing 1.8** linux\_cpu\_inference/README.md

```

1 # Linux CPU Quantized Inference
2
3 这是第三册贯穿项目的第一阶段切片：Linux 本地 CPU 量化线性层和最小推理流程。第三
  ↳ 册的贯穿目标不是这个 MLP 本身，而是一台能推理开源 7B 左右 decoder-only 大
  ↳ 模型的本地引擎，并逐步覆盖 tokenizer、模型加载、Transformer 层、KV cache
  ↳ 、CPU/GPU 后端、正确性报告和性能对标。
4
5 目标平台是 x86-64 Linux 实验环境。完整性能报告应记录 CPU 型号、内核版本、编译器
  ↳ 版本、是否能使用 PMU、NUMA 和线程亲和能力。若运行环境缺少 perf、NUMA、线
  ↳ 程亲和性或硬件性能计数器权限，报告必须把相关结论降级为“未验证”，不能把源
  ↳ 码路径存在当作实测性能证据。后续 GPU 后端会作为独立 backend 接入，不改变
  ↳ 当前切片的语义合同。
6
7 当前版本支持：
8
9 - 教学文本模型格式 `LCQI_MODEL_V1`
10 - 最小 tokenizer 合同格式 `LCQI_TOKENIZER_V1`
11 - safetensors header manifest 解析：metadata、tensor name、dtype、shape、data
  ↳ offsets、offset 边界和 dtype/shape 字节数校验

```

```

12 - 两层 MLP 教学切片
13 - int8 权重加 per-layer scale
14 - float 输入和激活
15 - 量化线性层、ReLU、分类 argmax
16 - int8 linear scalar baseline、packed layout、AVX2 路径和 shape sweep benchmark
17 - tiny reference decoder: RMSNorm、float linear、RoPE、GQA KV cache、decode ↵
    ↵ attention、SwiGLU、lm head
18 - GPT-2 reference path: byte-level BPE tokenizer、`config.json`、F32/F16/BF16 ↵
    ↵ safetensors 张量读取、HuggingFace GPT-2 name mapping、LayerNorm、绝对位置 ↵
    ↵ 嵌入、packed QKV、causal attention、GELU、tied lm head、full-prefix ↵
    ↵ greedy generation、KV cache greedy generation 和分阶段 benchmark
19 - reference trace CSV: 从 prompt/tokenizer 合同到 logits/sampler/token, 逐 ↵
    ↵ checkpoint 输出 shape、stride、dtype、layout、checksum、max_abs、values ↵
    ↵ 摘要
20 - 7B ledger CSV: 输出典型 7B shape 的 FLOPs、KV 容量、权重/KV 字节和 bandwidth ↵
    ↵ only tokens/s 上限
21 - serving SLO probe CSV: 输出 admission、batch state、KV 预算、queue wait、TTFT ↵
    ↵ 、TPOT、取消回收和拒绝率
22 - 自研 GPT-2 冒烟: 加载 `openai-community/gpt2` 的 `config.json`、`vocab.json` ↵
    ↵ 、`merges.txt`、`model.safetensors`, 在 LCQI C++ reference path 上生成 1 ↵
    ↵ 个 token, 并把文件 hash、prompt ids、generated ids 和文本输出写入报告
23 - 自研 GPT-2 benchmark 对比: 同一个 GPT-2 F32 safetensors 模型分别跑 full ↵
    ↵ prefix 和 cached-KV, 并在可用时附带 llama.cpp/SmolLM2 Q4_K_M 作为成熟引擎 ↵
    ↵ 参考
24 - 真实 small 开源模型冒烟: 通过 llama.cpp 加载 SmolLM2-135M-Instruct 的 GGUF 量 ↵
    ↵ 化文件, 在本地 CPU 上生成一段 assistant 文本, 并把模型 hash、命令、输出和 ↵
    ↵ 性能行写入报告
25 - CLI 推理和简单 benchmark
26 - CTest 正确性测试
27
28 后续扩展方向:
29
30 - 真实 7B 级模型 metadata、tokenizer 和量化权重导入
31 - SentencePiece tokenizer、GGUF 权重导入、safetensors 分片加载和更多模型方言 ↵
    ↵ name mapping
32 - GPT-2 reference path 的 Transformers/PyTorch golden logits 对齐报告
33 - KV cache 分页、内存规划和可选量化
34 - CPU packed weight、SIMD、线程和 NUMA 优化
35 - GPU backend adapter 和 device kernel
36 - 与现有框架的 prefill/decode/内存/误差对标报告
37
38 构建:
39
40 ```bash
41 cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -↵
    ↵ DCMMAKE_BUILD_TYPE=Debug
42 cmake --build books/cpu-volume-3/build/lcqi-debug
43 ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
44 ```
45
46 运行:
47

```

```

48 ```bash
49 books/cpu-volume-3/build/lcqi-debug/labs/linux_cpu_inference/lcqi_cli \
50   books/cpu-volume-3/labs/linux_cpu_inference/models/tiny_mlp_i8.txt \
51   1.0 2.0 -1.0 0.5 1000
52 ```
53
54 kernel benchmark:
55
56 ```bash
57 books/cpu-volume-3/build/lcqi-debug/labs/linux_cpu_inference/lcqi_bench 200
58 ```
59
60 输出 CSV 字段:
61
62 ```text
63 target_arch,input_size,output_size,output_block,repeat,backend,available,↵
64   ↵ average_us,max_abs_diff,checksum
65 ```
66 当前 benchmark 按 backend 逐行输出: `scalar`、`packed_scalar`、`avx2`。x86-64 ↵
67   ↵ 构建会编译 AVX2 路径; 不支持 AVX2/FMA 的机器会把该 backend 标为 `↵
68   ↵ available=0`, 避免把源码存在误报成当前环境实测性能。
69
70 这还不是生产级高性能 kernel。当前 benchmark 建立 scalar baseline、packed layout↵
71   ↵ 、AVX2 基线正确性和 shape sweep 口径。要达到完整 AI Infra 证据, 还必须继↵
72   ↵ 续补反汇编分析、AVX-512、perf counter、oneDNN/OpenBLAS/llama.cpp/ggml 对↵
73   ↵ 照、sanitizer、fuzz 和 CI。
74
75 safetensors manifest gate:
76
77 `lcqi_tests` 会运行时生成一个最小 safetensors fixture, 验证 header 长度、`↵
78   ↵ __metadata__`、tensor name、dtype、shape、data offsets、payload 边界和 ↵
79   ↵ dtype/shape 推导出的字节数。它还会生成 offset 超出 payload、dtype/shape ↵
80   ↵ 字节数不匹配的坏文件, 确认加载器会拒绝。这个 gate 只覆盖模型资产目录表, ↵
81   ↵ 不等于已经支持完整 LLaMA/Qwen 权重映射、分片合并或 mmap 执行; 后续真实模↵
82   ↵ 型加载必须在这个 manifest 上继续接 name mapping、shape verifier、quant ↵
83   ↵ descriptor 和 first token trace。
84
85 reference decoder trace:
86
87 ```bash
88 books/cpu-volume-3/build/lcqi-debug/labs/linux_cpu_inference/lcqi_trace
89 ```
90
91 输出 CSV 字段:
92
93 ```text
94 name,token_position,layer_id,shape,stride,dtype,layout,checksum,max_abs,values
95 ```
96
97 这条 trace 固定 prompt fixture `tok1 tok2 tok3`, tokenizer 输出完整 `tokens↵
98   ↵ =[0,1,2,3,4]`, 其中 `0/4` 是 BOS/EOS; decoder trace 会剥离 BOS/EOS 后进入↵

```

```

    ↪ tiny decoder。这样报告能同时固定 tokenizer 合同和 decoder 输入合同，避免
    ↪ 把二者混成一个不可解释的 token 数组。随后 trace 把 embedding、RMSNorm、Q/
    ↪ K/V、RoPE、KV cache slot、attention、SwiGLU、layer output、final norm、
    ↪ logits、argmax sampler 和 generated token 串成可回归检查的硬链路。`
    ↪ lcqi_tests` 会检查 tokenizer contract hash、关键 checkpoint 的 shape、
    ↪ stride、dtype、layout 和 logits golden，CTest 也会直接运行 `lcqi_trace`。
87
88 当前 tokenizer fixture 的关键输出：
89
90 ```text
91 tokenizer,0,-1,5,1,i32,contiguous,10,4,0;1;2;3;4
92 tokenizer_contract,0,-1,scalar,scalar,u64,fnv1a,13452902845388333734,0,tiny-
    ↪ chat-template-v1
93 ```
94
95 GPT-2 自研 reference smoke:
96
97 ```bash
98 books/cpu-volume-3/build/lcqi-debug/labs/linux_cpu_inference/lcqi_gpt2
99 make cpu3-gpt2-smoke
100 make cpu3-gpt2-benchmark-compare
101 ```
102
103 第一条命令不需要下载模型，直接运行仓库内 tiny GPT-2 fixture，输出固定 prompt
    ↪ ids、logits、predicted token 和 greedy generated ids。`lcqi_tests` 会覆盖
    ↪ tiny GPT-2 forward/generation、byte-level BPE 编码解码、F32/F16/BF16
    ↪ safetensors 读取、HuggingFace 风格 tiny GPT-2 checkpoint 加载和坏 shape
    ↪ 拒绝。
104
105 第二条命令显式依赖网络，只下载 `openai-community/gpt2` 的 `config.json`、`vocab-
    ↪ .json`、`merges.txt`、`model.safetensors` 到 `~/cache/lcqi-gpt2-smoke/`
    ↪ `openai-community--gpt2`，模型文件不进入 Git。脚本会校验四个文件的 bytes
    ↪ 和 SHA256，构建 `lcqi_gpt2`，运行：
106
107 ```text
108 Hello, my name is
109 ```
110
111 当前验证报告写到：
112
113 ```text
114 books/cpu-volume-3/results/lcqi-gpt2-smoke.txt
115 ```
116
117 最近一次本机报告的关键行是：
118
119 ```text
120 prompt_ids 15496 11 616 1438 318
121 generated_ids 15496 11 616 1438 318 1757
122 generated_text Hello, my name is John
123 validation=PASS
124 ```

```

```

125
126 这说明 LCQI 自研 C++ reference path 已经能跑标准 GPT-2 safetensors 的 tokenizer↵
    ↪ 、权重导入、forward 和 greedy generation。默认真实模型路径现在使用 `↵
    ↪ cached_kv`，也可以用 `--engine full` 复现旧的 full-prefix 重算路径：
127
128 ```bash
129 books/cpu-volume-3/build/lcqi-release/labs/linux_cpu_inference/lcqi_gpt2 \
130 --benchmark --engine cached \
131 ~/.cache/lcqi-gpt2-smoke/openai-community--gpt2 \
132 "Hello, my name is" 4
133 ```
134
135 `make cpu3-gpt2-benchmark-compare` 会生成：
136
137 ```text
138 books/cpu-volume-3/results/lcqi-gpt2-benchmark-compare.txt
139 ```
140
141 最近一次本机 aarch64/Qualcomm 报告中，LCQI full-prefix 在 GPT-2 F32 上生成 4 个↵
    ↪ token 的生成阶段为 `2522.81 ms`，约 `1.586 token/s`；LCQI cached-KV 同样↵
    ↪ 输出 `Hello, my name is John. I'm`，生成阶段为 `826.152 ms`，约 `4.842 ↵
    ↪ token/s`。cached-KV 的第一个新 token 来自 prefill 最后一步，后续 decode ↵
    ↪ 实际执行 3 个 forward step，耗时 `309.524 ms`，约 `9.692 step/s`。同机 ↵
    ↪ llama.cpp 使用 SmolLM2-135M-Instruct Q4_K_M 的 `llama-bench -p 5 -n 4` 作↵
    ↪ 为成熟引擎参照，prefill `76.41 token/s`、decode `27.83 token/s`。这个对照↵
    ↪ 不是同模型质量排名：LCQI 跑的是 GPT-2 F32 reference path，llama.cpp 跑的↵
    ↪ 是 GGUF 量化 small instruct 模型。它揭示的是工程差距：LCQI 已经消除了“每↵
    ↪ 步重算前缀”的主要算法错误，但还缺 packed weight、量化 matmul、SIMD/SVE/↵
    ↪ NEON kernel、线程并行、mmap/lazy load、采样器和成熟的内存调度。
142
143 它仍然不是生产级推理引擎：当前 GPT-2 路径还没有把 logits 和 Transformers/↵
    ↪ PyTorch 逐数值对齐成外部 golden 报告，也没有 GGUF 自研导入、分页 KV cache↵
    ↪ 、量化权重执行和高性能多线程 kernel。
144
145 7B ledger：
146
147 ```bash
148 books/cpu-volume-3/build/lcqi-debug/labs/linux_cpu_inference/lcqi_ledger \
149 > books/cpu-volume-3/results/lcqi-sevenb-ledger-2026-06-24.csv
150 ```
151
152 输出 CSV 字段：
153
154 ```text
155 section,item,value,unit,notes
156 ```
157
158 这份账本固定 `L=32,H=4096,n_q=32,n_kv=8,head_dim=128,context=4096`，报告 decode↵
    ↪ FLOPs、KV 容量、fp16/int8/int4 每 token 字节和不同有效带宽下的 tokens/s ↵
    ↪ 上限。CTest 会检查 `int4_at_100_gbps=23.602 tokens/s` 这一行，防止账本公↵
    ↪ 式漂移。
159

```

```

160 serving SLO probe:
161
162 ```bash
163 books/cpu-volume-3/build/lcqi-debug/labs/linux_cpu_inference/lcqi_serving_probe↵
164 ↵ \
164 > books/cpu-volume-3/results/lcqi-serving-slo-probe-2026-06-24.csv
165 ```
166
167 输出 CSV 字段:
168
169 ```text
170 row_type,request_id,accepted,rejection_reason,prompt_tokens,reserved_tokens,↵
171 ↵ kv_reserved_mib,queue_wait_ms,prefill_ms,ttft_ms,decode_step_p95_ms,↵
172 ↵ tpot_ms,completion_tokens,finish_reason,cancel_reclaim_ms,metric_name,↵
173 ↵ metric_value
174 ```
175
176 probe 固定四个请求：短请求、RAG 请求、取消请求和超长请求。它演示 admission 先按↵
177 ↵ `prompt_tokens + max_new_tokens` 预留 KV，取消请求释放预算，超长请求返回↵
178 ↵ `memory_budget_exceeded`，并输出 `ttft_p95_ms`、`queue_wait_p95_ms`、↵
179 ↵ `rejection_rate` 等服务化指标。
180
181 真实 small 开源模型 smoke:
182
183 ```bash
184 make cpu3-smolllm2-smoke
185 ```
186
187 这条命令显式依赖网络和外部构建，因此不放进默认 CTest。脚本会把 llama.cpp 和模型↵
188 ↵ 文件缓存到 `~/.cache/lcqi-smolllm2-small-smoke`，模型不进入 Git。固定模型↵
189 ↵ 是 `HuggingFaceTB/SmolLM2-135M-Instruct` 的 GGUF 量化版本，来源仓库为 `↵
190 ↵ bartowski/SmolLM2-135M-Instruct-GGUF`，文件为 `SmolLM2-135M-Instruct-↵
191 ↵ Q4_K_M.gguf`。脚本会校验文件大小 `105454432` 字节和 SHA256:
192
193 ```text
194 2e8040ceae7815abe0dcb3540b9995ead1fa0d2ca9e797d0a635ae4433c68c2d
195 ```
196
197 脚本构建固定 commit 的 `llama-simple-chat`，用 `-ngl 0` 强制 CPU 路径，输入短↵
198 ↵ prompt，检查 assistant response 非空，并生成:
199
200 ```text
201 books/cpu-volume-3/results/lcqi-smolllm2-small-model-smoke.txt
202 ```
203
204 这份报告只能证明真实 GGUF small 模型已经完成加载、tokenizer/chat template、↵
205 ↵ prefill、decode 和文本输出。它不能证明 LCQI 自研 C++ 引擎已经支持完整↵
206 ↵ GGUF、不能证明回答质量正确，也不能替代 tiny fixture 的逐层 golden trace。↵
207 ↵ tiny fixture 继续负责可手算语义和回归定位；SmolLM2 smoke 负责把“真实开源↵
208 ↵ 模型能在本机跑起来”变成可复查证据。

```

tiny MLP 模型文件是最短闭环的资产。它决定输入维度、hidden 维度、输出维度、权重量化

值、scale 和 bias；测试中的 logits golden 依赖这份文件。

**Listing 1.9** models/tiny\_mlp\_i8.txt

```

1 LCQI_MODEL_V1
2 input_size 4
3 hidden_size 3
4 output_size 2
5 hidden_scale 0.1
6 hidden_weights
7 5 -3 2 1
8 -4 6 1 -2
9 3 2 -5 4
10 hidden_bias
11 0.1 -0.2 0.0
12 output_scale 0.05
13 output_weights
14 4 -6 2
15 -3 5 8
16 output_bias
17 -0.5 0.4

```

tiny tokenizer fixture 固定 BOS、EOS、UNK、词表和模板 hash。tokenizer 合同如果漂移，后面的 decoder trace 和测试都会被污染。

**Listing 1.10** models/tiny\_tokenizer.txt

```

1 LCQI_TOKENIZER_V1
2 bos_id 0
3 eos_id 4
4 unk_id -1
5 chat_template_hash tiny-chat-template-v1
6 vocab_size 5
7 vocab
8 0 <bos>
9 1 tok1
10 2 tok2
11 3 tok3
12 4 <eos>

```

## 为什么要先讲调用链

实践与代码卷如果只把 **\*.cpp** 和 **\*.hpp** 全部贴出来，读者仍然不知道从哪里下手。LCQI 的代码同时包含 tiny MLP、int8 kernel、tokenizer、safetensors、reference decoder、GPT-2、benchmark、trace、ledger、serving probe 和外部模型脚本；这些文件不是平铺关系，而是几条调用链。先讲调用链，再给完整源码，读者才能把每个文件放到正确位置。

### 核心思想

读 LCQI 的顺序不是“哪个文件长先读哪个”，而是“入口、合同、资产、执行、证据”。入口说明目标怎样构建，合同说明类型怎样连接，资产说明模型从哪里来，执行说明 token 或张量怎样流动，证据说明测试和报告守住了哪些边界。



## 一、全工程入口：哪些文件决定能不能复现

| 文件         | 负责什么                                                                         | 读源码时要确认什么                                                                                         |
|------------|------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------|
| 第三册README  | 第三册 AI 计算工程路线、当前 LCQI 能力、GPT-2 smoke 和结果报告位置。                                | 能力声明必须能在 <code>labs/linux_cpu_inference</code> 、 <code>tools</code> 和 <code>results</code> 中找到证据。 |
| 第三册CMake   | 第三册根 CMake，负责把 <code>linux_cpu_inference</code> 接入。                          | 根 CMake 应保持薄层；真正 target 定义必须在 lab 内部，避免入口文件承担实现细节。                                                |
| 本卷CMake    | 本卷复用同一份 LCQI lab。                                                            | 本卷不复制源码，只把同一工程接入自己的测试构建，防止两份代码漂移。                                                                 |
| 本卷Makefile | 本卷的书籍构建入口。                                                                   | 它构建 PDF/EPUB，也提供 <code>test</code> 目标去构建第三册 LCQI；实践与代码卷不能只排版不验证。                                  |
| LCQICMake  | <code>lcqi</code> 库、CLI、benchmark、trace、GPT-2、ledger、serving probe、CTest 目标。 | 读这里能看到哪些文件进入库，哪些是可执行程序，哪些依赖可选外部库。                                                                 |
| LCQIREADME | 本地构建、tiny 推理、benchmark、trace、GPT-2、SmolLM2 smoke 的命令。                        | 先按 README 跑命令，再读源码；否则不知道某个函数是否真的能从命令行触达。                                                          |

上表里的短标签对应下面这些真实路径：

- `books/cpu-volume-3/README.md`
- `books/cpu-volume-3/CMakeLists.txt`
- `books/cpu-volume-3-practice/CMakeLists.txt`
- `books/cpu-volume-3-practice/Makefile`
- `books/cpu-volume-3/labs/linux_cpu_inference/CMakeLists.txt`
- `books/cpu-volume-3/labs/linux_cpu_inference/README.md`

这些入口文件回答“代码在哪里”和“怎样运行”。读者进入本卷代码走读章时，先不要直接跳到 `gpt2_reference.cpp`。先确认 `lcqi` 静态库包含 `model.cpp`、`inference.cpp`、`int8_kernels.cpp`、`reference_decoder.cpp`、`tokenizer.cpp`、`safetensors.cpp` 和 `gpt2_reference.cpp`；再确认 `lcqi_cli`、`lcqi_bench`、`lcqi_trace`、`lcqi_gpt2`、`lcqi_ledger`、`lcqi_serving_probe` 和 `lcqi_tests` 分别从哪条路径进入。

## 二、Tiny MLP 调用链：最短闭环怎样跑通

Tiny MLP 是最小推理闭环，用来证明“模型资产、shape 检查、int8 linear、激活函数、logits、argmax、CLI 输出”可以完整连起来。它不是大模型能力本身，但它是后面所有复杂路径的可手算底座。

| 文件                                  | 关键类型或函数                                                                                   | 这段代码的意思                                                   |
|-------------------------------------|-------------------------------------------------------------------------------------------|-----------------------------------------------------------|
| <code>models/tiny_mlp_i8.txt</code> | 输入维度、hidden 维度、输出维度、int8 权重、scale、bias。                                                   | 这是模型资产。测试中的 <code>golden logits</code> 来自这里，资产改动必须带着测试改动。 |
| <code>include/lcqi/model.hpp</code> | <code>QuantizedLinearLayer</code> 、 <code>TinyMlpModel</code> 、 <code>load_model</code> 。 | 定义 tiny MLP 如何在内存中表示；只负责模型结构和加载入口，不做推理。                   |
| <code>src/model.cpp</code>          | <code>read_i8_vector</code> 、 <code>validate_layer</code> 、 <code>load_model</code> 。     | 解析文本模型，拒绝坏 key、坏维度、权重数量不匹配和非法 scale，把错误挡在执行前。             |



|                                         |                                                                                                                  |                                                                                                |
|-----------------------------------------|------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------|
| <code>include/lcqi/inference.hpp</code> | <code>InferenceResult</code> 、 <code>linear_i8</code> 、 <code>relu_inplace</code> 、 <code>run_inference</code> 。 | 定义推理输出要包含 <code>hidden</code> 、 <code>logits</code> 和 <code>predicted class</code> ，方便测试定位错误层。 |
| <code>src/inference.cpp</code>          | 输入 <code>shape</code> 检查、两层 <code>int8 linear</code> 、 <code>ReLU</code> 、 <code>argmax</code> 。                 | 最短执行路径：模型加输入变成 <code>logits</code> 。这里的代码应保持清楚，不能混入 <code>CLI</code> 或文件解析。                    |
| <code>src/main.cpp</code>               | 解析命令行输入并打印结果。                                                                                                    | <code>CLI</code> 是薄层；它只接模型路径和输入，不复制 <code>run_inference</code> 逻辑。                             |

这条链路按下面顺序读：`main()` 取得模型路径和输入向量，调用 `load_model()` 把 `tiny_mlp_i8.txt` 变成 `TinyMlpModel`，再调用 `run_inference()`。在 `run_inference()` 内部，第一层 `linear_i8()` 把 `float` 输入和 `int8` 权重相乘，乘以 `scale` 后加 `bias`；`relu_inplace()` 把负 `hidden` 截断；第二层 `linear_i8()` 得到 `logits`；最后 `argmax()` 给出 `predicted class`。

读这个路径时要抓住边界条件。输入长度不等于 `input_size` 要立即拒绝；权重数量不是 `rows*columns` 要在加载时拒绝；`logits` 为空不能取 `argmax`。这些错误如果拖到 `SIMD kernel` 或 `benchmark` 里才暴露，定位成本会高很多。

### 三、Int8 Kernel 调用链：reference、packed、AVX2 和 benchmark 怎样对照

| 文件                                         | 关键代码                                                                                                                                              | 这段代码的意思                                                                                                                                             |
|--------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>include/lcqi/int8_kernels.hpp</code> | <code>PackedLinearI8</code> 、 <code>KernelBenchmarkCase</code> 、 <code>benchmark_linear_i8_case</code> 。                                          | 定义 <code>kernel</code> 对标口径： <code>shape</code> 、 <code>repeat</code> 、 <code>backend</code> 、平均耗时、 <code>checksum</code> 、 <code>max diff</code> 。 |
| <code>src/int8_kernels.cpp</code>          | <code>pack_linear_i8</code> 、 <code>linear_i8_packed</code> 、 <code>linear_i8_packed_with_workspace</code> 、 <code>deterministic fixture</code> 。 | <code>scalar</code> 和 <code>packed</code> 路径是正确性基线， <code>workspace</code> 路径减少重复分配， <code>fixture</code> 让 <code>benchmark</code> 和测试可复现。          |
| <code>src/int8_kernels_avx2.cpp</code>     | <code>linear_i8_packed_avx2_available</code> 、 <code>linear_i8_packed_avx2</code> 。                                                               | 平台优化路径。必须先判断机器能力，再执行 <code>AVX2/FMA</code> ；不可用时报告 <code>unavailable</code> ，而不是假装测过。                                                               |
| <code>src/bench.cpp</code>                 | <code>shape</code> 列表、 <code>repeat</code> 参数、 <code>backend</code> 输出。                                                                           | <code>benchmark</code> 程序只负责测 <code>kernel</code> 入口。性能结论必须和 <code>max_abs_diff</code> 、 <code>CPU</code> 、编译器、构建类型一起报告。                            |

`kernel` 文件要按“可信基线到优化路径”的顺序读。先看 `deterministic layer` 和 `input` 怎么生成，保证每次运行输入相同；再看 `linear_i8_packed()` 怎样把 `int8` 权重、`scale`、`bias` 和 `float` 输入算成输出；再看 `workspace` 路径怎样复用临时缓冲；最后看 `AVX2` 路径怎样在支持平台上替换内层循环。测试用 `packed` 与 `scalar` 的 `max_abs_diff` 对照，`benchmark` 再比较耗时。

### 四、Tokenizer、Safetensors 与 Reference Decoder：真实模型前的三道门

| 文件                                     | 关键类型或函数                                                                           | 这段代码的意思                                                                            |
|----------------------------------------|-----------------------------------------------------------------------------------|------------------------------------------------------------------------------------|
| <code>models/tiny_tokenizer.txt</code> | <code>BOS</code> 、 <code>EOS</code> 、 <code>UNK</code> 、词表、模板 <code>hash</code> 。 | 固定 <code>tokenizer</code> 合同，避免 <code>prompt</code> 到 <code>token id</code> 的规则漂移。 |

|                                                                                          |                                                                                                                                 |                                                                        |
|------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------|
| <code>include/lcqi/tokenizer.hpp</code> 与 <code>src/tokenizer.cpp</code>                 | <code>TokenizerModel</code> 、 <code>load_tokenizer</code> 、 <code>encode_prompt</code> 、 <code>tokenizer_contract_hash</code> 。 | 把 prompt 切成稳定 token id；未知 token 是否允许，必须由 UNK 合同决定。                     |
| <code>include/lcqi/safetensors.hpp</code> 与 <code>src/safetensors.cpp</code>             | <code>SafeTensorManifest</code> 、 <code>SafeTensorFile</code> 、manifest parser、dtype 转换。                                        | 解析模型字节之前先验证 header、dtype、shape、offset 和 payload 边界，防止坏资产进入数学层。         |
| <code>include/lcqi/reference_decoder.hpp</code> 与 <code>src/reference_decoder.cpp</code> | embedding、RMSNorm、RoPE、KV cache、attention、SwiGLU、trace checkpoint。                                                              | Tiny decoder-only reference path。它用 checkpoint 解释每一步张量状态，不只返回最终 token。 |
| <code>src/trace.cpp</code>                                                               | trace CSV 输出。                                                                                                                   | 把 reference decoder 的 checkpoint 打印成可复查表格，定位错误发生在哪个阶段。                 |

这三道门分别解决三个问题。Tokenizer 解决“人写的 prompt 如何变成 token id”；Safetensors 解决“磁盘里的模型字节如何变成带 shape 的张量”；Reference Decoder 解决“token 和权重如何经过 decoder-only 层生成 logits”。任何一层含糊，后面的 GPT-2 路径都会变成看似能跑、实则无法定位的黑盒。

## 五、GPT-2 调用链：从 HuggingFace 资产到生成 token

GPT-2 路径是本卷代码走读里最长的一条链，因为它要处理真实模型格式。读 `gpt2_reference.cpp` 时按模块读，不要从第一行一路滚到最后。

| 文件                                               | 关键代码                                                                                                                                     | 这段代码的意思                                                                               |
|--------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------|
| <code>include/lcqi/gpt2_reference.hpp</code>     | <code>Gpt2Config</code> 、 <code>Gpt2ReferenceModel</code> 、 <code>Gpt2KvCache</code> 、 <code>Gpt2Tokenizer</code> 、forward/generate API。 | 公开 GPT-2 reference path 的合同：配置、权重、KV cache、tokenizer、full-prefix 和 cached generation。 |
| <code>src/gpt2_reference.cpp</code>              | JSON parser、byte-level BPE、safetensors loader、LayerNorm、QKV、causal attention、GELU、lm head、KV cache。                                      | 正确性优先实现。它把 HuggingFace GPT-2 权重名、Conv1D 转置、tokenizer 和 forward 规则全部显式写出来。             |
| <code>src/gpt2_cli.cpp</code>                    | tiny mode、真实模型目录模式、generation、benchmark 输出。                                                                                              | 命令行入口。无参数跑 tiny fixture；传模型目录时加载真实 GPT-2 资产；benchmark 对比 full-prefix 和 cached-KV。     |
| <code>tools/run_gpt2_smoke.py</code>             | 下载标准 GPT-2 资产，调用 <code>lcqi_gpt2</code> 。                                                                                                | 证明自研 safetensors、BPE、forward、KV cache 和 greedy generation 能跑真实开源模型。                   |
| <code>tools/run_gpt2_benchmark_compare.py</code> | 运行 LCQI full-prefix、cached-KV，并可附带成熟引擎参考。                                                                                                | 用同一模型口径比较算法差距和 kernel 差距，不把不同模型的速度混成一个结论。                                             |

GPT-2 代码按五段读。第一段是配置和权重加载：`load_gpt2_config()` 读 `config.json`，`load_gpt2_reference_model()` 从 safetensors 中取 embedding、每层 LayerNorm、QKV、MLP 和 final norm。HuggingFace GPT-2 的 Conv1D 权重形状和普通线性层方向不同，所以代码显式调用转置函数，不能跳过。

第二段是 tokenizer: `load_gpt2_tokenizer()` 读 `vocab.json` 和 `merges.txt`, `gpt2_encode()` 执行 byte-level BPE, `gpt2_decode()` 把 id 还原为文本。这里的重点是字节到 Unicode 映射、BPE merge rank 和 unknown token 边界。

第三段是 full-prefix forward: `run_gpt2_forward()` 每生成一个新 token 都用完整 token 前缀重新跑所有层。它慢, 但容易验证。流程是 token embedding 加 position embedding, 逐层执行 LayerNorm、packed QKV、causal attention、残差、MLP、残差, 最后 final LayerNorm 和 tied lm head 得到 logits。

第四段是 cached-KV forward: `Gpt2KvCache` 保存每层每个 position 的 key/value, `run_gpt2_forward_cached()` 每次只处理最新 token, 并让 attention 读取过去 cache。这里要特别看 `validate_written()`: 读取未写入 cache slot 必须报错, 否则生成错误会很难定位。

第五段是 greedy generation: `gpt2_generate_greedy()` 用 full-prefix 路径, `gpt2_generate_greedy_cached()` 用 KV cache 路径。两个函数都要检查 prompt 非空、`max_new_tokens` 非负、不会超过 `max_positions`, 并在生成 EOS 时停止。

六、报告、服务探针、外部对标和测试

| 文件                                             | 关键代码                                                                                          | 这段代码的意思                                                  |
|------------------------------------------------|-----------------------------------------------------------------------------------------------|----------------------------------------------------------|
| <code>src/ledger.cpp</code>                    | 7B 参数量、KV cache、bandwidth-only 上限。                                                            | 输出的是工程账本, 不是实测速度; 它帮助读者估算权重和 KV 内存压力。                    |
| <code>src/serving_probe.cpp</code>             | 短请求、RAG 请求、取消请求、超长请求, TTFT/TPOT 字段。                                                           | 把服务层边界变成字段: 哪些请求接受, 哪些拒绝, 取消后资源是否回收。                     |
| <code>src/onednn_bench.cpp</code>              | oneDNN matmul benchmark。                                                                      | 可选外部库对标入口。没有 oneDNN 环境时不能声称已验证。                          |
| <code>tools/run_smolllm2_small_smoke.py</code> | 下载 SmolLM2-135M-Instruct GGUF, 构建固定 commit 的 llama.cpp smoke。                                 | 这是成熟外部引擎上的真实 small 模型 smoke, 用来建立外部参照, 不代表 LCQI 已具备同等能力。 |
| <code>tools/run_lcqi_benchmark.sh</code>       | 收集 LCQI 本地 benchmark。                                                                         | 在固定机器上重复运行同一组 shape, 形成可复查 CSV 或文本报告。                    |
| <code>tests/lcqi_tests.cpp</code>              | tiny MLP、tokenizer、safetensors、GPT-2 tiny、HF-style checkpoint、bad shape、decoder、trace、kernel。 | 这是本卷最重要的代码证据。新增功能如果没有进入这里, 就不能写成已经守住的能力。                 |

读测试文件时, 不要只看最后是否返回 0。它里面的 fixture 本身就是源码讲解的一部分: 测试会手写 safetensors header, 构造 F16/BF16/F32 payload, 生成 tiny GPT-2 权重, 验证 bad shape 会被拒绝, 比较 packed int8 kernel 和 scalar 路径。这些测试告诉读者“代码真正承诺了什么”。如果一个能力只出现在 README 或 CLI 输出里, 却没有进入测试或 smoke 报告, 就只能算实验入口, 不能算稳定能力。

源码阅读任务

读完本章后, 再去读相邻主题中的完整源码。每看到一个文件, 都把它归到五类之一: 入口、合同、资产解析、执行路径、证据。归不进去的文件, 就说明工程边界需要重新解释或重构。

## 第2章实践：张量布局、地址流与第一个 MatVec

本章对应原书中的第三册“张量布局、地址流与第一个矩阵向量内核”。不要把它当作概念复述，本章要把 TensorView、dtype、shape、stride、row-major、地址公式、bias、ReLU、checksum 放到同一个可运行项目里检查。贯穿代码仍然是 `books/cpu-volume-3/labs/linux_cpu_inference`；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `linear_i8`，核心命令或测试入口是 `lcqi_tests`。

**Listing 2.1** 第2章实践：张量布局、地址流与第一个 MatVec 的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  <- DCMMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_tests / linear_i8
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：手算一个 2 行 4 列权重矩阵的地址偏移。读者要先手写预期结果，再运行项目观察输出。动作是：把 `weight[out,k]` 展开成 `base+(out*K+k)*sizeof(i8)`，再对应到 `linear_i8` 的循环。如果输出里没有 `logit` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：张量不是裸数组；每个 kernel 入口都要能解释 dtype、shape、stride、scale 和输出缓冲这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 `books/cpu-volume-3/labs/linux_cpu_inference/README.md`、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交一页地址流账本：输入向量、hidden 权重、输出 logits 各自的 shape、连续性、读写次数和错误边界。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。



## 第 3 章实践：GEMM、packing 与 SIMD 证据

本章对应原书中的第三册“从矩阵向量乘走向 GEMM、packing 与 SIMD”。不要把它当作概念复述，本章要把 packed weight、scalar baseline、AVX2 backend、output block、shape sweep、反汇编和降级放到同一个可运行项目里检查。贯穿代码仍然是 [books/cpu-volume-3/labs/linux\\_cpu\\_inference](#)；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `output_block`，核心命令或测试入口是 `lcqi_bench`。

**Listing 3.1** 第 3 章实践：GEMM、packing 与 SIMD 证据的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  <- DCMMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_bench / output_block
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：同一组输入同时跑 `scalar`、`packed_scalar` 和 `avx2`。读者要先手写预期结果，再运行项目观察输出。动作是：比较 CSV 中 `backend`、`available`、`average_us`、`max_abs_diff` 和 `checksum`。如果输出里没有 `max_abs_diff` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：SIMD 路径只有在结果和 baseline 对齐后才有资格谈速度；`available=0` 也是证据，不是失败这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 [books/cpu-volume-3/labs/linux\\_cpu\\_inference/README.md](#)、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交一个 shape sweep 表，至少解释一个 SIMD 有收益、一个收益不明显、一个当前环境不可验证的场景。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。

## 第4章实践：int8 量化、累加器与重新量化

本章对应原书中的第三册“量化系统总览”和“int8 量化、累加器与重新量化”。不要把它当作概念复述，本章要把 int8 weight、float activation、scale、accumulator、requantize、oracle tolerance 放到同一个可运行项目里检查。贯穿代码仍然是 `books/cpu-volume-3/labs/linux_cpu_inference`；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `TinyMlpModel`，核心命令或测试入口是 `lcqi_tests`。

Listing 4.1 第4章实践：int8 量化、累加器与重新量化的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  <- DCMMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_tests / TinyMlpModel
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：构造一行 int8 权重 `[2, -1, 3, 0]` 和输入 `[1, 2, -1, 0.5]`。读者要先手写预期结果，再运行项目观察输出。动作是：先手算 int32/float accumulator，再和 `linear_i8` 输出比较。如果输出里没有 `scale` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：量化不是把 float 强转成 int8；必须记录 scale 作用位置、累加宽度、误差阈值和 golden 输出这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 `books/cpu-volume-3/labs/linux_cpu_inference/README.md`、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交量化合同卡：每个 tensor 的 dtype、scale 粒度、accumulator 类型、误差门限和拒绝条件。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。

# 代码走读：公共合同、Tiny MLP 与 Int8 Kernel

## 为什么先读头文件

头文件是系统对外承诺的最小合同。实现可以重写，kernel 可以替换，benchmark 可以补充平台路径，但头文件里暴露的数据结构、函数入口和返回字段决定了测试、CLI、工具脚本和章节讲解怎样调用这套工程。

### 核心思想

读 LCQI 时不要从最长的 `gpt2_reference.cpp` 开始。先读公共头文件，明确“模型如何表示”“推理返回什么”“trace 记录什么”“KV cache 怎样寻址”“tokenizer 的坏输入怎样处理”，再进入实现。这样读到长循环和 shape 检查时，知道它们在维护哪份合同。

## Tiny MLP 与推理入口

`model.hpp` 定义 tiny MLP 的模型对象、权重、bias、scale 和加载入口。它的责任是把文本模型资产解释成 C++ 对象，不负责执行推理。

Listing 4.2 `include/lcqi/model.hpp`

```

1 #pragma once
2
3 #include <stdint>
4 #include <filesystem>
5 #include <string>
6 #include <vector>
7
8 namespace lcqi {
9
10 struct QuantizedLinearLayer {
11     std::int32_t input_size = 0;
12     std::int32_t output_size = 0;
13     float scale = 1.0F;
14     std::vector<std::int8_t> weights;
15     std::vector<float> bias;
16 };
17
18 struct TinyMlpModel {
19     std::int32_t input_size = 0;
20     std::int32_t hidden_size = 0;
21     std::int32_t output_size = 0;
22     QuantizedLinearLayer hidden;
23     QuantizedLinearLayer output;
24 };
25
26 TinyMlpModel load_model(const std::filesystem::path& path);
27
28 } // namespace lcqi

```

`inference.hpp` 定义推理结果和 `run_inference` 入口。返回值包含 hidden、logits 和 pre-

dicted class，测试可以用这些字段定位是加载错、kernel 错，还是 argmax 错。

**Listing 4.3** include/lcqi/inference.hpp

```

1  #pragma once
2
3  #include <cstdint>
4  #include <span>
5  #include <vector>
6
7  #include <lcqi/model.hpp>
8
9  namespace lcqi {
10
11  struct InferenceResult {
12      std::vector<float> logits;
13      std::int32_t predicted_class = 0;
14  };
15
16  void linear_i8(const QuantizedLinearLayer& layer,
17               std::span<const float> input,
18               std::span<float> output);
19
20  void relu_inplace(std::span<float> values);
21
22  InferenceResult run_inference(const TinyMlpModel& model,
23                               std::span<const float> input);
24
25 } // namespace lcqi

```

## Int8 Kernel、Tokenizer 与 Safetensors

int8 kernel 头文件定义 packed linear、deterministic fixture、scalar/workspace/AVX2 路径和误差比较。注意 AVX2 是运行时能力，不是源码存在就等于当前机器已验证。

**Listing 4.4** include/lcqi/int8\_kernels.hpp

```

1  #pragma once
2
3  #include <cstdint>
4  #include <span>
5  #include <vector>
6
7  #include <lcqi/model.hpp>
8
9  namespace lcqi {
10
11  struct PackedLinearI8 {
12      std::int32_t input_size = 0;
13      std::int32_t output_size = 0;
14      std::int32_t output_block_size = 0;
15      float scale = 1.0F;
16      std::vector<std::int8_t> packed_weights;

```



```

17     std::vector<float> bias;
18 };
19
20 struct KernelBenchmarkCase {
21     std::int32_t input_size = 0;
22     std::int32_t output_size = 0;
23     std::int32_t output_block_size = 0;
24     std::int32_t repeat = 0;
25 };
26
27 struct KernelBackendBenchmark {
28     const char* name = "";
29     bool available = false;
30     double average_us = 0.0;
31     float max_abs_diff = 0.0F;
32     float checksum = 0.0F;
33 };
34
35 struct KernelBenchmarkResult {
36     KernelBenchmarkCase benchmark_case;
37     std::vector<KernelBackendBenchmark> backends;
38 };
39
40 PackedLinearI8 pack_linear_i8(const QuantizedLinearLayer& layer,
41                               std::int32_t output_block_size);
42
43 void linear_i8_packed(const PackedLinearI8& layer,
44                      std::span<const float> input,
45                      std::span<float> output);
46
47 void linear_i8_packed_with_workspace(const PackedLinearI8& layer,
48                                     std::span<const float> input,
49                                     std::span<float> output,
50                                     std::span<float> accumulator_workspace);
51
52 bool linear_i8_packed_avx2_available() noexcept;
53
54 void linear_i8_packed_avx2(const PackedLinearI8& layer,
55                            std::span<const float> input,
56                            std::span<float> output);
57
58 QuantizedLinearLayer make_deterministic_i8_layer(std::int32_t input_size,
59                                                  std::int32_t output_size,
60                                                  float scale);
61
62 std::vector<float> make_deterministic_input(std::int32_t input_size);
63
64 float max_abs_diff(std::span<const float> lhs, std::span<const float> rhs);
65
66 KernelBenchmarkResult benchmark_linear_i8_case(const KernelBenchmarkCase& ↵
67     ↵ benchmark_case);

```

```
68 } // namespace lcqi
```

tokenizer 头文件固定 BOS/EOS/UNK、词表和 template hash。prompt 进入模型前，必须先变成稳定 token id；未知 token 行为也必须被写进合同。

**Listing 4.5** include/lcqi/tokenizer.hpp

```
1 #pragma once
2
3 #include <cstdint>
4 #include <filesystem>
5 #include <string>
6 #include <string_view>
7 #include <vector>
8
9 namespace lcqi {
10
11 struct TokenEntry {
12     std::string text;
13     std::int32_t id = 0;
14 };
15
16 struct TokenizerModel {
17     std::int32_t bos_id = -1;
18     std::int32_t eos_id = -1;
19     std::int32_t unk_id = -1;
20     std::string chat_template_hash;
21     std::vector<TokenEntry> vocab;
22 };
23
24 [[nodiscard]] TokenizerModel load_tokenizer(const std::filesystem::path& path);
25
26 [[nodiscard]] std::vector<std::int32_t> encode_prompt(
27     const TokenizerModel& tokenizer,
28     std::string_view prompt);
29
30 [[nodiscard]] std::uint64_t tokenizer_contract_hash(const TokenizerModel& ↵
31     ↵ tokenizer);
32 } // namespace lcqi
```

safetensors 头文件固定 dtype、shape、offset 和 payload 边界。它负责模型字节怎样被解释，不应该和数学推理逻辑混在一起。

**Listing 4.6** include/lcqi/safetensors.hpp

```
1 #pragma once
2
3 #include <cstdint>
4 #include <filesystem>
5 #include <span>
6 #include <string>
7 #include <string_view>
```

```

8 #include <utility>
9 #include <vector>
10
11 namespace lcqi {
12
13 struct SafeTensorEntry {
14     std::string name;
15     std::string dtype;
16     std::vector<std::int64_t> shape;
17     std::uint64_t data_begin = 0;
18     std::uint64_t data_end = 0;
19
20     [[nodiscard]] std::uint64_t byte_size() const;
21 };
22
23 struct SafeTensorManifest {
24     std::uint64_t header_size = 0;
25     std::uint64_t data_start_offset = 0;
26     std::vector<std::pair<std::string, std::string>> metadata;
27     std::vector<SafeTensorEntry> tensors;
28
29     [[nodiscard]] const SafeTensorEntry* find_tensor(std::string_view name) ↵
        ↵ const;
30 };
31
32 struct SafeTensorFile {
33     std::filesystem::path path;
34     SafeTensorManifest manifest;
35     std::vector<std::uint8_t> bytes;
36 };
37
38 [[nodiscard]] SafeTensorManifest load_safetensors_manifest(
39     const std::filesystem::path& path);
40
41 [[nodiscard]] SafeTensorFile load_safetensors_file(const std::filesystem::path& ↵
        ↵ path);
42
43 [[nodiscard]] std::vector<float> read_safetensor_f32_tensor(
44     const SafeTensorFile& file,
45     std::string_view name);
46
47 [[nodiscard]] std::span<const std::uint8_t> read_safetensor_tensor_bytes(
48     const SafeTensorFile& file,
49     std::string_view name);
50
51 } // namespace lcqi

```

## Reference Decoder 与 GPT-2 合同

reference decoder 头文件定义张量、层配置、trace checkpoint、采样结果和 decode 函数。它是定位推理错误的主要合同，不只返回最终 token。

Listing 4.7 include/lcqi/reference\_decoder.hpp

```

1  #pragma once
2
3  #include <cstdint>
4  #include <span>
5  #include <string>
6  #include <vector>
7
8  namespace lcqi {
9
10 struct LinearWeightsF32 {
11     std::int32_t input_size = 0;
12     std::int32_t output_size = 0;
13     std::vector<float> weights;
14     std::vector<float> bias;
15 };
16
17 struct DecoderConfig {
18     std::int32_t hidden_size = 0;
19     std::int32_t query_heads = 0;
20     std::int32_t kv_heads = 0;
21     std::int32_t head_dim = 0;
22     std::int32_t intermediate_size = 0;
23     std::int32_t vocab_size = 0;
24     std::int32_t max_context = 0;
25     float rms_epsilon = 1.0e-5F;
26     float rope_base = 10000.0F;
27 };
28
29 struct DecoderLayerWeightsF32 {
30     std::vector<float> rms_attention_weight;
31     LinearWeightsF32 wq;
32     LinearWeightsF32 wk;
33     LinearWeightsF32 wv;
34     LinearWeightsF32 wo;
35     std::vector<float> rms_mlp_weight;
36     LinearWeightsF32 w_gate;
37     LinearWeightsF32 w_up;
38     LinearWeightsF32 w_down;
39 };
40
41 struct ReferenceDecoderModel {
42     DecoderConfig config;
43     std::vector<float> token_embedding;
44     std::vector<DecoderLayerWeightsF32> layers;
45     std::vector<float> final_norm_weight;
46     LinearWeightsF32 lm_head;
47 };
48
49 struct ReferenceDecodeResult {
50     std::vector<float> logits;
51     std::int32_t predicted_token = 0;

```

```

52 };
53
54 struct ReferenceTensorCheckpoint {
55     std::string name;
56     std::int32_t token_position = 0;
57     std::int32_t layer_id = 0;
58     std::vector<std::int32_t> shape;
59     std::vector<std::int32_t> stride;
60     std::string dtype;
61     std::string layout;
62     float checksum = 0.0F;
63     float max_abs = 0.0F;
64     std::vector<float> values;
65 };
66
67 struct ReferenceDecodeTraceResult {
68     ReferenceDecodeResult result;
69     std::vector<ReferenceTensorCheckpoint> checkpoints;
70 };
71
72 class ReferenceKVCache {
73 public:
74     ReferenceKVCache(const DecoderConfig& config, std::int32_t layer_count);
75
76     void append(std::int32_t layer_id,
77                std::int32_t model_position,
78                std::span<const float> key,
79                std::span<const float> value);
80
81     [[nodiscard]] std::span<const float> key(std::int32_t layer_id,
82                                              std::int32_t model_position,
83                                              std::int32_t kv_head) const;
84
85     [[nodiscard]] std::span<const float> value(std::int32_t layer_id,
86                                              std::int32_t model_position,
87                                              std::int32_t kv_head) const;
88
89     [[nodiscard]] std::int32_t layer_count() const noexcept;
90     [[nodiscard]] std::int32_t filled_tokens() const noexcept;
91
92 private:
93     [[nodiscard]] std::size_t base_offset(std::int32_t layer_id,
94                                           std::int32_t model_position,
95                                           std::int32_t kv_head) const;
96     void validate_address(std::int32_t layer_id,
97                          std::int32_t model_position,
98                          std::int32_t kv_head) const;
99
100     DecoderConfig config_;
101     std::int32_t layer_count_ = 0;
102     std::int32_t filled_tokens_ = 0;
103     std::vector<float> keys_;

```

```

104     std::vector<float> values_;
105     std::vector<std::uint8_t> written_;
106 };
107
108 void rms_norm(std::span<const float> input,
109              std::span<const float> weight,
110              float epsilon,
111              std::span<float> output);
112
113 void linear_f32(const LinearWeightsF32& layer,
114               std::span<const float> input,
115               std::span<float> output);
116
117 void apply_rope(std::span<float> heads,
118                std::int32_t head_count,
119                std::int32_t head_dim,
120                std::int32_t model_position,
121                float rope_base);
122
123 void attention_decode(const DecoderConfig& config,
124                     const ReferenceKVCache& cache,
125                     std::int32_t layer_id,
126                     std::int32_t model_position,
127                     std::span<const float> query,
128                     std::span<float> output);
129
130 ReferenceDecodeResult run_reference_decode(const ReferenceDecoderModel& model,
131                                         std::span<const std::int32_t> token_ids) {
132     ↪ token_ids);
133
134 ReferenceDecodeTraceResult run_reference_decode_with_trace(
135     const ReferenceDecoderModel& model,
136     std::span<const std::int32_t> token_ids);
137
138 ReferenceDecoderModel make_tiny_reference_decoder_model();
139 } // namespace lcqi

```

GPT-2 reference 头文件定义 GPT-2 配置、byte-level BPE tokenizer、HuggingFace safetensors 加载、forward、cached forward 和 greedy generation。它单独存在，是因为 GPT-2 使用 LayerNorm、绝对位置嵌入、packed QKV、GELU 和 tied lm head，不能直接塞进 LLaMA-like reference decoder。

**Listing 4.8** include/lcqi/gpt2\_reference.hpp

```

1 #pragma once
2
3 #include <cstdint>
4 #include <filesystem>
5 #include <span>
6 #include <string>
7 #include <string_view>
8 #include <unordered_map>
9 #include <vector>

```

```

10
11 namespace lcqi {
12
13 struct Gpt2Config {
14     std::int32_t hidden_size = 0;
15     std::int32_t head_count = 0;
16     std::int32_t layer_count = 0;
17     std::int32_t max_positions = 0;
18     std::int32_t vocab_size = 0;
19     std::int32_t intermediate_size = 0;
20     std::int32_t bos_token_id = -1;
21     std::int32_t eos_token_id = -1;
22     float layer_norm_epsilon = 1.0e-5F;
23     std::string activation_function = "gelu_new";
24 };
25
26 struct Gpt2LinearF32 {
27     std::int32_t input_size = 0;
28     std::int32_t output_size = 0;
29     std::vector<float> weights;
30     std::vector<float> bias;
31 };
32
33 struct Gpt2LayerWeightsF32 {
34     std::vector<float> ln_1_weight;
35     std::vector<float> ln_1_bias;
36     Gpt2LinearF32 c_attn;
37     Gpt2LinearF32 c_proj;
38     std::vector<float> ln_2_weight;
39     std::vector<float> ln_2_bias;
40     Gpt2LinearF32 c_fc;
41     Gpt2LinearF32 mlp_c_proj;
42 };
43
44 struct Gpt2ReferenceModel {
45     Gpt2Config config;
46     std::vector<float> token_embedding;
47     std::vector<float> position_embedding;
48     std::vector<Gpt2LayerWeightsF32> layers;
49     std::vector<float> final_ln_weight;
50     std::vector<float> final_ln_bias;
51     std::vector<float> lm_head_weight;
52     bool tie_lm_head_to_embedding = true;
53 };
54
55 struct Gpt2ForwardResult {
56     std::vector<float> logits;
57     std::int32_t predicted_token = 0;
58 };
59
60 class Gpt2KvCache {
61 public:

```



```

62     explicit Gpt2KvCache(const Gpt2Config& config);
63
64     void append(std::int32_t layer_id,
65                std::int32_t model_position,
66                std::span<const float> key,
67                std::span<const float> value);
68
69     [[nodiscard]] std::span<const float> key(std::int32_t layer_id,
70                                              std::int32_t model_position,
71                                              std::int32_t head) const;
72
73     [[nodiscard]] std::span<const float> value(std::int32_t layer_id,
74                                                std::int32_t model_position,
75                                                std::int32_t head) const;
76
77     [[nodiscard]] std::int32_t filled_tokens() const noexcept;
78     [[nodiscard]] std::size_t byte_size() const noexcept;
79
80 private:
81     [[nodiscard]] std::size_t base_offset(std::int32_t layer_id,
82                                           std::int32_t model_position,
83                                           std::int32_t head) const;
84     void validate_address(std::int32_t layer_id,
85                          std::int32_t model_position,
86                          std::int32_t head) const;
87     void validate_written(std::int32_t layer_id,
88                          std::int32_t model_position) const;
89
90     Gpt2Config config_;
91     std::int32_t filled_tokens_ = 0;
92     std::vector<float> keys_;
93     std::vector<float> values_;
94     std::vector<std::uint8_t> written_;
95 };
96
97 struct Gpt2Tokenizer {
98     std::int32_t bos_token_id = -1;
99     std::int32_t eos_token_id = -1;
100     std::unordered_map<std::string, std::int32_t> vocab;
101     std::vector<std::string> id_to_token;
102     std::unordered_map<std::string, std::int32_t> bpe_ranks;
103 };
104
105 [[nodiscard]] Gpt2Tokenizer load_gpt2_tokenizer(
106     const std::filesystem::path& vocab_json,
107     const std::filesystem::path& merges_txt,
108     std::int32_t bos_token_id,
109     std::int32_t eos_token_id);
110
111 [[nodiscard]] std::vector<std::int32_t> gpt2_encode(
112     const Gpt2Tokenizer& tokenizer,
113     std::string_view text);

```

```

114
115 [[nodiscard]] std::string gpt2_decode(
116     const Gpt2Tokenizer& tokenizer,
117     std::span<const std::int32_t> token_ids);
118
119 [[nodiscard]] Gpt2Config load_gpt2_config(const std::filesystem::path& ↵
    ↵ config_json);
120
121 [[nodiscard]] Gpt2ReferenceModel load_gpt2_reference_model(
122     const std::filesystem::path& config_json,
123     const std::filesystem::path& safetensors_path);
124
125 [[nodiscard]] Gpt2ReferenceModel load_gpt2_from_directory(
126     const std::filesystem::path& model_directory);
127
128 [[nodiscard]] Gpt2ForwardResult run_gpt2_forward(
129     const Gpt2ReferenceModel& model,
130     std::span<const std::int32_t> token_ids);
131
132 [[nodiscard]] Gpt2ForwardResult run_gpt2_forward_cached(
133     const Gpt2ReferenceModel& model,
134     Gpt2KvCache& cache,
135     std::int32_t token_id);
136
137 [[nodiscard]] std::vector<std::int32_t> gpt2_generate_greedy(
138     const Gpt2ReferenceModel& model,
139     std::span<const std::int32_t> prompt_token_ids,
140     std::int32_t max_new_tokens);
141
142 [[nodiscard]] std::vector<std::int32_t> gpt2_generate_greedy_cached(
143     const Gpt2ReferenceModel& model,
144     std::span<const std::int32_t> prompt_token_ids,
145     std::int32_t max_new_tokens);
146
147 [[nodiscard]] Gpt2ReferenceModel make_tiny_gpt2_reference_model();
148
149 } // namespace lcqi

```

## 最短闭环：从模型文件到 logits

tiny MLP 路径不是为了证明 LCQI 已经是完整大模型引擎，而是为了建立最短的、可手算的、可测试的推理闭环。读者要把模型加载、输入检查、int8 linear、ReLU、第二层 logits、argmax 和 CLI 输出连起来看。

## 模型加载与基础推理

**model.cpp** 负责解析 tiny 文本模型、权重、scale、bias 和 shape。坏格式、坏维度和坏数值都应该在这里被拒绝，不能拖到 kernel 内部才崩溃。

**Listing 4.9** src/model.cpp

```

1 #include <lcqi/model.hpp>
2

```

```

3  #include <fstream>
4  #include <limits>
5  #include <stdexcept>
6  #include <string>
7
8  namespace lcqi {
9  namespace {
10
11  constexpr std::int32_t LCQI_INT8_MIN_VALUE = -128;
12  constexpr std::int32_t LCQI_INT8_MAX_VALUE = 127;
13
14  std::string read_key(std::istream& input) {
15      std::string key;
16      if (!(input >> key)) {
17          throw std::runtime_error("unexpected end of model file");
18      }
19      return key;
20  }
21
22  void expect_key(std::istream& input, const std::string& expected) {
23      const std::string key = read_key(input);
24      if (key != expected) {
25          throw std::runtime_error("expected key '" + expected + "', got '" + key +
↳ + "'");
26      }
27  }
28
29  std::int32_t read_i32(std::istream& input, const std::string& key) {
30      expect_key(input, key);
31      std::int32_t value = 0;
32      if (!(input >> value)) {
33          throw std::runtime_error("invalid int32 value for key '" + key + "'");
34      }
35      return value;
36  }
37
38  float read_float(std::istream& input, const std::string& key) {
39      expect_key(input, key);
40      float value = 0.0F;
41      if (!(input >> value)) {
42          throw std::runtime_error("invalid float value for key '" + key + "'");
43      }
44      return value;
45  }
46
47  std::vector<std::int8_t> read_i8_vector(std::istream& input,
48                                          const std::string& key,
49                                          std::int32_t count) {
50      expect_key(input, key);
51      if (count < 0) {
52          throw std::runtime_error("negative vector size for key '" + key + "'");
53      }

```

[illegible]

```

102     if (rows <= 0 || columns <= 0) {
103         throw std::runtime_error(name + " layer has invalid shape");
104     }
105     const std::int64_t count =
106         static_cast<std::int64_t>(rows) * static_cast<std::int64_t>(columns);
107     if (count > static_cast<std::int64_t>(std::numeric_limits<std::int32_t>::max()
↵ max())) {
108         throw std::runtime_error(name + " layer element count is too large");
109     }
110     return static_cast<std::int32_t>(count);
111 }
112
113 } // namespace
114
115 TinyMlpModel load_model(const std::filesystem::path& path) {
116     std::ifstream input(path);
117     if (!input) {
118         throw std::runtime_error("cannot open model file: " + path.string());
119     }
120
121     expect_key(input, "LCQI_MODEL_V1");
122
123     TinyMlpModel model;
124     model.input_size = read_i32(input, "input_size");
125     model.hidden_size = read_i32(input, "hidden_size");
126     model.output_size = read_i32(input, "output_size");
127
128     const std::int32_t hidden_weight_count =
129         checked_element_count(model.input_size, model.hidden_size, "hidden");
130     const std::int32_t output_weight_count =
131         checked_element_count(model.hidden_size, model.output_size, "output");
132
133     model.hidden.input_size = model.input_size;
134     model.hidden.output_size = model.hidden_size;
135     model.hidden.scale = read_float(input, "hidden_scale");
136     model.hidden.weights = read_i8_vector(
137         input,
138         "hidden_weights",
139         hidden_weight_count);
140     model.hidden.bias = read_float_vector(input, "hidden_bias", model.↵
↵ hidden_size);
141
142     model.output.input_size = model.hidden_size;
143     model.output.output_size = model.output_size;
144     model.output.scale = read_float(input, "output_scale");
145     model.output.weights = read_i8_vector(
146         input,
147         "output_weights",
148         output_weight_count);
149     model.output.bias = read_float_vector(input, "output_bias", model.↵
↵ output_size);
150

```

```

151     validate_layer(model.hidden, "hidden");
152     validate_layer(model.output, "output");
153     return model;
154 }
155
156 } // namespace lcqi

```

`inference.cpp` 执行 int8 linear、ReLU、第二层 logits 和 predicted class。它是系统中最短的“模型资产到输出”路径。

Listing 4.10 `src/inference.cpp`

```

1  #include <lcqi/inference.hpp>
2
3  #include <algorithm>
4  #include <cmath>
5  #include <stdexcept>
6
7  namespace lcqi {
8  namespace {
9
10 void validate_input_size(std::span<const float> input, std::int32_t expected) {
11     if (input.size() != static_cast<std::size_t>(expected)) {
12         throw std::runtime_error("input size does not match model shape");
13     }
14 }
15
16 std::int32_t argmax(std::span<const float> values) {
17     if (values.empty()) {
18         throw std::runtime_error("cannot take argmax of empty logits");
19     }
20     std::int32_t best_index = 0;
21     float best_value = values[0];
22     for (std::int32_t i = 1; i < static_cast<std::int32_t>(values.size()); ++i)↵
↵ {
23         const float candidate = values[static_cast<std::size_t>(i)];
24         if (candidate > best_value) {
25             best_value = candidate;
26             best_index = i;
27         }
28     }
29     return best_index;
30 }
31
32 } // namespace
33
34 void linear_i8(const QuantizedLinearLayer& layer,
35               std::span<const float> input,
36               std::span<float> output) {
37     if (input.size() != static_cast<std::size_t>(layer.input_size)) {
38         throw std::runtime_error("linear_i8 input size mismatch");
39     }
40     if (output.size() != static_cast<std::size_t>(layer.output_size)) {

```

```

41     throw std::runtime_error("linear_i8 output size mismatch");
42 }
43
44 for (std::int32_t out = 0; out < layer.output_size; ++out) {
45     const std::size_t row_base =
46         static_cast<std::size_t>(out) * static_cast<std::size_t>(layer.↵
↵ input_size);
47     float acc = layer.bias[static_cast<std::size_t>(out)];
48     for (std::int32_t in = 0; in < layer.input_size; ++in) {
49         const std::int8_t weight =
50             layer.weights[row_base + static_cast<std::size_t>(in)];
51         acc += static_cast<float>(weight) * layer.scale *
52             input[static_cast<std::size_t>(in)];
53     }
54     output[static_cast<std::size_t>(out)] = acc;
55 }
56 }
57
58 void relu_inplace(std::span<float> values) {
59     for (float& value : values) {
60         value = std::max(0.0F, value);
61     }
62 }
63
64 InferenceResult run_inference(const TinyMlpModel& model,
65                               std::span<const float> input) {
66     validate_input_size(input, model.input_size);
67
68     std::vector<float> hidden(static_cast<std::size_t>(model.hidden_size), 0.0F↵
↵ );
69     std::vector<float> logits(static_cast<std::size_t>(model.output_size), 0.0F↵
↵ );
70
71     linear_i8(model.hidden, input, hidden);
72     relu_inplace(hidden);
73     linear_i8(model.output, hidden, logits);
74
75     InferenceResult result;
76     result.predicted_class = argmax(logits);
77     result.logits = std::move(logits);
78     return result;
79 }
80
81 } // namespace lcqi

```

`main.cpp` 把模型路径和输入接到 `run_inference`，并打印结果。CLI 应保持薄层，不复制模型解析和推理逻辑。

Listing 4.11 src/main.cpp

```

1 #include <lcqi/inference.hpp>
2 #include <lcqi/model.hpp>
3

```

```

4 #include <chrono>
5 #include <cstdint>
6 #include <cstdlib>
7 #include <filesystem>
8 #include <iostream>
9 #include <span>
10 #include <stdexcept>
11 #include <string>
12 #include <vector>
13
14 namespace {
15
16 constexpr std::int32_t LCQI_DEFAULT_REPEAT = 1000;
17 constexpr int LCQI_EXPECTED_ARGC_MIN = 2;
18 constexpr double LCQI_SECONDS_TO_MICROSECONDS = 1000000.0;
19
20 std::vector<float> parse_input(int argc, char** argv, std::int32_t ↵
    ↵ expected_size) {
21     if (argc < LCQI_EXPECTED_ARGC_MIN + expected_size) {
22         throw std::runtime_error("usage: lcqi_cli <model.txt> <input floats...>↵
    ↵ [repeat]");
23     }
24     std::vector<float> input(static_cast<std::size_t>(expected_size));
25     for (std::int32_t i = 0; i < expected_size; ++i) {
26         input[static_cast<std::size_t>(i)] =
27             std::stof(argv[LCQI_EXPECTED_ARGC_MIN + i]);
28     }
29     return input;
30 }
31
32 std::int32_t parse_repeat(int argc, char** argv, std::int32_t input_size) {
33     const int repeat_index = LCQI_EXPECTED_ARGC_MIN + input_size;
34     if (argc <= repeat_index) {
35         return LCQI_DEFAULT_REPEAT;
36     }
37     const std::int32_t repeat = static_cast<std::int32_t>(std::stoi(argv[↵
    ↵ repeat_index]));
38     if (repeat <= 0) {
39         throw std::runtime_error("repeat must be positive");
40     }
41     return repeat;
42 }
43
44 } // namespace
45
46 int main(int argc, char** argv) {
47     try {
48         if (argc < LCQI_EXPECTED_ARGC_MIN) {
49             throw std::runtime_error("usage: lcqi_cli <model.txt> <input floats↵
    ↵ ...> [repeat]");
50         }
51

```



```

52     const std::filesystem::path model_path = argv[1];
53     const lcqi::TinyMlpModel model = lcqi::load_model(model_path);
54     const std::vector<float> input = parse_input(argc, argv, model.↵
↵ input_size);
55     const std::int32_t repeat = parse_repeat(argc, argv, model.input_size);
56
57     const lcqi::InferenceResult result = lcqi::run_inference(model, input);
58     std::cout << "predicted_class " << result.predicted_class << "\n";
59     std::cout << "logits";
60     for (const float value : result.logits) {
61         std::cout << ' ' << value;
62     }
63     std::cout << "\n";
64
65     const auto begin = std::chrono::steady_clock::now();
66     std::int32_t checksum = 0;
67     for (std::int32_t i = 0; i < repeat; ++i) {
68         checksum += lcqi::run_inference(model, input).predicted_class;
69     }
70     const auto end = std::chrono::steady_clock::now();
71     const std::chrono::duration<double> elapsed = end - begin;
72     const double average_us =
73         elapsed.count() * LCQI_SECONDS_TO_MICROSECONDS / static_cast<double>↵
↵ >(repeat);
74     std::cout << "repeat " << repeat << "\n";
75     std::cout << "average_us " << average_us << "\n";
76     std::cout << "checksum " << checksum << "\n";
77     return EXIT_SUCCESS;
78 } catch (const std::exception& error) {
79     std::cerr << "error: " << error.what() << "\n";
80     return EXIT_FAILURE;
81 }
82 }

```

## Int8 Kernel 基础实现

`int8_kernels.cpp` 包含 scalar、packed、workspace、deterministic fixture 和误差比较。它是 AVX2 路径的 reference 对照，也是 benchmark 的正确性底座。

Listing 4.12 `src/int8_kernels.cpp`

```

1 #include <lcqi/int8_kernels.hpp>
2
3 #include <lcqi/inference.hpp>
4
5 #include <algorithm>
6 #include <chrono>
7 #include <cmath>
8 #include <stdexcept>
9 #include <string>
10
11 namespace lcqi {

```

```

12 namespace {
13
14 constexpr std::int32_t LCQI_I8_PATTERN_MODULUS = 23;
15 constexpr std::int32_t LCQI_I8_PATTERN_CENTER = 11;
16 constexpr std::int32_t LCQI_INPUT_PATTERN_MODULUS = 17;
17 constexpr std::int32_t LCQI_INPUT_PATTERN_CENTER = 8;
18 constexpr float LCQI_INPUT_PATTERN_SCALE = 0.125F;
19 constexpr double LCQI_SECONDS_TO_MICROSECONDS = 1000000.0;
20 constexpr std::int32_t LCQI_SIMD_OUTPUT_BLOCK = 8;
21 constexpr std::size_t LCQI_BENCHMARK_BACKEND_COUNT = 3;
22 constexpr const char* LCQI_BACKEND_SCALAR = "scalar";
23 constexpr const char* LCQI_BACKEND_PACKED_SCALAR = "packed_scalar";
24 constexpr const char* LCQI_BACKEND_AVX2 = "avx2";
25
26 void require_positive(std::int32_t value, const char* name) {
27     if (value <= 0) {
28         throw std::runtime_error(std::string(name) + " must be positive");
29     }
30 }
31
32 std::size_t checked_size(std::int32_t value, const char* name) {
33     if (value < 0) {
34         throw std::runtime_error(std::string(name) + " cannot be negative");
35     }
36     return static_cast<std::size_t>(value);
37 }
38
39 void validate_quantized_layer(const QuantizedLinearLayer& layer) {
40     require_positive(layer.input_size, "input_size");
41     require_positive(layer.output_size, "output_size");
42     const std::size_t expected_weights =
43         checked_size(layer.input_size, "input_size") *
44         checked_size(layer.output_size, "output_size");
45     if (layer.weights.size() != expected_weights) {
46         throw std::runtime_error("quantized layer weight size mismatch");
47     }
48     if (layer.bias.size() != checked_size(layer.output_size, "output_size")) {
49         throw std::runtime_error("quantized layer bias size mismatch");
50     }
51 }
52
53 std::int32_t round_up_to_block(std::int32_t value, std::int32_t block_size) {
54     require_positive(value, "value");
55     require_positive(block_size, "block_size");
56     return ((value + block_size - 1) / block_size) * block_size;
57 }
58
59 float checksum(std::span<const float> values) {
60     float sum = 0.0F;
61     for (const float value : values) {
62         sum += value;
63     }

```

```

64     return sum;
65 }
66
67 void add_unavailable_backend(KernelBenchmarkResult& result, const char* name) {
68     result.backends.push_back(KernelBackendBenchmark{
69         .name = name,
70         .available = false,
71         .average_us = 0.0,
72         .max_abs_diff = 0.0F,
73         .checksum = 0.0F,
74     });
75 }
76
77 template <typename Callable>
78 double measure_average_us(std::int32_t repeat, Callable&& callable) {
79     require_positive(repeat, "repeat");
80     const auto begin = std::chrono::steady_clock::now();
81     for (std::int32_t index = 0; index < repeat; ++index) {
82         callable();
83     }
84     const auto end = std::chrono::steady_clock::now();
85     const std::chrono::duration<double> elapsed = end - begin;
86     return elapsed.count() * LCQI_SECONDS_TO_MICROSECONDS / static_cast<double>←
    ↪ >(repeat);
87 }
88
89 } // namespace
90
91 PackedLinearI8 pack_linear_i8(const QuantizedLinearLayer& layer,
92                               std::int32_t output_block_size) {
93     validate_quantized_layer(layer);
94     require_positive(output_block_size, "output_block_size");
95
96     PackedLinearI8 packed;
97     packed.input_size = layer.input_size;
98     packed.output_size = layer.output_size;
99     packed.output_block_size = output_block_size;
100     packed.scale = layer.scale;
101     packed.bias = layer.bias;
102
103     const std::int32_t padded_outputs =
104         round_up_to_block(layer.output_size, output_block_size);
105     packed.packed_weights.assign(
106         checked_size(padded_outputs, "padded_outputs") *
107         checked_size(layer.input_size, "input_size"),
108         0);
109
110     std::size_t packed_index = 0;
111     for (std::int32_t output_block = 0; output_block < padded_outputs;
112          output_block += output_block_size) {
113         for (std::int32_t input_index = 0; input_index < layer.input_size; ++←
    ↪ input_index) {

```

```

114         for (std::int32_t offset = 0; offset < output_block_size; ++offset)↵
115         ↵ {
116             const std::int32_t output_index = output_block + offset;
117             if (output_index < layer.output_size) {
118                 const std::size_t source =
119                     checked_size(output_index, "output_index") *
120                     checked_size(layer.input_size, "input_size") +
121                     checked_size(input_index, "input_index");
122                 packed.packed_weights[packed_index] = layer.weights[source]↵
123                 ↵ ];
124             }
125             ++packed_index;
126         }
127     }
128     return packed;
129 }
130 void linear_i8_packed(const PackedLinearI8& layer,
131                      std::span<const float> input,
132                      std::span<float> output) {
133     std::vector<float> accumulators(
134         checked_size(layer.output_block_size, "output_block_size"),
135         0.0F);
136     linear_i8_packed_with_workspace(layer, input, output, accumulators);
137 }
138
139 void linear_i8_packed_with_workspace(const PackedLinearI8& layer,
140                                     std::span<const float> input,
141                                     std::span<float> output,
142                                     std::span<float> accumulator_workspace) {
143     require_positive(layer.input_size, "packed input_size");
144     require_positive(layer.output_size, "packed output_size");
145     require_positive(layer.output_block_size, "packed output_block_size");
146     if (input.size() != checked_size(layer.input_size, "input_size")) {
147         throw std::runtime_error("linear_i8_packed input size mismatch");
148     }
149     if (output.size() != checked_size(layer.output_size, "output_size")) {
150         throw std::runtime_error("linear_i8_packed output size mismatch");
151     }
152     if (layer.bias.size() != checked_size(layer.output_size, "output_size")) {
153         throw std::runtime_error("linear_i8_packed bias size mismatch");
154     }
155
156     const std::int32_t padded_outputs =
157         round_up_to_block(layer.output_size, layer.output_block_size);
158     const std::size_t expected_packed_size =
159         checked_size(padded_outputs, "padded_outputs") *
160         checked_size(layer.input_size, "input_size");
161     if (layer.packed_weights.size() != expected_packed_size) {
162         throw std::runtime_error("linear_i8_packed packed weight size mismatch"↵
163         ↵ );

```

```

163     }
164     if (accumulator_workspace.size() != checked_size(layer.output_block_size, "↵
↵ output_block_size")) {
165         throw std::runtime_error("linear_i8_packed accumulator workspace size ↵
↵ mismatch");
166     }
167
168     std::size_t packed_index = 0;
169     for (std::int32_t output_block = 0; output_block < padded_outputs;
170         output_block += layer.output_block_size) {
171         std::fill(accumulator_workspace.begin(), accumulator_workspace.end(), ↵
↵ 0.0F);
172         for (std::int32_t offset = 0; offset < layer.output_block_size; ++↵
↵ offset) {
173             const std::int32_t output_index = output_block + offset;
174             if (output_index < layer.output_size) {
175                 accumulator_workspace[checked_size(offset, "offset")] =
176                     layer.bias[checked_size(output_index, "output_index")];
177             }
178         }
179
180         for (std::int32_t input_index = 0; input_index < layer.input_size; ++↵
↵ input_index) {
181             const float input_value = input[checked_size(input_index, "↵
↵ input_index")];
182             for (std::int32_t offset = 0; offset < layer.output_block_size; ++↵
↵ offset) {
183                 const std::int32_t output_index = output_block + offset;
184                 const std::int8_t weight = layer.packed_weights[packed_index↵
↵ ++];
185                 if (output_index < layer.output_size) {
186                     accumulator_workspace[checked_size(offset, "offset")] +=
187                         static_cast<float>(weight) * layer.scale * input_value;
188                 }
189             }
190         }
191
192         for (std::int32_t offset = 0; offset < layer.output_block_size; ++↵
↵ offset) {
193             const std::int32_t output_index = output_block + offset;
194             if (output_index < layer.output_size) {
195                 output[checked_size(output_index, "output_index")] =
196                     accumulator_workspace[checked_size(offset, "offset")];
197             }
198         }
199     }
200 }
201
202 QuantizedLinearLayer make_deterministic_i8_layer(std::int32_t input_size,
203                                                  std::int32_t output_size,
204                                                  float scale) {
205     require_positive(input_size, "input_size");

```

```

206     require_positive(output_size, "output_size");
207     QuantizedLinearLayer layer;
208     layer.input_size = input_size;
209     layer.output_size = output_size;
210     layer.scale = scale;
211     const std::size_t weight_count =
212         checked_size(input_size, "input_size") * checked_size(output_size, "↵
↵ output_size");
213     layer.weights.resize(weight_count);
214     layer.bias.resize(checked_size(output_size, "output_size"));
215     for (std::int32_t output_index = 0; output_index < output_size; ++↵
↵ output_index) {
216         layer.bias[checked_size(output_index, "output_index")] =
217             static_cast<float>((output_index % LCQI_INPUT_PATTERN_MODULUS) -
218                               LCQI_INPUT_PATTERN_CENTER) *
219             LCQI_INPUT_PATTERN_SCALE;
220         for (std::int32_t input_index = 0; input_index < input_size; ++↵
↵ input_index) {
221             const std::int32_t raw =
222                 ((output_index * input_size + input_index) % ↵
↵ LCQI_I8_PATTERN_MODULUS) -
223                 LCQI_I8_PATTERN_CENTER;
224             layer.weights[checked_size(output_index, "output_index") *
225                           checked_size(input_size, "input_size") +
226                           checked_size(input_index, "input_index")] =
227                 static_cast<std::int8_t>(raw);
228         }
229     }
230     return layer;
231 }
232
233 std::vector<float> make_deterministic_input(std::int32_t input_size) {
234     require_positive(input_size, "input_size");
235     std::vector<float> input(checked_size(input_size, "input_size"));
236     for (std::int32_t input_index = 0; input_index < input_size; ++input_index)↵
↵ {
237         input[checked_size(input_index, "input_index")] =
238             static_cast<float>((input_index % LCQI_INPUT_PATTERN_MODULUS) -
239                               LCQI_INPUT_PATTERN_CENTER) *
240             LCQI_INPUT_PATTERN_SCALE;
241     }
242     return input;
243 }
244
245 float max_abs_diff(std::span<const float> lhs, std::span<const float> rhs) {
246     if (lhs.size() != rhs.size()) {
247         throw std::runtime_error("max_abs_diff size mismatch");
248     }
249     float diff = 0.0F;
250     for (std::size_t index = 0; index < lhs.size(); ++index) {
251         diff = std::max(diff, std::fabs(lhs[index] - rhs[index]));
252     }

```

```

253     return diff;
254 }
255
256 KernelBenchmarkResult benchmark_linear_i8_case(const KernelBenchmarkCase& ←
    ↪ benchmark_case) {
257     require_positive(benchmark_case.input_size, "benchmark input_size");
258     require_positive(benchmark_case.output_size, "benchmark output_size");
259     require_positive(benchmark_case.output_block_size, "benchmark ↪
    ↪ output_block_size");
260     require_positive(benchmark_case.repeat, "benchmark repeat");
261
262     const QuantizedLinearLayer layer =
263         make_deterministic_i8_layer(benchmark_case.input_size,
264                                     benchmark_case.output_size,
265                                     0.015625F);
266     const PackedLinearI8 packed = pack_linear_i8(layer, benchmark_case.↪
    ↪ output_block_size);
267     const PackedLinearI8 packed_simd = pack_linear_i8(layer, ↪
    ↪ LCQI_SIMD_OUTPUT_BLOCK);
268     const std::vector<float> input = make_deterministic_input(benchmark_case.↪
    ↪ input_size);
269     std::vector<float> scalar_output(checked_size(benchmark_case.output_size, "↪
    ↪ output_size"),
270                                     0.0F);
271     std::vector<float> packed_output(scalar_output.size(), 0.0F);
272     std::vector<float> simd_output(scalar_output.size(), 0.0F);
273     std::vector<float> packed_workspace(
274         checked_size(benchmark_case.output_block_size, "output_block_size"),
275         0.0F);
276
277     linear_i8(layer, input, scalar_output);
278     linear_i8_packed_with_workspace(packed, input, packed_output, ↪
    ↪ packed_workspace);
279
280     KernelBenchmarkResult result;
281     result.benchmark_case = benchmark_case;
282     result.backends.reserve(LCQI_BENCHMARK_BACKEND_COUNT);
283     result.backends.push_back(KernelBackendBenchmark{
284         .name = LCQI_BACKEND_SCALAR,
285         .available = true,
286         .average_us = measure_average_us(benchmark_case.repeat,
287                                         [&]() { linear_i8(layer, input, ↪
    ↪ scalar_output); }),
288         .max_abs_diff = 0.0F,
289         .checksum = checksum(scalar_output),
290     });
291     result.backends.push_back(KernelBackendBenchmark{
292         .name = LCQI_BACKEND_PACKED_SCALAR,
293         .available = true,
294         .average_us = measure_average_us(
295             benchmark_case.repeat,
296             [&]() { linear_i8_packed_with_workspace(packed,

```

```

297                                     input,
298                                     packed_output,
299                                     packed_workspace); }},
300     .max_abs_diff = max_abs_diff(scalar_output, packed_output),
301     .checksum = checksum(packed_output),
302 });
303 if (linear_i8_packed_avx2_available()) {
304     linear_i8_packed_avx2(packed_simd, input, simd_output);
305     result.backends.push_back(KernelBackendBenchmark{
306         .name = LCQI_BACKEND_AVX2,
307         .available = true,
308         .average_us = measure_average_us(
309             benchmark_case.repeat,
310             [&]() { linear_i8_packed_avx2(packed_simd, input, simd_output); ←
    ↪ }),
311         .max_abs_diff = max_abs_diff(scalar_output, simd_output),
312         .checksum = checksum(simd_output),
313     });
314 } else {
315     add_unavailable_backend(result, LCQI_BACKEND_AVX2);
316 }
317 return result;
318 }
319
320 } // namespace lcqi
321
322 #if !defined(LCQI_ENABLE_AVX2)
323 namespace lcqi {
324
325 bool linear_i8_packed_avx2_available() noexcept {
326     return false;
327 }
328
329 void linear_i8_packed_avx2(const PackedLinearI8&, std::span<const float>, std::span<float>) {
    ↪ span<float>) {
330     throw std::runtime_error("AVX2 packed kernel is not enabled in this build") ←
    ↪ ;
331 }
332
333 } // namespace lcqi
334 #endif

```

`int8_kernels_avx2.cpp` 是平台优化路径。读者要先看可用性判断，再看向量化实现；如果机器不支持 AVX2/FMA，报告必须写 `unavailable`。

**Listing 4.13** `src/int8_kernels_avx2.cpp`

```

1 #include <lcqi/int8_kernels.hpp>
2
3 #if defined(LCQI_ENABLE_AVX2)
4
5 #include <immintrin.h>
6

```



```

7  #include <algorithm>
8  #include <cstdint>
9  #include <cstring>
10 #include <stdexcept>
11 #include <string>
12
13 namespace lcqi {
14 namespace {
15
16 constexpr std::int32_t LCQI_AVX2_OUTPUT_BLOCK = 8;
17
18 void require_positive(std::int32_t value, const char* name) {
19     if (value <= 0) {
20         throw std::runtime_error(std::string(name) + " must be positive");
21     }
22 }
23
24 std::size_t checked_size(std::int32_t value, const char* name) {
25     if (value < 0) {
26         throw std::runtime_error(std::string(name) + " cannot be negative");
27     }
28     return static_cast<std::size_t>(value);
29 }
30
31 std::int32_t round_up_to_block(std::int32_t value, std::int32_t block_size) {
32     require_positive(value, "value");
33     require_positive(block_size, "block_size");
34     return ((value + block_size - 1) / block_size) * block_size;
35 }
36
37 void validate_avx2_contract(const PackedLinearI8& layer,
38                             std::span<const float> input,
39                             std::span<float> output) {
40     require_positive(layer.input_size, "packed input_size");
41     require_positive(layer.output_size, "packed output_size");
42     if (layer.output_block_size != LCQI_AVX2_OUTPUT_BLOCK) {
43         throw std::runtime_error("AVX2 packed kernel requires output_block_size ←
44     ↪ == 8");
45     }
46     if (input.size() != checked_size(layer.input_size, "input_size")) {
47         throw std::runtime_error("AVX2 packed kernel input size mismatch");
48     }
49     if (output.size() != checked_size(layer.output_size, "output_size")) {
50         throw std::runtime_error("AVX2 packed kernel output size mismatch");
51     }
52     if (layer.bias.size() != checked_size(layer.output_size, "output_size")) {
53         throw std::runtime_error("AVX2 packed kernel bias size mismatch");
54     }
55     const std::int32_t padded_outputs =
56         round_up_to_block(layer.output_size, layer.output_block_size);
57     const std::size_t expected_packed_size =
58         checked_size(padded_outputs, "padded_outputs") *

```

```

58     checked_size(layer.input_size, "input_size");
59     if (layer.packed_weights.size() != expected_packed_size) {
60         throw std::runtime_error("AVX2 packed kernel packed weight size ↵
↵ mismatch");
61     }
62 }
63
64 } // namespace
65
66 bool linear_i8_packed_avx2_available() noexcept {
67     #if defined(__GNUC__) || defined(__clang__)
68         __builtin_cpu_init();
69         return __builtin_cpu_supports("avx2") && __builtin_cpu_supports("fma");
70     #else
71         return true;
72     #endif
73 }
74
75 void linear_i8_packed_avx2(const PackedLinearI8& layer,
76                             std::span<const float> input,
77                             std::span<float> output) {
78     if (!linear_i8_packed_avx2_available()) {
79         throw std::runtime_error("AVX2/FMA is not available on this CPU");
80     }
81     validate_avx2_contract(layer, input, output);
82
83     const std::int32_t padded_outputs =
84         round_up_to_block(layer.output_size, layer.output_block_size);
85     std::size_t packed_index = 0;
86     alignas(32) float block_output[LCQI_AVX2_OUTPUT_BLOCK] = {};
87
88     for (std::int32_t output_block = 0; output_block < padded_outputs;
89         output_block += LCQI_AVX2_OUTPUT_BLOCK) {
90         __m256 accumulator = _mm256_setzero_ps();
91
92         for (std::int32_t offset = 0; offset < LCQI_AVX2_OUTPUT_BLOCK; ++offset ↵
↵ ) {
93             const std::int32_t output_index = output_block + offset;
94             block_output[checked_size(offset, "offset")] =
95                 output_index < layer.output_size
96                 ? layer.bias[checked_size(output_index, "output_index")]
97                 : 0.0F;
98         }
99         accumulator = _mm256_load_ps(block_output);
100
101         for (std::int32_t input_index = 0; input_index < layer.input_size; ++ ↵
↵ input_index) {
102             std::uint64_t packed_bytes = 0;
103             std::memcpy(&packed_bytes,
104                         layer.packed_weights.data() + packed_index,
105                         sizeof(packed_bytes));
106             const __m128i packed_i8 =

```

```

107         _mm_cvtsi64_si128(static_cast<long long>(packed_bytes));
108         packed_index += LCQI_AVX2_OUTPUT_BLOCK;
109         const __m256i weights_i32 = _mm256_cvtepi8_epi32(packed_i8);
110         const __m256 weights_f32 = _mm256_cvtepi32_ps(weights_i32);
111         const __m256 input_scaled =
112             _mm256_set1_ps(input[checked_size(input_index, "input_index")] ←
↪ * layer.scale);
113         accumulator = _mm256_fmadd_ps(weights_f32, input_scaled, ←
↪ accumulator);
114     }
115
116     _mm256_store_ps(block_output, accumulator);
117     for (std::int32_t offset = 0; offset < LCQI_AVX2_OUTPUT_BLOCK; ++offset↪
↪ ) {
118         const std::int32_t output_index = output_block + offset;
119         if (output_index < layer.output_size) {
120             output[checked_size(output_index, "output_index")] =
121                 block_output[checked_size(offset, "offset")];
122         }
123     }
124 }
125 }
126
127 } // namespace lcqi
128
129 #endif

```

第零部分

# Transformer、KV Cache 与服务运行时

## 第 5 章实践：Transformer 切片与 KV Cache 状态

本章对应原书中的第三册“AI 推理原理、KV cache 与计算账本”和“KV cache 实现细节”。不要把它当作概念复述，本章要把 RMSNorm、Q/K/V、RoPE、GQA、KV slot、position、state edge 放到同一个可运行项目里检查。贯穿代码仍然是 `books/cpu-volume-3/labs/linux_cpu_inference`；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `kv_cache`，核心命令或测试入口是 `lcqi_trace`。

Listing 5.1 第 5 章实践：Transformer 切片与 KV Cache 状态的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  ↪ DCMMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_trace / kv_cache
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：固定 prompt `tok1tok2tok3`，观察第一个 decode step。读者要先手写预期结果，再运行项目观察输出。动作是：把 trace 中 tokenizer、q/k/v、`kv_cache_slot`、attention 和 logits 串成状态表。如果输出里没有 `token_position` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：KV cache 是请求状态，不是临时 activation；position、layer、head、layout 和生命周期必须进入 trace 这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 `books/cpu-volume-3/labs/linux_cpu_inference/README.md`、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交 KV cache 状态表：每层每个 token 写入哪里、何时读取、何时禁止复用。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。

## 第 6 章实践：Reference Decoder、logits 与采样

本章对应原书中的第三册“Reference decoder 实现”和“推理算法：logits、softmax、采样”。不要把它当作概念复述，本章要把 reference decoder、logits golden、argmax sampler、tokenizer contract hash、BOS/EOS 放到同一个可运行项目里检查。贯穿代码仍然是 [books/cpu-volume-3/labs/linux\\_cpu\\_inference](#)；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 **argmax**，核心命令或测试入口是 **lcqi\_trace**。

**Listing 6.1** 第 6 章实践：Reference Decoder、logits 与采样的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  ↪ DCMMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_trace / argmax
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：让 trace 固定输出 tokenizer contract 和 logits golden。读者要先手写预期结果，再运行项目观察输出。动作是：比较 generated token、logits 数组和 tokenizer hash。如果输出里没有 **tokenizer\_contract** 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：采样前必须先证明 logits 可复查；sampler 的随机性、温度和 top-k 只能建立在 reference 通过之后这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 [books/cpu-volume-3/labs/linux\\_cpu\\_inference/README.md](#)、相关 **include/lcqi** 头文件和一个 **src** 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交采样报告：argmax、top-k、temperature 三种策略各自需要哪些额外字段才能复现。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。

## 第7章实践：服务调度、SLO 与资源预算

本章对应原书中的第三册“业务服务实战：请求、调度、队列、取消与资源预算”。不要把它当作概念复述，本章要把 admission、KV budget、queue wait、TTFT、TPOT、cancel、reject、p95 放到同一个可运行项目里检查。贯穿代码仍然是 `books/cpu-volume-3/labs/linux_cpu_inference`；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `ttft_ms`，核心命令或测试入口是 `lcqi_serving_probe`。

Listing 7.1 第7章实践：服务调度、SLO 与资源预算的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  <- DCMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_serving_probe / ttft_ms
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：运行 serving probe 中短请求、RAG 请求、取消请求和超长请求。读者要先手写预期结果，再运行项目观察输出。动作是：解释 `memory_budget_exceeded`、cancel reclaim、`queue_wait_p95_ms` 和 `rejection_rate`。如果输出里没有 `rejection_rate` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：推理服务不是只看 tokens/s；是否接收请求、如何预留 KV、取消后是否回收，决定系统是否可运营这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 `books/cpu-volume-3/labs/linux_cpu_inference/README.md`、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交 SLO 表：每类请求的 prompt/new token、预算、是否接受、TTFT、TPOT、拒绝原因和未验证边界。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。



## 第 8 章实践：RAG、Agent 与状态图边界

本章对应原书中的第三册“上层 AI 应用编排：RAG、Agent、工具和状态图”。不要把它当作概念复述，本章要把 retrieval、tool call、prompt assembly、state graph、budget、trace id 放到同一个可运行项目里检查。贯穿代码仍然是 `books/cpu-volume-3/labs/linux_cpu_inference`；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `req_rag`，核心命令或测试入口是 `lcqi_serving_probe`。

Listing 8.1 第 8 章实践：RAG、Agent 与状态图边界的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  ↳ DCMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_serving_probe / req_rag
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：把 RAG 请求拆成 retrieve、rerank、prompt build、prefill、decode 五段。读者要先手写预期结果，再运行项目观察输出。动作是：为每段写输入输出和失败后能否重试。如果输出里没有 `prompt_tokens` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：上层编排不能把外部工具结果和模型状态混成字符串；每个边界都要有身份、预算和可重放记录这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 `books/cpu-volume-3/labs/linux_cpu_inference/README.md`、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交一个 RAG 状态图：节点、边、超时、重试、缓存键、prompt 版本和审计字段。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。



第零部分

# 模型导入、AI 编译器与内存计划

## 第 9 章实践：模型导入、manifest 与 Shape Verifier

本章对应原书中的第三册“模型导入、manifest 与 shape verifier”。不要把它当作概念复述，本章要把 safetensors header、metadata、tensor name、dtype、shape、`data_offsets`、payload boundary 放到同一个可运行项目里检查。贯穿代码仍然是 `books/cpu-volume-3/labs/linux_cpu_inference`；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `load_safetensors_manifest`，核心命令或测试入口是 `lcqi_tests`。

**Listing 9.1** 第 9 章实践：模型导入、manifest 与 Shape Verifier 的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  ↪ DCMMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_tests / load_safetensors_manifest
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：加载测试临时 safetensors fixture。读者要先手写预期结果，再运行项目观察输出。动作是：故意构造 offset 越界和 dtype/shape 字节数不匹配。如果输出里没有 `data_offsets` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：模型导入的第一职责是早拒绝：名字、shape、dtype、offset 和 payload 都必须在 kernel 运行前被验证这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 `books/cpu-volume-3/labs/linux_cpu_inference/README.md`、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交 manifest gate 报告：合法 fixture 和两个坏 fixture 的字段、错误原因和拒绝位置。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。

## 第 10 章实践：真实模型加载闭环与首 Token Trace

本章对应原书中的第三册“真实模型加载闭环：Tokenizer、权重格式与首 token trace”。不要把它当作概念复述，本章要把 tokenizer、weight mapping、chat template、first token trace、checksum 放到同一个可运行项目里检查。贯穿代码仍然是 [books/cpu-volume-3/labs/linux\\_cpu\\_inference](#)；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

本章新增一个明确边界：tiny fixture 负责手算和逐层 oracle；LCQI GPT-2 reference path 负责证明自研代码已经能加载标准 GPT-2 safetensors、执行 byte-level BPE、跑 GPT-2 forward 并 greedy 生成；SmolLM2 small smoke 负责证明外部 llama.cpp 路径能加载 GGUF small 模型。三者不能互相冒充。tiny MLP 算对，不等于 GPT-2 完成；GPT-2 reference 能出一个 token，不等于高性能 KV cache runtime 完成；SmolLM2 能输出文本，也不等于 LCQI 自研 GGUF 后端完成。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里有三组核心对象。第一组是 [tiny\\_tokenizer.txt](#)、reference decoder 和 [lcqi\\_trace](#)，用于手算 token 合同和逐 checkpoint 追踪。第二组是 [gpt2\\_reference.hpp/.cpp](#)、[lcqi\\_gpt2](#) 和 [tools/run\\_gpt2\\_smoke.py](#)，用于自研 GPT-2 safetensors 路径。第三组是 [tools/run\\_smolllm2\\_small\\_smoke.py](#)，用于下载并校验 SmolLM2-135M-Instruct 的 GGUF 量化文件，再通过 llama.cpp 在 CPU 上实际生成 assistant 文本。

Listing 10.1 第 10 章实践：真实模型加载闭环与首 Token Trace 的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  ↪ DCMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_trace / tiny_tokenizer.txt
books/cpu-volume-3/build/lcqi-debug/labs/linux_cpu_inference/lcqi_gpt2
make cpu3-gpt2-smoke
make cpu3-smolllm2-smoke
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：从 tokenizer fixture 开始，不跳到真实 7B。读者要先手写预期结果，再运行项目观察输出。动作是：解释 BOS/EOS、UNK、template hash 和 decoder 输入 tokens 的边界。如果输出里没有 `tokens=[0,1,2,3,4]` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：真实模型加载不是下载权重；必须先让 tokenizer 合同、权重 role、shape verifier 和首 token trace 形成闭环。这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、自研 GPT-2 reference smoke：先把标准 safetensors 跑通

GPT-2 不是 LLaMA。它使用 LayerNorm，不是 RMSNorm；使用 learned absolute position embedding，不是 RoPE；attention 里的 Q/K/V 在 HuggingFace 权重中是 packed Conv1D 权重；MLP 使用 GELU，不是 SwiGLU；`lm_head` 通常和 token embedding 共享权重。把 GPT-2 硬塞进原来的 LLaMA-like tiny decoder，会让概念混乱，所以 LCQI 新增了独立的 `gpt2_reference` 模块。

仓库内最小验收分两层。第一层不下载模型，运行 tiny GPT-2 fixture：

**Listing 10.2** 仓库内 tiny GPT-2 数学闭环

```
books/cpu-volume-3/build/lcqi-debug/labs/linux_cpu_inference/lcqi_gpt2
```

它固定 prompt ids 为 123，检查 logits、predicted token 和 greedy generated ids。对应测试还会临时写出一个 HuggingFace 风格 `model.safetensors`，再用 loader 读回来，覆盖 tensor name mapping、shape verifier、Conv1D 转置和 tied lm head。坏 shape 必须在 loader 阶段被拒绝。

第二层下载真实 GPT-2 资产：

**Listing 10.3** LCQI 自研 GPT-2 smoke

```
make cpu3-gpt2-smoke
```

脚本下载 `openai-community/gpt2` 的 `config.json`、`vocab.json`、`merges.txt` 和 `model.safetensors` 到用户缓存目录，不提交模型文件。最近一次报告写到：

**Listing 10.4** GPT-2 smoke 报告关键字段

```
books/cpu-volume-3/results/lcqi-gpt2-smoke.txt
prompt_ids 15496 11 616 1438 318
generated_ids 15496 11 616 1438 318 1757
generated_text Hello, my name is John
validation=PASS
```

这说明自研 LCQI C++ reference path 已经能跑标准 GPT-2 safetensors 的 tokenizer、权重导入、forward 和 greedy generation。边界也要写清：当前实现为了可读和可验证，会重跑完整前缀，没有 KV cache、没有 packed weight、没有多线程，也没有把 logits 与 PyTorch/Transformers 的逐数值 golden 报告纳入仓库。因此它是正确性闭环，不是性能结论。

#### 四、真实 small 模型 smoke：不用 tiny，先让开源模型实际跑起来

本章的真实模型 smoke 固定为 `SmolLM2-135M-Instruct`，规模是 135M 参数，不是 tiny 随机夹具。脚本使用 GGUF 文件 `SmolLM2-135M-Instruct-Q4_K_M.gguf`，来源仓库是 `bartowski/SmolLM2-135M-Instruct-GGUF`，原模型组织是 `HuggingFaceTB`，license 字段为 `apache-2.0`。脚本不会把模型文件提交到仓库，而是缓存到用户目录；验收时检查文件大小 105454432 字节和 SHA256：

**Listing 10.5** SmolLM2 small GGUF 文件身份

```
model_source_id=HuggingFaceTB/SmolLM2-135M-Instruct
model_gguf_repo_id=bartowski/SmolLM2-135M-Instruct-GGUF
model_file=SmolLM2-135M-Instruct-Q4_K_M.gguf
model_expected_bytes=105454432
model_expected_sha256=2↵
↵ e8040ceae7815abe0dcb3540b9995eaa1fa0d2ca9e797d0a635ae4433c68c2d
```

运行命令是：

**Listing 10.6** 真实 small 模型 CPU 冒烟

```
make cpu3-smolllm2-smoke
```

这条命令会做四步。第一，下载或复用 GGUF 文件，并校验 SHA256，防止“文件名对了但内容不对”。第二，拉取固定 commit 的 `llama.cpp` 并构建 `llama-simple-chat`，避免依赖本机临时安装。第三，用 `-ngl0` 强制 CPU 路径，输入一个短 prompt，确认程序能走过 tokenizer、chat template、prefill、decode 和 sampler。第四，把报告写到：

Listing 10.7 真实模型 smoke 报告路径

```
books/cpu-volume-3/results/lcqi-smolllm2-small-model-smoke.txt
```

报告里必须至少出现这些字段：`model_expected_sha256`、`llama_cpp_commit`、`command`、`exit_code=0`、`validation=PASS` 和 `assistant_response`。如果只有下载日志，没有 assistant response，不能算通过；如果 assistant response 内容答错了，也不能改写成“质量通过”。本章只验证真实模型资产可加载、可 tokenize、可 decode、可生成非空文本。

## 五、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 `books/cpu-volume-3/labs/linux_cpu_inference/README.md`、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：运行 GPT-2 smoke。报告不许只贴“生成了 John”，必须解释 `config.json`、`vocab.json`、`merges.txt`、`model.safetensors` 各自对应哪段代码，为什么 GPT-2 需要 byte-level BPE，为什么 HuggingFace Conv1D 权重要转置，为什么 `lm_head` 可以 tied 到 embedding。

任务四：运行真实 small 模型 smoke。报告不许只贴“命令执行成功”，必须解释脚本为什么校验模型 hash、为什么固定 `llama.cpp` commit、为什么使用 `-ngl0`，以及为什么 assistant response 非空只是 smoke，不是质量评测。

任务五：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。GPT-2 反例可以是：tensor name 缺失、shape 和 config 不一致、vocab 中缺少 BPE merge 后 token、`max_new_tokens` 超过 context。真实模型 smoke 的反例可以是：模型 SHA256 不匹配、网络不可用、`cmake` 缺失、assistant response 为空、命令超时。每个反例都要写清期望处理方式。

## 六、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。新增 GPT-2 smoke 后，最低验收还包括：真实 GPT-2 文件身份有 hash，prompt ids 可解释，generated ids 可解释，文本反解码可解释，并明确声明这不是高性能推理结论。新增 SmolLM2 smoke 后，最低验收还包括：模型身份有 hash，模型规模不是 tiny，CPU 路径实际运行，报告里有非空 assistant response，并明确声明这不是自研 LCQI 完整 GGUF 支持。

递进大作业：提交 GPT-2 reference path 的下一阶段迁移清单：KV cache、prefill/decode 分离、packed weight、线程并行、Transformers golden logits、更多模型方言 name mapping、分片 safetensors 和性能报告。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。

# 代码走读: Reference Decoder、Tokenizer、Safetensors 与 GPT-2

## 从 tiny MLP 走向 decoder-only 模型

Reference decoder 和 GPT-2 reference path 是 LCQI 从教学闭环走向真实模型导入的关键层。这里的代码更长，是因为真实模型不只有矩阵乘法：它还需要 tokenizer、位置、归一化、QKV 拆分、attention mask、KV cache、激活函数、lm head、权重命名映射、dtype 转换和坏资产拒绝。

## Tiny Reference Decoder

`reference_decoder.cpp` 实现 embedding、RMSNorm、Q/K/V、RoPE、KV cache、attention、SviGLU、final norm、logits 和 sampler，并在关键位置生成 trace checkpoint。读者应按 checkpoint 顺序读，而不是从最终 logits 倒推。

Listing 10.8 `src/reference_decoder.cpp`

```

1  #include <lcqi/reference_decoder.hpp>
2
3  #include <algorithm>
4  #include <array>
5  #include <cmath>
6  #include <limits>
7  #include <stdexcept>
8  #include <string>
9
10 namespace lcqi {
11 namespace {
12
13 constexpr float LCQI_SOFTMAX_NEGATIVE_INFINITY = -std::numeric_limits<float>::infinity();
14 constexpr std::int32_t LCQI_ROPE_PAIR_WIDTH = 2;
15 constexpr const char* LCQI_TRACE_DTYPE_F32 = "f32";
16 constexpr const char* LCQI_TRACE_LAYOUT_CONTIGUOUS = "contiguous";
17 constexpr const char* LCQI_TRACE_LAYOUT_ROW_MAJOR = "row_major";
18 constexpr const char* LCQI_TRACE_LAYOUT_KV_CONTIGUOUS = "kv_contiguous";
19
20 void require_positive(std::int32_t value, const char* name) {
21     if (value <= 0) {
22         throw std::runtime_error(std::string(name) + " must be positive");
23     }
24 }
25
26 void validate_config(const DecoderConfig& config) {
27     require_positive(config.hidden_size, "hidden_size");
28     require_positive(config.query_heads, "query_heads");
29     require_positive(config.kv_heads, "kv_heads");
30     require_positive(config.head_dim, "head_dim");
31     require_positive(config.intermediate_size, "intermediate_size");
32     require_positive(config.vocab_size, "vocab_size");

```



```

33     require_positive(config.max_context, "max_context");
34     if (config.query_heads % config.kv_heads != 0) {
35         throw std::runtime_error("query_heads must be divisible by kv_heads");
36     }
37     if (config.hidden_size != config.query_heads * config.head_dim) {
38         throw std::runtime_error("hidden_size must equal query_heads * head_dim↵
↵ ");
39     }
40     if (config.head_dim % LCQI_ROPE_PAIR_WIDTH != 0) {
41         throw std::runtime_error("head_dim must be even for RoPE");
42     }
43     if (config.rope_base <= 1.0F) {
44         throw std::runtime_error("rope_base must be greater than 1");
45     }
46 }
47
48 std::size_t checked_size(std::int32_t value, const char* name) {
49     if (value < 0) {
50         throw std::runtime_error(std::string(name) + " cannot be negative");
51     }
52     return static_cast<std::size_t>(value);
53 }
54
55 std::int32_t kv_width(const DecoderConfig& config) {
56     return config.kv_heads * config.head_dim;
57 }
58
59 void validate_linear(const LinearWeightsF32& layer, const char* name) {
60     require_positive(layer.input_size, "linear input_size");
61     require_positive(layer.output_size, "linear output_size");
62     const std::size_t expected_weights =
63         checked_size(layer.input_size, name) * checked_size(layer.output_size, ↵
↵ name);
64     if (layer.weights.size() != expected_weights) {
65         throw std::runtime_error(std::string(name) + " weight size mismatch");
66     }
67     if (!layer.bias.empty() &&
68         layer.bias.size() != checked_size(layer.output_size, name)) {
69         throw std::runtime_error(std::string(name) + " bias size mismatch");
70     }
71 }
72
73 void validate_span_size(std::size_t actual, std::int32_t expected, const char* ↵
↵ name) {
74     if (actual != checked_size(expected, name)) {
75         throw std::runtime_error(std::string(name) + " size mismatch");
76     }
77 }
78
79 std::int32_t argmax(std::span<const float> values) {
80     if (values.empty()) {
81         throw std::runtime_error("cannot take argmax of empty values");

```



```

82     }
83     std::int32_t best_index = 0;
84     float best_value = values[0];
85     for (std::int32_t index = 1; index < static_cast<std::int32_t>(values.size()
↵    )); ++index) {
86         const float candidate = values[static_cast<std::size_t>(index)];
87         if (candidate > best_value) {
88             best_value = candidate;
89             best_index = index;
90         }
91     }
92     return best_index;
93 }
94
95 float silu(float value) {
96     return value / (1.0F + std::exp(-value));
97 }
98
99 std::span<const float> row_span(std::span<const float> values,
100                                std::int32_t row,
101                                std::int32_t row_width) {
102     const std::size_t begin = checked_size(row, "row") * checked_size(row_width,
↵    , "row_width");
103     return values.subspan(begin, checked_size(row_width, "row_width"));
104 }
105
106 void add_inplace(std::span<float> target, std::span<const float> source) {
107     if (target.size() != source.size()) {
108         throw std::runtime_error("residual add size mismatch");
109     }
110     for (std::size_t index = 0; index < target.size(); ++index) {
111         target[index] += source[index];
112     }
113 }
114
115 void swiglu_mlp(const DecoderLayerWeightsF32& layer,
116                std::span<const float> input,
117                std::span<float> output) {
118     std::vector<float> gate(checked_size(layer.w_gate.output_size, "gate"), 0.0F
↵    );
119     std::vector<float> up(checked_size(layer.w_up.output_size, "up"), 0.0F);
120     std::vector<float> hidden(gate.size(), 0.0F);
121     linear_f32(layer.w_gate, input, gate);
122     linear_f32(layer.w_up, input, up);
123     if (gate.size() != up.size()) {
124         throw std::runtime_error("SwiGLU gate/up size mismatch");
125     }
126     for (std::size_t index = 0; index < hidden.size(); ++index) {
127         hidden[index] = silu(gate[index]) * up[index];
128     }
129     linear_f32(layer.w_down, hidden, output);
130 }

```

```

131
132 float checksum(std::span<const float> values) {
133     float sum = 0.0F;
134     for (const float value : values) {
135         sum += value;
136     }
137     return sum;
138 }
139
140 float max_abs(std::span<const float> values) {
141     float result = 0.0F;
142     for (const float value : values) {
143         result = std::max(result, std::fabs(value));
144     }
145     return result;
146 }
147
148 std::vector<std::int32_t> contiguous_stride(std::span<const std::int32_t> shape ←
    ↪ ) {
149     std::vector<std::int32_t> stride(shape.size(), 1);
150     std::int32_t running = 1;
151     for (std::int32_t index = static_cast<std::int32_t>(shape.size()) - 1; ←
    ↪ index >= 0; --index) {
152         stride[checked_size(index, "stride index")] = running;
153         running *= shape[checked_size(index, "shape index")];
154     }
155     return stride;
156 }
157
158 void add_checkpoint(std::vector<ReferenceTensorCheckpoint>* checkpoints,
159                    const char* name,
160                    std::int32_t token_position,
161                    std::int32_t layer_id,
162                    std::span<const std::int32_t> shape,
163                    const char* layout,
164                    std::span<const float> values) {
165     if (checkpoints == nullptr) {
166         return;
167     }
168     ReferenceTensorCheckpoint checkpoint;
169     checkpoint.name = name;
170     checkpoint.token_position = token_position;
171     checkpoint.layer_id = layer_id;
172     checkpoint.shape.assign(shape.begin(), shape.end());
173     checkpoint.stride = contiguous_stride(shape);
174     checkpoint.dtype = LCQI_TRACE_DTYPE_F32;
175     checkpoint.layout = layout;
176     checkpoint.checksum = checksum(values);
177     checkpoint.max_abs = max_abs(values);
178     checkpoint.values.assign(values.begin(), values.end());
179     checkpoints->push_back(std::move(checkpoint));
180 }

```

```

181
182 void swiglu_mlp_with_trace(const DecoderLayerWeightsF32& layer,
183                             std::span<const float> input,
184                             std::span<float> output,
185                             std::int32_t token_position,
186                             std::int32_t layer_id,
187                             std::vector<ReferenceTensorCheckpoint>* checkpoints)↵
    ↵ {
188     std::vector<float> gate(check_size(layer.w_gate.output_size, "gate"), 0.0↵
    ↵ F);
189     std::vector<float> up(check_size(layer.w_up.output_size, "up"), 0.0F);
190     std::vector<float> hidden(gate.size(), 0.0F);
191     linear_f32(layer.w_gate, input, gate);
192     linear_f32(layer.w_up, input, up);
193     if (gate.size() != up.size()) {
194         throw std::runtime_error("SwiGLU gate/up size mismatch");
195     }
196     const std::array<std::int32_t, 1> intermediate_shape{layer.w_gate.↵
    ↵ output_size};
197     add_checkpoint(checkpoints,
198                     "mlp_gate",
199                     token_position,
200                     layer_id,
201                     intermediate_shape,
202                     LCQI_TRACE_LAYOUT_CONTIGUOUS,
203                     gate);
204     add_checkpoint(checkpoints,
205                     "mlp_up",
206                     token_position,
207                     layer_id,
208                     intermediate_shape,
209                     LCQI_TRACE_LAYOUT_CONTIGUOUS,
210                     up);
211     for (std::size_t index = 0; index < hidden.size(); ++index) {
212         hidden[index] = silu(gate[index]) * up[index];
213     }
214     add_checkpoint(checkpoints,
215                     "mlp_hidden",
216                     token_position,
217                     layer_id,
218                     intermediate_shape,
219                     LCQI_TRACE_LAYOUT_CONTIGUOUS,
220                     hidden);
221     linear_f32(layer.w_down, hidden, output);
222 }
223
224 void forward_layer(const DecoderConfig& config,
225                   const DecoderLayerWeightsF32& layer,
226                   std::int32_t layer_id,
227                   std::int32_t model_position,
228                   ReferenceKVCache& cache,
229                   std::span<float> hidden,

```

```

230         std::vector<ReferenceTensorCheckpoint>* checkpoints = ↵
↵ nullptr) {
231     validate_span_size(hidden.size(), config.hidden_size, "hidden");
232
233     std::vector<float> normed(check_size(config.hidden_size, "hidden_size"), ↵
↵ 0.0F);
234     std::vector<float> query(check_size(config.hidden_size, "hidden_size"), ↵
↵ 0.0F);
235     std::vector<float> key(check_size(kv_width(config), "kv_width"), 0.0F);
236     std::vector<float> value(check_size(kv_width(config), "kv_width"), 0.0F);
237     std::vector<float> attention(check_size(config.hidden_size, "hidden_size"↵
↵ ), 0.0F);
238     std::vector<float> attention_projected(check_size(config.hidden_size, "↵
↵ hidden_size"), 0.0F);
239     std::vector<float> mlp(check_size(config.hidden_size, "hidden_size"), 0.0↵
↵ F);
240
241     const std::array<std::int32_t, 1> hidden_shape{config.hidden_size};
242     const std::array<std::int32_t, 2> query_shape{config.query_heads, config.↵
↵ head_dim};
243     const std::array<std::int32_t, 2> kv_shape{config.kv_heads, config.head_dim↵
↵ };
244     const std::array<std::int32_t, 4> kv_slot_shape{
245         1, 1, config.kv_heads, config.head_dim};
246
247     rms_norm(hidden, layer.rms_attention_weight, config.rms_epsilon, normed);
248     add_checkpoint(checkpoints,
249         "attn_norm",
250         model_position,
251         layer_id,
252         hidden_shape,
253         LCQI_TRACE_LAYOUT_CONTIGUOUS,
254         normed);
255     linear_f32(layer.wq, normed, query);
256     linear_f32(layer.wk, normed, key);
257     linear_f32(layer.wv, normed, value);
258     add_checkpoint(checkpoints,
259         "q_before_rope",
260         model_position,
261         layer_id,
262         query_shape,
263         LCQI_TRACE_LAYOUT_CONTIGUOUS,
264         query);
265     add_checkpoint(checkpoints,
266         "k_before_rope",
267         model_position,
268         layer_id,
269         kv_shape,
270         LCQI_TRACE_LAYOUT_CONTIGUOUS,
271         key);
272     add_checkpoint(checkpoints,
273         "v",

```

```

274         model_position,
275         layer_id,
276         kv_shape,
277         LCQI_TRACE_LAYOUT_CONTIGUOUS,
278         value);
279     apply_rope(query, config.query_heads, config.head_dim, model_position, ↵
↵ config.rope_base);
280     apply_rope(key, config.kv_heads, config.head_dim, model_position, config.↵
↵ rope_base);
281     add_checkpoint(checkpoints,
282         "q_after_rope",
283         model_position,
284         layer_id,
285         query_shape,
286         LCQI_TRACE_LAYOUT_CONTIGUOUS,
287         query);
288     add_checkpoint(checkpoints,
289         "k_after_rope",
290         model_position,
291         layer_id,
292         kv_shape,
293         LCQI_TRACE_LAYOUT_CONTIGUOUS,
294         key);
295
296     cache.append(layer_id, model_position, key, value);
297     add_checkpoint(checkpoints,
298         "kv_cache_key_slot",
299         model_position,
300         layer_id,
301         kv_slot_shape,
302         LCQI_TRACE_LAYOUT_KV_CONTIGUOUS,
303         key);
304     attention_decode(config, cache, layer_id, model_position, query, attention)↵
↵ ;
305     add_checkpoint(checkpoints,
306         "attention_out",
307         model_position,
308         layer_id,
309         hidden_shape,
310         LCQI_TRACE_LAYOUT_CONTIGUOUS,
311         attention);
312     linear_f32(layer.wo, attention, attention_projected);
313     add_inplace(hidden, attention_projected);
314     add_checkpoint(checkpoints,
315         "post_attention_residual",
316         model_position,
317         layer_id,
318         hidden_shape,
319         LCQI_TRACE_LAYOUT_CONTIGUOUS,
320         hidden);
321
322     rms_norm(hidden, layer.rms_mlp_weight, config.rms_epsilon, normed);

```

```

323     add_checkpoint(checkpoints,
324                     "mlp_norm",
325                     model_position,
326                     layer_id,
327                     hidden_shape,
328                     LCQI_TRACE_LAYOUT_CONTIGUOUS,
329                     normed);
330     if (checkpoints == nullptr) {
331         swiglu_mlp(layer, normed, mlp);
332     } else {
333         swiglu_mlp_with_trace(layer, normed, mlp, model_position, layer_id, ↵
↵ checkpoints);
334     }
335     add_inplace(hidden, mlp);
336     add_checkpoint(checkpoints,
337                     "layer_out",
338                     model_position,
339                     layer_id,
340                     hidden_shape,
341                     LCQI_TRACE_LAYOUT_CONTIGUOUS,
342                     hidden);
343 }
344
345 LinearWeightsF32 make_linear(std::int32_t input_size,
346                               std::int32_t output_size,
347                               std::vector<float> weights,
348                               std::vector<float> bias = {}) {
349     LinearWeightsF32 layer;
350     layer.input_size = input_size;
351     layer.output_size = output_size;
352     layer.weights = std::move(weights);
353     layer.bias = std::move(bias);
354     validate_linear(layer, "tiny linear");
355     return layer;
356 }
357
358 std::vector<float> zeros(std::int32_t count) {
359     return std::vector<float>(checked_size(count, "zero count"), 0.0F);
360 }
361
362 } // namespace
363
364 ReferenceDecodeTraceResult run_reference_decode_with_trace(
365     const ReferenceDecoderModel& model,
366     std::span<const std::int32_t> token_ids) {
367     validate_config(model.config);
368     if (token_ids.empty()) {
369         throw std::runtime_error("reference decode needs at least one token");
370     }
371     if (token_ids.size() > checked_size(model.config.max_context, "max_context" ↵
↵ )) {
372         throw std::runtime_error("token_ids exceed max_context");

```

```

373     }
374     if (model.token_embedding.size() !=
375         checked_size(model.config.vocab_size, "vocab_size") *
376         checked_size(model.config.hidden_size, "hidden_size")) {
377         throw std::runtime_error("token embedding size mismatch");
378     }
379     if (model.final_norm_weight.size() != checked_size(model.config.hidden_size,
380 ↪ , "hidden_size")) {
381         throw std::runtime_error("final norm weight size mismatch");
382     }
383     validate_linear(model.lm_head, "lm_head");
384
385     ReferenceDecodeTraceResult trace;
386     ReferenceKVCache cache(model.config, static_cast<std::int32_t>(model.layers.
387 ↪ .size()));
388     std::vector<float> hidden(checked_size(model.config.hidden_size, "
389 ↪ hidden_size"), 0.0F);
390     const std::array<std::int32_t, 1> hidden_shape{model.config.hidden_size};
391     for (std::int32_t position = 0; position < static_cast<std::int32_t>(
392 ↪ token_ids.size());
393         ++position) {
394         const std::int32_t token_id = token_ids[checked_size(position, "
395 ↪ position")];
396         if (token_id < 0 || token_id >= model.config.vocab_size) {
397             throw std::runtime_error("token id out of vocabulary");
398         }
399         const std::span<const float> embedding =
400             row_span(model.token_embedding, token_id, model.config.hidden_size)
401 ↪ ;
402         std::copy(embedding.begin(), embedding.end(), hidden.begin());
403         add_checkpoint(&trace.checkpoints,
404             "embedding",
405             position,
406             -1,
407             hidden_shape,
408             LCQI_TRACE_LAYOUT_ROW_MAJOR,
409             hidden);
410         for (std::int32_t layer_id = 0; layer_id < static_cast<std::int32_t>(
411 ↪ model.layers.size());
412             ++layer_id) {
413             forward_layer(model.config,
414                 model.layers[checked_size(layer_id, "layer_id")],
415                 layer_id,
416                 position,
417                 cache,
418                 hidden,
419                 &trace.checkpoints);
420         }
421     }
422
423     std::vector<float> normed(checked_size(model.config.hidden_size, "
424 ↪ hidden_size"), 0.0F);

```



```

417     std::vector<float> logits(checkered_size(model.config.vocab_size, "vocab_size" ←
    ↪ "), 0.0F);
418     const std::array<std::int32_t, 1> logits_shape{model.config.vocab_size};
419     const std::int32_t last_position = static_cast<std::int32_t>(token_ids.size ←
    ↪ ()) - 1;
420     rms_norm(hidden, model.final_norm_weight, model.config.rms_epsilon, normed) ←
    ↪ ;
421     add_checkpoint(&trace.checkpoints,
422                   "final_norm",
423                   last_position,
424                   -1,
425                   hidden_shape,
426                   LCQI_TRACE_LAYOUT_CONTIGUOUS,
427                   normed);
428     linear_f32(model.lm_head, normed, logits);
429     add_checkpoint(&trace.checkpoints,
430                   "logits",
431                   last_position,
432                   -1,
433                   logits_shape,
434                   LCQI_TRACE_LAYOUT_CONTIGUOUS,
435                   logits);
436
437     trace.result.predicted_token = argmax(logits);
438     trace.result.logits = std::move(logits);
439     return trace;
440 }
441
442 ReferenceKVCache::ReferenceKVCache(const DecoderConfig& config, std::int32_t ←
    ↪ layer_count)
443 : config_(config),
444   layer_count_(layer_count) {
445     validate_config(this->config_);
446     require_positive(this->layer_count_, "layer_count");
447     const std::size_t element_count =
448         checked_size(this->layer_count_, "layer_count") *
449         checked_size(this->config_.max_context, "max_context") *
450         checked_size(this->config_.kv_heads, "kv_heads") *
451         checked_size(this->config_.head_dim, "head_dim");
452     this->keys_.assign(element_count, 0.0F);
453     this->values_.assign(element_count, 0.0F);
454     this->written_.assign(
455         checked_size(this->layer_count_, "layer_count") *
456         checked_size(this->config_.max_context, "max_context"),
457         0);
458 }
459
460 void ReferenceKVCache::append(std::int32_t layer_id,
461                               std::int32_t model_position,
462                               std::span<const float> key,
463                               std::span<const float> value) {
464     validate_span_size(key.size(), kv_width(this->config_), "KV key");

```

```

465 validate_span_size(value.size(), kv_width(this->config_), "KV value");
466 this->validate_address(layer_id, model_position, 0);
467 for (std::int32_t kv_head = 0; kv_head < this->config_.kv_heads; ++kv_head) {
468     const std::size_t target = this->base_offset(layer_id, model_position,
469     kv_head);
469     const std::size_t source =
470         checked_size(kv_head, "kv_head") * checked_size(this->config_.
471     head_dim, "head_dim");
471     std::copy_n(key.begin() + static_cast<std::ptrdiff_t>(source),
472         this->config_.head_dim,
473         this->keys_.begin() + static_cast<std::ptrdiff_t>(target));
474     std::copy_n(value.begin() + static_cast<std::ptrdiff_t>(source),
475         this->config_.head_dim,
476         this->values_.begin() + static_cast<std::ptrdiff_t>(target));
477 }
478 const std::size_t written_index =
479     checked_size(layer_id, "layer_id") * checked_size(this->config_.
480     max_context, "max_context") +
481     checked_size(model_position, "model_position");
481 this->written_[written_index] = 1;
482 this->filled_tokens_ = std::max(this->filled_tokens_, model_position + 1);
483 }
484
485 std::span<const float> ReferenceKVCache::key(std::int32_t layer_id,
486     std::int32_t model_position,
487     std::int32_t kv_head) const {
488     this->validate_address(layer_id, model_position, kv_head);
489     const std::size_t offset = this->base_offset(layer_id, model_position,
490     kv_head);
491     return std::span<const float>(this->keys_.data() + offset,
492     checked_size(this->config_.head_dim, "
493     head_dim"));
494 }
495
496 std::span<const float> ReferenceKVCache::value(std::int32_t layer_id,
497     std::int32_t model_position,
498     std::int32_t kv_head) const {
499     this->validate_address(layer_id, model_position, kv_head);
500     const std::size_t offset = this->base_offset(layer_id, model_position,
501     kv_head);
502     return std::span<const float>(this->values_.data() + offset,
503     checked_size(this->config_.head_dim, "
504     head_dim"));
505 }
506
507 std::int32_t ReferenceKVCache::layer_count() const noexcept {
508     return this->layer_count_;
509 }
510
511 std::int32_t ReferenceKVCache::filled_tokens() const noexcept {

```

```

508     return this->filled_tokens_;
509 }
510
511 std::size_t ReferenceKVCache::base_offset(std::int32_t layer_id,
512                                           std::int32_t model_position,
513                                           std::int32_t kv_head) const {
514     return (((checked_size(layer_id, "layer_id") *
515                    checked_size(this->config.max_context, "max_context")) +
516            checked_size(model_position, "model_position")) *
517            checked_size(this->config.kv_heads, "kv_heads") +
518            checked_size(kv_head, "kv_head")) *
519            checked_size(this->config.head_dim, "head_dim");
520 }
521
522 void ReferenceKVCache::validate_address(std::int32_t layer_id,
523                                         std::int32_t model_position,
524                                         std::int32_t kv_head) const {
525     if (layer_id < 0 || layer_id >= this->layer_count_) {
526         throw std::runtime_error("KV layer_id out of range");
527     }
528     if (model_position < 0 || model_position >= this->config.max_context) {
529         throw std::runtime_error("KV model_position out of range");
530     }
531     if (kv_head < 0 || kv_head >= this->config.kv_heads) {
532         throw std::runtime_error("KV head out of range");
533     }
534 }
535
536 void rms_norm(std::span<const float> input,
537              std::span<const float> weight,
538              float epsilon,
539              std::span<float> output) {
540     if (input.size() != weight.size() || input.size() != output.size()) {
541         throw std::runtime_error("RMSNorm size mismatch");
542     }
543     if (input.empty()) {
544         throw std::runtime_error("RMSNorm input is empty");
545     }
546     float sum_square = 0.0F;
547     for (const float value : input) {
548         sum_square += value * value;
549     }
550     const float mean_square = sum_square / static_cast<float>(input.size());
551     const float scale = 1.0F / std::sqrt(mean_square + epsilon);
552     for (std::size_t index = 0; index < input.size(); ++index) {
553         output[index] = input[index] * scale * weight[index];
554     }
555 }
556
557 void linear_f32(const LinearWeightsF32& layer,
558               std::span<const float> input,
559               std::span<float> output) {

```

```

560 validate_linear(layer, "linear_f32");
561 validate_span_size(input.size(), layer.input_size, "linear input");
562 validate_span_size(output.size(), layer.output_size, "linear output");
563 for (std::int32_t out = 0; out < layer.output_size; ++out) {
564     const std::size_t row_base =
565         checked_size(out, "out") * checked_size(layer.input_size, "↵
↵ input_size");
566     float accumulator = layer.bias.empty() ? 0.0F : layer.bias[checked_size↵
↵ (out, "out")];
567     for (std::int32_t in = 0; in < layer.input_size; ++in) {
568         accumulator += layer.weights[row_base + checked_size(in, "in")] *
569             input[checked_size(in, "in")];
570     }
571     output[checked_size(out, "out")] = accumulator;
572 }
573 }
574
575 void apply_rope(std::span<float> heads,
576                std::int32_t head_count,
577                std::int32_t head_dim,
578                std::int32_t model_position,
579                float rope_base) {
580     require_positive(head_count, "head_count");
581     require_positive(head_dim, "head_dim");
582     validate_span_size(heads.size(), head_count * head_dim, "RoPE heads");
583     if (head_dim % LCQI_ROPE_PAIR_WIDTH != 0) {
584         throw std::runtime_error("RoPE head_dim must be even");
585     }
586     if (model_position < 0) {
587         throw std::runtime_error("RoPE model_position cannot be negative");
588     }
589     if (rope_base <= 1.0F) {
590         throw std::runtime_error("RoPE base must be greater than 1");
591     }
592     for (std::int32_t head = 0; head < head_count; ++head) {
593         for (std::int32_t offset = 0; offset < head_dim; offset += ↵
↵ LCQI_ROPE_PAIR_WIDTH) {
594             const float exponent = static_cast<float>(offset) / static_cast<↵
↵ float>(head_dim);
595             const float theta = 1.0F / std::pow(rope_base, exponent);
596             const float angle = static_cast<float>(model_position) * theta;
597             const float cosine = std::cos(angle);
598             const float sine = std::sin(angle);
599             const std::size_t first =
600                 checked_size(head, "head") * checked_size(head_dim, "head_dim")↵
↵ +
601                 checked_size(offset, "offset");
602             const float x0 = heads[first];
603             const float x1 = heads[first + 1];
604             heads[first] = x0 * cosine - x1 * sine;
605             heads[first + 1] = x0 * sine + x1 * cosine;
606         }

```

```

607     }
608 }
609
610 void attention_decode(const DecoderConfig& config,
611                     const ReferenceKVCache& cache,
612                     std::int32_t layer_id,
613                     std::int32_t model_position,
614                     std::span<const float> query,
615                     std::span<float> output) {
616     validate_config(config);
617     validate_span_size(query.size(), config.hidden_size, "attention query");
618     validate_span_size(output.size(), config.hidden_size, "attention output");
619     if (model_position < 0 || model_position >= cache.filled_tokens()) {
620         throw std::runtime_error("attention model_position has not been written");
621     }
622     std::fill(output.begin(), output.end(), 0.0F);
623
624     const std::int32_t group_size = config.query_heads / config.kv_heads;
625     const float score_scale = 1.0F / std::sqrt(static_cast<float>(config.head_dim));
626     std::vector<float> scores(check_size(model_position + 1, "score count"),
627                             LCQI_SOFTMAX_NEGATIVE_INFINITY);
628     std::vector<float> probabilities(scores.size(), 0.0F);
629
630     for (std::int32_t query_head = 0; query_head < config.query_heads; ++query_head) {
631         const std::int32_t kv_head = query_head / group_size;
632         const std::size_t query_base =
633             checked_size(query_head, "query_head") * checked_size(config.head_dim, "head_dim");
634         float max_score = LCQI_SOFTMAX_NEGATIVE_INFINITY;
635         for (std::int32_t position = 0; position <= model_position; ++position) {
636             const std::span<const float> key = cache.key(layer_id, position, kv_head);
637             float dot = 0.0F;
638             for (std::int32_t dim = 0; dim < config.head_dim; ++dim) {
639                 dot += query[query_base + checked_size(dim, "dim")] *
640                     key[checked_size(dim, "dim")];
641             }
642             const float score = dot * score_scale;
643             scores[checked_size(position, "position")] = score;
644             max_score = std::max(max_score, score);
645         }
646
647         float probability_sum = 0.0F;
648         for (std::int32_t position = 0; position <= model_position; ++position) {
649             const std::size_t index = checked_size(position, "position");
650             const float unnormalized = std::exp(scores[index] - max_score);
651             probabilities[index] = unnormalized;

```

```

652         probability_sum += unnormalized;
653     }
654     if (probability_sum <= 0.0F) {
655         throw std::runtime_error("attention probability sum is invalid");
656     }
657     for (std::int32_t position = 0; position <= model_position; ++position)↵
↵ {
658         const float probability =
659             probabilities[checked_size(position, "position")] / ↵
↵ probability_sum;
660         const std::span<const float> value = cache.value(layer_id, position↵
↵ , kv_head);
661         for (std::int32_t dim = 0; dim < config.head_dim; ++dim) {
662             output[query_base + checked_size(dim, "dim")] +=
663                 probability * value[checked_size(dim, "dim")];
664         }
665     }
666 }
667 }
668
669 ReferenceDecodeResult run_reference_decode(const ReferenceDecoderModel& model,
670                                           std::span<const std::int32_t> ↵
↵ token_ids) {
671     validate_config(model.config);
672     if (token_ids.empty()) {
673         throw std::runtime_error("reference decode needs at least one token");
674     }
675     if (token_ids.size() > checked_size(model.config.max_context, "max_context"↵
↵ )) {
676         throw std::runtime_error("token_ids exceed max_context");
677     }
678     if (model.token_embedding.size() !=
679         checked_size(model.config.vocab_size, "vocab_size") *
680         checked_size(model.config.hidden_size, "hidden_size")) {
681         throw std::runtime_error("token embedding size mismatch");
682     }
683     if (model.final_norm_weight.size() != checked_size(model.config.hidden_size↵
↵ , "hidden_size")) {
684         throw std::runtime_error("final norm weight size mismatch");
685     }
686     validate_linear(model.lm_head, "lm_head");
687
688     ReferenceKVCache cache(model.config, static_cast<std::int32_t>(model.layers↵
↵ .size()));
689     std::vector<float> hidden(checked_size(model.config.hidden_size, "↵
↵ hidden_size"), 0.0F);
690     for (std::int32_t position = 0; position < static_cast<std::int32_t>(↵
↵ token_ids.size());
691         ++position) {
692         const std::int32_t token_id = token_ids[checked_size(position, "↵
↵ position")];
693         if (token_id < 0 || token_id >= model.config.vocab_size) {

```

```

694         throw std::runtime_error("token id out of vocabulary");
695     }
696     const std::span<const float> embedding =
697         row_span(model.token_embedding, token_id, model.config.hidden_size)↵
↵ ;
698     std::copy(embedding.begin(), embedding.end(), hidden.begin());
699     for (std::int32_t layer_id = 0; layer_id < static_cast<std::int32_t>(↵
↵ model.layers.size());
700         ++layer_id) {
701         forward_layer(model.config,
702             model.layers[checked_size(layer_id, "layer_id")],
703             layer_id,
704             position,
705             cache,
706             hidden);
707     }
708 }
709
710 std::vector<float> normed(checked_size(model.config.hidden_size, "↵
↵ hidden_size"), 0.0F);
711 std::vector<float> logits(checked_size(model.config.vocab_size, "vocab_size↵
↵ "), 0.0F);
712 rms_norm(hidden, model.final_norm_weight, model.config.rms_epsilon, normed)↵
↵ ;
713 linear_f32(model.lm_head, normed, logits);
714
715 ReferenceDecodeResult result;
716 result.predicted_token = argmax(logits);
717 result.logits = std::move(logits);
718 return result;
719 }
720
721 ReferenceDecoderModel make_tiny_reference_decoder_model() {
722     ReferenceDecoderModel model;
723     model.config.hidden_size = 4;
724     model.config.query_heads = 2;
725     model.config.kv_heads = 1;
726     model.config.head_dim = 2;
727     model.config.intermediate_size = 6;
728     model.config.vocab_size = 5;
729     model.config.max_context = 8;
730     model.config.rms_epsilon = 1.0e-5F;
731     model.config.rope_base = 10000.0F;
732
733     model.token_embedding = {
734         0.20F, -0.10F, 0.05F, 0.30F,
735         -0.40F, 0.10F, 0.25F, -0.20F,
736         0.15F, 0.35F, -0.30F, 0.05F,
737         0.50F, -0.25F, 0.10F, -0.15F,
738         -0.05F, 0.20F, 0.40F, -0.35F,
739     };
740

```



```

741 DecoderLayerWeightsF32 layer;
742 layer.rms_attention_weight = {1.0F, 0.9F, 1.1F, 1.0F};
743 layer.wq = make_linear(4,
744                        4,
745                        {
746                            0.20F, -0.10F, 0.05F, 0.30F,
747                            -0.15F, 0.25F, 0.10F, -0.05F,
748                            0.05F, 0.10F, -0.20F, 0.15F,
749                            0.30F, -0.05F, 0.20F, 0.10F,
750                        });
751 layer.wk = make_linear(4,
752                        2,
753                        {
754                            0.10F, 0.20F, -0.10F, 0.05F,
755                            -0.05F, 0.15F, 0.25F, -0.20F,
756                        });
757 layer.wv = make_linear(4,
758                        2,
759                        {
760                            0.25F, -0.05F, 0.10F, 0.15F,
761                            -0.10F, 0.30F, 0.05F, 0.20F,
762                        });
763 layer.wo = make_linear(4,
764                        4,
765                        {
766                            0.20F, 0.10F, -0.05F, 0.15F,
767                            -0.10F, 0.25F, 0.20F, -0.05F,
768                            0.05F, -0.20F, 0.30F, 0.10F,
769                            0.15F, 0.05F, -0.10F, 0.25F,
770                        });
771 layer.rms_mlp_weight = {0.95F, 1.05F, 1.0F, 0.9F};
772 layer.w_gate = make_linear(4,
773                            6,
774                            {
775                                0.20F, -0.10F, 0.05F, 0.10F,
776                                -0.05F, 0.15F, 0.20F, -0.10F,
777                                0.10F, 0.05F, -0.15F, 0.25F,
778                                -0.20F, 0.10F, 0.15F, 0.05F,
779                                0.05F, -0.25F, 0.10F, 0.20F,
780                                0.15F, 0.05F, 0.05F, -0.15F,
781                            });
782 layer.w_up = make_linear(4,
783                            6,
784                            {
785                                0.10F, 0.20F, -0.05F, 0.05F,
786                                -0.15F, 0.05F, 0.25F, 0.10F,
787                                0.20F, -0.10F, 0.05F, 0.15F,
788                                0.05F, 0.10F, -0.20F, 0.25F,
789                                -0.05F, 0.15F, 0.10F, -0.10F,
790                                0.25F, -0.05F, 0.05F, 0.10F,
791                            });
792 layer.w_down = make_linear(6,

```

```

793         4,
794         {
795             0.10F, -0.05F, 0.15F, 0.20F, -0.10F, 0.05F,
796             -0.20F, 0.10F, 0.05F, -0.05F, 0.15F, 0.20F,
797             0.05F, 0.15F, -0.10F, 0.10F, 0.20F, -0.05F,
798             0.20F, 0.05F, 0.10F, -0.15F, -0.05F, 0.10F,
799         });
800     model.layers.push_back(std::move(layer));
801     model.final_norm_weight = {1.0F, 0.95F, 1.05F, 1.0F};
802     model.lm_head = make_linear(4,
803         5,
804         {
805             0.20F, -0.10F, 0.05F, 0.15F,
806             -0.05F, 0.25F, 0.10F, -0.20F,
807             0.15F, 0.05F, -0.25F, 0.10F,
808             -0.10F, 0.20F, 0.15F, 0.05F,
809             0.05F, -0.15F, 0.20F, 0.25F,
810         },
811         zeros(5));
812     return model;
813 }
814
815 } // namespace lcqi

```

## Tokenizer 与 Safetensors 实现

`tokenizer.cpp` 把 prompt 变成 token id, 并固定 unknown token 行为。读者应检查 BOS/EOS 是否稳定插入, UNK 是否按合同处理。

Listing 10.9 src/tokenizer.cpp

```

1  #include <lcqi/tokenizer.hpp>
2
3  #include <cstdint>
4  #include <fstream>
5  #include <stdexcept>
6  #include <string>
7  #include <string_view>
8
9  namespace lcqi {
10 namespace {
11
12 constexpr std::uint64_t LCQI_FNV_OFFSET_BASIS = 1469598103934665603ULL;
13 constexpr std::uint64_t LCQI_FNV_PRIME = 1099511628211ULL;
14 constexpr std::int32_t LCQI_INVALID_TOKEN_ID = -1;
15
16 std::string read_key(std::istream& input) {
17     std::string key;
18     if (!(input >> key)) {
19         throw std::runtime_error("unexpected end of tokenizer file");
20     }
21     return key;

```

```

22 }
23
24 void expect_key(std::istream& input, const std::string& expected) {
25     const std::string key = read_key(input);
26     if (key != expected) {
27         throw std::runtime_error("expected tokenizer key '" + expected + "', ←
↪ got '" + key + "'");
28     }
29 }
30
31 std::int32_t read_i32(std::istream& input, const std::string& key) {
32     expect_key(input, key);
33     std::int32_t value = 0;
34     if (!(input >> value)) {
35         throw std::runtime_error("invalid tokenizer int32 value for key '" + ←
↪ key + "'");
36     }
37     return value;
38 }
39
40 std::string read_string(std::istream& input, const std::string& key) {
41     expect_key(input, key);
42     std::string value;
43     if (!(input >> value)) {
44         throw std::runtime_error("invalid tokenizer string value for key '" + ←
↪ key + "'");
45     }
46     return value;
47 }
48
49 std::int32_t lookup_token_id(const TokenizerModel& tokenizer, std::string_view ←
↪ text) {
50     for (const TokenEntry& entry : tokenizer.vocab) {
51         if (entry.text == text) {
52             return entry.id;
53         }
54     }
55     return LCQI_INVALID_TOKEN_ID;
56 }
57
58 void validate_tokenizer(const TokenizerModel& tokenizer) {
59     if (tokenizer.vocab.empty()) {
60         throw std::runtime_error("tokenizer vocab is empty");
61     }
62     for (std::size_t i = 0; i < tokenizer.vocab.size(); ++i) {
63         const TokenEntry& left = tokenizer.vocab[i];
64         if (left.id < 0) {
65             throw std::runtime_error("tokenizer token id must be non-negative")↪
↪ ;
66         }
67         for (std::size_t j = i + 1; j < tokenizer.vocab.size(); ++j) {
68             const TokenEntry& right = tokenizer.vocab[j];

```

```

69         if (left.id == right.id) {
70             throw std::runtime_error("duplicate tokenizer token id");
71         }
72         if (left.text == right.text) {
73             throw std::runtime_error("duplicate tokenizer token text");
74         }
75     }
76 }
77 if (tokenizer.bos_id != LCQI_INVALID_TOKEN_ID &&
78     lookup_token_id(tokenizer, "<bos>") != tokenizer.bos_id) {
79     throw std::runtime_error("tokenizer BOS id does not match vocab");
80 }
81 if (tokenizer.eos_id != LCQI_INVALID_TOKEN_ID &&
82     lookup_token_id(tokenizer, "<eos>") != tokenizer.eos_id) {
83     throw std::runtime_error("tokenizer EOS id does not match vocab");
84 }
85 }
86
87 std::uint64_t hash_byte(std::uint64_t hash, unsigned char value) {
88     hash ^= static_cast<std::uint64_t>(value);
89     hash *= LCQI_FNV_PRIME;
90     return hash;
91 }
92
93 std::uint64_t hash_string(std::uint64_t hash, std::string_view text) {
94     for (const unsigned char value : text) {
95         hash = hash_byte(hash, value);
96     }
97     return hash;
98 }
99
100 std::uint64_t hash_i32(std::uint64_t hash, std::int32_t value) {
101     const std::uint32_t bits = static_cast<std::uint32_t>(value);
102     for (std::int32_t shift = 0; shift < 32; shift += 8) {
103         hash = hash_byte(hash, static_cast<unsigned char>((bits >> shift) & 0x
↵ xFFU));
104     }
105     return hash;
106 }
107
108 } // namespace
109
110 TokenizerModel load_tokenizer(const std::filesystem::path& path) {
111     std::ifstream input(path);
112     if (!input) {
113         throw std::runtime_error("cannot open tokenizer file: " + path.string()
↵ );
114     }
115
116     expect_key(input, "LCQI_TOKENIZER_V1");
117
118     TokenizerModel tokenizer;

```

```

119     tokenizer.bos_id = read_i32(input, "bos_id");
120     tokenizer.eos_id = read_i32(input, "eos_id");
121     tokenizer.unk_id = read_i32(input, "unk_id");
122     tokenizer.chat_template_hash = read_string(input, "chat_template_hash");
123     const std::int32_t vocab_size = read_i32(input, "vocab_size");
124     if (vocab_size <= 0) {
125         throw std::runtime_error("tokenizer vocab_size must be positive");
126     }
127
128     tokenizer.vocab.reserve(static_cast<std::size_t>(vocab_size));
129     expect_key(input, "vocab");
130     for (std::int32_t i = 0; i < vocab_size; ++i) {
131         TokenEntry entry;
132         if (!(input >> entry.id >> entry.text)) {
133             throw std::runtime_error("invalid tokenizer vocab entry");
134         }
135         tokenizer.vocab.push_back(entry);
136     }
137
138     validate_tokenizer(tokenizer);
139     return tokenizer;
140 }
141
142 std::vector<std::int32_t> encode_prompt(const TokenizerModel& tokenizer,
143                                         std::string_view prompt) {
144     std::vector<std::int32_t> ids;
145     if (tokenizer.bos_id != LCQI_INVALID_TOKEN_ID) {
146         ids.push_back(tokenizer.bos_id);
147     }
148
149     std::size_t begin = 0;
150     while (begin < prompt.size()) {
151         while (begin < prompt.size() && prompt[begin] == ' ') {
152             ++begin;
153         }
154         if (begin >= prompt.size()) {
155             break;
156         }
157         std::size_t end = begin;
158         while (end < prompt.size() && prompt[end] != ' ') {
159             ++end;
160         }
161         const std::string_view token = prompt.substr(begin, end - begin);
162         const std::int32_t token_id = lookup_token_id(tokenizer, token);
163         if (token_id == LCQI_INVALID_TOKEN_ID) {
164             if (tokenizer.unk_id == LCQI_INVALID_TOKEN_ID) {
165                 throw std::runtime_error("tokenizer encountered unknown token");
166             }
167             ids.push_back(tokenizer.unk_id);
168         } else {
169             ids.push_back(token_id);

```

```

170     }
171     begin = end;
172 }
173
174 if (tokenizer.eos_id != LCQI_INVALID_TOKEN_ID) {
175     ids.push_back(tokenizer.eos_id);
176 }
177 return ids;
178 }
179
180 std::uint64_t tokenizer_contract_hash(const TokenizerModel& tokenizer) {
181     std::uint64_t hash = LCQI_FNV_OFFSET_BASIS;
182     hash = hash_string(hash, tokenizer.chat_template_hash);
183     hash = hash_i32(hash, tokenizer.bos_id);
184     hash = hash_i32(hash, tokenizer.eos_id);
185     hash = hash_i32(hash, tokenizer.unk_id);
186     for (const TokenEntry& entry : tokenizer.vocab) {
187         hash = hash_string(hash, entry.text);
188         hash = hash_i32(hash, entry.id);
189     }
190     return hash;
191 }
192
193 } // namespace lcqi

```

**safetensors.cpp** 解析 header、metadata、dtype、shape 和 data offsets。坏 dtype、坏 shape、坏 offset 和 payload 太短都应该被拒绝。

**Listing 10.10** src/safetensors.cpp

```

1 #include <lcqi/safetensors.hpp>
2
3 #include <algorithm>
4 #include <cctype>
5 #include <cstdint>
6 #include <cmath>
7 #include <cstring>
8 #include <fstream>
9 #include <limits>
10 #include <stdexcept>
11 #include <string>
12 #include <string_view>
13 #include <vector>
14
15 namespace lcqi {
16 namespace {
17
18 constexpr std::size_t LCQI_SAFETENSORS_PREFIX_BYTES = 8;
19 constexpr std::uint64_t LCQI_SAFETENSORS_MAX_HEADER_BYTES = 16U * 1024U * 1024U↵
    ↵ ;
20 constexpr std::size_t LCQI_SAFETENSORS_OFFSET_COUNT = 2;
21 constexpr std::uint64_t LCQI_BYTE_BITS = 8;
22 constexpr std::uint64_t LCQI_DECIMAL_BASE = 10;

```

```

23 constexpr std::uint64_t LCQI_BYTE_MASK = 0xFFU;
24 constexpr unsigned char LCQI_JSON_CONTROL_CHAR_LIMIT = 0x20U;
25 constexpr std::uint64_t LCQI_DTYPE_BYTES_F64 = 8;
26 constexpr std::uint64_t LCQI_DTYPE_BYTES_F32 = 4;
27 constexpr std::uint64_t LCQI_DTYPE_BYTES_F16 = 2;
28 constexpr std::uint64_t LCQI_DTYPE_BYTES_I64 = 8;
29 constexpr std::uint64_t LCQI_DTYPE_BYTES_I32 = 4;
30 constexpr std::uint64_t LCQI_DTYPE_BYTES_I16 = 2;
31 constexpr std::uint64_t LCQI_DTYPE_BYTES_I8 = 1;
32 constexpr std::uint64_t LCQI_DTYPE_BYTES_U8 = 1;
33 constexpr std::uint32_t LCQI_F32_EXPONENT_MASK = 0x7F800000U;
34 constexpr std::uint32_t LCQI_F16_SIGN_MASK = 0x8000U;
35 constexpr std::uint32_t LCQI_F16_EXPONENT_MASK = 0x7C00U;
36 constexpr std::uint32_t LCQI_F16_MANTISSA_MASK = 0x03FFU;
37 constexpr std::uint32_t LCQI_F16_EXPONENT_BIAS = 15U;
38 constexpr std::uint32_t LCQI_F32_EXPONENT_BIAS = 127U;
39 constexpr std::uint32_t LCQI_F32_MANTISSA_BITS = 23U;
40 constexpr std::uint32_t LCQI_F16_MANTISSA_BITS = 10U;
41 constexpr std::uint32_t LCQI_F16_TO_F32_SIGN_SHIFT = 16U;
42
43 class HeaderCursor {
44 public:
45     explicit HeaderCursor(std::string_view text) : text_(text) {}
46
47     void expect(char expected) {
48         this->skip_ws();
49         if (this->eof() || this->text_[this->position_] != expected) {
50             throw std::runtime_error("invalid safetensors header syntax");
51         }
52         ++this->position_;
53     }
54
55     [[nodiscard]] bool consume(char expected) {
56         this->skip_ws();
57         if (!this->eof() && this->text_[this->position_] == expected) {
58             ++this->position_;
59             return true;
60         }
61         return false;
62     }
63
64     [[nodiscard]] std::string parse_string() {
65         this->skip_ws();
66         this->expect('\"');
67         std::string result;
68         while (!this->eof()) {
69             const char ch = this->text_[this->position_++];
70             if (ch == '\"') {
71                 return result;
72             }
73             if (static_cast<unsigned char>(ch) < LCQI_JSON_CONTROL_CHAR_LIMIT) ←
↪ {

```



```

74         throw std::runtime_error("control character in safetensors ↵
↵ string");
75     }
76     if (ch == '\\') {
77         result.push_back(this->parse_escape());
78     } else {
79         result.push_back(ch);
80     }
81 }
82 throw std::runtime_error("unterminated safetensors string");
83 }
84
85 [[nodiscard]] std::uint64_t parse_u64() {
86     this->skip_ws();
87     if (this->eof() ||
88 ↵ !std::isdigit(static_cast<unsigned char>(this->text_[this->↵
↵ position_]))) {
89         throw std::runtime_error("expected unsigned integer in safetensors ↵
↵ header");
90     }
91     std::uint64_t value = 0;
92     while (!this->eof() &&
93 ↵ std::isdigit(static_cast<unsigned char>(this->text_[this->↵
↵ position_]))) {
94         const std::uint64_t digit =
95             static_cast<std::uint64_t>(this->text_[this->position_] - '0');
96         if (value >
97 ↵ (std::numeric_limits<std::uint64_t>::max() - digit) / ↵
↵ LCQI_DECIMAL_BASE) {
98             throw std::runtime_error("safetensors integer overflow");
99         }
100         value = value * LCQI_DECIMAL_BASE + digit;
101         ++this->position_;
102     }
103     return value;
104 }
105
106 [[nodiscard]] std::vector<std::int64_t> parse_i64_array() {
107     std::vector<std::int64_t> values;
108     this->expect('[');
109     if (this->consume(' ')) {
110         return values;
111     }
112     while (true) {
113         const std::uint64_t raw = this->parse_u64();
114         if (raw > static_cast<std::uint64_t>(std::numeric_limits<std::↵
↵ int64_t>::max())) {
115             throw std::runtime_error("safetensors shape value is too large"↵
↵ );
116         }
117         values.push_back(static_cast<std::int64_t>(raw));
118         if (this->consume(' ')) {

```

```

119         return values;
120     }
121     this->expect(',');
122 }
123 }
124
125 [[nodiscard]] std::vector<std::uint64_t> parse_u64_array() {
126     std::vector<std::uint64_t> values;
127     this->expect('[');
128     if (this->consume(' ')) {
129         return values;
130     }
131     while (true) {
132         values.push_back(this->parse_u64());
133         if (this->consume(' ')) {
134             return values;
135         }
136         this->expect(',');
137     }
138 }
139
140 [[nodiscard]] bool eof_after_ws() {
141     this->skip_ws();
142     return this->eof();
143 }
144
145 private:
146     std::string_view text_;
147     std::size_t position_ = 0;
148
149     [[nodiscard]] bool eof() const {
150         return this->position_ >= this->text_.size();
151     }
152
153     void skip_ws() {
154         while (!this->eof() &&
155             std::isspace(static_cast<unsigned char>(this->text_[this->position_])) {
156             ++this->position_;
157         }
158     }
159
160     [[nodiscard]] char parse_escape() {
161         if (this->eof()) {
162             throw std::runtime_error("unterminated safetensors escape");
163         }
164         const char escaped = this->text_[this->position_++];
165         switch (escaped) {
166             case '"':
167             case '\\':
168             case '/':
169                 return escaped;

```

```

170         case 'b':
171             return '\\b';
172         case 'f':
173             return '\\f';
174         case 'n':
175             return '\\n';
176         case 'r':
177             return '\\r';
178         case 't':
179             return '\\t';
180         default:
181             throw std::runtime_error("unsupported safetensors string escape↵
↵ ");
182     }
183 }
184 };
185
186 std::uint64_t read_le_u64(std::span<const std::uint8_t> bytes) {
187     if (bytes.size() < LCQI_SAFETENSORS_PREFIX_BYTES) {
188         throw std::runtime_error("safetensors file is too small");
189     }
190     std::uint64_t value = 0;
191     for (std::size_t i = 0; i < LCQI_SAFETENSORS_PREFIX_BYTES; ++i) {
192         value |= static_cast<std::uint64_t>(bytes[i]) << (LCQI_BYTE_BITS * i);
193     }
194     return value;
195 }
196
197 std::string read_file_bytes(const std::filesystem::path& path) {
198     std::ifstream input(path, std::ios::binary);
199     if (!input) {
200         throw std::runtime_error("cannot open safetensors file: " + path.string↵
↵ ());
201     }
202     input.seekg(0, std::ios::end);
203     const std::streamoff size = input.tellg();
204     if (size < 0) {
205         throw std::runtime_error("cannot determine safetensors file size");
206     }
207     input.seekg(0, std::ios::beg);
208     std::string bytes(static_cast<std::size_t>(size), '\\0');
209     if (!bytes.empty() && !input.read(bytes.data(), static_cast<std::streamsize>↵
↵ >(bytes.size()))) {
210         throw std::runtime_error("cannot read safetensors file");
211     }
212     return bytes;
213 }
214
215 std::vector<std::uint8_t> read_file_u8(const std::filesystem::path& path) {
216     std::ifstream input(path, std::ios::binary);
217     if (!input) {
218         throw std::runtime_error("cannot open safetensors file: " + path.string↵

```

```

    ↪ ();
219 }
220 input.seekg(0, std::ios::end);
221 const std::streamoff size = input.tellg();
222 if (size < 0) {
223     throw std::runtime_error("cannot determine safetensors file size");
224 }
225 input.seekg(0, std::ios::beg);
226 std::vector<std::uint8_t> bytes(static_cast<std::size_t>(size), 0);
227 if (!bytes.empty() &&
228     !input.read(reinterpret_cast<char*>(bytes.data()),
229                 static_cast<std::streamsize>(bytes.size()))) {
230     throw std::runtime_error("cannot read safetensors file");
231 }
232 return bytes;
233 }
234
235 void parse_metadata_value(HeaderCursor& cursor,
236                           SafeTensorManifest& manifest,
237                           const std::string& key) {
238     if (key == "__metadata__") {
239         cursor.expect('{');
240         if (cursor.consume('}')) {
241             return;
242         }
243         while (true) {
244             const std::string metadata_key = cursor.parse_string();
245             cursor.expect(':');
246             const std::string metadata_value = cursor.parse_string();
247             manifest.metadata.emplace_back(metadata_key, metadata_value);
248             if (cursor.consume('}')) {
249                 return;
250             }
251             cursor.expect(',');
252         }
253     }
254     throw std::runtime_error("unexpected safetensors metadata value");
255 }
256
257 SafeTensorEntry parse_tensor_entry(HeaderCursor& cursor, const std::string& ↪
    ↪ name) {
258     SafeTensorEntry entry;
259     entry.name = name;
260     bool saw_dtype = false;
261     bool saw_shape = false;
262     bool saw_offsets = false;
263
264     cursor.expect('{');
265     if (cursor.consume('}')) {
266         throw std::runtime_error("empty safetensors tensor entry");
267     }
268     while (true) {

```

```

269     const std::string field = cursor.parse_string();
270     cursor.expect(':');
271     if (field == "dtype") {
272         entry.dtype = cursor.parse_string();
273         saw_dtype = true;
274     } else if (field == "shape") {
275         entry.shape = cursor.parse_i64_array();
276         saw_shape = true;
277     } else if (field == "data_offsets") {
278         const std::vector<std::uint64_t> offsets = cursor.parse_u64_array();
279         ↪ ;
280         if (offsets.size() != LCQI_SAFETENSORS_OFFSET_COUNT) {
281             throw std::runtime_error("safetensors data_offsets must contain ↪
282             ↪ two values");
283         }
284         entry.data_begin = offsets[0];
285         entry.data_end = offsets[1];
286         saw_offsets = true;
287     } else {
288         throw std::runtime_error("unsupported safetensors tensor field: " + ↪
289         ↪ field);
290     }
291     if (cursor.consume('}')) {
292         break;
293     }
294     cursor.expect(',');
295 }
296
297 if (!saw_dtype || !saw_shape || !saw_offsets) {
298     throw std::runtime_error("safetensors tensor entry is missing required ↪
299     ↪ fields");
300 }
301 if (entry.data_end < entry.data_begin) {
302     throw std::runtime_error("safetensors tensor offset range is invalid");
303 }
304 return entry;
305 }
306
307 void parse_header(std::string_view header, SafeTensorManifest& manifest) {
308     HeaderCursor cursor(header);
309     cursor.expect('{');
310     if (cursor.consume('}')) {
311         throw std::runtime_error("safetensors header is empty");
312     }
313
314     while (true) {
315         const std::string key = cursor.parse_string();
316         cursor.expect(':');
317         if (key == "__metadata__") {
318             parse_metadata_value(cursor, manifest, key);
319         } else {
320             manifest.tensors.push_back(parse_tensor_entry(cursor, key));

```

```

317     }
318     if (cursor.consume('}')) {
319         break;
320     }
321     cursor.expect(',');
322 }
323
324 if (!cursor.eof_after_ws()) {
325     throw std::runtime_error("trailing data after safetensors header");
326 }
327 }
328
329 std::uint64_t expected_tensor_bytes(const SafeTensorEntry& entry);
330
331 void validate_manifest(const SafeTensorManifest& manifest,
332                       std::uint64_t file_size) {
333     const std::uint64_t payload_size = file_size - manifest.data_start_offset;
334     for (std::size_t i = 0; i < manifest.tensors.size(); ++i) {
335         const SafeTensorEntry& left = manifest.tensors[i];
336         if (left.name.empty() || left.dtype.empty()) {
337             throw std::runtime_error("safetensors tensor has empty name or ↵
↵ dtype");
338         }
339         if (left.data_end > payload_size) {
340             throw std::runtime_error("safetensors tensor data_offsets exceed ↵
↵ payload size");
341         }
342         const std::uint64_t expected_bytes = expected_tensor_bytes(left);
343         if (left.byte_size() != expected_bytes) {
344             throw std::runtime_error("safetensors tensor byte size does not ↵
↵ match dtype and shape");
345         }
346         for (const std::int64_t dim : left.shape) {
347             if (dim < 0) {
348                 throw std::runtime_error("safetensors shape dimension is ↵
↵ negative");
349             }
350         }
351         for (std::size_t j = i + 1; j < manifest.tensors.size(); ++j) {
352             const SafeTensorEntry& right = manifest.tensors[j];
353             if (left.name == right.name) {
354                 throw std::runtime_error("duplicate safetensors tensor name");
355             }
356             const bool overlaps =
357                 left.data_begin < right.data_end && right.data_begin < left.↵
↵ data_end;
358             if (overlaps) {
359                 throw std::runtime_error("overlapping safetensors tensor data ↵
↵ range");
360             }
361         }
362     }

```

```

363 }
364
365 std::uint64_t dtype_byte_size(std::string_view dtype) {
366     if (dtype == "F64") {
367         return LCQI_DTYPE_BYTES_F64;
368     }
369     if (dtype == "F32") {
370         return LCQI_DTYPE_BYTES_F32;
371     }
372     if (dtype == "F16" || dtype == "BF16") {
373         return LCQI_DTYPE_BYTES_F16;
374     }
375     if (dtype == "I64" || dtype == "U64") {
376         return LCQI_DTYPE_BYTES_I64;
377     }
378     if (dtype == "I32" || dtype == "U32") {
379         return LCQI_DTYPE_BYTES_I32;
380     }
381     if (dtype == "I16" || dtype == "U16") {
382         return LCQI_DTYPE_BYTES_I16;
383     }
384     if (dtype == "I8") {
385         return LCQI_DTYPE_BYTES_I8;
386     }
387     if (dtype == "U8" || dtype == "B00L") {
388         return LCQI_DTYPE_BYTES_U8;
389     }
390     throw std::runtime_error("unsupported safetensors dtype: " + std::string(
391         dtype));
392 }
393
394 std::uint64_t expected_tensor_bytes(const SafeTensorEntry& entry) {
395     std::uint64_t element_count = 1;
396     for (const std::int64_t dim : entry.shape) {
397         if (dim < 0) {
398             throw std::runtime_error("safetensors shape dimension is negative");
399         }
400         const std::uint64_t extent = static_cast<std::uint64_t>(dim);
401         if (extent != 0 &&
402             element_count > std::numeric_limits<std::uint64_t>::max() / extent) {
403             throw std::runtime_error("safetensors shape element count overflow");
404         }
405         element_count *= extent;
406     }
407     const std::uint64_t dtype_bytes = dtype_byte_size(entry.dtype);
408     if (dtype_bytes != 0 &&
409         element_count > std::numeric_limits<std::uint64_t>::max() / dtype_bytes) {
410         throw std::runtime_error("safetensors tensor byte size overflow");
411     }
412 }

```

```

410     }
411     return element_count * dtype_bytes;
412 }
413
414 SafeTensorManifest parse_manifest_from_bytes(std::span<const std::uint8_t> ↵
    ↵ bytes) {
415     const std::uint64_t header_size = read_le_u64(bytes);
416     if (header_size == 0 || header_size > LCQI_SAFETENSORS_MAX_HEADER_BYTES) {
417         throw std::runtime_error("safetensors header size is invalid");
418     }
419     const std::uint64_t data_start_offset =
420         static_cast<std::uint64_t>(LCQI_SAFETENSORS_PREFIX_BYTES) + header_size ↵
    ↵ ;
421     if (data_start_offset > bytes.size()) {
422         throw std::runtime_error("safetensors header extends beyond file size") ↵
    ↵ ;
423     }
424
425     SafeTensorManifest manifest;
426     manifest.header_size = header_size;
427     manifest.data_start_offset = data_start_offset;
428     const char* header_begin =
429         reinterpret_cast<const char*>(bytes.data() + ↵
    ↵ LCQI_SAFETENSORS_PREFIX_BYTES);
430     const std::string_view header(header_begin, static_cast<std::size_t>(↵
    ↵ header_size));
431     parse_header(header, manifest);
432     validate_manifest(manifest, static_cast<std::uint64_t>(bytes.size()));
433     return manifest;
434 }
435
436 std::uint16_t read_le_u16(std::span<const std::uint8_t> bytes, std::size_t ↵
    ↵ offset) {
437     return static_cast<std::uint16_t>(
438         static_cast<std::uint16_t>(bytes[offset]) |
439         (static_cast<std::uint16_t>(bytes[offset + 1]) << LCQI_BYTE_BITS));
440 }
441
442 std::uint32_t read_le_u32(std::span<const std::uint8_t> bytes, std::size_t ↵
    ↵ offset) {
443     std::uint32_t value = 0;
444     for (std::size_t i = 0; i < LCQI_DTYPE_BYTES_F32; ++i) {
445         value |= static_cast<std::uint32_t>(bytes[offset + i]) << (↵
    ↵ LCQI_BYTE_BITS * i);
446     }
447     return value;
448 }
449
450 float bits_to_float(std::uint32_t bits) {
451     float value = 0.0F;
452     std::memcpy(&value, &bits, sizeof(value));
453     return value;

```



```

454 }
455
456 float f16_to_f32(std::uint16_t bits) {
457     const std::uint32_t sign = (static_cast<std::uint32_t>(bits) & ↵
    ↵ LCQI_F16_SIGN_MASK)
458         << LCQI_F16_TO_F32_SIGN_SHIFT;
459     const std::uint32_t exponent = static_cast<std::uint32_t>(bits) & ↵
    ↵ LCQI_F16_EXPONENT_MASK;
460     const std::uint32_t mantissa = static_cast<std::uint32_t>(bits) & ↵
    ↵ LCQI_F16_MANTISSA_MASK;
461
462     if (exponent == 0U) {
463         if (mantissa == 0U) {
464             return bits_to_float(sign);
465         }
466         float value = static_cast<float>(mantissa) /
467             static_cast<float>(1U << LCQI_F16_MANTISSA_BITS);
468         value = std::ldexp(value, 1 - static_cast<int>(LCQI_F16_EXPONENT_BIAS)); ↵
    ↵ ;
469         return (sign == 0U) ? value : -value;
470     }
471
472     if (exponent == LCQI_F16_EXPONENT_MASK) {
473         const std::uint32_t f32_mantissa =
474             mantissa << (LCQI_F32_MANTISSA_BITS - LCQI_F16_MANTISSA_BITS);
475         return bits_to_float(sign | LCQI_F32_EXPONENT_MASK | f32_mantissa);
476     }
477
478     const std::uint32_t f32_exponent =
479         ((exponent >> LCQI_F16_MANTISSA_BITS) - LCQI_F16_EXPONENT_BIAS +
480         LCQI_F32_EXPONENT_BIAS)
481         << LCQI_F32_MANTISSA_BITS;
482     const std::uint32_t f32_mantissa =
483         mantissa << (LCQI_F32_MANTISSA_BITS - LCQI_F16_MANTISSA_BITS);
484     return bits_to_float(sign | f32_exponent | f32_mantissa);
485 }
486
487 float bf16_to_f32(std::uint16_t bits) {
488     return bits_to_float(static_cast<std::uint32_t>(bits) << 16U);
489 }
490
491 std::span<const std::uint8_t> tensor_payload_span(const SafeTensorFile& file,
492     const SafeTensorEntry& entry) ↵
    ↵ {
493     const std::uint64_t begin = file.manifest.data_start_offset + entry. ↵
    ↵ data_begin;
494     const std::uint64_t end = file.manifest.data_start_offset + entry.data_end;
495     if (begin > end || end > file.bytes.size()) {
496         throw std::runtime_error("safetensors tensor byte range is invalid");
497     }
498     return std::span<const std::uint8_t>(
499         file.bytes.data() + static_cast<std::size_t>(begin),

```

```

500     static_cast<std::size_t>(end - begin));
501 }
502
503 } // namespace
504
505 std::uint64_t SafeTensorEntry::byte_size() const {
506     return this->data_end - this->data_begin;
507 }
508
509 const SafeTensorEntry* SafeTensorManifest::find_tensor(std::string_view name) ←
    ↪ const {
510     const auto found = std::find_if(
511         this->tensors.begin(),
512         this->tensors.end(),
513         [name](const SafeTensorEntry& entry) {
514             return entry.name == name;
515         });
516     if (found == this->tensors.end()) {
517         return nullptr;
518     }
519     return &(*found);
520 }
521
522 SafeTensorManifest load_safetensors_manifest(const std::filesystem::path& path) ←
    ↪ {
523     const std::string bytes = read_file_bytes(path);
524     return parse_manifest_from_bytes(
525         std::span<const std::uint8_t>(
526             reinterpret_cast<const std::uint8_t*>(bytes.data()),
527             bytes.size()));
528 }
529
530 SafeTensorFile load_safetensors_file(const std::filesystem::path& path) {
531     SafeTensorFile file;
532     file.path = path;
533     file.bytes = read_file_u8(path);
534     file.manifest = parse_manifest_from_bytes(file.bytes);
535     return file;
536 }
537
538 std::span<const std::uint8_t> read_safetensor_tensor_bytes(
539     const SafeTensorFile& file,
540     std::string_view name) {
541     const SafeTensorEntry* entry = file.manifest.find_tensor(name);
542     if (entry == nullptr) {
543         throw std::runtime_error("missing safetensors tensor: " + std::string(←
    ↪ name));
544     }
545     return tensor_payload_span(file, *entry);
546 }
547
548 std::vector<float> read_safetensor_f32_tensor(const SafeTensorFile& file,

```

```

549         std::string_view name) {
550     const SafeTensorEntry* entry = file.manifest.find_tensor(name);
551     if (entry == nullptr) {
552         throw std::runtime_error("missing safetensors tensor: " + std::string(↵
↵ name));
553     }
554     const std::span<const std::uint8_t> bytes = tensor_payload_span(file, *↵
↵ entry);
555     std::vector<float> values;
556     if (entry->dtype == "F32") {
557         values.resize(bytes.size() / LCQI_DTYPE_BYTES_F32);
558         for (std::size_t index = 0; index < values.size(); ++index) {
559             values[index] = bits_to_float(
560                 read_le_u32(bytes, index * LCQI_DTYPE_BYTES_F32));
561         }
562         return values;
563     }
564     if (entry->dtype == "F16") {
565         values.resize(bytes.size() / LCQI_DTYPE_BYTES_F16);
566         for (std::size_t index = 0; index < values.size(); ++index) {
567             values[index] = f16_to_f32(
568                 read_le_u16(bytes, index * LCQI_DTYPE_BYTES_F16));
569         }
570         return values;
571     }
572     if (entry->dtype == "BF16") {
573         values.resize(bytes.size() / LCQI_DTYPE_BYTES_F16);
574         for (std::size_t index = 0; index < values.size(); ++index) {
575             values[index] = bf16_to_f32(
576                 read_le_u16(bytes, index * LCQI_DTYPE_BYTES_F16));
577         }
578         return values;
579     }
580     throw std::runtime_error("safetensors tensor is not float-compatible: " + ↵
↵ entry->name);
581 }
582
583 } // namespace lcqi

```

## GPT-2 Reference Path

`gpt2_reference.cpp` 包括 JSON 配置读取、byte-level BPE、safetensors F32/F16/BF16 张量读取、HuggingFace 权重名前缀解析、LayerNorm、packed QKV、causal attention、GELU、tied lm head、full-prefix greedy generation 和 cached-KV generation。它是正确性优先的 reference path，暂不追求吞吐。

**Listing 10.11** `src/gpt2_reference.cpp`

```

1 #include <lcqi/gpt2_reference.hpp>
2
3 #include <lcqi/safetensors.hpp>
4

```

```

5 #include <algorithm>
6 #include <array>
7 #include <cmath>
8 #include <cstdint>
9 #include <fstream>
10 #include <limits>
11 #include <optional>
12 #include <sstream>
13 #include <stdexcept>
14 #include <string>
15 #include <string_view>
16 #include <unordered_set>
17 #include <utility>
18
19 namespace lcqi {
20 namespace {
21
22 constexpr std::int32_t LCQI_GPT2_DEFAULT_BOS_EOS = 50256;
23 constexpr std::int32_t LCQI_GPT2_ATTENTION_PROJECTIONS = 3;
24 constexpr std::int32_t LCQI_GPT2_PAIR_TOKEN_COUNT = 2;
25 constexpr std::int32_t LCQI_GPT2_BYTE_COUNT = 256;
26 constexpr std::int32_t LCQI_GPT2_UNICODE_PRIVATE_BASE = 256;
27 constexpr std::int32_t LCQI_GPT2_ASCII_EXCLAMATION = 33;
28 constexpr std::int32_t LCQI_GPT2_ASCII_TILDE = 126;
29 constexpr std::int32_t LCQI_GPT2_LATIN_INVERTED_EXCLAMATION = 161;
30 constexpr std::int32_t LCQI_GPT2_LATIN_NOT_SIGN = 172;
31 constexpr std::int32_t LCQI_GPT2_LATIN_REGISTERED = 174;
32 constexpr std::int32_t LCQI_GPT2_LATIN_Y_DIAERESIS = 255;
33 constexpr std::int32_t LCQI_GPT2_UTF8_CONTINUATION_MASK = 0x3F;
34 constexpr std::int32_t LCQI_GPT2_UTF8_CONTINUATION_BITS = 0x80;
35 constexpr std::int32_t LCQI_GPT2_UTF8_TWO_BYTE_HEAD = 0xC0;
36 constexpr std::int32_t LCQI_GPT2_UTF8_THREE_BYTE_HEAD = 0xE0;
37 constexpr std::int32_t LCQI_GPT2_UTF8_FOUR_BYTE_HEAD = 0xF0;
38 constexpr std::int32_t LCQI_GPT2_UTF8_ONE_BYTE_LIMIT = 0x80;
39 constexpr std::int32_t LCQI_GPT2_UTF8_TWO_BYTE_LIMIT = 0x800;
40 constexpr std::int32_t LCQI_GPT2_UTF8_THREE_BYTE_LIMIT = 0x10000;
41 constexpr unsigned char LCQI_GPT2_ASCII_APOSTROPHE = 0x27U;
42 constexpr float LCQI_GPT2_GELU_SQRT_2_OVER_PI = 0.7978845608028654F;
43 constexpr float LCQI_GPT2_GELU_CUBIC_COEFF = 0.044715F;
44 constexpr float LCQI_GPT2_SOFTMAX_NEGATIVE_INFINITY =
45     -std::numeric_limits<float>::infinity();
46
47 class JsonCursor {
48 public:
49     explicit JsonCursor(std::string_view text) : text_(text) {}
50
51     void expect(char expected) {
52         this->skip_ws();
53         if (this->eof() || this->text_[this->position_] != expected) {
54             throw std::runtime_error("invalid JSON syntax");
55         }
56         ++this->position_;

```

```

57     }
58
59     [[nodiscard]] bool consume(char expected) {
60         this->skip_ws();
61         if (!this->eof() && this->text_[this->position_] == expected) {
62             ++this->position_;
63             return true;
64         }
65         return false;
66     }
67
68     [[nodiscard]] std::string parse_string() {
69         this->skip_ws();
70         this->expect('"');
71         std::string result;
72         while (!this->eof()) {
73             const char ch = this->text_[this->position_++];
74             if (ch == '"') {
75                 return result;
76             }
77             if (ch == '\\') {
78                 result += this->parse_escape();
79             } else {
80                 result.push_back(ch);
81             }
82         }
83         throw std::runtime_error("unterminated JSON string");
84     }
85
86     [[nodiscard]] std::int64_t parse_i64() {
87         this->skip_ws();
88         bool negative = false;
89         if (this->consume('-')) {
90             negative = true;
91         }
92         if (this->eof() || this->text_[this->position_] < '0' ||
93             this->text_[this->position_] > '9') {
94             throw std::runtime_error("expected JSON integer");
95         }
96         std::int64_t value = 0;
97         while (!this->eof() && this->text_[this->position_] >= '0' &&
98             this->text_[this->position_] <= '9') {
99             const std::int64_t digit = this->text_[this->position_] - '0';
100             if (value >
101                 (std::numeric_limits<std::int64_t>::max() - digit) / 10) {
102                 throw std::runtime_error("JSON integer overflow");
103             }
104             value = value * 10 + digit;
105             ++this->position_;
106         }
107         return negative ? -value : value;
108     }

```

```

109
110 [[nodiscard]] double parse_number() {
111     this->skip_ws();
112     const std::size_t begin = this->position_;
113     if (!this->eof() && this->text_[this->position_] == '-') {
114         ++this->position_;
115     }
116     while (!this->eof() && this->text_[this->position_] >= '0' &&
117           this->text_[this->position_] <= '9') {
118         ++this->position_;
119     }
120     if (!this->eof() && this->text_[this->position_] == '.') {
121         ++this->position_;
122         while (!this->eof() && this->text_[this->position_] >= '0' &&
123               this->text_[this->position_] <= '9') {
124             ++this->position_;
125         }
126     }
127     if (!this->eof() &&
128         (this->text_[this->position_] == 'e' || this->text_[this->position_↵
↵ ] == 'E')) {
129         ++this->position_;
130         if (!this->eof() &&
131             (this->text_[this->position_] == '+' || this->text_[this->↵
↵ position_] == '-')) {
132             ++this->position_;
133         }
134         while (!this->eof() && this->text_[this->position_] >= '0' &&
135               this->text_[this->position_] <= '9') {
136             ++this->position_;
137         }
138     }
139     if (begin == this->position_) {
140         throw std::runtime_error("expected JSON number");
141     }
142     return std::stod(std::string(this->text_.substr(begin, this->position_ ↵
↵ - begin)));
143 }
144
145 void skip_value() {
146     this->skip_ws();
147     if (this->eof()) {
148         throw std::runtime_error("unexpected end of JSON");
149     }
150     const char ch = this->text_[this->position_];
151     if (ch == '"') {
152         static_cast<void>(this->parse_string());
153         return;
154     }
155     if (ch == '{') {
156         this->expect('{');
157         if (this->consume('}')) {

```

```

158         return;
159     }
160     while (true) {
161         static_cast<void>(this->parse_string());
162         this->expect(':');
163         this->skip_value();
164         if (this->consume('}')) {
165             return;
166         }
167         this->expect(',');
168     }
169 }
170 if (ch == '[') {
171     this->expect('[');
172     if (this->consume(']')) {
173         return;
174     }
175     while (true) {
176         this->skip_value();
177         if (this->consume(']')) {
178             return;
179         }
180         this->expect(',');
181     }
182 }
183 if (ch == 't') {
184     this->expect_literal("true");
185     return;
186 }
187 if (ch == 'f') {
188     this->expect_literal("false");
189     return;
190 }
191 if (ch == 'n') {
192     this->expect_literal("null");
193     return;
194 }
195 static_cast<void>(this->parse_number());
196 }
197
198 [[nodiscard]] bool eof_after_ws() {
199     this->skip_ws();
200     return this->eof();
201 }
202
203 private:
204     std::string_view text_;
205     std::size_t position_ = 0;
206
207     [[nodiscard]] bool eof() const {
208         return this->position_ >= this->text_.size();
209     }

```

```

210
211 void skip_ws() {
212     while (!this->eof() &&
213            (this->text_[this->position_] == ' ' ||
214             this->text_[this->position_] == '\n' ||
215             this->text_[this->position_] == '\r' ||
216             this->text_[this->position_] == '\t')) {
217         ++this->position_;
218     }
219 }
220
221 void expect_literal(std::string_view literal) {
222     for (const char expected : literal) {
223         if (this->eof() || this->text_[this->position_] != expected) {
224             throw std::runtime_error("invalid JSON literal");
225         }
226         ++this->position_;
227     }
228 }
229
230 [[nodiscard]] std::string parse_escape() {
231     if (this->eof()) {
232         throw std::runtime_error("unterminated JSON escape");
233     }
234     const char escaped = this->text_[this->position_++];
235     switch (escaped) {
236         case '"':
237         case '\\':
238         case '/':
239             return std::string(1, escaped);
240         case 'b':
241             return "\b";
242         case 'f':
243             return "\f";
244         case 'n':
245             return "\n";
246         case 'r':
247             return "\r";
248         case 't':
249             return "\t";
250         case 'u':
251             return this->parse_unicode_escape();
252         default:
253             throw std::runtime_error("unsupported JSON escape");
254     }
255 }
256
257 [[nodiscard]] std::string parse_unicode_escape() {
258     std::uint32_t value = 0;
259     for (std::int32_t i = 0; i < 4; ++i) {
260         if (this->eof()) {
261             throw std::runtime_error("unterminated JSON unicode escape");

```



```

262     }
263     const char ch = this->text_[this->position_++];
264     value <= 4U;
265     if (ch >= '0' && ch <= '9') {
266         value += static_cast<std::uint32_t>(ch - '0');
267     } else if (ch >= 'a' && ch <= 'f') {
268         value += static_cast<std::uint32_t>(10 + ch - 'a');
269     } else if (ch >= 'A' && ch <= 'F') {
270         value += static_cast<std::uint32_t>(10 + ch - 'A');
271     } else {
272         throw std::runtime_error("invalid JSON unicode escape");
273     }
274 }
275 std::string encoded;
276 if (value < LCQI_GPT2_UTF8_ONE_BYTE_LIMIT) {
277     encoded.push_back(static_cast<char>(value));
278 } else if (value < LCQI_GPT2_UTF8_TWO_BYTE_LIMIT) {
279     encoded.push_back(static_cast<char>(
280         LCQI_GPT2_UTF8_TWO_BYTE_HEAD | (value >> 6U)));
281     encoded.push_back(static_cast<char>(
282         LCQI_GPT2_UTF8_CONTINUATION_BITS |
283         (value & LCQI_GPT2_UTF8_CONTINUATION_MASK)));
284 } else {
285     encoded.push_back(static_cast<char>(
286         LCQI_GPT2_UTF8_THREE_BYTE_HEAD | (value >> 12U)));
287     encoded.push_back(static_cast<char>(
288         LCQI_GPT2_UTF8_CONTINUATION_BITS |
289         ((value >> 6U) & LCQI_GPT2_UTF8_CONTINUATION_MASK)));
290     encoded.push_back(static_cast<char>(
291         LCQI_GPT2_UTF8_CONTINUATION_BITS |
292         (value & LCQI_GPT2_UTF8_CONTINUATION_MASK)));
293 }
294 return encoded;
295 }
296 };
297
298 std::string read_text_file(const std::filesystem::path& path) {
299     std::ifstream input(path);
300     if (!input) {
301         throw std::runtime_error("cannot open text file: " + path.string());
302     }
303     std::ostringstream buffer;
304     buffer << input.rdbuf();
305     return buffer.str();
306 }
307
308 std::size_t checked_size(std::int64_t value, const char* name) {
309     if (value < 0) {
310         throw std::runtime_error(std::string(name) + " cannot be negative");
311     }
312     return static_cast<std::size_t>(value);
313 }

```

```

314
315 std::string encode_codepoint(std::uint32_t value) {
316     std::string encoded;
317     if (value < LCQI_GPT2_UTF8_ONE_BYTE_LIMIT) {
318         encoded.push_back(static_cast<char>(value));
319     } else if (value < LCQI_GPT2_UTF8_TWO_BYTE_LIMIT) {
320         encoded.push_back(static_cast<char>(
321             LCQI_GPT2_UTF8_TWO_BYTE_HEAD | (value >> 6U)));
322         encoded.push_back(static_cast<char>(
323             LCQI_GPT2_UTF8_CONTINUATION_BITS |
324             (value & LCQI_GPT2_UTF8_CONTINUATION_MASK)));
325     } else if (value < LCQI_GPT2_UTF8_THREE_BYTE_LIMIT) {
326         encoded.push_back(static_cast<char>(
327             LCQI_GPT2_UTF8_THREE_BYTE_HEAD | (value >> 12U)));
328         encoded.push_back(static_cast<char>(
329             LCQI_GPT2_UTF8_CONTINUATION_BITS |
330             ((value >> 6U) & LCQI_GPT2_UTF8_CONTINUATION_MASK)));
331         encoded.push_back(static_cast<char>(
332             LCQI_GPT2_UTF8_CONTINUATION_BITS |
333             (value & LCQI_GPT2_UTF8_CONTINUATION_MASK)));
334     } else {
335         encoded.push_back(static_cast<char>(
336             LCQI_GPT2_UTF8_FOUR_BYTE_HEAD | (value >> 18U)));
337         encoded.push_back(static_cast<char>(
338             LCQI_GPT2_UTF8_CONTINUATION_BITS |
339             ((value >> 12U) & LCQI_GPT2_UTF8_CONTINUATION_MASK)));
340         encoded.push_back(static_cast<char>(
341             LCQI_GPT2_UTF8_CONTINUATION_BITS |
342             ((value >> 6U) & LCQI_GPT2_UTF8_CONTINUATION_MASK)));
343         encoded.push_back(static_cast<char>(
344             LCQI_GPT2_UTF8_CONTINUATION_BITS |
345             (value & LCQI_GPT2_UTF8_CONTINUATION_MASK)));
346     }
347     return encoded;
348 }
349
350 void require_positive(std::int32_t value, const char* name) {
351     if (value <= 0) {
352         throw std::runtime_error(std::string(name) + " must be positive");
353     }
354 }
355
356 void validate_config(const Gpt2Config& config) {
357     require_positive(config.hidden_size, "hidden_size");
358     require_positive(config.head_count, "head_count");
359     require_positive(config.layer_count, "layer_count");
360     require_positive(config.max_positions, "max_positions");
361     require_positive(config.vocab_size, "vocab_size");
362     require_positive(config.intermediate_size, "intermediate_size");
363     if (config.hidden_size % config.head_count != 0) {
364         throw std::runtime_error("hidden_size must be divisible by head_count")↵
↵ ;

```

```

365     }
366     if (config.layer_norm_epsilon <= 0.0F) {
367         throw std::runtime_error("layer_norm_epsilon must be positive");
368     }
369 }
370
371 std::int32_t head_dim(const Gpt2Config& config) {
372     return config.hidden_size / config.head_count;
373 }
374
375 std::size_t matrix_size(std::int32_t rows, std::int32_t columns) {
376     return checked_size(rows, "rows") * checked_size(columns, "columns");
377 }
378
379 void validate_vector_size(std::size_t actual, std::int32_t expected, const char↵
↵ * name) {
380     if (actual != checked_size(expected, name)) {
381         throw std::runtime_error(std::string(name) + " size mismatch");
382     }
383 }
384
385 void validate_linear(const Gpt2LinearF32& linear, const char* name) {
386     require_positive(linear.input_size, "linear input_size");
387     require_positive(linear.output_size, "linear output_size");
388     if (linear.weights.size() != matrix_size(linear.output_size, linear.↵
↵ input_size)) {
389         throw std::runtime_error(std::string(name) + " weight size mismatch");
390     }
391     if (!linear.bias.empty() &&
392         linear.bias.size() != checked_size(linear.output_size, "linear ↵
↵ output_size")) {
393         throw std::runtime_error(std::string(name) + " bias size mismatch");
394     }
395 }
396
397 void validate_model(const Gpt2ReferenceModel& model) {
398     validate_config(model.config);
399     if (model.layers.size() != checked_size(model.config.layer_count, "↵
↵ layer_count")) {
400         throw std::runtime_error("GPT-2 layer count mismatch");
401     }
402     if (model.token_embedding.size() !=
403         matrix_size(model.config.vocab_size, model.config.hidden_size)) {
404         throw std::runtime_error("GPT-2 token embedding size mismatch");
405     }
406     if (model.position_embedding.size() !=
407         matrix_size(model.config.max_positions, model.config.hidden_size)) {
408         throw std::runtime_error("GPT-2 position embedding size mismatch");
409     }
410     validate_vector_size(model.final_ln_weight.size(), model.config.hidden_size↵
↵ , "ln_f weight");
411     validate_vector_size(model.final_ln_bias.size(), model.config.hidden_size, ↵

```

```

    ↪ "ln_f bias");
412 if (model.tie_lm_head_to_embedding) {
413     if (!model.lm_head_weight.empty()) {
414         throw std::runtime_error("tied GPT-2 lm_head should not carry ↪
    ↪ separate weights");
415     }
416 } else if (model.lm_head_weight.size() !=
417     matrix_size(model.config.vocab_size, model.config.hidden_size)) ↪
    ↪ {
418     throw std::runtime_error("GPT-2 lm_head size mismatch");
419 }
420 for (const Gpt2LayerWeightsF32& layer : model.layers) {
421     validate_vector_size(layer.ln_1_weight.size(), model.config.hidden_size ↪
    ↪ , "ln_1 weight");
422     validate_vector_size(layer.ln_1_bias.size(), model.config.hidden_size, ↪
    ↪ "ln_1 bias");
423     validate_linear(layer.c_attn, "c_attn");
424     validate_linear(layer.c_proj, "c_proj");
425     validate_vector_size(layer.ln_2_weight.size(), model.config.hidden_size ↪
    ↪ , "ln_2 weight");
426     validate_vector_size(layer.ln_2_bias.size(), model.config.hidden_size, ↪
    ↪ "ln_2 bias");
427     validate_linear(layer.c_fc, "c_fc");
428     validate_linear(layer.mlp_c_proj, "mlp_c_proj");
429 }
430 }
431
432 std::span<const float> row_span(std::span<const float> values,
433     std::int32_t row,
434     std::int32_t row_width) {
435     return values.subspan(check_size(row, "row") * checked_size(row_width, "↪
    ↪ row_width"),
436         checked_size(row_width, "row_width"));
437 }
438
439 void add_inplace(std::span<float> target, std::span<const float> source) {
440     if (target.size() != source.size()) {
441         throw std::runtime_error("GPT-2 residual add size mismatch");
442     }
443     for (std::size_t index = 0; index < target.size(); ++index) {
444         target[index] += source[index];
445     }
446 }
447
448 void layer_norm(std::span<const float> input,
449     std::span<const float> weight,
450     std::span<const float> bias,
451     float epsilon,
452     std::span<float> output) {
453     if (input.empty() || input.size() != weight.size() || input.size() != bias.↪
    ↪ size() ||
454         input.size() != output.size()) {

```

```

455     throw std::runtime_error("GPT-2 LayerNorm size mismatch");
456 }
457 float mean = 0.0F;
458 for (const float value : input) {
459     mean += value;
460 }
461 mean /= static_cast<float>(input.size());
462
463 float variance = 0.0F;
464 for (const float value : input) {
465     const float centered = value - mean;
466     variance += centered * centered;
467 }
468 variance /= static_cast<float>(input.size());
469 const float scale = 1.0F / std::sqrt(variance + epsilon);
470
471 for (std::size_t index = 0; index < input.size(); ++index) {
472     output[index] = (input[index] - mean) * scale * weight[index] + bias[↵
↵ index];
473 }
474 }
475
476 void linear_f32(const Gpt2LinearF32& linear,
477                 std::span<const float> input,
478                 std::span<float> output) {
479     validate_linear(linear, "GPT-2 linear");
480     if (input.size() != checked_size(linear.input_size, "linear input") ||
481         output.size() != checked_size(linear.output_size, "linear output")) {
482         throw std::runtime_error("GPT-2 linear input/output size mismatch");
483     }
484     for (std::int32_t out = 0; out < linear.output_size; ++out) {
485         float sum = linear.bias.empty() ? 0.0F : linear.bias[checked_size(out, ↵
↵ "out")];
486         const std::size_t row_base = checked_size(out, "out") *
487                                     checked_size(linear.input_size, "linear ↵
↵ input");
488         for (std::int32_t in = 0; in < linear.input_size; ++in) {
489             sum += linear.weights[row_base + checked_size(in, "in")] *
490                   input[checked_size(in, "in")];
491         }
492         output[checked_size(out, "out")] = sum;
493     }
494 }
495
496 float gelu_new(float value) {
497     const float cubic = value * value * value;
498     return 0.5F * value *
499           (1.0F + std::tanh(LCQI_GPT2_GELU_SQRT_2_OVER_PI *
500                             (value + LCQI_GPT2_GELU_CUBIC_COEFF * cubic)));
501 }
502
503 std::int32_t argmax(std::span<const float> values) {

```

```

504     if (values.empty()) {
505         throw std::runtime_error("cannot take argmax of empty logits");
506     }
507     std::int32_t best_index = 0;
508     float best_value = values.front();
509     for (std::int32_t index = 1; index < static_cast<std::int32_t>(values.size()
↵    )); ++index) {
510         const float value = values[checked_size(index, "argmax index")];
511         if (value > best_value) {
512             best_value = value;
513             best_index = index;
514         }
515     }
516     return best_index;
517 }
518
519 void project_qkv(const Gpt2Config& config,
520                 const Gpt2LayerWeightsF32& layer,
521                 std::span<const float> hidden,
522                 std::span<float> q,
523                 std::span<float> k,
524                 std::span<float> v) {
525     std::vector<float> packed(matrix_size(LCQI_GPT2_ATTENTION_PROJECTIONS,
526                                           config.hidden_size),
527                               0.0F);
528     linear_f32(layer.c_attn, hidden, packed);
529     const std::size_t width = checked_size(config.hidden_size, "hidden_size");
530     std::copy_n(packed.begin(), config.hidden_size, q.begin());
531     std::copy_n(packed.begin() + static_cast<std::ptrdiff_t>(width),
532               config.hidden_size,
533               k.begin());
534     std::copy_n(packed.begin() + static_cast<std::ptrdiff_t>(width * 2),
535               config.hidden_size,
536               v.begin());
537 }
538
539 void attention_cached_position(const Gpt2Config& config,
540                              const Gpt2KvCache& cache,
541                              std::int32_t layer_id,
542                              std::int32_t model_position,
543                              std::span<const float> query,
544                              std::span<float> output) {
545     const std::int32_t local_head_dim = head_dim(config);
546     const float score_scale = 1.0F / std::sqrt(static_cast<float>(
↵    local_head_dim));
547     std::fill(output.begin(), output.end(), 0.0F);
548     std::vector<float> scores(checked_size(model_position + 1, "model_position"
↵    ),
549                             LCQI_GPT2_SOFTMAX_NEGATIVE_INFINITY);
550
551     for (std::int32_t head = 0; head < config.head_count; ++head) {
552         float max_score = LCQI_GPT2_SOFTMAX_NEGATIVE_INFINITY;

```

```

553     for (std::int32_t past = 0; past <= model_position; ++past) {
554         const std::span<const float> key = cache.key(layer_id, past, head);
555         float dot = 0.0F;
556         for (std::int32_t dim = 0; dim < local_head_dim; ++dim) {
557             const std::size_t query_index =
558                 matrix_size(head, local_head_dim) + checked_size(dim, "head"
↵ dim");
559             dot += query[query_index] * key[checked_size(dim, "head dim")];
560         }
561         const float score = dot * score_scale;
562         scores[checked_size(past, "past")] = score;
563         max_score = std::max(max_score, score);
564     }
565
566     float denominator = 0.0F;
567     for (std::int32_t past = 0; past <= model_position; ++past) {
568         const std::size_t index = checked_size(past, "past");
569         scores[index] = std::exp(scores[index] - max_score);
570         denominator += scores[index];
571     }
572     if (denominator <= 0.0F) {
573         throw std::runtime_error("GPT-2 cached attention denominator is
↵ invalid");
574     }
575
576     for (std::int32_t past = 0; past <= model_position; ++past) {
577         const float probability = scores[checked_size(past, "past")] /
↵ denominator;
578         const std::span<const float> value = cache.value(layer_id, past,
↵ head);
579         for (std::int32_t dim = 0; dim < local_head_dim; ++dim) {
580             const std::size_t output_index =
581                 matrix_size(head, local_head_dim) + checked_size(dim, "head"
↵ dim");
582             output[output_index] += probability * value[checked_size(dim, "
↵ head dim")];
583         }
584     }
585 }
586 }
587
588 void attention_full_sequence(const Gpt2Config& config,
589                             std::span<const float> q_by_pos,
590                             std::span<const float> k_by_pos,
591                             std::span<const float> v_by_pos,
592                             std::int32_t sequence_length,
593                             std::span<float> output_by_pos) {
594     const std::int32_t local_head_dim = head_dim(config);
595     const float score_scale = 1.0F / std::sqrt(static_cast<float>(<
↵ local_head_dim));
596     std::fill(output_by_pos.begin(), output_by_pos.end(), 0.0F);
597     std::vector<float> scores(checked_size(sequence_length, "sequence_length"),

```



```

598         LCQI_GPT2_SOFTMAX_NEGATIVE_INFINITY);
599
600     for (std::int32_t position = 0; position < sequence_length; ++position) {
601         for (std::int32_t head = 0; head < config.head_count; ++head) {
602             float max_score = LCQI_GPT2_SOFTMAX_NEGATIVE_INFINITY;
603             for (std::int32_t past = 0; past <= position; ++past) {
604                 float dot = 0.0F;
605                 for (std::int32_t dim = 0; dim < local_head_dim; ++dim) {
606                     const std::size_t q_index =
607                         matrix_size(position, config.hidden_size) +
608                         matrix_size(head, local_head_dim) + checked_size(dim, "↵
↵ head dim");
609                     const std::size_t k_index =
610                         matrix_size(past, config.hidden_size) +
611                         matrix_size(head, local_head_dim) + checked_size(dim, "↵
↵ head dim");
612                     dot += q_by_pos[q_index] * k_by_pos[k_index];
613                 }
614                 const float score = dot * score_scale;
615                 scores[checked_size(past, "past")] = score;
616                 max_score = std::max(max_score, score);
617             }
618
619             float denominator = 0.0F;
620             for (std::int32_t past = 0; past <= position; ++past) {
621                 const std::size_t index = checked_size(past, "past");
622                 scores[index] = std::exp(scores[index] - max_score);
623                 denominator += scores[index];
624             }
625             if (denominator <= 0.0F) {
626                 throw std::runtime_error("GPT-2 attention denominator is ↵
↵ invalid");
627             }
628             for (std::int32_t past = 0; past <= position; ++past) {
629                 const float probability = scores[checked_size(past, "past")] / ↵
↵ denominator;
630                 for (std::int32_t dim = 0; dim < local_head_dim; ++dim) {
631                     const std::size_t output_index =
632                         matrix_size(position, config.hidden_size) +
633                         matrix_size(head, local_head_dim) + checked_size(dim, "↵
↵ head dim");
634                     const std::size_t value_index =
635                         matrix_size(past, config.hidden_size) +
636                         matrix_size(head, local_head_dim) + checked_size(dim, "↵
↵ head dim");
637                     output_by_pos[output_index] += probability * v_by_pos[↵
↵ value_index];
638                 }
639             }
640         }
641     }
642 }

```



```

643
644 void forward_gpt2_layer(const Gpt2Config& config,
645                         const Gpt2LayerWeightsF32& layer,
646                         std::int32_t sequence_length,
647                         std::vector<float>& hidden_by_pos) {
648     std::vector<float> normed(checkered_size(config.hidden_size, "hidden_size"), ←
    ↪ 0.0F);
649     std::vector<float> q_by_pos(matrix_size(sequence_length, config.hidden_size, ←
    ↪ ), 0.0F);
650     std::vector<float> k_by_pos(q_by_pos.size(), 0.0F);
651     std::vector<float> v_by_pos(q_by_pos.size(), 0.0F);
652     std::vector<float> attention_by_pos(q_by_pos.size(), 0.0F);
653     std::vector<float> projected(checkered_size(config.hidden_size, "hidden_size" ←
    ↪ ), 0.0F);
654     std::vector<float> mlp_fc(checkered_size(config.intermediate_size, " ←
    ↪ intermediate_size"), 0.0F);
655     std::vector<float> mlp_out(checkered_size(config.hidden_size, "hidden_size"), ←
    ↪ 0.0F);
656
657     for (std::int32_t position = 0; position < sequence_length; ++position) {
658         const std::span<float> hidden = std::span<float>(
659             hidden_by_pos.data() + matrix_size(position, config.hidden_size),
660             checked_size(config.hidden_size, "hidden_size"));
661         layer_norm(hidden,
662                 layer.ln_1_weight,
663                 layer.ln_1_bias,
664                 config.layer_norm_epsilon,
665                 normed);
666         std::span<float> q = std::span<float>(
667             q_by_pos.data() + matrix_size(position, config.hidden_size),
668             checked_size(config.hidden_size, "hidden_size"));
669         std::span<float> k = std::span<float>(
670             k_by_pos.data() + matrix_size(position, config.hidden_size),
671             checked_size(config.hidden_size, "hidden_size"));
672         std::span<float> v = std::span<float>(
673             v_by_pos.data() + matrix_size(position, config.hidden_size),
674             checked_size(config.hidden_size, "hidden_size"));
675         project_qkv(config, layer, normed, q, k, v);
676     }
677
678     attention_full_sequence(config, q_by_pos, k_by_pos, v_by_pos, ←
    ↪ sequence_length, attention_by_pos);
679
680     for (std::int32_t position = 0; position < sequence_length; ++position) {
681         const std::span<float> hidden = std::span<float>(
682             hidden_by_pos.data() + matrix_size(position, config.hidden_size),
683             checked_size(config.hidden_size, "hidden_size"));
684         const std::span<const float> attention = std::span<const float>(
685             attention_by_pos.data() + matrix_size(position, config.hidden_size) ←
    ↪ ,
686             checked_size(config.hidden_size, "hidden_size"));
687         linear_f32(layer.c_proj, attention, projected);

```

```

688     add_inplace(hidden, projected);
689
690     layer_norm(hidden,
691                 layer.ln_2_weight,
692                 layer.ln_2_bias,
693                 config.layer_norm_epsilon,
694                 normed);
695     linear_f32(layer.c_fc, normed, mlp_fc);
696     for (float& value : mlp_fc) {
697         value = gelu_new(value);
698     }
699     linear_f32(layer.mlp_c_proj, mlp_fc, mlp_out);
700     add_inplace(hidden, mlp_out);
701 }
702 }
703
704 void forward_gpt2_layer_cached(const Gpt2Config& config,
705                                const Gpt2LayerWeightsF32& layer,
706                                std::int32_t layer_id,
707                                std::int32_t model_position,
708                                Gpt2KvCache& cache,
709                                std::vector<float>& hidden) {
710     std::vector<float> normed(check_size(config.hidden_size, "hidden_size"), ←
711                               ↪ 0.0F);
712     std::vector<float> query(check_size(config.hidden_size, "hidden_size"), ←
713                              ↪ 0.0F);
714     std::vector<float> key(check_size(config.hidden_size, "hidden_size"), 0.0F);
715     std::vector<float> value(check_size(config.hidden_size, "hidden_size"), ←
716                              ↪ 0.0F);
717     std::vector<float> attention(check_size(config.hidden_size, "hidden_size"), ←
718                                  ↪ 0.0F);
719     std::vector<float> projected(check_size(config.hidden_size, "hidden_size"), ←
720                                  ↪ 0.0F);
721     std::vector<float> mlp_fc(check_size(config.intermediate_size, "intermediate_size"), ←
722                               ↪ 0.0F);
723     std::vector<float> mlp_out(check_size(config.hidden_size, "hidden_size"), ←
724                                ↪ 0.0F);
725
726     layer_norm(hidden,
727                 layer.ln_1_weight,
728                 layer.ln_1_bias,
729                 config.layer_norm_epsilon,
730                 normed);
731     project_qkv(config, layer, normed, query, key, value);
732     cache.append(layer_id, model_position, key, value);
733     attention_cached_position(config, cache, layer_id, model_position, query, ←
734                               ↪ attention);
735
736     linear_f32(layer.c_proj, attention, projected);
737     add_inplace(hidden, projected);
738 }

```

```

731     layer_norm(hidden,
732                 layer.ln_2_weight,
733                 layer.ln_2_bias,
734                 config.layer_norm_epsilon,
735                 normed);
736     linear_f32(layer.c_fc, normed, mlp_fc);
737     for (float& value_ref : mlp_fc) {
738         value_ref = gelu_new(value_ref);
739     }
740     linear_f32(layer.mlp_c_proj, mlp_fc, mlp_out);
741     add_inplace(hidden, mlp_out);
742 }
743
744 void compute_logits(const Gpt2ReferenceModel& model,
745                    std::span<const float> hidden,
746                    std::span<float> logits) {
747     if (logits.size() != checked_size(model.config.vocab_size, "vocab_size")) {
748         throw std::runtime_error("GPT-2 logits size mismatch");
749     }
750     const std::span<const float> weight =
751         model.tie_lm_head_to_embedding ? std::span<const float>(model.↵
↵ token_embedding)
752                                         : std::span<const float>(model.↵
↵ lm_head_weight);
753     for (std::int32_t token = 0; token < model.config.vocab_size; ++token) {
754         float sum = 0.0F;
755         const std::span<const float> row = row_span(weight, token, model.config.↵
↵ .hidden_size);
756         for (std::int32_t dim = 0; dim < model.config.hidden_size; ++dim) {
757             sum += row[checked_size(dim, "dim")] * hidden[checked_size(dim, "↵
↵ dim")];
758         }
759         logits[checked_size(token, "token")] = sum;
760     }
761 }
762
763 std::vector<std::string> gpt2_bytes_to_unicode() {
764     std::vector<std::int32_t> bytes;
765     for (std::int32_t value = LCQI_GPT2_ASCII_EXCLAMATION;
766          value <= LCQI_GPT2_ASCII_TILDE;
767          ++value) {
768         bytes.push_back(value);
769     }
770     for (std::int32_t value = LCQI_GPT2_LATIN_INVERTED_EXCLAMATION;
771          value <= LCQI_GPT2_LATIN_NOT_SIGN;
772          ++value) {
773         bytes.push_back(value);
774     }
775     for (std::int32_t value = LCQI_GPT2_LATIN_REGISTERED;
776          value <= LCQI_GPT2_LATIN_Y_DIAERESIS;
777          ++value) {
778         bytes.push_back(value);

```

```

779     }
780
781     std::unordered_set<std::int32_t> byte_set(bytes.begin(), bytes.end());
782     std::vector<std::int32_t> chars = bytes;
783     std::int32_t private_offset = 0;
784     for (std::int32_t byte = 0; byte < LCQI_GPT2_BYTE_COUNT; ++byte) {
785         if (!byte_set.contains(byte)) {
786             bytes.push_back(byte);
787             chars.push_back(LCQI_GPT2_UNICODE_PRIVATE_BASE + private_offset);
788             ++private_offset;
789         }
790     }
791
792     std::vector<std::string> mapping(LCQI_GPT2_BYTE_COUNT);
793     for (std::size_t index = 0; index < bytes.size(); ++index) {
794         mapping[checked_size(bytes[index], "byte unicode index")] =
795             encode_codepoint(static_cast<std::uint32_t>(chars[index]));
796     }
797     return mapping;
798 }
799
800 std::unordered_map<std::string, unsigned char> gpt2_unicode_to_bytes(
801     const std::vector<std::string>& byte_encoder) {
802     std::unordered_map<std::string, unsigned char> result;
803     for (std::size_t byte = 0; byte < byte_encoder.size(); ++byte) {
804         result.emplace(byte_encoder[byte], static_cast<unsigned char>(byte));
805     }
806     return result;
807 }
808
809 std::string bytes_to_bpe_alphabet(std::string_view text) {
810     static const std::vector<std::string> BYTE_ENCODER = gpt2_bytes_to_unicode(
811         ↵ ↵ ();
812     std::string result;
813     for (const unsigned char byte : text) {
814         result += BYTE_ENCODER[byte];
815     }
816     return result;
817 }
818
819 std::vector<std::string> utf8_symbols(std::string_view text) {
820     std::vector<std::string> symbols;
821     std::size_t offset = 0;
822     while (offset < text.size()) {
823         const unsigned char lead = static_cast<unsigned char>(text[offset]);
824         std::size_t width = 1;
825         if ((lead & 0xE0U) == LCQI_GPT2_UTF8_TWO_BYTE_HEAD) {
826             width = 2;
827         } else if ((lead & 0xF0U) == LCQI_GPT2_UTF8_THREE_BYTE_HEAD) {
828             width = 3;
829         } else if ((lead & 0xF8U) == LCQI_GPT2_UTF8_FOUR_BYTE_HEAD) {
830             width = 4;

```

```

830     }
831     if (offset + width > text.size()) {
832         throw std::runtime_error("invalid UTF-8/BPE symbol boundary");
833     }
834     symbols.emplace_back(text.substr(offset, width));
835     offset += width;
836 }
837 return symbols;
838 }
839
840 std::string pair_key(std::string_view left, std::string_view right) {
841     std::string key(left);
842     key.push_back('\n');
843     key += right;
844     return key;
845 }
846
847 std::vector<std::string> bpe_apply(const Gpt2Tokenizer& tokenizer, std::string_view token) {
848     std::vector<std::string> word = utf8_symbols(token);
849     if (word.size() <= 1) {
850         return word;
851     }
852
853     while (true) {
854         std::optional<std::int32_t> best_rank;
855         std::string best_left;
856         std::string best_right;
857         for (std::size_t index = 0; index + 1 < word.size(); ++index) {
858             const auto rank = tokenizer.bpe_ranks.find(pair_key(word[index],
859                 word[index + 1]));
860             if (rank != tokenizer.bpe_ranks.end() &&
861                 (!best_rank.has_value() || rank->second < *best_rank)) {
862                 best_rank = rank->second;
863                 best_left = word[index];
864                 best_right = word[index + 1];
865             }
866         }
867         if (!best_rank.has_value()) {
868             return word;
869         }
870
871         std::vector<std::string> merged;
872         std::size_t index = 0;
873         while (index < word.size()) {
874             if (index + 1 < word.size() && word[index] == best_left &&
875                 word[index + 1] == best_right) {
876                 merged.push_back(best_left + best_right);
877                 index += LCQI_GPT2_PAIR_TOKEN_COUNT;
878             } else {
879                 merged.push_back(word[index]);
880                 ++index;
881             }
882         }
883     }
884 }

```

```

880     }
881 }
882 word = std::move(merged);
883 if (word.size() == 1) {
884     return word;
885 }
886 }
887 }
888
889 bool is_ascii_letter(unsigned char value) {
890     return (value >= 'A' && value <= 'Z') || (value >= 'a' && value <= 'z');
891 }
892
893 bool is_ascii_digit(unsigned char value) {
894     return value >= '0' && value <= '9';
895 }
896
897 bool is_ascii_space(unsigned char value) {
898     return value == ' ' || value == '\t' || value == '\n' ||
899         value == '\r' || value == '\f' || value == '\v';
900 }
901
902 bool is_gpt2_punctuation(unsigned char value) {
903     return !is_ascii_letter(value) && !is_ascii_digit(value) &&
904         !is_ascii_space(value);
905 }
906
907 std::size_t consume_utf8_codepoint(std::string_view text, std::size_t offset) {
908     const unsigned char lead = static_cast<unsigned char>(text[offset]);
909     std::size_t width = 1;
910     if ((lead & 0xE0U) == LCQI_GPT2_UTF8_TWO_BYTE_HEAD) {
911         width = 2;
912     } else if ((lead & 0xF0U) == LCQI_GPT2_UTF8_THREE_BYTE_HEAD) {
913         width = 3;
914     } else if ((lead & 0xF8U) == LCQI_GPT2_UTF8_FOUR_BYTE_HEAD) {
915         width = 4;
916     }
917     if (offset + width > text.size()) {
918         throw std::runtime_error("invalid UTF-8 text for GPT-2 tokenizer");
919     }
920     return offset + width;
921 }
922
923 bool starts_with_contraction(std::string_view text,
924                             std::size_t offset,
925                             std::string_view contraction) {
926     if (offset + contraction.size() > text.size()) {
927         return false;
928     }
929     for (std::size_t index = 0; index < contraction.size(); ++index) {
930         char left = text[offset + index];
931         const char right = contraction[index];

```

```

932     if (left >= 'A' && left <= 'Z') {
933         left = static_cast<char>(left - 'A' + 'a');
934     }
935     if (left != right) {
936         return false;
937     }
938 }
939 return true;
940 }
941
942 std::vector<std::string> gpt2_pretokenize(std::string_view text) {
943     std::vector<std::string> pieces;
944     std::size_t offset = 0;
945     const std::array<std::string_view, 7> contractions{
946         "s", "t", "re", "ve", "m", "ll", "d",
947     };
948     while (offset < text.size()) {
949         bool emitted_contraction = false;
950         for (const std::string_view contraction : contractions) {
951             if (starts_with_contraction(text, offset, contraction)) {
952                 pieces.emplace_back(text.substr(offset, contraction.size()));
953                 offset += contraction.size();
954                 emitted_contraction = true;
955                 break;
956             }
957         }
958         if (emitted_contraction) {
959             continue;
960         }
961
962         const std::size_t begin = offset;
963         bool has_prefix_space = false;
964         if (text[offset] == ' ' && offset + 1 < text.size() &&
965             !is_ascii_space(static_cast<unsigned char>(text[offset + 1]))) {
966             has_prefix_space = true;
967             ++offset;
968         }
969
970         if (offset >= text.size()) {
971             pieces.emplace_back(text.substr(begin, offset - begin));
972             continue;
973         }
974
975         const unsigned char ch = static_cast<unsigned char>(text[offset]);
976         if (is_ascii_letter(ch)) {
977             while (offset < text.size() &&
978                 is_ascii_letter(static_cast<unsigned char>(text[offset]))) {
979                 ++offset;
980             }
981         } else if (is_ascii_digit(ch)) {
982             while (offset < text.size() &&
983                 is_ascii_digit(static_cast<unsigned char>(text[offset]))) {

```

```

984         ++offset;
985     }
986     } else if (is_gpt2_punctuation(ch) && ch != LCQI_GPT2_ASCII_APOSTROPHE) ←
    ↪ {
987         while (offset < text.size() &&
988             is_gpt2_punctuation(static_cast<unsigned char>(text[offset])) ←
    ↪ ) &&
989             static_cast<unsigned char>(text[offset]) != ←
    ↪ LCQI_GPT2_ASCII_APOSTROPHE) {
990             offset = consume_utf8_codepoint(text, offset);
991         }
992     } else if (is_ascii_space(ch)) {
993         while (offset < text.size() &&
994             is_ascii_space(static_cast<unsigned char>(text[offset]))) {
995             ++offset;
996         }
997     } else {
998         offset = consume_utf8_codepoint(text, offset);
999     }
1000
1001     if (has_prefix_space || offset > begin) {
1002         pieces.emplace_back(text.substr(begin, offset - begin));
1003     } else {
1004         ++offset;
1005     }
1006 }
1007 return pieces;
1008 }
1009
1010 std::int32_t parse_i32_field(JsonCursor& cursor) {
1011     const std::int64_t value = cursor.parse_i64();
1012     if (value < std::numeric_limits<std::int32_t>::min() ||
1013         value > std::numeric_limits<std::int32_t>::max()) {
1014         throw std::runtime_error("JSON int32 field out of range");
1015     }
1016     return static_cast<std::int32_t>(value);
1017 }
1018
1019 std::unordered_map<std::string, std::int32_t> parse_vocab_json(std::string_view ←
    ↪ text) {
1020     JsonCursor cursor(text);
1021     std::unordered_map<std::string, std::int32_t> vocab;
1022     cursor.expect('{');
1023     if (cursor.consume('}')) {
1024         return vocab;
1025     }
1026     while (true) {
1027         const std::string key = cursor.parse_string();
1028         cursor.expect(':');
1029         const std::int32_t id = parse_i32_field(cursor);
1030         vocab.emplace(key, id);
1031         if (cursor.consume('}')) {

```



```

1032         break;
1033     }
1034     cursor.expect(',');
1035 }
1036 if (!cursor.eof_after_ws()) {
1037     throw std::runtime_error("trailing data after vocab JSON");
1038 }
1039 return vocab;
1040 }
1041
1042 std::vector<std::string> build_id_to_token(
1043     const std::unordered_map<std::string, std::int32_t>& vocab) {
1044     std::int32_t max_id = -1;
1045     for (const auto& [token, id] : vocab) {
1046         if (id < 0) {
1047             throw std::runtime_error("GPT-2 vocab id cannot be negative");
1048         }
1049         max_id = std::max(max_id, id);
1050     }
1051     std::vector<std::string> result(check_size(max_id + 1, "max vocab id"));
1052     for (const auto& [token, id] : vocab) {
1053         std::string& slot = result[checked_size(id, "vocab id")];
1054         if (!slot.empty()) {
1055             throw std::runtime_error("duplicate GPT-2 vocab id");
1056         }
1057         slot = token;
1058     }
1059     return result;
1060 }
1061
1062 std::unordered_map<std::string, std::int32_t> parse_merges(std::string_view ↵
    ↵ text) {
1063     std::unordered_map<std::string, std::int32_t> ranks;
1064     std::istringstream input{std::string(text)};
1065     std::string line;
1066     std::int32_t rank = 0;
1067     while (std::getline(input, line)) {
1068         if (line.empty() || line[0] == '#') {
1069             continue;
1070         }
1071         std::istringstream line_stream(line);
1072         std::string left;
1073         std::string right;
1074         if (!(line_stream >> left >> right)) {
1075             throw std::runtime_error("invalid GPT-2 merges line");
1076         }
1077         ranks.emplace(pair_key(left, right), rank);
1078         ++rank;
1079     }
1080     return ranks;
1081 }
1082

```

```

1083 void transpose_hf_conv1d(std::vector<float>& values,
1084                          std::int32_t input_size,
1085                          std::int32_t output_size) {
1086     if (values.size() != matrix_size(input_size, output_size)) {
1087         throw std::runtime_error("GPT-2 Conv1D tensor size mismatch");
1088     }
1089     std::vector<float> transposed(matrix_size(output_size, input_size), 0.0F);
1090     for (std::int32_t input = 0; input < input_size; ++input) {
1091         for (std::int32_t output = 0; output < output_size; ++output) {
1092             transposed[matrix_size(output, input_size) + checked_size(input, "↵
↵ input")] =
1093                 values[matrix_size(input, output_size) + checked_size(output, "↵
↵ output")];
1094         }
1095     }
1096     values = std::move(transposed);
1097 }
1098
1099 std::vector<float> require_tensor(const SafeTensorFile& file,
1100                                  std::string_view name,
1101                                  std::span<const std::int64_t> expected_shape)↵
↵ {
1102     const SafeTensorEntry* entry = file.manifest.find_tensor(name);
1103     if (entry == nullptr) {
1104         throw std::runtime_error("missing GPT-2 tensor: " + std::string(name));
1105     }
1106     if (entry->shape.size() != expected_shape.size()) {
1107         throw std::runtime_error("GPT-2 tensor rank mismatch: " + std::string(↵
↵ name));
1108     }
1109     for (std::size_t index = 0; index < expected_shape.size(); ++index) {
1110         if (entry->shape[index] != expected_shape[index]) {
1111             throw std::runtime_error("GPT-2 tensor shape mismatch: " + std::↵
↵ string(name));
1112         }
1113     }
1114     return read_safetensor_f32_tensor(file, name);
1115 }
1116
1117 std::vector<float> require_first_tensor(const SafeTensorFile& file,
1118                                         std::span<const std::string_view> names↵
↵ ,
1119                                         std::span<const std::int64_t> ↵
↵ expected_shape) {
1120     for (const std::string_view name : names) {
1121         if (file.manifest.find_tensor(name) != nullptr) {
1122             return require_tensor(file, name, expected_shape);
1123         }
1124     }
1125     throw std::runtime_error("missing GPT-2 tensor: " + std::string(names.front↵
↵ ()));
1126 }

```

```

1127
1128 Gpt2LinearF32 make_linear(std::int32_t input_size,
1129                             std::int32_t output_size,
1130                             std::vector<float> weights,
1131                             std::vector<float> bias = {}) {
1132     Gpt2LinearF32 linear;
1133     linear.input_size = input_size;
1134     linear.output_size = output_size;
1135     linear.weights = std::move(weights);
1136     linear.bias = std::move(bias);
1137     validate_linear(linear, "make GPT-2 linear");
1138     return linear;
1139 }
1140
1141 std::vector<float> zeros(std::int32_t count) {
1142     return std::vector<float>(checked_size(count, "zero count"), 0.0F);
1143 }
1144
1145 std::vector<float> ones(std::int32_t count) {
1146     return std::vector<float>(checked_size(count, "one count"), 1.0F);
1147 }
1148
1149 std::string tensor_name(std::int32_t layer, std::string_view suffix) {
1150     return "transformer.h." + std::to_string(layer) + "." + std::string(suffix) ←
    ↪ ;
1151 }
1152
1153 std::string bare_tensor_name(std::int32_t layer, std::string_view suffix) {
1154     return "h." + std::to_string(layer) + "." + std::string(suffix);
1155 }
1156
1157 std::array<std::string_view, 2> embedding_names(std::string_view bare) {
1158     if (bare == "wte.weight") {
1159         return {"transformer.wte.weight", "wte.weight"};
1160     }
1161     if (bare == "wpe.weight") {
1162         return {"transformer.wpe.weight", "wpe.weight"};
1163     }
1164     if (bare == "ln_f.weight") {
1165         return {"transformer.ln_f.weight", "ln_f.weight"};
1166     }
1167     if (bare == "ln_f.bias") {
1168         return {"transformer.ln_f.bias", "ln_f.bias"};
1169     }
1170     throw std::runtime_error("unsupported GPT-2 embedding tensor alias");
1171 }
1172
1173 std::array<std::string, 2> layer_names(std::int32_t layer, std::string_view ←
    ↪ suffix) {
1174     return {tensor_name(layer, suffix), bare_tensor_name(layer, suffix)};
1175 }
1176

```

```

1177 std::vector<float> require_layer_tensor(const SafeTensorFile& file,
1178                                         std::int32_t layer,
1179                                         std::string_view suffix,
1180                                         std::span<const std::int64_t> ↵
1181                                         ↵ expected_shape) {
1182     const std::array<std::string, 2> names = layer_names(layer, suffix);
1183     const std::array<std::string_view, 2> views{names[0], names[1]};
1184     return require_first_tensor(file, views, expected_shape);
1185 }
1186
1187 std::filesystem::path find_gpt2_weight_file(const std::filesystem::path& ↵
1188 ↵ model_directory) {
1189     const std::array<const char*, 2> candidates{
1190         "model.safetensors",
1191         "pytorch_model.safetensors",
1192     };
1193     for (const char* candidate : candidates) {
1194         const std::filesystem::path path = model_directory / candidate;
1195         if (std::filesystem::exists(path)) {
1196             return path;
1197         }
1198     }
1199     throw std::runtime_error("cannot find model.safetensors in GPT-2 directory" ↵
1200     ↵ );
1201 }
1202 // namespace
1203
1204 Gpt2Tokenizer load_gpt2_tokenizer(const std::filesystem::path& vocab_json,
1205                                   const std::filesystem::path& merges_txt,
1206                                   std::int32_t bos_token_id,
1207                                   std::int32_t eos_token_id) {
1208     Gpt2Tokenizer tokenizer;
1209     tokenizer.bos_token_id = bos_token_id;
1210     tokenizer.eos_token_id = eos_token_id;
1211     tokenizer.vocab = parse_vocab_json(read_text_file(vocab_json));
1212     tokenizer.id_to_token = build_id_to_token(tokenizer.vocab);
1213     tokenizer.bpe_ranks = parse_merges(read_text_file(merges_txt));
1214     if (tokenizer.vocab.empty() || tokenizer.bpe_ranks.empty()) {
1215         throw std::runtime_error("GPT-2 tokenizer assets are empty");
1216     }
1217     return tokenizer;
1218 }
1219
1220 std::vector<std::int32_t> gpt2_encode(const Gpt2Tokenizer& tokenizer,
1221                                       std::string_view text) {
1222     std::vector<std::int32_t> ids;
1223     for (const std::string& piece : gpt2_pretokenize(text)) {
1224         const std::string byte_piece = bytes_to_bpe_alphabet(piece);
1225         const std::vector<std::string> bpe_tokens = bpe_apply(tokenizer, ↵
1226         ↵ byte_piece);
1227         for (const std::string& token : bpe_tokens) {

```

```

1225         const auto found = tokenizer.vocab.find(token);
1226         if (found == tokenizer.vocab.end()) {
1227             throw std::runtime_error("GPT-2 BPE token is missing from vocab");
1228         }
1229         ids.push_back(found->second);
1230     }
1231 }
1232 return ids;
1233 }
1234
1235 std::string gpt2_decode(const Gpt2Tokenizer& tokenizer,
1236                        std::span<const std::int32_t> token_ids) {
1237     static const std::vector<std::string> BYTE_ENCODER = gpt2_bytes_to_unicode();
1238     static const std::unordered_map<std::string, unsigned char> BYTE_DECODER =
1239         gpt2_unicode_to_bytes(BYTE_ENCODER);
1240
1241     std::string token_text;
1242     for (const std::int32_t token_id : token_ids) {
1243         if (token_id < 0 || token_id >= static_cast<std::int32_t>(tokenizer.id_to_token.size())) {
1244             throw std::runtime_error("GPT-2 token id out of range");
1245         }
1246         token_text += tokenizer.id_to_token[checked_size(token_id, "token id")];
1247     }
1248
1249     std::string result;
1250     const std::vector<std::string> symbols = utf8_symbols(token_text);
1251     for (const std::string& symbol : symbols) {
1252         const auto found = BYTE_DECODER.find(symbol);
1253         if (found == BYTE_DECODER.end()) {
1254             throw std::runtime_error("GPT-2 token cannot be decoded to byte");
1255         }
1256         result.push_back(static_cast<char>(found->second));
1257     }
1258     return result;
1259 }
1260
1261 Gpt2Config load_gpt2_config(const std::filesystem::path& config_json) {
1262     const std::string config_text = read_text_file(config_json);
1263     JsonCursor cursor(config_text);
1264     Gpt2Config config;
1265     config.bos_token_id = LCQI_GPT2_DEFAULT_BOS_EOS;
1266     config.eos_token_id = LCQI_GPT2_DEFAULT_BOS_EOS;
1267
1268     cursor.expect('{');
1269     if (cursor.consume('}') {
1270         throw std::runtime_error("GPT-2 config is empty");
1271     }
1272     while (true) {

```

```

1273     const std::string key = cursor.parse_string();
1274     cursor.expect(':');
1275     if (key == "n_embd") {
1276         config.hidden_size = parse_i32_field(cursor);
1277     } else if (key == "n_head") {
1278         config.head_count = parse_i32_field(cursor);
1279     } else if (key == "n_layer") {
1280         config.layer_count = parse_i32_field(cursor);
1281     } else if (key == "n_positions" || key == "n_ctx") {
1282         const std::int32_t value = parse_i32_field(cursor);
1283         config.max_positions = std::max(config.max_positions, value);
1284     } else if (key == "vocab_size") {
1285         config.vocab_size = parse_i32_field(cursor);
1286     } else if (key == "n_inner") {
1287         config.intermediate_size = parse_i32_field(cursor);
1288     } else if (key == "layer_norm_epsilon") {
1289         config.layer_norm_epsilon = static_cast<float>(cursor.parse_number←
↪ ());
1290     } else if (key == "activation_function") {
1291         config.activation_function = cursor.parse_string();
1292     } else if (key == "bos_token_id") {
1293         config.bos_token_id = parse_i32_field(cursor);
1294     } else if (key == "eos_token_id") {
1295         config.eos_token_id = parse_i32_field(cursor);
1296     } else {
1297         cursor.skip_value();
1298     }
1299     if (cursor.consume('}')) {
1300         break;
1301     }
1302     cursor.expect(',');
1303 }
1304 if (!cursor.eof_after_ws()) {
1305     throw std::runtime_error("trailing data after GPT-2 config JSON");
1306 }
1307 if (config.intermediate_size == 0 && config.hidden_size > 0) {
1308     config.intermediate_size = config.hidden_size * 4;
1309 }
1310 validate_config(config);
1311 return config;
1312 }
1313
1314 Gpt2ReferenceModel load_gpt2_reference_model(const std::filesystem::path& ←
↪ config_json,
1315                                             const std::filesystem::path& ←
↪ safetensors_path) {
1316     Gpt2ReferenceModel model;
1317     model.config = load_gpt2_config(config_json);
1318     const SafeTensorFile file = load_safetensors_file(safetensors_path);
1319     const Gpt2Config& config = model.config;
1320     const std::array<std::int64_t, 2> wte_shape{config.vocab_size, config.←
↪ hidden_size};

```

```

1321     const std::array<std::int64_t, 2> wpe_shape{config.max_positions, config.↵
↵ hidden_size};
1322     const std::array<std::int64_t, 1> hidden_shape{config.hidden_size};
1323
1324     const std::array<std::string_view, 2> wte_names = embedding_names("wte.↵
↵ weight");
1325     const std::array<std::string_view, 2> wpe_names = embedding_names("wpe.↵
↵ weight");
1326     model.token_embedding = require_first_tensor(file, wte_names, wte_shape);
1327     model.position_embedding = require_first_tensor(file, wpe_names, wpe_shape)↵
↵ ;
1328     model.layers.resize(check_size(config.layer_count, "layer_count"));
1329
1330     for (std::int32_t layer_id = 0; layer_id < config.layer_count; ++layer_id) ↵
↵ {
1331         Gpt2LayerWeightsF32& layer = model.layers[check_size(layer_id, "↵
↵ layer_id")];
1332         layer.ln_1_weight = require_layer_tensor(file, layer_id, "ln_1.weight",↵
↵ hidden_shape);
1333         layer.ln_1_bias = require_layer_tensor(file, layer_id, "ln_1.bias", ↵
↵ hidden_shape);
1334         layer.ln_2_weight = require_layer_tensor(file, layer_id, "ln_2.weight",↵
↵ hidden_shape);
1335         layer.ln_2_bias = require_layer_tensor(file, layer_id, "ln_2.bias", ↵
↵ hidden_shape);
1336
1337         const std::array<std::int64_t, 2> c_attn_shape{
1338             config.hidden_size,
1339             config.hidden_size * LCQI_GPT2_ATTENTION_PROJECTIONS};
1340         std::vector<float> c_attn_weight =
1341             require_layer_tensor(file, layer_id, "attn.c_attn.weight", ↵
↵ c_attn_shape);
1342         transpose_hf_conv1d(c_attn_weight,
1343                             config.hidden_size,
1344                             config.hidden_size * ↵
↵ LCQI_GPT2_ATTENTION_PROJECTIONS);
1345         const std::array<std::int64_t, 1> c_attn_bias_shape{
1346             config.hidden_size * LCQI_GPT2_ATTENTION_PROJECTIONS};
1347         layer.c_attn = make_linear(config.hidden_size,
1348                                   config.hidden_size * ↵
↵ LCQI_GPT2_ATTENTION_PROJECTIONS,
1349                                   std::move(c_attn_weight),
1350                                   require_layer_tensor(file,
1351                                                         layer_id,
1352                                                         "attn.c_attn.bias",
1353                                                         c_attn_bias_shape));
1354
1355         const std::array<std::int64_t, 2> c_proj_shape{config.hidden_size, ↵
↵ config.hidden_size};
1356         std::vector<float> c_proj_weight =
1357             require_layer_tensor(file, layer_id, "attn.c_proj.weight", ↵
↵ c_proj_shape);

```

```

1358     transpose_hf_conv1d(c_proj_weight, config.hidden_size, config.↵
↵ hidden_size);
1359     layer.c_proj = make_linear(config.hidden_size,
1360                               config.hidden_size,
1361                               std::move(c_proj_weight),
1362                               require_layer_tensor(file,
1363                                                    layer_id,
1364                                                    "attn.c_proj.bias",
1365                                                    hidden_shape));
1366
1367     const std::array<std::int64_t, 2> c_fc_shape{
1368         config.hidden_size,
1369         config.intermediate_size};
1370     std::vector<float> c_fc_weight =
1371         require_layer_tensor(file, layer_id, "mlp.c_fc.weight", c_fc_shape)↵
↵ ;
1372     transpose_hf_conv1d(c_fc_weight, config.hidden_size, config.↵
↵ intermediate_size);
1373     const std::array<std::int64_t, 1> intermediate_shape{config.↵
↵ intermediate_size};
1374     layer.c_fc = make_linear(config.hidden_size,
1375                              config.intermediate_size,
1376                              std::move(c_fc_weight),
1377                              require_layer_tensor(file,
1378                                                    layer_id,
1379                                                    "mlp.c_fc.bias",
1380                                                    intermediate_shape));
1381
1382     const std::array<std::int64_t, 2> mlp_proj_shape{
1383         config.intermediate_size,
1384         config.hidden_size};
1385     std::vector<float> mlp_proj_weight =
1386         require_layer_tensor(file, layer_id, "mlp.c_proj.weight", ↵
↵ mlp_proj_shape);
1387     transpose_hf_conv1d(mlp_proj_weight, config.intermediate_size, config.↵
↵ hidden_size);
1388     layer.mlp_c_proj = make_linear(config.intermediate_size,
1389                                    config.hidden_size,
1390                                    std::move(mlp_proj_weight),
1391                                    require_layer_tensor(file,
1392                                                         layer_id,
1393                                                         "mlp.c_proj.bias",
1394                                                         hidden_shape));
1395 }
1396
1397 const std::array<std::string_view, 2> ln_weight_names = embedding_names("↵
↵ ln_f.weight");
1398 const std::array<std::string_view, 2> ln_bias_names = embedding_names("ln_f↵
↵ .bias");
1399 model.final_ln_weight = require_first_tensor(file, ln_weight_names, ↵
↵ hidden_shape);
1400 model.final_ln_bias = require_first_tensor(file, ln_bias_names, ↵

```



```

    ↪ hidden_shape);
1401 if (file.manifest.find_tensor("lm_head.weight") != nullptr) {
1402     model.tie_lm_head_to_embedding = false;
1403     model.lm_head_weight = require_tensor(file, "lm_head.weight", wte_shape ↪
    ↪ );
1404 } else {
1405     model.tie_lm_head_to_embedding = true;
1406 }
1407 validate_model(model);
1408 return model;
1409 }
1410
1411 Gpt2ReferenceModel load_gpt2_from_directory(const std::filesystem::path& ↪
    ↪ model_directory) {
1412     return load_gpt2_reference_model(model_directory / "config.json",
1413                                     find_gpt2_weight_file(model_directory));
1414 }
1415
1416 Gpt2KvCache::Gpt2KvCache(const Gpt2Config& config) : config_(config) {
1417     validate_config(this->config_);
1418     const std::size_t element_count =
1419         checked_size(this->config_.layer_count, "layer_count") *
1420         checked_size(this->config_.max_positions, "max_positions") *
1421         checked_size(this->config_.head_count, "head_count") *
1422         checked_size(head_dim(this->config_), "head_dim");
1423     this->keys_.assign(element_count, 0.0F);
1424     this->values_.assign(element_count, 0.0F);
1425     this->written_.assign(
1426         checked_size(this->config_.layer_count, "layer_count") *
1427         checked_size(this->config_.max_positions, "max_positions"),
1428         0);
1429 }
1430
1431 void Gpt2KvCache::append(std::int32_t layer_id,
1432                          std::int32_t model_position,
1433                          std::span<const float> key,
1434                          std::span<const float> value) {
1435     const std::size_t hidden_size = checked_size(this->config_.hidden_size, " ↪
    ↪ hidden_size");
1436     if (key.size() != hidden_size || value.size() != hidden_size) {
1437         throw std::runtime_error("GPT-2 KV key/value size mismatch");
1438     }
1439     this->validate_address(layer_id, model_position, 0);
1440     const std::int32_t local_head_dim = head_dim(this->config_);
1441     for (std::int32_t head = 0; head < this->config_.head_count; ++head) {
1442         const std::size_t source =
1443             checked_size(head, "head") * checked_size(local_head_dim, "head_dim ↪
    ↪ ");
1444         const std::size_t target = this->base_offset(layer_id, model_position, ↪
    ↪ head);
1445         std::copy_n(key.begin() + static_cast<std::ptrdiff_t>(source),
1446                    local_head_dim,

```

```

1447         this->keys_.begin() + static_cast<std::ptrdiff_t>(target));
1448     std::copy_n(value.begin() + static_cast<std::ptrdiff_t>(source),
1449               local_head_dim,
1450               this->values_.begin() + static_cast<std::ptrdiff_t>(target)↵
1451     ↵ );
1452 }
1453 const std::size_t written_index =
1454     checked_size(layer_id, "layer_id") *
1455     checked_size(this->config_.max_positions, "max_positions") +
1456     checked_size(model_position, "model_position");
1457 this->written_[written_index] = 1;
1458 this->filled_tokens_ = std::max(this->filled_tokens_, model_position + 1);
1459 }
1460 std::span<const float> Gpt2KvCache::key(std::int32_t layer_id,
1461                                       std::int32_t model_position,
1462                                       std::int32_t head) const {
1463     this->validate_address(layer_id, model_position, head);
1464     this->validate_written(layer_id, model_position);
1465     const std::size_t offset = this->base_offset(layer_id, model_position, head)↵
1466     ↵ );
1467     return std::span<const float>(this->keys_.data() + offset,
1468                                checked_size(head_dim(this->config_), "↵
1469     ↵ head_dim"));
1470 }
1471 std::span<const float> Gpt2KvCache::value(std::int32_t layer_id,
1472                                       std::int32_t model_position,
1473                                       std::int32_t head) const {
1474     this->validate_address(layer_id, model_position, head);
1475     this->validate_written(layer_id, model_position);
1476     const std::size_t offset = this->base_offset(layer_id, model_position, head)↵
1477     ↵ );
1478     return std::span<const float>(this->values_.data() + offset,
1479                                checked_size(head_dim(this->config_), "↵
1480     ↵ head_dim"));
1481 }
1482 std::int32_t Gpt2KvCache::filled_tokens() const noexcept {
1483     return this->filled_tokens_;
1484 }
1485 std::size_t Gpt2KvCache::byte_size() const noexcept {
1486     return (this->keys_.size() + this->values_.size()) * sizeof(float) +
1487            this->written_.size() * sizeof(std::uint8_t);
1488 }
1489 std::size_t Gpt2KvCache::base_offset(std::int32_t layer_id,
1490                                       std::int32_t model_position,
1491                                       std::int32_t head) const {
1492     return (((checked_size(layer_id, "layer_id") *
1493               checked_size(this->config_.max_positions, "max_positions")) +

```

```

1494         checked_size(model_position, "model_position")) *
1495         checked_size(this->config_.head_count, "head_count") +
1496         checked_size(head, "head")) *
1497         checked_size(head_dim(this->config_), "head_dim");
1498     }
1499
1500     void Gpt2KvCache::validate_address(std::int32_t layer_id,
1501                                       std::int32_t model_position,
1502                                       std::int32_t head) const {
1503         if (layer_id < 0 || layer_id >= this->config_.layer_count) {
1504             throw std::runtime_error("GPT-2 KV layer_id out of range");
1505         }
1506         if (model_position < 0 || model_position >= this->config_.max_positions) {
1507             throw std::runtime_error("GPT-2 KV model_position out of range");
1508         }
1509         if (head < 0 || head >= this->config_.head_count) {
1510             throw std::runtime_error("GPT-2 KV head out of range");
1511         }
1512     }
1513
1514     void Gpt2KvCache::validate_written(std::int32_t layer_id,
1515                                       std::int32_t model_position) const {
1516         const std::size_t written_index =
1517             checked_size(layer_id, "layer_id") *
1518             checked_size(this->config_.max_positions, "max_positions") +
1519             checked_size(model_position, "model_position");
1520         if (written_index >= this->written_.size() || this->written_[written_index] ←
1521             == 0) {
1522             throw std::runtime_error("GPT-2 KV slot has not been written");
1523         }
1524     }
1525
1526     Gpt2ForwardResult run_gpt2_forward(const Gpt2ReferenceModel& model,
1527                                       std::span<const std::int32_t> token_ids) {
1528         validate_model(model);
1529         if (token_ids.empty()) {
1530             throw std::runtime_error("GPT-2 forward needs at least one token");
1531         }
1532         if (token_ids.size() > checked_size(model.config.max_positions, " ←
1533             max_positions")) {
1534             throw std::runtime_error("GPT-2 token_ids exceed max_positions");
1535         }
1536
1537         const std::int32_t sequence_length = static_cast<std::int32_t>(token_ids. ←
1538             size());
1539         std::vector<float> hidden_by_pos(matrix_size(sequence_length, model.config. ←
1540             hidden_size), 0.0F);
1541         for (std::int32_t position = 0; position < sequence_length; ++position) {
1542             const std::int32_t token_id = token_ids[checked_size(position, " ←
1543                 position")];
1544             if (token_id < 0 || token_id >= model.config.vocab_size) {
1545                 throw std::runtime_error("GPT-2 token id out of range");
1546             }
1547         }
1548     }

```

```

1541     }
1542     const std::span<const float> token =
1543         row_span(model.token_embedding, token_id, model.config.hidden_size)↵
1544     ↵ ;
1545     const std::span<const float> position_embedding =
1546         row_span(model.position_embedding, position, model.config.↵
1547     ↵ hidden_size);
1548     std::span<float> hidden = std::span<float>(
1549         hidden_by_pos.data() + matrix_size(position, model.config.↵
1550     ↵ hidden_size),
1551         checked_size(model.config.hidden_size, "hidden_size"));
1552     for (std::int32_t dim = 0; dim < model.config.hidden_size; ++dim) {
1553         hidden[checked_size(dim, "dim")] =
1554             token[checked_size(dim, "dim")] + position_embedding[↵
1555     ↵ checked_size(dim, "dim")];
1556     }
1557 }
1558
1559 for (const Gpt2LayerWeightsF32& layer : model.layers) {
1560     forward_gpt2_layer(model.config, layer, sequence_length, hidden_by_pos)↵
1561     ↵ ;
1562 }
1563
1564 std::vector<float> normed(checked_size(model.config.hidden_size, "↵
1565     ↵ hidden_size"), 0.0F);
1566 const std::span<const float> final_hidden = std::span<const float>(
1567     hidden_by_pos.data() + matrix_size(sequence_length - 1, model.config.↵
1568     ↵ hidden_size),
1569     checked_size(model.config.hidden_size, "hidden_size"));
1570 layer_norm(final_hidden,
1571     model.final_ln_weight,
1572     model.final_ln_bias,
1573     model.config.layer_norm_epsilon,
1574     normed);
1575
1576 Gpt2ForwardResult result;
1577 result.logits.assign(checked_size(model.config.vocab_size, "vocab_size"), ↵
1578     ↵ 0.0F);
1579 compute_logits(model, normed, result.logits);
1580 result.predicted_token = argmax(result.logits);
1581 return result;
1582 }
1583
1584 Gpt2ForwardResult run_gpt2_forward_cached(const Gpt2ReferenceModel& model,
1585     Gpt2KvCache& cache,
1586     std::int32_t token_id) {
1587     validate_model(model);
1588     if (token_id < 0 || token_id >= model.config.vocab_size) {
1589         throw std::runtime_error("GPT-2 token id out of range");
1590     }
1591     const std::int32_t model_position = cache.filled_tokens();
1592     if (model_position >= model.config.max_positions) {

```

```

1585     throw std::runtime_error("GPT-2 cached forward would exceed ↵
↵ max_positions");
1586 }
1587
1588 std::vector<float> hidden(check_size(model.config.hidden_size, "↵
↵ hidden_size"), 0.0F);
1589 const std::span<const float> token =
1590     row_span(model.token_embedding, token_id, model.config.hidden_size);
1591 const std::span<const float> position_embedding =
1592     row_span(model.position_embedding, model_position, model.config.↵
↵ hidden_size);
1593 for (std::int32_t dim = 0; dim < model.config.hidden_size; ++dim) {
1594     hidden[check_size(dim, "dim")] =
1595         token[check_size(dim, "dim")] + position_embedding[check_size(↵
↵ dim, "dim")];
1596 }
1597
1598 for (std::int32_t layer_id = 0;
1599     layer_id < static_cast<std::int32_t>(model.layers.size()));
1600     ++layer_id) {
1601     forward_gpt2_layer_cached(model.config,
1602                               model.layers[check_size(layer_id, "layer_id↵
↵ ")]),
1603                               layer_id,
1604                               model_position,
1605                               cache,
1606                               hidden);
1607 }
1608
1609 std::vector<float> normed(check_size(model.config.hidden_size, "↵
↵ hidden_size"), 0.0F);
1610 layer_norm(hidden,
1611            model.final_ln_weight,
1612            model.final_ln_bias,
1613            model.config.layer_norm_epsilon,
1614            normed);
1615
1616 Gpt2ForwardResult result;
1617 result.logits.assign(check_size(model.config.vocab_size, "vocab_size"), ↵
↵ 0.0F);
1618 compute_logits(model, normed, result.logits);
1619 result.predicted_token = argmax(result.logits);
1620 return result;
1621 }
1622
1623 std::vector<std::int32_t> gpt2_generate_greedy(const Gpt2ReferenceModel& model,
1624                                                std::span<const std::int32_t> ↵
↵ prompt_token_ids,
1625                                                std::int32_t max_new_tokens) {
1626     if (max_new_tokens < 0) {
1627         throw std::runtime_error("max_new_tokens cannot be negative");
1628     }

```

```

1629     std::vector<std::int32_t> tokens(prompt_token_ids.begin(), prompt_token_ids↵
↵ .end());
1630     for (std::int32_t step = 0; step < max_new_tokens; ++step) {
1631         if (tokens.size() >= checked_size(model.config.max_positions, "↵
↵ max_positions")) {
1632             throw std::runtime_error("GPT-2 generation would exceed ↵
↵ max_positions");
1633         }
1634         const Gpt2ForwardResult result = run_gpt2_forward(model, tokens);
1635         tokens.push_back(result.predicted_token);
1636         if (model.config.eos_token_id >= 0 && result.predicted_token == model.↵
↵ config.eos_token_id) {
1637             break;
1638         }
1639     }
1640     return tokens;
1641 }
1642
1643 std::vector<std::int32_t> gpt2_generate_greedy_cached(
1644     const Gpt2ReferenceModel& model,
1645     std::span<const std::int32_t> prompt_token_ids,
1646     std::int32_t max_new_tokens) {
1647     if (prompt_token_ids.empty()) {
1648         throw std::runtime_error("GPT-2 generation needs at least one prompt ↵
↵ token");
1649     }
1650     if (max_new_tokens < 0) {
1651         throw std::runtime_error("max_new_tokens cannot be negative");
1652     }
1653     if (prompt_token_ids.size() > checked_size(model.config.max_positions, "↵
↵ max_positions")) {
1654         throw std::runtime_error("GPT-2 prompt exceeds max_positions");
1655     }
1656
1657     Gpt2KvCache cache(model.config);
1658     std::vector<std::int32_t> tokens(prompt_token_ids.begin(), prompt_token_ids↵
↵ .end());
1659     if (max_new_tokens == 0) {
1660         return tokens;
1661     }
1662     Gpt2ForwardResult result;
1663     for (const std::int32_t token_id : prompt_token_ids) {
1664         result = run_gpt2_forward_cached(model, cache, token_id);
1665     }
1666     for (std::int32_t step = 0; step < max_new_tokens; ++step) {
1667         if (tokens.size() >= checked_size(model.config.max_positions, "↵
↵ max_positions")) {
1668             throw std::runtime_error("GPT-2 generation would exceed ↵
↵ max_positions");
1669         }
1670         tokens.push_back(result.predicted_token);
1671         if (model.config.eos_token_id >= 0 && result.predicted_token == model.↵

```

```

    ↪ config.eos_token_id) {
1672         break;
1673     }
1674     if (step + 1 < max_new_tokens) {
1675         result = run_gpt2_forward_cached(model, cache, result.↪
    ↪ predicted_token);
1676     }
1677 }
1678 return tokens;
1679 }
1680
1681 Gpt2ReferenceModel make_tiny_gpt2_reference_model() {
1682     Gpt2ReferenceModel model;
1683     model.config.hidden_size = 4;
1684     model.config.head_count = 2;
1685     model.config.layer_count = 1;
1686     model.config.max_positions = 8;
1687     model.config.vocab_size = 6;
1688     model.config.intermediate_size = 8;
1689     model.config.bos_token_id = 0;
1690     model.config.eos_token_id = 5;
1691     model.config.layer_norm_epsilon = 1.0e-5F;
1692     model.config.activation_function = "gelu_new";
1693
1694     model.token_embedding = {
1695         0.20F, -0.10F, 0.00F, 0.10F,
1696         -0.10F, 0.30F, 0.20F, -0.20F,
1697         0.05F, 0.10F, -0.30F, 0.25F,
1698         0.40F, -0.20F, 0.15F, 0.05F,
1699         -0.30F, 0.05F, 0.35F, -0.15F,
1700         0.10F, 0.20F, -0.10F, 0.30F,
1701     };
1702     model.position_embedding = {
1703         0.01F, 0.02F, 0.03F, 0.04F,
1704         -0.02F, 0.01F, 0.00F, 0.03F,
1705         0.03F, -0.01F, 0.02F, 0.00F,
1706         0.00F, 0.02F, -0.02F, 0.01F,
1707         0.02F, 0.00F, 0.01F, -0.01F,
1708         -0.01F, 0.03F, 0.00F, 0.02F,
1709         0.01F, -0.02F, 0.03F, 0.00F,
1710         0.00F, 0.01F, -0.01F, 0.02F,
1711     };
1712
1713     Gpt2LayerWeightsF32 layer;
1714     layer.ln_1_weight = ones(4);
1715     layer.ln_1_bias = zeros(4);
1716     layer.ln_2_weight = {0.9F, 1.1F, 1.0F, 0.95F};
1717     layer.ln_2_bias = {0.01F, -0.02F, 0.00F, 0.03F};
1718     layer.c_attn = make_linear(
1719         4,
1720         12,
1721         {

```

```

1722         0.10F, -0.20F, 0.05F, 0.15F,
1723         -0.05F, 0.12F, 0.20F, -0.10F,
1724         0.18F, 0.02F, -0.08F, 0.11F,
1725         -0.12F, 0.06F, 0.14F, 0.04F,
1726         0.07F, 0.16F, -0.09F, 0.03F,
1727         -0.15F, 0.05F, 0.13F, 0.10F,
1728         0.04F, -0.11F, 0.19F, -0.02F,
1729         0.09F, 0.08F, -0.06F, 0.17F,
1730         0.20F, -0.04F, 0.01F, 0.06F,
1731         -0.08F, 0.15F, 0.07F, -0.03F,
1732         0.05F, 0.10F, -0.14F, 0.12F,
1733         0.11F, -0.07F, 0.16F, 0.02F,
1734     },
1735     {
1736         0.01F, -0.02F, 0.00F, 0.03F,
1737         -0.01F, 0.02F, 0.01F, -0.02F,
1738         0.00F, 0.01F, -0.01F, 0.02F,
1739     });
1740     layer.c_proj = make_linear(
1741         4,
1742         4,
1743         {
1744             0.10F, 0.05F, -0.08F, 0.12F,
1745             -0.04F, 0.11F, 0.09F, -0.02F,
1746             0.07F, -0.10F, 0.13F, 0.05F,
1747             0.02F, 0.14F, -0.06F, 0.08F,
1748         },
1749         {0.01F, 0.00F, -0.01F, 0.02F});
1750     layer.c_fc = make_linear(
1751         4,
1752         8,
1753         {
1754             0.20F, -0.10F, 0.05F, 0.03F,
1755             -0.05F, 0.14F, 0.12F, -0.08F,
1756             0.09F, 0.07F, -0.11F, 0.16F,
1757             -0.12F, 0.05F, 0.18F, 0.02F,
1758             0.04F, -0.09F, 0.13F, 0.10F,
1759             0.11F, 0.03F, -0.07F, 0.15F,
1760             -0.02F, 0.17F, 0.06F, -0.04F,
1761             0.08F, -0.06F, 0.10F, 0.12F,
1762         },
1763         {0.01F, -0.02F, 0.00F, 0.03F, -0.01F, 0.02F, 0.01F, 0.00F});
1764     layer.mlp_c_proj = make_linear(
1765         8,
1766         4,
1767         {
1768             0.10F, -0.08F, 0.06F, 0.12F, -0.04F, 0.07F, 0.03F, 0.11F,
1769             -0.05F, 0.13F, 0.02F, -0.09F, 0.15F, 0.04F, -0.03F, 0.08F,
1770             0.07F, 0.01F, -0.12F, 0.10F, 0.06F, -0.05F, 0.14F, 0.02F,
1771             0.03F, 0.09F, 0.11F, -0.06F, 0.05F, 0.12F, -0.08F, 0.04F,
1772         },
1773         {0.00F, 0.01F, -0.01F, 0.02F});

```



```

1774     model.layers.push_back(std::move(layer));
1775     model.final_ln_weight = {1.0F, 0.95F, 1.05F, 1.0F};
1776     model.final_ln_bias = {0.0F, 0.01F, -0.01F, 0.02F};
1777     model.tie_lm_head_to_embedding = true;
1778     validate_model(model);
1779     return model;
1780 }
1781
1782 } // namespace lcqi

```

`gpt2_cli.cpp` 负责把模型目录、prompt、`max_new_tokens` 和 benchmark 选项接到 GPT-2 reference path。无参数时运行 tiny GPT-2 fixture；传入 HuggingFace 模型目录时加载真实资产。

Listing 10.12 `src/gpt2_cli.cpp`

```

1  #include <lcqi/gpt2_reference.hpp>
2
3  #include <chrono>
4  #include <cstdlib>
5  #include <exception>
6  #include <filesystem>
7  #include <iostream>
8  #include <span>
9  #include <stdexcept>
10 #include <string>
11 #include <string_view>
12 #include <vector>
13
14 namespace {
15
16 constexpr int LCQI_GPT2_TINY_ARGC = 1;
17 constexpr int LCQI_GPT2_MIN_REAL_ARGC = 3;
18 constexpr std::int32_t LCQI_GPT2_DEFAULT_MAX_NEW_TOKENS = 8;
19 constexpr std::int32_t LCQI_GPT2_TINY_PROMPT_0 = 1;
20 constexpr std::int32_t LCQI_GPT2_TINY_PROMPT_1 = 2;
21 constexpr std::int32_t LCQI_GPT2_TINY_PROMPT_2 = 3;
22
23 enum class Gpt2EngineMode {
24     full_prefix,
25     cached_kv,
26 };
27
28 struct Gpt2CliOptions {
29     Gpt2EngineMode engine = Gpt2EngineMode::cached_kv;
30     bool benchmark = false;
31     std::filesystem::path model_dir;
32     std::string prompt;
33     std::int32_t max_new_tokens = LCQI_GPT2_DEFAULT_MAX_NEW_TOKENS;
34 };
35
36 struct Gpt2BenchmarkTimings {
37     double load_ms = 0.0;
38     double tokenize_ms = 0.0;

```

```

39     double prefill_ms = 0.0;
40     double decode_ms = 0.0;
41     double generate_ms = 0.0;
42     double total_ms = 0.0;
43     std::size_t prefill_steps = 0;
44     std::size_t decode_steps = 0;
45     std::size_t kv_cache_bytes = 0;
46 };
47
48 using Clock = std::chrono::steady_clock;
49
50 void print_ids(std::span<const std::int32_t> ids) {
51     for (std::size_t index = 0; index < ids.size(); ++index) {
52         if (index != 0) {
53             std::cout << ' ';
54         }
55         std::cout << ids[index];
56     }
57 }
58
59 [[nodiscard]] double elapsed_ms(Clock::time_point start, Clock::time_point end)↵
    ↵ {
60     return std::chrono::duration<double, std::milli>(end - start).count();
61 }
62
63 [[nodiscard]] std::size_t checked_size(std::int64_t value, const char* name) {
64     if (value < 0) {
65         throw std::runtime_error(std::string(name) + " cannot be negative");
66     }
67     return static_cast<std::size_t>(value);
68 }
69
70 [[nodiscard]] const char* engine_name(Gpt2EngineMode engine) {
71     switch (engine) {
72         case Gpt2EngineMode::full_prefix:
73             return "full_prefix";
74         case Gpt2EngineMode::cached_kv:
75             return "cached_kv";
76     }
77     return "unknown";
78 }
79
80 [[nodiscard]] Gpt2EngineMode parse_engine(std::string_view value) {
81     if (value == "full" || value == "full_prefix") {
82         return Gpt2EngineMode::full_prefix;
83     }
84     if (value == "cached" || value == "cached_kv") {
85         return Gpt2EngineMode::cached_kv;
86     }
87     throw std::runtime_error("unknown --engine value, expected cached or full")↵
    ↵ ;
88 }

```

```

89
90 [[nodiscard]] std::int32_t parse_non_negative_i32(const std::string& text,
91                                                    const char* name) {
92     const std::int32_t value = static_cast<std::int32_t>(std::stoi(text));
93     if (value < 0) {
94         throw std::runtime_error(std::string(name) + " cannot be negative");
95     }
96     return value;
97 }
98
99 [[nodiscard]] Gpt2CliOptions parse_options(int argc, char** argv) {
100     Gpt2CliOptions options;
101     std::vector<std::string> positional;
102     for (int index = 1; index < argc; ++index) {
103         const std::string arg = argv[index];
104         if (arg == "--benchmark") {
105             options.benchmark = true;
106         } else if (arg == "--engine") {
107             if (index + 1 >= argc) {
108                 throw std::runtime_error("--engine requires a value");
109             }
110             ++index;
111             options.engine = parse_engine(argv[index]);
112         } else if (arg.rfind("--engine=", 0) == 0) {
113             options.engine = parse_engine(std::string_view(arg).substr(std::string(
114 ↪ "--engine=").size()));
115         } else if (arg == "--full") {
116             options.engine = Gpt2EngineMode::full_prefix;
117         } else if (arg == "--cached") {
118             options.engine = Gpt2EngineMode::cached_kv;
119         } else if (!arg.empty() && arg.front() == '-') {
120             throw std::runtime_error("unknown option: " + arg);
121         } else {
122             positional.push_back(arg);
123         }
124     }
125     if (positional.size() < 2 || positional.size() > 3) {
126         throw std::runtime_error(
127             "usage: lcqi_gpt2 [--engine cached|full] [--benchmark] "
128             "model_dir prompt [max_new_tokens]");
129     }
130     options.model_dir = positional[0];
131     options.prompt = positional[1];
132     if (positional.size() == 3) {
133         options.max_new_tokens = parse_non_negative_i32(positional[2], "
134 ↪ max_new_tokens");
135     }
136     return options;
137 }
138 [[nodiscard]] std::vector<std::int32_t> generate_full_prefix(

```

```

139     const lcqi::Gpt2ReferenceModel& model,
140     std::span<const std::int32_t> prompt_ids,
141     std::int32_t max_new_tokens,
142     Gpt2BenchmarkTimings& timings) {
143     const Clock::time_point generate_start = Clock::now();
144     std::vector<std::int32_t> tokens(prompt_ids.begin(), prompt_ids.end());
145     for (std::int32_t step = 0; step < max_new_tokens; ++step) {
146         if (tokens.size() >= checked_size(model.config.max_positions, "↵
↵ max_positions")) {
147             throw std::runtime_error("GPT-2 generation would exceed ↵
↵ max_positions");
148         }
149         const lcqi::Gpt2ForwardResult result = lcqi::run_gpt2_forward(model, ↵
↵ tokens);
150         ++timings.decode_steps;
151         tokens.push_back(result.predicted_token);
152         if (model.config.eos_token_id >= 0 && result.predicted_token == model.↵
↵ config.eos_token_id) {
153             break;
154         }
155     }
156     const Clock::time_point generate_end = Clock::now();
157     timings.generate_ms = elapsed_ms(generate_start, generate_end);
158     timings.decode_ms = timings.generate_ms;
159     return tokens;
160 }

161 [[nodiscard]] std::vector<std::int32_t> generate_cached(
162     const lcqi::Gpt2ReferenceModel& model,
163     std::span<const std::int32_t> prompt_ids,
164     std::int32_t max_new_tokens,
165     Gpt2BenchmarkTimings& timings) {
166     if (prompt_ids.empty()) {
167         throw std::runtime_error("GPT-2 generation needs at least one prompt ↵
↵ token");
168     }
169     lcqi::Gpt2KvCache cache(model.config);
170     timings.kv_cache_bytes = cache.byte_size();
171     std::vector<std::int32_t> tokens(prompt_ids.begin(), prompt_ids.end());
172     if (max_new_tokens == 0) {
173         return tokens;
174     }
175
176     const Clock::time_point generate_start = Clock::now();
177     lcqi::Gpt2ForwardResult result;
178     const Clock::time_point prefill_start = Clock::now();
179     for (const std::int32_t token_id : prompt_ids) {
180         result = lcqi::run_gpt2_forward_cached(model, cache, token_id);
181         ++timings.prefill_steps;
182     }
183     const Clock::time_point prefill_end = Clock::now();
184     timings.prefill_ms = elapsed_ms(prefill_start, prefill_end);
185

```

```

186
187     const Clock::time_point decode_start = Clock::now();
188     for (std::int32_t step = 0; step < max_new_tokens; ++step) {
189         if (tokens.size() >= checked_size(model.config.max_positions, "↵
↵ max_positions")) {
190             throw std::runtime_error("GPT-2 generation would exceed ↵
↵ max_positions");
191         }
192         tokens.push_back(result.predicted_token);
193         if (model.config.eos_token_id >= 0 && result.predicted_token == model.↵
↵ config.eos_token_id) {
194             break;
195         }
196         if (step + 1 < max_new_tokens) {
197             result = lcqi::run_gpt2_forward_cached(model, cache, result.↵
↵ predicted_token);
198             ++timings.decode_steps;
199         }
200     }
201     const Clock::time_point decode_end = Clock::now();
202     const Clock::time_point generate_end = Clock::now();
203     timings.decode_ms = elapsed_ms(decode_start, decode_end);
204     timings.generate_ms = elapsed_ms(generate_start, generate_end);
205     return tokens;
206 }
207
208 void print_benchmark(const Gpt2BenchmarkTimings& timings,
209                     std::size_t prompt_tokens,
210                     std::size_t generated_tokens) {
211     const double new_tokens = static_cast<double>(generated_tokens);
212     const double generated_tokens_per_second =
213         timings.generate_ms > 0.0 ? new_tokens * 1000.0 / timings.generate_ms :↵
↵ 0.0;
214     const double decode_tokens_per_second =
215         timings.decode_ms > 0.0
216         ? static_cast<double>(timings.decode_steps) * 1000.0 / timings.↵
↵ decode_ms
217         : 0.0;
218
219     std::cout << "benchmark_load_ms " << timings.load_ms << "\n";
220     std::cout << "benchmark_tokenize_ms " << timings.tokenize_ms << "\n";
221     std::cout << "benchmark_prefill_ms " << timings.prefill_ms << "\n";
222     std::cout << "benchmark_decode_ms " << timings.decode_ms << "\n";
223     std::cout << "benchmark_generate_ms " << timings.generate_ms << "\n";
224     std::cout << "benchmark_total_ms " << timings.total_ms << "\n";
225     std::cout << "benchmark_prompt_tokens " << prompt_tokens << "\n";
226     std::cout << "benchmark_generated_tokens " << generated_tokens << "\n";
227     std::cout << "benchmark_prefill_steps " << timings.prefill_steps << "\n";
228     std::cout << "benchmark_decode_steps " << timings.decode_steps << "\n";
229     std::cout << "benchmark_generate_tokens_per_second "
230         << generated_tokens_per_second << "\n";
231     std::cout << "benchmark_decode_tokens_per_second "

```

```

232         << decode_tokens_per_second << "\n";
233     std::cout << "benchmark_kv_cache_bytes " << timings.kv_cache_bytes << "\n";
234 }
235
236 int run_tiny_mode() {
237     const lcqi::Gpt2ReferenceModel model = lcqi::make_tiny_gpt2_reference_model(
238         ↵ );
239     const std::vector<std::int32_t> prompt{
240         LCQI_GPT2_TINY_PROMPT_0,
241         LCQI_GPT2_TINY_PROMPT_1,
242         LCQI_GPT2_TINY_PROMPT_2,
243     };
244     const lcqi::Gpt2ForwardResult result = lcqi::run_gpt2_forward(model, prompt(
245         ↵ );
246     const std::vector<std::int32_t> generated =
247         lcqi::gpt2_generate_greedy(model, prompt, 2);
248     std::cout << "mode tiny\n";
249     std::cout << "prompt_ids ";
250     print_ids(prompt);
251     std::cout << "\n";
252     std::cout << "predicted_token " << result.predicted_token << "\n";
253     std::cout << "logits";
254     for (const float value : result.logits) {
255         std::cout << ' ' << value;
256     }
257     std::cout << "\n";
258     std::cout << "generated_ids ";
259     print_ids(generated);
260     std::cout << "\n";
261     return EXIT_SUCCESS;
262 }
263
264 int run_real_model_mode(int argc, char** argv) {
265     if (argc < LCQI_GPT2_MIN_REAL_ARGC) {
266         throw std::runtime_error(
267             "usage: lcqi_gpt2 [--engine cached|full] [--benchmark] "
268             "model_dir prompt [max_new_tokens]");
269     }
270     const Clock::time_point total_start = Clock::now();
271     const Gpt2CliOptions options = parse_options(argc, argv);
272
273     Gpt2BenchmarkTimings timings;
274     const Clock::time_point load_start = Clock::now();
275     const lcqi::Gpt2ReferenceModel model = lcqi::load_gpt2_from_directory(
276         ↵ options.model_dir);
277     const lcqi::Gpt2Tokenizer tokenizer = lcqi::load_gpt2_tokenizer(
278         options.model_dir / "vocab.json",
279         options.model_dir / "merges.txt",
280         model.config.bos_token_id,
281         model.config.eos_token_id);
282     timings.load_ms = elapsed_ms(load_start, Clock::now());

```

```

281     const Clock::time_point tokenize_start = Clock::now();
282     const std::vector<std::int32_t> prompt_ids = lcqi::gpt2_encode(tokenizer, ↵
↵ options.prompt);
283     timings.tokenize_ms = elapsed_ms(tokenize_start, Clock::now());
284
285     std::vector<std::int32_t> generated;
286     if (options.engine == Gpt2EngineMode::full_prefix) {
287         generated = generate_full_prefix(model, prompt_ids, options.↵
↵ max_new_tokens, timings);
288     } else {
289         generated = generate_cached(model, prompt_ids, options.max_new_tokens, ↵
↵ timings);
290     }
291     timings.total_ms = elapsed_ms(total_start, Clock::now());
292
293     std::cout << "mode gpt2_directory\n";
294     std::cout << "engine " << engine_name(options.engine) << "\n";
295     std::cout << "model_dir " << options.model_dir.string() << "\n";
296     std::cout << "prompt_ids ";
297     print_ids(prompt_ids);
298     std::cout << "\n";
299     std::cout << "generated_ids ";
300     print_ids(generated);
301     std::cout << "\n";
302     std::cout << "generated_text " << lcqi::gpt2_decode(tokenizer, generated) ↵
↵ << "\n";
303     if (options.benchmark) {
304         const std::size_t generated_tokens = generated.size() - prompt_ids.size()↵
↵ ();
305         print_benchmark(timings, prompt_ids.size(), generated_tokens);
306     }
307     return EXIT_SUCCESS;
308 }
309
310 } // namespace
311
312 int main(int argc, char** argv) {
313     try {
314         if (argc == LCQI_GPT2_TINY_ARGC) {
315             return run_tiny_mode();
316         }
317         return run_real_model_mode(argc, argv);
318     } catch (const std::exception& error) {
319         std::cerr << "error: " << error.what() << "\n";
320         return EXIT_FAILURE;
321     }
322 }

```



## 第 11 章实践：运行时内存计划、liveness 与 Workspace

本章对应原书中的第三册“推理运行时与内存规划”。不要把它当作概念复述，本章要把 activation arena、workspace、liveness、state edge、debug dump、buffer reuse 放到同一个可运行项目里检查。贯穿代码仍然是 [books/cpu-volume-3/labs/linux\\_cpu\\_inference](#)；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `layout`，核心命令或测试入口是 `lcqi_trace`。

**Listing 11.1** 第 11 章实践：运行时内存计划、liveness 与 Workspace 的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  ↪ DCMMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_trace / layout
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：比较普通 activation 和 KV cache 的生命周期。读者要先手写预期结果，再运行项目观察输出。动作是：说明为什么 KV cache 不能被 activation planner 当作临时缓冲复用。如果输出里没有 `checksum` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：内存复用只在生命周期不重叠时成立；debug/oracle/materialization 会延长生命周期这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 [books/cpu-volume-3/labs/linux\\_cpu\\_inference/README.md](#)、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交 liveness 表：每个 checkpoint 的 `defined_at`、`last_used_at`、是否 state、是否允许复用。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。



## 第 12 章实践：IR、Pass、Lowering 与 Executable Plan

本章对应原书中的第三册“AI 编译器细化：IR、pass、lowering 与 executable plan”。不要把它当作概念复述，本章要把 typed graph、pass、fusion、layout propagation、backend capability、fallback reason 放到同一个可运行项目里检查。贯穿代码仍然是 [books/cpu-volume-3/labs/linux\\_cpu\\_inference](#)；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `run_inference`，核心命令或测试入口是 `lcqi_cli`。

**Listing 12.1** 第 12 章实践：IR、Pass、Lowering 与 Executable Plan 的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  <- DCMMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_cli / run_inference
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：把 tiny MLP 写成三节点 IR：linear、relu、linear。读者要先手写预期结果，再运行项目观察输出。动作是：说明哪些 pass 可以融合，哪些必须保留为 oracle checkpoint。如果输出里没有 `backend` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：IR 的价值不是画图，而是让 shape、layout、量化、后端选择和 fallback 诊断有稳定承载物这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 [books/cpu-volume-3/labs/linux\\_cpu\\_inference/README.md](#)、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交 executable plan 草图：节点、输入输出 tensor、layout、kernel、fallback reason 和 profiler event。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。

第零部分

# 后端优化、工业专项与验收

## 第 13 章实践：CPU 后端、微内核与性能证据

本章对应原书中的第三册“CPU 后端优化细讲”和“CPU decode 实战”。不要把它当作概念复述，本章要把 scalar baseline、packed weight、AVX2、objdump、perf boundary、shape sweep 放到同一个可运行项目里检查。贯穿代码仍然是 [books/cpu-volume-3/labs/linux\\_cpu\\_inference](#)；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 **avx2**，核心命令或测试入口是 **lcqi\_bench**。

**Listing 13.1** 第 13 章实践：CPU 后端、微内核与性能证据的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  ↳ DCMMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_bench / avx2
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：对同一 kernel 保存 benchmark CSV 和反汇编摘要。读者要先手写预期结果，再运行项目观察输出。动作是：解释速度差来自地址流、SIMD lane、packing、load 或 reduction 哪一项。如果输出里没有 **average\_us** 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：CPU 后端优化必须同时有 oracle、反汇编、shape sweep 和降级边界；只有耗时表不够这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 [books/cpu-volume-3/labs/linux\\_cpu\\_inference/README.md](#)、相关 **include/lcqi** 头文件和一个 **src** 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交 CPU kernel evidence：命令、机器、编译器、输入 shape、backend、checksum、反汇编摘要和未验证 PMU 边界。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。

## 第 14 章实践：GPU 后端边界、copy 与同步成本

本章对应原书中的第三册“GPU 硬件基础”和“GPU 后端优化细讲”。不要把它当作概念复述，本章要把 device buffer、stream、event、kernel launch、copy、sync、fallback 放到同一个可运行项目里检查。贯穿代码仍然是 [books/cpu-volume-3/labs/linux\\_cpu\\_inference](#)；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `StepTrace`，核心命令或测试入口是 `lcqi_trace`。

**Listing 14.1** 第 14 章实践：GPU 后端边界、copy 与同步成本的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  <- DCMMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_trace / StepTrace
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：当前项目没有 GPU kernel，也要写清 adapter 合同。读者要先手写预期结果，再运行项目观察输出。动作是：把 CPU trace 字段扩展成 GPU StepTrace 需要哪些字段。如果输出里没有 `copy_time_ns` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：GPU 快不快不能只看 kernel；copy、sync、sampler、fallback 和 batch wait 必须进入端到端 trace 这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 [books/cpu-volume-3/labs/linux\\_cpu\\_inference/README.md](#)、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交 GPU adapter 设计：buffer owner、stream、event、H2D/D2H 字节、sync 点、fallback 和 oracle 对照。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。

## 第 15 章实践：工业 LLM 专项能力对照

本章对应原书中的第三册“工业 LLM 专项：投机解码、MoE、LoRA、多 GPU 与源码对照”。不要把它当作概念复述，本章要把 speculative decoding、MoE routing、LoRA adapter、multi-GPU、source reading 放到同一个可运行项目里检查。贯穿代码仍然是 [books/cpu-volume-3/labs/linux\\_cpu\\_inference](#)；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `bandwidth_limit`，核心命令或测试入口是 `lcqi_ledger`。

Listing 15.1 第 15 章实践：工业 LLM 专项能力对照的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
  <- DCMMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: lcqi_ledger / bandwidth_limit
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：选择一个专项，不写泛泛介绍。读者要先手写预期结果，再运行项目观察输出。动作是：把它接到现有 trace、memory、backend 或 serving 字段上。如果输出里没有 `tokens/s` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：工业专项不是功能清单；必须说明它购买什么能力、增加什么状态、破坏哪些假设这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 [books/cpu-volume-3/labs/linux\\_cpu\\_inference/README.md](#)、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交专项设计卡：动机、状态新增、trace 字段、回归测试、性能账本和不能验证的部分。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。



## 第 16 章实践：工程验收、回归门禁与最终项目

本章对应原书中的第三册“工程验收与回归体系”。不要把它当作概念复述，本章要把 reference、unit test、trace golden、benchmark gate、SLO gate、coverage、CI 放到同一个可运行项目里检查。贯穿代码仍然是 `books/cpu-volume-3/labs/linux_cpu_inference`；如果某个实验在当前机器上不能验证，报告必须写清降级原因，不能把源码存在当作实测证据。

### 一、对应原书问题：概念必须落到对象和命令

原书讲的是系统边界，本章实践要回答三个问题。第一，哪个代码对象承载这个概念；第二，哪个命令能产生可复查输出；第三，哪个坏输入或缺失字段能证明边界不是空话。这里的核心对象包括 `lcqi_tests`，核心命令或测试入口是 `ctest`。

Listing 16.1 第 16 章实践：工程验收、回归门禁与最终项目的最小命令入口

```
cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -<
↳ DCMAKE_BUILD_TYPE=Debug
cmake --build books/cpu-volume-3/build/lcqi-debug
ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure
# chapter focus: ctest / lcqi_tests
```

### 二、从特例推到规律：先抓一个能手算的切片

本章先看特例：把所有前面报告合成最终验收包。读者要先手写预期结果，再运行项目观察输出。动作是：运行 `ctest` 并说明每个 test 保护哪条合同。如果输出里没有 `100% tests passed` 或对应字段无法解释，就不要继续优化；先回到 reference、输入合同和 trace 字段。

从这个特例推出的规律是：实践卷最终目标是长期维护：结果正确、trace 稳定、性能可解释、服务指标可回归、失败边界不伪装这条规律要写进报告，不要只写“运行成功”。运行成功只说明某条路径没崩，解释清楚输入、输出、错误路径和不可验证边界，才说明学会了这个知识点。

### 三、实践任务：把观察固定成报告字段

任务一：定位源码。至少阅读 `books/cpu-volume-3/labs/linux_cpu_inference/README.md`、相关 `include/lcqi` 头文件和一个 `src` 实现文件。报告要写出函数入口、输入对象、输出对象和错误处理方式。

任务二：运行命令。保存构建命令、测试命令、核心输出和环境字段。若命令依赖 AVX2、perf、GPU、NUMA 或外部库，必须记录当前机器是否支持；不支持时把结论降级为“接口已设计，性能未验证”。

任务三：构造反例。围绕本章知识点设计一个坏输入、缺失字段或不满足前提的场景，说明系统应拒绝、降级、fallback 还是继续执行。只给正常路径不合格。

### 四、验收和递进大作业

验收标准：报告能从一个具体输入推到工程规则；命令可以在仓库中复跑；输出字段能对应到源码；失败路径说得清；未验证边界不伪装成结论。

递进大作业：提交最终项目：README、命令记录、报告索引、测试结果、benchmark 表、trace 摘要、风险清单和下一阶段计划。这个作业会被后续章节继续引用，命名、目录和字段不要随意变化。

## 代码走读：工具、Smoke、Benchmark 与回归测试

### 工具不是附属品

LCQI 的工程证据不只在 C++ 源码里。smoke 脚本决定真实模型怎样下载、校验和运行；benchmark 程序决定性能口径；trace、ledger 和 serving probe 决定报告字段；CTest 决定哪些行为能长期回归。读者必须把这些文件当作工程合同的一部分。

### Benchmark、Trace、Ledger 与 Serving Probe

`bench.cpp` 输出 int8 kernel benchmark。它服务性能观察，但必须和 max diff 对照一起看。

Listing 16.2 src/bench.cpp

```

1 #include <lcqi/int8_kernels.hpp>
2
3 #include <algorithm>
4 #include <cstdlib>
5 #include <iostream>
6 #include <stdexcept>
7 #include <string>
8 #include <vector>
9
10 namespace {
11
12 constexpr const char* LCQI_TARGET_ARCH_X86_64 = "x86_64";
13 constexpr const char* LCQI_TARGET_ARCH_I386 = "i386";
14 constexpr const char* LCQI_TARGET_ARCH_UNKNOWN = "unknown";
15 constexpr std::int32_t LCQI_DEFAULT_REPEAT = 200;
16 constexpr int LCQI_REPEAT_ARG_INDEX = 1;
17 constexpr std::int32_t LCQI_LARGE_CASE_REPEAT_DIVISOR = 4;
18 constexpr std::int32_t LCQI_MIN_LARGE_CASE_REPEAT = 1;
19 constexpr std::int32_t LCQI_DECODE_SHAPE_SIZE = 4096;
20
21 const char* target_arch() noexcept {
22     #if defined(__x86_64__) || defined(_M_X64)
23         return LCQI_TARGET_ARCH_X86_64;
24     #elif defined(__i386__) || defined(_M_IX86)
25         return LCQI_TARGET_ARCH_I386;
26     #else
27         return LCQI_TARGET_ARCH_UNKNOWN;
28     #endif
29 }
30
31 std::int32_t parse_repeat(int argc, char** argv) {
32     if (argc <= LCQI_REPEAT_ARG_INDEX) {
33         return LCQI_DEFAULT_REPEAT;
34     }
35     const std::int32_t repeat = static_cast<std::int32_t>(std::stoi(argv[↵
↵ LCQI_REPEAT_ARG_INDEX]));
36     if (repeat <= 0) {

```

```

37     throw std::runtime_error("repeat must be positive");
38 }
39 return repeat;
40 }
41
42 std::vector<lcqi::KernelBenchmarkCase> make_cases(std::int32_t repeat) {
43     return {
44         {128, 128, 16, repeat},
45         {256, 256, 16, repeat},
46         {512, 512, 16, repeat},
47         {513, 257, 16, repeat},
48         {1024, 4096, 16, std::max(LCQI_MIN_LARGE_CASE_REPEAT,
49                                     repeat / LCQI_LARGE_CASE_REPEAT_DIVISOR)},
50         {LCQI_DECODE_SHAPE_SIZE, LCQI_DECODE_SHAPE_SIZE, 16,
51          std::max(LCQI_MIN_LARGE_CASE_REPEAT, repeat / ↵
52 ↵ LCQI_LARGE_CASE_REPEAT_DIVISOR)}},
53     };
54 }
55 } // namespace
56
57 int main(int argc, char** argv) {
58     try {
59         const std::int32_t repeat = parse_repeat(argc, argv);
60         std::cout
61             << "target_arch,input_size,output_size,output_block,repeat,backend,↵
62 ↵ available,"
63             << "average_us,max_abs_diff,checksum\n";
64         for (const lcqi::KernelBenchmarkCase& benchmark_case : make_cases(↵
65 ↵ repeat)) {
66             const lcqi::KernelBenchmarkResult result =
67                 lcqi::benchmark_linear_i8_case(benchmark_case);
68             for (const lcqi::KernelBackendBenchmark& backend : result.backends)↵
69 ↵ {
70                 std::cout << target_arch() << ','
71                     << result.benchmark_case.input_size << ','
72                     << result.benchmark_case.output_size << ','
73                     << result.benchmark_case.output_block_size << ','
74                     << result.benchmark_case.repeat << ','
75                     << backend.name << ','
76                     << (backend.available ? 1 : 0) << ','
77                     << backend.average_us << ','
78                     << backend.max_abs_diff << ','
79                     << backend.checksum << '\n';
80             }
81         }
82         return EXIT_SUCCESS;
83     } catch (const std::exception& error) {
84         std::cerr << "error: " << error.what() << '\n';
85         return EXIT_FAILURE;
86     }
87 }

```



`trace.cpp` 输出 reference decoder checkpoint。它用于定位错误边界，不应该被简化成只打印最终 token。

Listing 16.3 `src/trace.cpp`

```

1  #include <lcqi/reference_decoder.hpp>
2  #include <lcqi/tokenizer.hpp>
3
4  #include <algorithm>
5  #include <array>
6  #include <cstdlib>
7  #include <exception>
8  #include <filesystem>
9  #include <iostream>
10 #include <span>
11 #include <sstream>
12 #include <string>
13 #include <string_view>
14
15 namespace {
16
17 constexpr std::array<std::int32_t, 3> LCQI_TRACE_TOKEN_IDS{1, 2, 3};
18 constexpr std::string_view LCQI_TRACE_PROMPT = "tok1 tok2 tok3";
19 constexpr std::int32_t LCQI_TRACE_METADATA_LAYER = -1;
20 constexpr std::int32_t LCQI_TRACE_BOS_ID = 0;
21 constexpr std::int32_t LCQI_TRACE_EOS_ID = 4;
22 constexpr std::size_t LCQI_TRACE_VALUE_LIMIT = 8;
23
24 std::filesystem::path tokenizer_fixture_path() {
25     return std::filesystem::path(__FILE__).parent_path().parent_path() /
26         "models" / "tiny_tokenizer.txt";
27 }
28
29 std::vector<std::int32_t> decoder_tokens_from_prompt(const lcqi::TokenizerModel&
    ↪ & tokenizer,
30                                                     std::string_view prompt) {
31     std::vector<std::int32_t> full_tokens = lcqi::encode_prompt(tokenizer,
    ↪ prompt);
32     if (full_tokens.size() != LCQI_TRACE_TOKEN_IDS.size() + 2 ||
33         full_tokens.front() != LCQI_TRACE_BOS_ID ||
34         full_tokens.back() != LCQI_TRACE_EOS_ID) {
35         throw std::runtime_error("unexpected tokenizer fixture ids");
36     }
37
38     std::vector<std::int32_t> decoder_tokens(
39         full_tokens.begin() + 1,
40         full_tokens.end() - 1);
41     if (!std::equal(decoder_tokens.begin(),
42                    decoder_tokens.end(),
43                    LCQI_TRACE_TOKEN_IDS.begin(),
44                    LCQI_TRACE_TOKEN_IDS.end())) {
45         throw std::runtime_error("unexpected decoder token ids");
46     }

```

```

47     return decoder_tokens;
48 }
49
50 std::string join_ints(std::span<const std::int32_t> values) {
51     std::ostringstream stream;
52     for (std::size_t index = 0; index < values.size(); ++index) {
53         if (index != 0) {
54             stream << 'x';
55         }
56         stream << values[index];
57     }
58     return stream.str();
59 }
60
61 std::string join_int_values(std::span<const std::int32_t> values) {
62     std::ostringstream stream;
63     for (std::size_t index = 0; index < values.size(); ++index) {
64         if (index != 0) {
65             stream << ';';
66         }
67         stream << values[index];
68     }
69     return stream.str();
70 }
71
72 std::int32_t checksum_ints(std::span<const std::int32_t> values) {
73     std::int32_t result = 0;
74     for (const std::int32_t value : values) {
75         result += value;
76     }
77     return result;
78 }
79
80 std::int32_t max_abs_ints(std::span<const std::int32_t> values) {
81     std::int32_t result = 0;
82     for (const std::int32_t value : values) {
83         result = std::max(result, std::abs(value));
84     }
85     return result;
86 }
87
88 std::string join_floats(std::span<const float> values) {
89     std::ostringstream stream;
90     const std::size_t count = std::min(values.size(), LCQI_TRACE_VALUE_LIMIT);
91     for (std::size_t index = 0; index < count; ++index) {
92         if (index != 0) {
93             stream << ';';
94         }
95         stream << values[index];
96     }
97     if (values.size() > LCQI_TRACE_VALUE_LIMIT) {
98         stream << "...";

```

```

99     }
100     return stream.str();
101 }
102
103 void print_trace_csv(const lcqi::ReferenceDecodeTraceResult& trace,
104                    const lcqi::TokenizerModel& tokenizer,
105                    std::span<const std::int32_t> full_tokens) {
106     std::cout << "name,token_position,layer_id,shape,stride,dtype,layout,↵
↵ checksum,max_abs,values\n";
107     std::cout << "prompt,0," << LCQI_TRACE_METADATA_LAYER
108         << ",scalar,scalar,utf8,text,0,0," << LCQI_TRACE_PROMPT << '\n';
109     std::cout << "tokenizer,0," << LCQI_TRACE_METADATA_LAYER
110         << ', ' << full_tokens.size() << ",1,i32,contiguous,"
111         << checksum_ints(full_tokens) << ', '
112         << max_abs_ints(full_tokens) << ', '
113         << join_int_values(full_tokens) << '\n';
114     std::cout << "tokenizer_contract,0," << LCQI_TRACE_METADATA_LAYER
115         << ",scalar,scalar,u64,fnv1a,"
116         << tokenizer_contract_hash(tokenizer) << ",0,"
117         << tokenizer.chat_template_hash << '\n';
118     for (const lcqi::ReferenceTensorCheckpoint& checkpoint : trace.checkpoints)↵
↵     {
119         std::cout << checkpoint.name << ', '
120             << checkpoint.token_position << ', '
121             << checkpoint.layer_id << ', '
122             << join_ints(checkpoint.shape) << ', '
123             << join_ints(checkpoint.stride) << ', '
124             << checkpoint.dtype << ', '
125             << checkpoint.layout << ', '
126             << checkpoint.checksum << ', '
127             << checkpoint.max_abs << ', '
128             << join_floats(checkpoint.values) << '\n';
129     }
130     std::cout << "sampler_argmax,"
131         << LCQI_TRACE_TOKEN_IDS.size() - 1 << ', '
132         << LCQI_TRACE_METADATA_LAYER
133         << ', ' << trace.result.logits.size() << ",1,f32,argmax,"
134         << "0,0,token=" << trace.result.predicted_token << '\n';
135     std::cout << "generated_token,"
136         << LCQI_TRACE_TOKEN_IDS.size() << ', '
137         << LCQI_TRACE_METADATA_LAYER
138         << ",scalar,scalar,i32,generated,"
139         << trace.result.predicted_token << ', '
140         << trace.result.predicted_token << ', '
141         << trace.result.predicted_token << '\n';
142 }
143
144 } // namespace
145
146 int main() {
147     try {
148         const lcqi::TokenizerModel tokenizer =

```

```

149         lcqi::load_tokenizer(tokenizer_fixture_path());
150         const std::vector<std::int32_t> full_tokens =
151             lcqi::encode_prompt(tokenizer, LCQI_TRACE_PROMPT);
152         const std::vector<std::int32_t> decoder_tokens =
153             decoder_tokens_from_prompt(tokenizer, LCQI_TRACE_PROMPT);
154         const lcqi::ReferenceDecoderModel model = lcqi::↵
↵ make_tiny_reference_decoder_model();
155         const lcqi::ReferenceDecodeTraceResult trace =
156             lcqi::run_reference_decode_with_trace(model, decoder_tokens);
157         print_trace_csv(trace, tokenizer, full_tokens);
158         return EXIT_SUCCESS;
159     } catch (const std::exception& error) {
160         std::cerr << "error: " << error.what() << '\n';
161         return EXIT_FAILURE;
162     }
163 }

```

`ledger.cpp` 输出 7B 规模账本。它是估算，不是实测吞吐；读者要区分 FLOPs、KV 容量、权重/KV 字节和 bandwidth-only tokens/s 上限。

Listing 16.4 src/ledger.cpp

```

1  #include <cstdlib>
2  #include <cstdint>
3  #include <exception>
4  #include <iomanip>
5  #include <iostream>
6  #include <string_view>
7
8  namespace {
9
10 constexpr double LCQI_BYTES_PER_GB = 1000.0 * 1000.0 * 1000.0;
11 constexpr double LCQI_BYTES_PER_MIB = 1024.0 * 1024.0;
12 constexpr double LCQI_FLOPS_PER_MAC = 2.0;
13 constexpr std::int32_t LCQI_SEVENB_LAYER_COUNT = 32;
14 constexpr std::int32_t LCQI_SEVENB_HIDDEN_SIZE = 4096;
15 constexpr std::int32_t LCQI_SEVENB_QUERY_HEADS = 32;
16 constexpr std::int32_t LCQI_SEVENB_KV_HEADS = 8;
17 constexpr std::int32_t LCQI_SEVENB_HEAD_DIM = 128;
18 constexpr std::int32_t LCQI_SEVENB_INTERMEDIATE_SIZE = 11008;
19 constexpr std::int32_t LCQI_SEVENB_CONTEXT_LENGTH = 4096;
20
21 struct DTypeLedger {
22     std::string_view name;
23     double weight_gb_per_token = 0.0;
24     double kv_element_bytes = 0.0;
25 };
26
27 constexpr DTypeLedger LCQI_DTYPE_LEDGERS[] = {
28     {"fp16", 14.0, 2.0},
29     {"int8", 7.0, 2.0},
30     {"int4", 3.7, 2.0},
31 };

```

```

32
33 constexpr double LCQI_BANDWIDTH_GB_PER_S[] = {
34     50.0,
35     100.0,
36     200.0,
37     600.0,
38     1000.0,
39     1600.0,
40 };
41
42 double linear_macs_per_layer() noexcept {
43     const double hidden = LCQI_SEVENB_HIDDEN_SIZE;
44     const double kv_width = LCQI_SEVENB_KV_HEADS * LCQI_SEVENB_HEAD_DIM;
45     const double qkv = hidden * (hidden + 2.0 * kv_width);
46     const double output_projection = hidden * hidden;
47     const double mlp = 3.0 * hidden * LCQI_SEVENB_INTERMEDIATE_SIZE;
48     return qkv + output_projection + mlp;
49 }
50
51 double attention_macs_per_layer() noexcept {
52     return 2.0 * LCQI_SEVENB_QUERY_HEADS * LCQI_SEVENB_CONTEXT_LENGTH *
53           LCQI_SEVENB_HEAD_DIM;
54 }
55
56 double kv_bytes_per_token_all_layers(double element_bytes) noexcept {
57     return static_cast<double>(LCQI_SEVENB_LAYER_COUNT) * 2.0 *
58           LCQI_SEVENB_KV_HEADS * LCQI_SEVENB_HEAD_DIM * element_bytes;
59 }
60
61 double kv_read_gb_per_decode_token(double element_bytes) noexcept {
62     const double bytes = static_cast<double>(LCQI_SEVENB_CONTEXT_LENGTH) *
63           kv_bytes_per_token_all_layers(element_bytes);
64     return bytes / LCQI_BYTES_PER_GB;
65 }
66
67 double kv_capacity_mib(double element_bytes, std::int32_t context_length) ↵
68     ↵ noexcept {
69     const double bytes = static_cast<double>(context_length) *
70           kv_bytes_per_token_all_layers(element_bytes);
71     return bytes / LCQI_BYTES_PER_MIB;
72 }
73
74 void print_model_rows() {
75     const double linear_macs = linear_macs_per_layer();
76     const double attention_macs = attention_macs_per_layer();
77     const double total_flops =
78         (linear_macs + attention_macs) * LCQI_SEVENB_LAYER_COUNT * ↵
79         ↵ LCQI_FLOPS_PER_MAC;
80     std::cout << "section,item,value,unit,notes\n";
81     std::cout << "model,layers," << LCQI_SEVENB_LAYER_COUNT << ",count,7B ↵
82         ↵ fixture\n";
83     std::cout << "model,hidden," << LCQI_SEVENB_HIDDEN_SIZE << ",elements,H\n";

```

```

81     std::cout << "model,query_heads," << LCQI_SEVENB_QUERY_HEADS << ",count,GQA" <<
    << query_heads << "\n";
82     std::cout << "model,kv_heads," << LCQI_SEVENB_KV_HEADS << ",count,GQA KV" <<
    << kv_heads << "\n";
83     std::cout << "model,head_dim," << LCQI_SEVENB_HEAD_DIM << ",elements,per" <<
    << head_dim << "\n";
84     std::cout << "compute,linear_macs_per_layer," << linear_macs << ",macs," <<
    << decode_token << "\n";
85     std::cout << "compute,attention_macs_per_layer," << attention_macs
86         << ",macs,context=4096\n";
87     std::cout << "compute,total_decode_flops_per_token," << total_flops
88         << ",flops,linear plus attention\n";
89 }
90
91 void print_kv_rows() {
92     for (const std::int32_t context : {1024, 2048, 4096, 8192}) {
93         std::cout << "kv,fp16_capacity_context_" << context << ', '
94             << kv_capacity_mib(2.0, context)
95             << ",MiB,single session\n";
96     }
97 }
98
99 void print_dtype_rows() {
100     for (const DTypeLedger& dtype : LCQI_DTYPE_LEDGERS) {
101         const double kv_read_gb = kv_read_gb_per_decode_token(dtype,
    << kv_element_bytes);
102         const double total_gb = dtype.weight_gb_per_token + kv_read_gb;
103         std::cout << "decode_bytes," << dtype.name << "_weight_gb,"
104             << dtype.weight_gb_per_token << ",GB/token,weight stream\n";
105         std::cout << "decode_bytes," << dtype.name << "_kv_read_gb,"
106             << kv_read_gb << ",GB/token,context=4096\n";
107         std::cout << "decode_bytes," << dtype.name << "_total_gb,"
108             << total_gb << ",GB/token,weight plus KV\n";
109         for (const double bandwidth : LCQI_BANDWIDTH_GB_PER_S) {
110             std::cout << "bandwidth_limit," << dtype.name << "_at_"
111                 << static_cast<std::int32_t>(bandwidth) << "_gbps,"
112                 << bandwidth / total_gb
113                 << ",tokens/s,bandwidth-only upper bound\n";
114         }
115     }
116 }
117
118 } // namespace
119
120 int main() {
121     try {
122         std::cout << std::fixed << std::setprecision(3);
123         print_model_rows();
124         print_kv_rows();
125         print_dtype_rows();
126         return EXIT_SUCCESS;
127     } catch (const std::exception& error) {

```

```

128         std::cerr << "error: " << error.what() << '\n';
129         return EXIT_FAILURE;
130     }
131 }

```

`serving_probe.cpp` 固定短请求、RAG 请求、取消请求和超长请求。它把 TTFT、TPOT、拒绝和取消回收变成可观察字段。

**Listing 16.5** `src/serving_probe.cpp`

```

1  #include <algorithm>
2  #include <cstdlib>
3  #include <cstdint>
4  #include <exception>
5  #include <cmath>
6  #include <iomanip>
7  #include <iostream>
8  #include <string_view>
9  #include <vector>
10
11 namespace {
12
13 constexpr double LCQI_BYTES_PER_MIB = 1024.0 * 1024.0;
14 constexpr std::int32_t LCQI_PROBE_LAYER_COUNT = 32;
15 constexpr std::int32_t LCQI_PROBE_KV_HEADS = 8;
16 constexpr std::int32_t LCQI_PROBE_HEAD_DIM = 128;
17 constexpr double LCQI_PROBE_KV_ELEMENT_BYTES = 2.0;
18 constexpr double LCQI_PROBE_KV_BUDGET_MIB = 512.0;
19 constexpr double LCQI_PROBE_WORKSPACE_MIB = 64.0;
20 constexpr double LCQI_PROBE_ARENA_MIB = 32.0;
21 constexpr double LCQI_PROBE_TRACE_MIB = 4.0;
22 constexpr double LCQI_PROBE_QUEUE_WAIT_PER_REQUEST_MS = 8.0;
23 constexpr double LCQI_PROBE_PREFILL_MS_PER_TOKEN = 0.08;
24 constexpr double LCQI_PROBE_DECODE_STEP_MS = 14.0;
25 constexpr double LCQI_PROBE_CANCEL_RECLAIM_MS = 2.0;
26 constexpr std::size_t LCQI_PERCENTILE_MIN_INDEX = 1;
27
28 struct ProbeRequest {
29     std::string_view request_id;
30     std::int32_t prompt_tokens = 0;
31     std::int32_t max_new_tokens = 0;
32     bool cancel_after_prefill = false;
33 };
34
35 struct ProbeResult {
36     std::string_view request_id;
37     std::int32_t accepted = 0;
38     std::string_view rejection_reason;
39     std::int32_t prompt_tokens = 0;
40     std::int32_t reserved_tokens = 0;
41     double kv_reserved_mib = 0.0;
42     double queue_wait_ms = 0.0;
43     double prefill_ms = 0.0;

```

```

44     double ttft_ms = 0.0;
45     double decode_step_p95_ms = 0.0;
46     double tpot_ms = 0.0;
47     std::int32_t completion_tokens = 0;
48     std::string_view finish_reason;
49     double cancel_reclaim_ms = 0.0;
50 };
51
52 const std::vector<ProbeRequest> LCQI_PROBE_REQUESTS{
53     {"req_short", 32, 4, false},
54     {"req_rag", 512, 8, false},
55     {"req_cancel", 128, 16, true},
56     {"req_too_long", 4096, 512, false},
57 };
58
59 double kv_mib_for_tokens(std::int32_t tokens) noexcept {
60     const double bytes = static_cast<double>(tokens) * LCQI_PROBE_LAYER_COUNT *
        ↪ 2.0 *
61         LCQI_PROBE_KV_HEADS * LCQI_PROBE_HEAD_DIM *
62         LCQI_PROBE_KV_ELEMENT_BYTES;
63     return bytes / LCQI_BYTES_PER_MIB;
64 }
65
66 double percentile(std::vector<double> values, double fraction) {
67     if (values.empty()) {
68         return 0.0;
69     }
70     std::sort(values.begin(), values.end());
71     const std::size_t max_index = values.size() - LCQI_PERCENTILE_MIN_INDEX;
72     const std::size_t index =
73         std::min(max_index,
74             static_cast<std::size_t>(
75                 std::ceil(fraction * static_cast<double>(max_index))));
76     return values[index];
77 }
78
79 std::vector<ProbeResult> run_probe() {
80     std::vector<ProbeResult> results;
81     double reserved_kv_mib = 0.0;
82     std::int32_t accepted_before = 0;
83     for (const ProbeRequest& request : LCQI_PROBE_REQUESTS) {
84         ProbeResult result;
85         result.request_id = request.request_id;
86         result.prompt_tokens = request.prompt_tokens;
87         result.reserved_tokens = request.prompt_tokens + request.max_new_tokens
        ↪ ;
88         result.kv_reserved_mib = kv_mib_for_tokens(result.reserved_tokens);
89         const double static_session_mib =
90             result.kv_reserved_mib + LCQI_PROBE_WORKSPACE_MIB +
91             LCQI_PROBE_ARENA_MIB + LCQI_PROBE_TRACE_MIB;
92         if (reserved_kv_mib + result.kv_reserved_mib > LCQI_PROBE_KV_BUDGET_MIB
        ↪ ) {

```



```

93         result.accepted = 0;
94         result.rejection_reason = "memory_budget_exceeded";
95         result.finish_reason = "rejected";
96         results.push_back(result);
97         continue;
98     }
99
100     reserved_kv_mib += result.kv_reserved_mib;
101     result.accepted = 1;
102     result.rejection_reason = "none";
103     result.queue_wait_ms = static_cast<double>(accepted_before) *
104                             LCQI_PROBE_QUEUE_WAIT_PER_REQUEST_MS;
105     result.prefill_ms = static_cast<double>(request.prompt_tokens) *
106                             LCQI_PROBE_PREFILL_MS_PER_TOKEN;
107     result.ttft_ms = result.queue_wait_ms + result.prefill_ms + ↵
↵ LCQI_PROBE_DECODE_STEP_MS;
108     result.decode_step_p95_ms = LCQI_PROBE_DECODE_STEP_MS +
109                                 static_session_mib / 512.0;
110     result.tpot_ms = result.decode_step_p95_ms;
111     if (request.cancel_after_prefill) {
112         result.completion_tokens = 0;
113         result.finish_reason = "cancelled";
114         result.cancel_reclaim_ms = LCQI_PROBE_CANCEL_RECLAIM_MS;
115         reserved_kv_mib -= result.kv_reserved_mib;
116     } else {
117         result.completion_tokens = request.max_new_tokens;
118         result.finish_reason = "max_tokens";
119     }
120     ++accepted_before;
121     results.push_back(result);
122 }
123 return results;
124 }
125
126 void print_request_rows(const std::vector<ProbeResult>& results) {
127     std::cout << "row_type,request_id,accepted,rejection_reason,prompt_tokens,↵
↵ reserved_tokens,"
128                 "kv_reserved_mib,queue_wait_ms,prefill_ms,ttft_ms,↵
↵ decode_step_p95_ms,"
129                 "tpot_ms,completion_tokens,finish_reason,cancel_reclaim_ms,↵
↵ metric_name,"
130                 "metric_value\n";
131     for (const ProbeResult& result : results) {
132         std::cout << "request," << result.request_id << ',' <<
133                     << result.accepted << ',' <<
134                     << result.rejection_reason << ',' <<
135                     << result.prompt_tokens << ',' <<
136                     << result.reserved_tokens << ',' <<
137                     << result.kv_reserved_mib << ',' <<
138                     << result.queue_wait_ms << ',' <<
139                     << result.prefill_ms << ',' <<
140                     << result.ttft_ms << ',' <<

```

```

141         << result.decode_step_p95_ms << ', '
142         << result.tpot_ms << ', '
143         << result.completion_tokens << ', '
144         << result.finish_reason << ', '
145         << result.cancel_reclaim_ms << ",,\n";
146     }
147 }
148
149 void print_summary_row(const std::vector<ProbeResult>& results) {
150     std::vector<double> ttft_values;
151     std::vector<double> queue_values;
152     double accepted = 0.0;
153     double rejected = 0.0;
154     double cancelled = 0.0;
155     double reserved_peak_mib = 0.0;
156     double running_reserved_mib = 0.0;
157     for (const ProbeResult& result : results) {
158         if (result.accepted == 0) {
159             rejected += 1.0;
160             continue;
161         }
162         accepted += 1.0;
163         ttft_values.push_back(result.ttft_ms);
164         queue_values.push_back(result.queue_wait_ms);
165         running_reserved_mib += result.kv_reserved_mib;
166         reserved_peak_mib = std::max(reserved_peak_mib, running_reserved_mib);
167         if (result.finish_reason == "cancelled") {
168             cancelled += 1.0;
169             running_reserved_mib -= result.kv_reserved_mib;
170         }
171     }
172     const double total = accepted + rejected;
173     const double rejection_rate = total == 0.0 ? 0.0 : rejected / total;
174     std::cout << "summary,all,0,summary,0,0,0,0,0,0,0,0,summary,0,"
175               << "accepted_count," << accepted << '\n';
176     std::cout << "summary,all,0,summary,0,0,0,0,0,0,0,0,summary,"
177               << LCQI_PROBE_CANCEL_RECLAIM_MS << "cancel_count," << cancelled <<
178               << '\n';
179     std::cout << "summary,all,0,summary,0,0,0,0,0,0,0,0,summary,0,"
180               << "rejected_count," << rejected << '\n';
181     std::cout << "summary,all,0,summary,0,0,0,0,0,0,0,0,summary,0,"
182               << "rejection_rate," << rejection_rate << '\n';
183     std::cout << "summary,all,0,summary,0,0,0,0,0,0,0,0,summary,0,"
184               << "kv_reserved_peak_mib," << reserved_peak_mib << '\n';
185     std::cout << "summary,all,0,summary,0,0,0,0,0,0,0,0,summary,0,"
186               << "queue_wait_p95_ms," << percentile(queue_values, 0.95) << '\n' <<
187               << ;
188     std::cout << "summary,all,0,summary,0,0,0,0,0,0,0,0,summary,0,"
189               << "ttft_p95_ms," << percentile(ttft_values, 0.95) << '\n';
190 } // namespace

```

```

191
192 int main() {
193     try {
194         std::cout << std::fixed << std::setprecision(3);
195         const std::vector<ProbeResult> results = run_probe();
196         print_request_rows(results);
197         print_summary_row(results);
198         return EXIT_SUCCESS;
199     } catch (const std::exception& error) {
200         std::cerr << "error: " << error.what() << '\n';
201         return EXIT_FAILURE;
202     }
203 }

```

`onednn_bench.cpp` 是可选外部库对标入口。它只在环境安装 oneDNN 时构建，不能把不可用平台写成已验证性能。

**Listing 16.6** `src/onednn_bench.cpp`

```

1 #include <dnnl.hpp>
2
3 #include <chrono>
4 #include <cstdint>
5 #include <cstdlib>
6 #include <iostream>
7 #include <numeric>
8 #include <stdexcept>
9 #include <string>
10 #include <unordered_map>
11 #include <vector>
12
13 namespace {
14
15 constexpr std::int64_t LCQI_ONEDNN_BATCH = 1;
16 constexpr std::int64_t LCQI_ONEDNN_K = 4096;
17 constexpr std::int64_t LCQI_ONEDNN_O = 4096;
18 constexpr std::int32_t LCQI_ONEDNN_DEFAULT_REPEAT = 20;
19 constexpr std::int32_t LCQI_ONEDNN_WARMUP = 3;
20 constexpr int LCQI_REPEAT_ARG_INDEX = 1;
21 constexpr double LCQI_SECONDS_TO_MICROSECONDS = 1000000.0;
22 constexpr float LCQI_INPUT_SCALE = 0.125F;
23 constexpr std::int32_t LCQI_INPUT_MODULUS = 17;
24 constexpr std::int32_t LCQI_INPUT_CENTER = 8;
25 constexpr std::int32_t LCQI_WEIGHT_MODULUS = 23;
26 constexpr std::int32_t LCQI_WEIGHT_CENTER = 11;
27 constexpr float LCQI_WEIGHT_SCALE = 0.03125F;
28
29 std::int32_t parse_repeat(int argc, char** argv) {
30     if (argc <= LCQI_REPEAT_ARG_INDEX) {
31         return LCQI_ONEDNN_DEFAULT_REPEAT;
32     }
33     const std::int32_t repeat = static_cast<std::int32_t>(std::stoi(argv[↵
↵ LCQI_REPEAT_ARG_INDEX]));

```

```

34     if (repeat <= 0) {
35         throw std::runtime_error("repeat must be positive");
36     }
37     return repeat;
38 }
39
40 std::vector<float> make_input() {
41     std::vector<float> input(static_cast<std::size_t>(LCQI_ONEDNN_K));
42     for (std::int64_t index = 0; index < LCQI_ONEDNN_K; ++index) {
43         input[static_cast<std::size_t>(index)] =
44             static_cast<float>((index % LCQI_INPUT_MODULUS) - LCQI_INPUT_CENTER) *
45             LCQI_INPUT_SCALE;
46     }
47     return input;
48 }
49
50 std::vector<float> make_weights_kn() {
51     std::vector<float> weights(static_cast<std::size_t>(LCQI_ONEDNN_K * LCQI_ONEDNN_O));
52     for (std::int64_t k = 0; k < LCQI_ONEDNN_K; ++k) {
53         for (std::int64_t out = 0; out < LCQI_ONEDNN_O; ++out) {
54             const std::int64_t row_major_out_k = out * LCQI_ONEDNN_K + k;
55             const float value =
56                 static_cast<float>((row_major_out_k % LCQI_WEIGHT_MODULUS) -
57                                     LCQI_WEIGHT_CENTER) *
58                 LCQI_WEIGHT_SCALE;
59             weights[static_cast<std::size_t>(k * LCQI_ONEDNN_O + out)] = value;
60         }
61     }
62     return weights;
63 }
64
65 float checksum(const std::vector<float>& values) {
66     return std::accumulate(values.begin(), values.end(), 0.0F);
67 }
68
69 } // namespace
70
71 int main(int argc, char** argv) {
72     try {
73         const std::int32_t repeat = parse_repeat(argc, argv);
74         std::vector<float> input = make_input();
75         std::vector<float> weights = make_weights_kn();
76         std::vector<float> output(static_cast<std::size_t>(LCQI_ONEDNN_BATCH * LCQI_ONEDNN_O),
77                                0.0F);
78
79         dnnl::engine engine(dnnl::engine::kind::cpu, 0);
80         dnnl::stream stream(engine);
81
82         const dnnl::memory::dims source_dims = {LCQI_ONEDNN_BATCH, LCQI_ONEDNN_K, LCQI_ONEDNN_O};

```

```

↪ LCQI_ONEDNN_K};
83     const dnnl::memory::dims weight_dims = {LCQI_ONEDNN_K, LCQI_ONEDNN_0};
84     const dnnl::memory::dims output_dims = {LCQI_ONEDNN_BATCH, ↪
↪ LCQI_ONEDNN_0};

85
86     const dnnl::memory::desc source_desc(
87         source_dims, dnnl::memory::data_type::f32, dnnl::memory::format_tag↪
↪ ::ab);
88     const dnnl::memory::desc weight_desc(
89         weight_dims, dnnl::memory::data_type::f32, dnnl::memory::format_tag↪
↪ ::ab);
90     const dnnl::memory::desc output_desc(
91         output_dims, dnnl::memory::data_type::f32, dnnl::memory::format_tag↪
↪ ::ab);
92
93     dnnl::memory source_memory(source_desc, engine, input.data());
94     dnnl::memory weight_memory(weight_desc, engine, weights.data());
95     dnnl::memory output_memory(output_desc, engine, output.data());
96
97     dnnl::matmul::primitive_desc matmul_desc(engine, source_desc, ↪
↪ weight_desc, output_desc);
98     dnnl::matmul matmul(matmul_desc);
99     const std::unordered_map<int, dnnl::memory> arguments = {
100         {DNNL_ARG_SRC, source_memory},
101         {DNNL_ARG_WEIGHTS, weight_memory},
102         {DNNL_ARG_DST, output_memory},
103     };
104
105     for (std::int32_t index = 0; index < LCQI_ONEDNN_WARMUP; ++index) {
106         matmul.execute(stream, arguments);
107         stream.wait();
108     }
109
110     const auto begin = std::chrono::steady_clock::now();
111     for (std::int32_t index = 0; index < repeat; ++index) {
112         matmul.execute(stream, arguments);
113         stream.wait();
114     }
115     const auto end = std::chrono::steady_clock::now();
116     const std::chrono::duration<double> elapsed = end - begin;
117     const double average_us =
118         elapsed.count() * LCQI_SECONDS_TO_MICROSECONDS / static_cast<double>↪
↪ >(repeat);
119
120     std::cout << "library,input_size,output_size,dtype,repeat,average_us,↪
↪ checksum\n";
121     std::cout << "onednn," << LCQI_ONEDNN_K << ',' << LCQI_ONEDNN_0
122         << ",f32," << repeat << ',' << average_us << ','
123         << checksum(output) << '\n';
124     return EXIT_SUCCESS;
125 } catch (const std::exception& error) {
126     std::cerr << "error: " << error.what() << '\n';

```

```

127         return EXIT_FAILURE;
128     }
129 }

```

## 外部模型 Smoke 与对比脚本

`run_smollm2_small_smoke.py` 下载并校验 SmolLM2-135M-Instruct 的 GGUF 量化文件，构建固定 commit 的 `llama-simple-chat`，在 CPU 上运行一次真实开源 small 模型生成，并输出可复查报告。它是外部真实模型 smoke，不替代 LCQI tiny golden trace。

**Listing 16.7** `tools/run_smollm2_small_smoke.py`

```

1  #!/usr/bin/env python3
2  from __future__ import annotations
3
4  import argparse
5  import datetime as dt
6  import hashlib
7  import os
8  import re
9  import shlex
10 import shutil
11 import subprocess
12 import sys
13 from dataclasses import dataclass
14 from pathlib import Path
15
16
17 MODEL_SOURCE_ID = "HuggingFaceTB/SmolLM2-135M-Instruct"
18 MODEL_GGUF_REPO_ID = "bartowski/SmolLM2-135M-Instruct-GGUF"
19 MODEL_FILE_NAME = "SmolLM2-135M-Instruct-Q4_K_M.gguf"
20 MODEL_URL = (
21     "https://huggingface.co/bartowski/SmolLM2-135M-Instruct-GGUF"
22     f"/resolve/main/{MODEL_FILE_NAME}"
23 )
24 MODEL_EXPECTED_BYTES = 105454432
25 MODEL_EXPECTED_SHA256 = (
26     "2e8040ceae7815abe0dcb3540b9995eaa1fa0d2ca9e797d0a635ae4433c68c2d"
27 )
28 LLAMA_CPP_REPO_URL = "https://github.com/ggml-org/llama.cpp.git"
29 LLAMA_CPP_COMMIT = "b3fed31b99f9bd37725833674252bccb429bb183"
30 LLAMA_TARGET = "llama-simple-chat"
31 DEFAULT_PROMPT = "What is 2+3? Answer in one short sentence."
32 DEFAULT_CONTEXT_TOKENS = 512
33 DEFAULT_TIMEOUT_SECONDS = 180
34 DEFAULT_BUILD_JOBS = 2
35
36
37 @dataclass(frozen=True)
38 class CommandResult:
39     args: list[str]
40     returncode: int

```

```

41     stdout: str
42     stderr: str
43
44
45 def script_path() -> Path:
46     return Path(__file__).resolve()
47
48
49 def volume_dir() -> Path:
50     return script_path().parents[1]
51
52
53 def repo_dir() -> Path:
54     return script_path().parents[3]
55
56
57 def default_cache_dir() -> Path:
58     explicit_cache = os.environ.get("LCQI_SMOLLM2_CACHE_DIR")
59     if explicit_cache:
60         return Path(explicit_cache).expanduser()
61     xdg_cache = os.environ.get("XDG_CACHE_HOME")
62     cache_root = Path(xdg_cache).expanduser() if xdg_cache else Path.home() / "↵
↵ .cache"
63     return cache_root / "lcqi-smolllm2-small-smoke"
64
65
66 def default_report_path() -> Path:
67     return volume_dir() / "results" / "lcqi-smolllm2-small-model-smoke.txt"
68
69
70 def command_line(args: list[str]) -> str:
71     return " ".join(shlex.quote(arg) for arg in args)
72
73
74 def require_tool(name: str) -> str:
75     path = shutil.which(name)
76     if path is None:
77         raise RuntimeError(f"required tool not found in PATH: {name}")
78     return path
79
80
81 def run_command(
82     args: list[str],
83     *,
84     cwd: Path | None = None,
85     input_text: str | None = None,
86     timeout_seconds: int | None = None,
87 ) -> CommandResult:
88     try:
89         completed = subprocess.run(
90             args,
91             cwd=str(cwd) if cwd else None,

```

```

92         input=input_text,
93         text=True,
94         capture_output=True,
95         timeout=timeout_seconds,
96         check=False,
97     )
98 except subprocess.TimeoutExpired as exc:
99     stdout = exc.stdout if isinstance(exc.stdout, str) else ""
100     stderr = exc.stderr if isinstance(exc.stderr, str) else ""
101     return CommandResult(args=args, returncode=124, stdout=stdout, stderr=↵
↵ stderr)
102
103 return CommandResult(
104     args=args,
105     returncode=completed.returncode,
106     stdout=completed.stdout,
107     stderr=completed.stderr,
108 )
109
110
111 def ensure_success(result: CommandResult, action: str) -> None:
112     if result.returncode == 0:
113         return
114     raise RuntimeError(
115         f"{action} failed with exit code {result.returncode}\n"
116         f"command: {command_line(result.args)}\n"
117         f"stdout tail:\n{tail(result.stdout)}\n"
118         f"stderr tail:\n{tail(result.stderr)}"
119     )
120
121
122 def tail(text: str, max_chars: int = 4000) -> str:
123     if len(text) <= max_chars:
124         return text
125     return text[-max_chars:]
126
127
128 def sha256_file(path: Path) -> str:
129     digest = hashlib.sha256()
130     with path.open("rb") as handle:
131         for block in iter(lambda: handle.read(1024 * 1024), b''):
132             digest.update(block)
133     return digest.hexdigest()
134
135
136 def verify_model_file(path: Path) -> tuple[bool, str, int]:
137     if not path.exists():
138         return False, "", 0
139     actual_bytes = path.stat().st_size
140     actual_sha = sha256_file(path)
141     return (
142         actual_bytes == MODEL_EXPECTED_BYTES

```



```

143         and actual_sha == MODEL_EXPECTED_SHA256,
144         actual_sha,
145         actual_bytes,
146     )
147
148
149 def download_model(model_path: Path, *, force_download: bool, no_download: bool ↵
    ↵ ) -> None:
150     model_path.parent.mkdir(parents=True, exist_ok=True)
151
152     if force_download and model_path.exists():
153         model_path.unlink()
154
155     is_valid, actual_sha, actual_bytes = verify_model_file(model_path)
156     if is_valid:
157         print(f"[lcqi-smoke] model cache hit: {model_path}")
158         return
159
160     if model_path.exists():
161         print(
162             "[lcqi-smoke] cached model hash mismatch, redownloading: "
163             f"bytes={actual_bytes}, sha256={actual_sha}"
164         )
165         model_path.unlink()
166
167     if no_download:
168         raise RuntimeError(
169             f"model file is missing or invalid and --no-download was set: {↵
    ↵ model_path}"
170         )
171
172     require_tool("curl")
173     part_path = model_path.with_suffix(model_path.suffix + ".part")
174     if part_path.exists():
175         part_path.unlink()
176
177     print(f"[lcqi-smoke] downloading small GGUF model: {MODEL_GGUF_REPO_ID}/{↵
    ↵ MODEL_FILE_NAME}")
178     result = run_command(
179         [
180             "curl",
181             "-L",
182             "--fail",
183             "--retry",
184             "3",
185             "-o",
186             str(part_path),
187             MODEL_URL,
188         ]
189     )
190     ensure_success(result, "download SmolLM2 GGUF")
191

```

```

192     downloaded_ok, downloaded_sha, downloaded_bytes = verify_model_file(↵
↵ part_path)
193     if not downloaded_ok:
194         raise RuntimeError(
195             "downloaded model did not match expected identity: "
196             f"bytes={downloaded_bytes}, sha256={downloaded_sha}"
197         )
198     part_path.replace(model_path)
199     print(f"[lcqi-smoke] model verified: sha256={MODEL_EXPECTED_SHA256}")
200
201
202 def ensure_llama_simple_chat(cache_dir: Path, *, jobs: int, skip_build: bool) ↵
↵ -> Path:
203     cached_binary = cache_dir / "llama.cpp" / "build" / "bin" / LLAMA_TARGET
204     source_dir = cache_dir / "llama.cpp"
205     build_dir = source_dir / "build"
206     if cached_binary.exists() and os.access(cached_binary, os.X_OK) and ↵
↵ is_pinned_llama_checkout(source_dir):
207         print(f"[lcqi-smoke] llama.cpp binary cache hit: {cached_binary}")
208         return cached_binary
209
210     if skip_build:
211         raise RuntimeError(f"llama.cpp binary is missing and --skip-build was ↵
↵ set: {cached_binary}")
212
213     require_tool("git")
214     require_tool("cmake")
215
216     cache_dir.mkdir(parents=True, exist_ok=True)
217
218     if not (source_dir / ".git").exists():
219         print(f"[lcqi-smoke] cloning llama.cpp into cache: {source_dir}")
220         result = run_command(
221             [
222                 "git",
223                 "clone",
224                 "--filter=blob:none",
225                 "--no-checkout",
226                 LLAMA_CPP_REPO_URL,
227                 str(source_dir),
228             ]
229         )
230         ensure_success(result, "clone llama.cpp")
231
232     print(f"[lcqi-smoke] checking out pinned llama.cpp commit: {↵
↵ LLAMA_CPP_COMMIT}")
233     fetch = run_command(
234         ["git", "-C", str(source_dir), "fetch", "--depth", "1", "origin", ↵
↵ LLAMA_CPP_COMMIT]
235     )
236     ensure_success(fetch, "fetch pinned llama.cpp commit")
237     checkout = run_command(["git", "-C", str(source_dir), "checkout", "--detach↵

```

```

↪ ", LLAMA_CPP_COMMIT])
238 ensure_success(checkout, "checkout pinned llama.cpp commit")
239
240 print(f"[lcqi-smoke] configuring llama.cpp build: {build_dir}")
241 configure = run_command(
242     [
243         "cmake",
244         "-S",
245         str(source_dir),
246         "-B",
247         str(build_dir),
248         "-DCMAKE_BUILD_TYPE=Release",
249         "-DLLAMA_BUILD_TESTS=OFF",
250         "-DLLAMA_BUILD_EXAMPLES=ON",
251         "-DLLAMA_BUILD_SERVER=OFF",
252         "-DGGML_NATIVE=OFF",
253     ]
254 )
255 ensure_success(configure, "configure llama.cpp")
256
257 print(f"[lcqi-smoke] building {LLAMA_TARGET} with -j{jobs}")
258 build = run_command(
259     ["cmake", "--build", str(build_dir), "--target", LLAMA_TARGET, "-j", ↪
↪ str(jobs)]
260 )
261 ensure_success(build, f"build {LLAMA_TARGET}")
262
263 if not cached_binary.exists():
264     raise RuntimeError(f"expected binary was not produced: {cached_binary}"↪
↪ )
265 return cached_binary
266
267
268 def is_pinned_llama_checkout(source_dir: Path) -> bool:
269     if not (source_dir / ".git").exists():
270         return False
271     result = run_command(["git", "-C", str(source_dir), "rev-parse", "HEAD"])
272     if result.returncode != 0:
273         return False
274     return result.stdout.strip() == LLAMA_CPP_COMMIT
275
276
277 def strip_ansi(text: str) -> str:
278     return re.sub(r"\x1b\[([0-?]*[ -/]*[@-~])", "", text)
279
280
281 def extract_response(stdout: str) -> str:
282     cleaned = strip_ansi(stdout).replace("\r", "")
283     chunks = [chunk.strip() for chunk in re.split(r"(?:^|\n)> ?", cleaned) if ↪
↪ chunk.strip()]
284     if not chunks:
285         return ""

```

```

286     return chunks[0]
287
288
289 def run_model_smoke(
290     llama_binary: Path,
291     model_path: Path,
292     *,
293     prompt: str,
294     context_tokens: int,
295     timeout_seconds: int,
296 ) -> tuple[CommandResult, str]:
297     command = [
298         str(llama_binary),
299         "-m",
300         str(model_path),
301         "-ngl",
302         "0",
303         "-c",
304         str(context_tokens),
305     ]
306     print("[lcqi-smoke] running real small model generation")
307     result = run_command(
308         command,
309         input_text=prompt.rstrip("\n") + "\n\n",
310         timeout_seconds=timeout_seconds,
311     )
312     ensure_success(result, "run small model smoke")
313
314     response = extract_response(result.stdout)
315     if not response:
316         raise RuntimeError(
317             "small model smoke exited successfully but produced no assistant ←
↵ text\n"
318             f"stdout:\n{result.stdout}\n"
319             f"stderr:\n{result.stderr}"
320         )
321     return result, response
322
323
324 def current_repo_commit() -> str:
325     result = run_command(["git", "-C", str(repo_dir()), "rev-parse", "--short",
↵ "HEAD"])
326     if result.returncode != 0:
327         return "unknown"
328     return result.stdout.strip()
329
330
331 def current_uname() -> str:
332     result = run_command(["uname", "-a"])
333     if result.returncode != 0:
334         return "unknown"
335     return result.stdout.strip()

```

```

336
337
338 def write_report(
339     report_path: Path,
340     *,
341     cache_dir: Path,
342     llama_binary: Path,
343     model_path: Path,
344     prompt: str,
345     result: CommandResult,
346     response: str,
347 ) -> None:
348     report_path.parent.mkdir(parents=True, exist_ok=True)
349     clean_stdout = strip_ansi(result.stdout).strip()
350     report = "\n".join(
351         [
352             "LCQI SmolLM2 Small Model Smoke",
353             "",
354             f"timestamp_utc={dt.datetime.now(dt.UTC).isoformat()}",
355             f"repo_commit={current_repo_commit()}",
356             f"host={current_uname()}",
357             f"python={sys.version.split()[0]}",
358             f"cache_dir={cache_dir}",
359             "",
360             f"model_source_id={MODEL_SOURCE_ID}",
361             f"model_gguf_repo_id={MODEL_GGUF_REPO_ID}",
362             f"model_file={MODEL_FILE_NAME}",
363             f"model_url={MODEL_URL}",
364             f"model_expected_bytes={MODEL_EXPECTED_BYTES}",
365             f"model_expected_sha256={MODEL_EXPECTED_SHA256}",
366             f"model_path={model_path}",
367             "",
368             f"llama_cpp_repo={LLAMA_CPP_REPO_URL}",
369             f"llama_cpp_commit={LLAMA_CPP_COMMIT}",
370             f"llama_target={LLAMA_TARGET}",
371             f"llama_binary={llama_binary}",
372             "",
373             f"prompt={prompt}",
374             f"command={command_line(result.args)}",
375             f"exit_code={result.returncode}",
376             "",
377             "validation=PASS: model hash matched, llama-simple-chat exited 0, "
378             "assistant response was non-empty",
379             "",
380             "assistant_response:",
381             response,
382             "",
383             "stdout_clean_tail:",
384             tail(clean_stdout, 2000),
385             "",
386             "stderr_tail:",
387             tail(result.stderr.strip(), 2000),

```

```

388         "",
389     ]
390 )
391 report_path.write_text(report, encoding="utf-8")
392
393
394 def parse_args() -> argparse.Namespace:
395     parser = argparse.ArgumentParser(
396         description=(
397             "Run a real local small open-source model smoke for CPU Volume 3. "
398             "The script caches llama.cpp and the GGUF model outside the Git ↵
↵ repository."
399         )
400     )
401     parser.add_argument("--cache-dir", type=Path, default=default_cache_dir())
402     parser.add_argument("--model-path", type=Path, default=None)
403     parser.add_argument("--llama-bin", type=Path, default=None)
404     parser.add_argument("--report", type=Path, default=default_report_path())
405     parser.add_argument("--prompt", default=DEFAULT_PROMPT)
406     parser.add_argument("--ctx", type=int, default=DEFAULT_CONTEXT_TOKENS)
407     parser.add_argument("--timeout", type=int, default=DEFAULT_TIMEOUT_SECONDS)
408     parser.add_argument("--jobs", type=int, default=DEFAULT_BUILD_JOBS)
409     parser.add_argument("--force-download", action="store_true")
410     parser.add_argument("--no-download", action="store_true")
411     parser.add_argument("--skip-build", action="store_true")
412     return parser.parse_args()
413
414
415 def main() -> int:
416     args = parse_args()
417     cache_dir = args.cache_dir.expanduser().resolve()
418     model_path = (
419         args.model_path.expanduser().resolve()
420         if args.model_path
421         else cache_dir / "models" / MODEL_FILE_NAME
422     )
423     report_path = args.report.expanduser().resolve()
424
425     try:
426         download_model(
427             model_path,
428             force_download=args.force_download,
429             no_download=args.no_download,
430         )
431         if args.llama_bin:
432             llama_binary = args.llama_bin.expanduser().resolve()
433             if not llama_binary.exists() or not os.access(llama_binary, os.X_OK↵
↵ ):
434                 raise RuntimeError(f"--llama-bin is not executable: {↵
↵ llama_binary}")
435         else:
436             llama_binary = ensure_llama_simple_chat(

```

```

437         cache_dir,
438         jobs=args.jobs,
439         skip_build=args.skip_build,
440     )
441
442     result, response = run_model_smoke(
443         llama_binary,
444         model_path,
445         prompt=args.prompt,
446         context_tokens=args.ctx,
447         timeout_seconds=args.timeout,
448     )
449     write_report(
450         report_path,
451         cache_dir=cache_dir,
452         llama_binary=llama_binary,
453         model_path=model_path,
454         prompt=args.prompt,
455         result=result,
456         response=response,
457     )
458     except RuntimeError as exc:
459         print(f"[lcqi-smoke] ERROR: {exc}", file=sys.stderr)
460         return 1
461
462     print(f"report={report_path}")
463     print(f"model_sha256={MODEL_EXPECTED_SHA256}")
464     print(f"assistant_response={response}")
465     return 0
466
467
468 if __name__ == "__main__":
469     raise SystemExit(main())

```

`run_gpt2_smoke.py` 下载标准 GPT-2 的 `config.json`、`vocab.json`、`merges.txt` 和 `model.safetensors`，再调用 LCQI 自研 `lcqi_gpt2` reference path 生成 token。它证明自研 `safetensors`、BPE、GPT-2 forward、KV cache decode 和 greedy generation 已连成闭环。

Listing 16.8 tools/run\_gpt2\_smoke.py

```

1  #!/usr/bin/env python3
2  from __future__ import annotations
3
4  import argparse
5  import datetime as dt
6  import hashlib
7  import os
8  import shlex
9  import subprocess
10 import sys
11 import urllib.request
12 from dataclasses import dataclass
13 from pathlib import Path

```

```
14
15
16 MODEL_REPO_ID = "openai-community/gpt2"
17 MODEL_REVISION = "main"
18 MODEL_FILE_IDENTITIES = {
19     "config.json": (
20         665,
21         "0daed7749b4f02b8f76240d5444551d7b08712dab4d0adb8239c56ba823bb7b4",
22     ),
23     "vocab.json": (
24         1042301,
25         "196139668be63f3b5d6574427317ae82f612a97c5d1cdaf36ed2256dbf636783",
26     ),
27     "merges.txt": (
28         456318,
29         "1ce1664773c50f3e0cc8842619a93edc4624525b728b188a9e0be33b7726adc5",
30     ),
31     "model.safetensors": (
32         548105171,
33         "248dfc3911869ec493c76e65bf2fcf7f615828b0254c12b473182f0f81d3a707",
34     ),
35 }
36 DEFAULT_PROMPT = "Hello, my name is"
37 DEFAULT_MAX_NEW_TOKENS = 1
38 DEFAULT_TIMEOUT_SECONDS = 300
39 DEFAULT_BUILD_JOBS = 2
40
41
42 @dataclass(frozen=True)
43 class CommandResult:
44     args: list[str]
45     returncode: int
46     stdout: str
47     stderr: str
48
49
50 def script_path() -> Path:
51     return Path(__file__).resolve()
52
53
54 def volume_dir() -> Path:
55     return script_path().parents[1]
56
57
58 def repo_dir() -> Path:
59     return script_path().parents[3]
60
61
62 def default_cache_dir() -> Path:
63     explicit_cache = os.environ.get("LCQI_GPT2_CACHE_DIR")
64     if explicit_cache:
65         return Path(explicit_cache).expanduser()
```



```

66     xdg_cache = os.environ.get("XDG_CACHE_HOME")
67     cache_root = Path(xdg_cache).expanduser() if xdg_cache else Path.home() / "↵
↵ .cache"
68     return cache_root / "lcqi-gpt2-smoke" / MODEL_REPO_ID.replace("/", "--")
69
70
71 def default_build_dir() -> Path:
72     return volume_dir() / "build" / "lcqi-release"
73
74
75 def default_report_path() -> Path:
76     return volume_dir() / "results" / "lcqi-gpt2-smoke.txt"
77
78
79 def command_line(args: list[str]) -> str:
80     return " ".join(shlex.quote(arg) for arg in args)
81
82
83 def run_command(
84     args: list[str],
85     *,
86     cwd: Path | None = None,
87     timeout_seconds: int | None = None,
88 ) -> CommandResult:
89     try:
90         completed = subprocess.run(
91             args,
92             cwd=str(cwd) if cwd else None,
93             text=True,
94             capture_output=True,
95             timeout=timeout_seconds,
96             check=False,
97         )
98     except subprocess.TimeoutExpired as exc:
99         stdout = exc.stdout if isinstance(exc.stdout, str) else ""
100         stderr = exc.stderr if isinstance(exc.stderr, str) else ""
101         return CommandResult(args=args, returncode=124, stdout=stdout, stderr=↵
↵ stderr)
102     return CommandResult(
103         args=args,
104         returncode=completed.returncode,
105         stdout=completed.stdout,
106         stderr=completed.stderr,
107     )
108
109
110 def ensure_success(result: CommandResult, action: str) -> None:
111     if result.returncode == 0:
112         return
113     raise RuntimeError(
114         f"{action} failed with exit code {result.returncode}\n"
115         f"command: {command_line(result.args)}\n"

```

```

116         f"stdout tail:\n{tail(result.stdout)}\n"
117         f"stderr tail:\n{tail(result.stderr)}"
118     )
119
120
121 def tail(text: str, max_chars: int = 4000) -> str:
122     if len(text) <= max_chars:
123         return text
124     return text[-max_chars:]
125
126
127 def sha256_file(path: Path) -> str:
128     digest = hashlib.sha256()
129     with path.open("rb") as handle:
130         for block in iter(lambda: handle.read(1024 * 1024), b""):
131             digest.update(block)
132     return digest.hexdigest()
133
134
135 def model_url(file_name: str) -> str:
136     return (
137         f"https://huggingface.co/{MODEL_REPO_ID}/resolve/"
138         f"{MODEL_REVISION}/{file_name}"
139     )
140
141
142 def download_file(url: str, target: Path, timeout_seconds: int) -> None:
143     part_path = target.with_suffix(target.suffix + ".part")
144     if part_path.exists():
145         part_path.unlink()
146     request = urllib.request.Request(url, headers={"User-Agent": "lcqi-gpt2-↵
↵ smoke"})
147     with urllib.request.urlopen(request, timeout=timeout_seconds) as response:
148         with part_path.open("wb") as output:
149             while True:
150                 block = response.read(1024 * 1024)
151                 if not block:
152                     break
153                 output.write(block)
154     part_path.replace(target)
155
156
157 def file_identity_matches(path: Path, expected_bytes: int, expected_sha256: str↵
↵ ) -> bool:
158     return (
159         path.exists()
160         and path.stat().st_size == expected_bytes
161         and sha256_file(path) == expected_sha256
162     )
163
164
165 def ensure_model_files(cache_dir: Path, *, no_download: bool, timeout_seconds: ↵

```

```

↪ int) -> None:
166 cache_dir.mkdir(parents=True, exist_ok=True)
167 missing = [
168     name
169     for name, (expected_bytes, expected_sha256) in MODEL_FILE_IDENTITIES.↪
↪ items()
170     if not file_identity_matches(cache_dir / name, expected_bytes, ↪
↪ expected_sha256)
171 ]
172 if not missing:
173     print(f"[lcqi-gpt2] model cache hit: {cache_dir}")
174     return
175 if no_download:
176     raise RuntimeError(
177         f"missing GPT-2 model files and --no-download was set: {missing}"
178     )
179 for file_name in missing:
180     target = cache_dir / file_name
181     if target.exists():
182         target.unlink()
183     print(f"[lcqi-gpt2] downloading {MODEL_REPO_ID}/{file_name}")
184     download_file(model_url(file_name), target, timeout_seconds)
185     expected_bytes, expected_sha256 = MODEL_FILE_IDENTITIES[file_name]
186     if not file_identity_matches(target, expected_bytes, expected_sha256):
187         raise RuntimeError(
188             f"downloaded GPT-2 file identity mismatch: {file_name}"
189         )
190
191
192 def build_lcqi_gpt2(build_dir: Path, jobs: int) -> Path:
193     configure = run_command(
194         [
195             "cmake",
196             "-S",
197             str(volume_dir()),
198             "-B",
199             str(build_dir),
200             "-DCMAKE_BUILD_TYPE=Release",
201         ],
202         timeout_seconds=DEFAULT_TIMEOUT_SECONDS,
203     )
204     ensure_success(configure, "configure LCQI")
205     build = run_command(
206         [
207             "cmake",
208             "--build",
209             str(build_dir),
210             "--target",
211             "lcqi_gpt2",
212             "-j",
213             str(jobs),
214         ],

```

```

215         timeout_seconds=DEFAULT_TIMEOUT_SECONDS,
216     )
217     ensure_success(build, "build lcqi_gpt2")
218     binary = build_dir / "labs" / "linux_cpu_inference" / "lcqi_gpt2"
219     if not binary.exists():
220         raise RuntimeError(f"lcqi_gpt2 binary not found: {binary}")
221     return binary
222
223
224 def write_report(
225     report_path: Path,
226     *,
227     cache_dir: Path,
228     binary: Path,
229     prompt: str,
230     max_new_tokens: int,
231     engine: str,
232     result: CommandResult,
233 ) -> None:
234     report_path.parent.mkdir(parents=True, exist_ok=True)
235     lines = [
236         f"timestamp_utc={dt.datetime.now(dt.UTC).isoformat()}",
237         f"model_repo_id={MODEL_REPO_ID}",
238         f"model_revision={MODEL_REVISION}",
239         f"model_cache_dir={cache_dir}",
240     ]
241     for file_name, (expected_bytes, expected_sha256) in MODEL_FILE_IDENTITIES.items():
242         path = cache_dir / file_name
243         lines.append(f"file.{file_name}.bytes={path.stat().st_size}")
244         lines.append(f"file.{file_name}.sha256={sha256_file(path)}")
245         lines.append(f"file.{file_name}.expected_bytes={expected_bytes}")
246         lines.append(f"file.{file_name}.expected_sha256={expected_sha256}")
247     lines.extend(
248         [
249             f"binary={binary}",
250             f"prompt={prompt}",
251             f"max_new_tokens={max_new_tokens}",
252             f"engine={engine}",
253             f"command={command_line(result.args)}",
254             f"exit_code={result.returncode}",
255             "stdout_begin",
256             result.stdout.rstrip(),
257             "stdout_end",
258             "stderr_begin",
259             tail(result.stderr).rstrip(),
260             "stderr_end",
261         ]
262     )
263     passed = result.returncode == 0 and "generated_text " in result.stdout
264     lines.append(f"validation={'PASS' if passed else 'FAIL'}")
265     report_path.write_text("\n".join(lines) + "\n", encoding="utf-8")

```

```

266
267
268 def parse_args() -> argparse.Namespace:
269     parser = argparse.ArgumentParser(
270         description="Run LCQI self-owned GPT-2 smoke on a HuggingFace ↵
↵ safetensors directory."
271     )
272     parser.add_argument("--cache-dir", type=Path, default=default_cache_dir())
273     parser.add_argument("--build-dir", type=Path, default=default_build_dir())
274     parser.add_argument("--report", type=Path, default=default_report_path())
275     parser.add_argument("--prompt", default=DEFAULT_PROMPT)
276     parser.add_argument("--max-new-tokens", type=int, default=↵
↵ DEFAULT_MAX_NEW_TOKENS)
277     parser.add_argument("--engine", choices=("cached", "full"), default="cached↵
↵ ")
278     parser.add_argument("--benchmark", action="store_true")
279     parser.add_argument("--timeout-seconds", type=int, default=↵
↵ DEFAULT_TIMEOUT_SECONDS)
280     parser.add_argument("--jobs", type=int, default=DEFAULT_BUILD_JOBS)
281     parser.add_argument("--no-download", action="store_true")
282     return parser.parse_args()
283
284
285 def main() -> int:
286     args = parse_args()
287     try:
288         if args.max_new_tokens < 0:
289             raise RuntimeError("--max-new-tokens cannot be negative")
290         ensure_model_files(
291             args.cache_dir,
292             no_download=args.no_download,
293             timeout_seconds=args.timeout_seconds,
294         )
295         binary = build_lcqi_gpt2(args.build_dir, args.jobs)
296         command = [
297             str(binary),
298             "--engine",
299             args.engine,
300             str(args.cache_dir),
301             args.prompt,
302             str(args.max_new_tokens),
303         ]
304         if args.benchmark:
305             command.insert(1, "--benchmark")
306         result = run_command(command, cwd=repo_dir(), timeout_seconds=args.↵
↵ timeout_seconds)
307         write_report(
308             args.report,
309             cache_dir=args.cache_dir,
310             binary=binary,
311             prompt=args.prompt,
312             max_new_tokens=args.max_new_tokens,

```

```

313         engine=args.engine,
314         result=result,
315     )
316     print(f"[lcqi-gpt2] report: {args.report}")
317     print(tail(result.stdout, max_chars=2000))
318     return 0 if result.returncode == 0 else result.returncode
319 except Exception as error:
320     print(f"[lcqi-gpt2] error: {error}", file=sys.stderr)
321     return 1
322
323
324 if __name__ == "__main__":
325     raise SystemExit(main())

```

`run_gpt2_benchmark_compare.py` 用同一个 GPT-2 F32 safetensors 模型分别运行 LCQI full-prefix 与 cached-KV 路径，并在本机存在 llama.cpp/SmolLM2 缓存时附带成熟引擎参考。它的报告说明算法差距、kernel 差距和模型口径差异。

**Listing 16.9** tools/run\_gpt2\_benchmark\_compare.py

```

1  #!/usr/bin/env python3
2  from __future__ import annotations
3
4  import argparse
5  import datetime as dt
6  import os
7  import re
8  import shlex
9  import subprocess
10 import sys
11 from dataclasses import dataclass
12 from pathlib import Path
13
14 import run_gpt2_smoke
15
16
17 DEFAULT_PROMPT = "Hello, my name is"
18 DEFAULT_MAX_NEW_TOKENS = 4
19 DEFAULT_BUILD_JOBS = 2
20 DEFAULT_TIMEOUT_SECONDS = 300
21 DEFAULT_SMOLLM2_CACHE_DIR = Path.home() / ".cache" / "lcqi-smolllm2-small-smoke"
22 SMOLLM2_GGUF = "SmolLM2-135M-Instruct-Q4_K_M.gguf"
23
24
25 @dataclass(frozen=True)
26 class EngineRun:
27     name: str
28     command: list[str]
29     returncode: int
30     stdout: str
31     stderr: str
32     metrics: dict[str, str]
33

```

```

34
35 def volume_dir() -> Path:
36     return Path(__file__).resolve().parents[1]
37
38
39 def repo_dir() -> Path:
40     return Path(__file__).resolve().parents[3]
41
42
43 def default_report_path() -> Path:
44     return volume_dir() / "results" / "lcqi-gpt2-benchmark-compare.txt"
45
46
47 def default_build_dir() -> Path:
48     return volume_dir() / "build" / "lcqi-release"
49
50
51 def command_line(args: list[str]) -> str:
52     return " ".join(shlex.quote(arg) for arg in args)
53
54
55 def tail(text: str, max_chars: int = 4000) -> str:
56     return text if len(text) <= max_chars else text[-max_chars:]
57
58
59 def run_command(args: list[str], *, timeout_seconds: int) -> tuple[int, str, ←
    ↪ str]:
60     try:
61         completed = subprocess.run(
62             args,
63             cwd=repo_dir(),
64             text=True,
65             capture_output=True,
66             timeout=timeout_seconds,
67             check=False,
68         )
69     except subprocess.TimeoutExpired as exc:
70         stdout = exc.stdout if isinstance(exc.stdout, str) else ""
71         stderr = exc.stderr if isinstance(exc.stderr, str) else ""
72         return 124, stdout, stderr
73     return completed.returncode, completed.stdout, completed.stderr
74
75
76 def parse_lcqi_metrics(stdout: str) -> dict[str, str]:
77     metrics: dict[str, str] = {}
78     for line in stdout.splitlines():
79         if not line:
80             continue
81         key, _, value = line.partition(" ")
82         if key in {
83             "engine",
84             "generated_text",

```

```

85         "benchmark_load_ms",
86         "benchmark_tokenize_ms",
87         "benchmark_prefill_ms",
88         "benchmark_decode_ms",
89         "benchmark_generate_ms",
90         "benchmark_total_ms",
91         "benchmark_prompt_tokens",
92         "benchmark_generated_tokens",
93         "benchmark_prefill_steps",
94         "benchmark_decode_steps",
95         "benchmark_generate_tokens_per_second",
96         "benchmark_decode_tokens_per_second",
97         "benchmark_kv_cache_bytes",
98     }:
99         metrics[key] = value
100     return metrics
101
102
103 def run_lcqi(
104     binary: Path,
105     model_dir: Path,
106     *,
107     engine: str,
108     prompt: str,
109     max_new_tokens: int,
110     timeout_seconds: int,
111 ) -> EngineRun:
112     command = [
113         str(binary),
114         "--benchmark",
115         "--engine",
116         engine,
117         str(model_dir),
118         prompt,
119         str(max_new_tokens),
120     ]
121     returncode, stdout, stderr = run_command(command, timeout_seconds=↵
↵ timeout_seconds)
122     return EngineRun(
123         name=f"lcqi_{engine}",
124         command=command,
125         returncode=returncode,
126         stdout=stdout,
127         stderr=stderr,
128         metrics=parse_lcqi_metrics(stdout),
129     )
130
131
132 def llama_bench_binary(cache_dir: Path) -> Path:
133     return cache_dir / "llama.cpp" / "build" / "bin" / "llama-bench"
134
135

```



```

136 def smollm2_model_path(cache_dir: Path) -> Path:
137     return cache_dir / "models" / SMOLLM2_GGUF
138
139
140 def parse_llama_bench_metrics(stdout: str) -> dict[str, str]:
141     metrics: dict[str, str] = {}
142     lines = [line.strip() for line in stdout.splitlines() if line.strip()]
143     for line in lines:
144         if not line.startswith("|"):
145             continue
146         columns = [column.strip() for column in line.strip("|").split("|")]
147         if len(columns) != 7 or columns[0] == "model" or columns[0].startswith(↵
↵ "-"):
148             continue
149         metrics["llama_model"] = columns[0]
150         metrics["llama_size"] = columns[1]
151         metrics["llama_params"] = columns[2]
152         metrics["llama_backend"] = columns[3]
153         metrics["llama_threads"] = columns[4]
154         test_name = columns[5]
155         tokens_per_second = re.sub(r"\s+.*$", "", columns[6])
156         if test_name.startswith("pp"):
157             metrics["llama_prefill_test"] = test_name
158             metrics["llama_prefill_tps"] = tokens_per_second
159         elif test_name.startswith("tg"):
160             metrics["llama_decode_test"] = test_name
161             metrics["llama_decode_tps"] = tokens_per_second
162     return metrics
163
164
165 def run_llama_bench(cache_dir: Path, *, timeout_seconds: int) -> EngineRun | ↵
↵ None:
166     binary = llama_bench_binary(cache_dir)
167     model = smollm2_model_path(cache_dir)
168     if not binary.exists() or not os.access(binary, os.X_OK) or not model.↵
↵ exists():
169         return None
170     command = [
171         str(binary),
172         "-m",
173         str(model),
174         "-ngl",
175         "0",
176         "-p",
177         "5",
178         "-n",
179         "4",
180         "-r",
181         "1",
182     ]
183     returncode, stdout, stderr = run_command(command, timeout_seconds=↵
↵ timeout_seconds)

```

```

184     return EngineRun(
185         name="llama_cpp_smolllm2_q4_k_m",
186         command=command,
187         returncode=returncode,
188         stdout=stdout,
189         stderr=stderr,
190         metrics=parse_llama_bench_metrics(stdout),
191     )
192
193
194 def write_report(path: Path, runs: list[EngineRun], *, prompt: str, ←
    ↪ max_new_tokens: int) -> None:
195     path.parent.mkdir(parents=True, exist_ok=True)
196     lines = [
197         "LCQI GPT-2 Benchmark Compare",
198         "",
199         f"timestamp_utc={dt.datetime.now(dt.UTC).isoformat()}",
200         f"repo_commit={current_commit()}",
201         f"working_tree_dirty={working_tree_dirty()}",
202         f"python={sys.version.split()[0]}",
203         f"prompt={prompt}",
204         f"max_new_tokens={max_new_tokens}",
205         "",
206         "scope=LCQI full_prefix and cached_kv run the same GPT-2 F32 ←
    ↪ safetensors model. "
207         "llama.cpp uses SmolLM2-135M-Instruct Q4_K_M as an external mature-←
    ↪ engine reference, "
208         "so its numbers are engineering context, not same-model quality ranking←
    ↪ .",
209         "",
210     ]
211     for run in runs:
212         lines.extend(
213             [
214                 f"[{run.name}]",
215                 f"returncode={run.returncode}",
216                 f"command={command_line(run.command)}",
217             ]
218         )
219         for key in sorted(run.metrics):
220             lines.append(f"{key}={run.metrics[key]}")
221         lines.extend(
222             [
223                 "stdout_tail:",
224                 tail(run.stdout.strip(), 2000),
225                 "stderr_tail:",
226                 tail(run.stderr.strip(), 2000),
227                 "",
228             ]
229         )
230     while lines and lines[-1] == "":
231         lines.pop()

```

```

232     path.write_text("\n".join(lines) + "\n", encoding="utf-8")
233
234
235 def current_commit() -> str:
236     result = subprocess.run(
237         ["git", "-C", str(repo_dir()), "rev-parse", "--short", "HEAD"],
238         text=True,
239         capture_output=True,
240         check=False,
241     )
242     return result.stdout.strip() if result.returncode == 0 else "unknown"
243
244
245 def working_tree_dirty() -> str:
246     result = subprocess.run(
247         ["git", "-C", str(repo_dir()), "status", "--porcelain"],
248         text=True,
249         capture_output=True,
250         check=False,
251     )
252     if result.returncode != 0:
253         return "unknown"
254     return "true" if result.stdout.strip() else "false"
255
256
257 def parse_args() -> argparse.Namespace:
258     parser = argparse.ArgumentParser(
259         description="Compare LCQI GPT-2 full-prefix and KV-cache paths, with ↵
↵ optional llama.cpp context."
260     )
261     parser.add_argument("--cache-dir", type=Path, default=run_gpt2_smoke.↵
↵ default_cache_dir())
262     parser.add_argument("--build-dir", type=Path, default=default_build_dir())
263     parser.add_argument("--report", type=Path, default=default_report_path())
264     parser.add_argument("--prompt", default=DEFAULT_PROMPT)
265     parser.add_argument("--max-new-tokens", type=int, default=↵
↵ DEFAULT_MAX_NEW_TOKENS)
266     parser.add_argument("--jobs", type=int, default=DEFAULT_BUILD_JOBS)
267     parser.add_argument("--timeout-seconds", type=int, default=↵
↵ DEFAULT_TIMEOUT_SECONDS)
268     parser.add_argument("--no-download", action="store_true")
269     parser.add_argument("--skip-llama", action="store_true")
270     parser.add_argument("--smollm2-cache-dir", type=Path, default=↵
↵ DEFAULT_SMOLLM2_CACHE_DIR)
271     return parser.parse_args()
272
273
274 def main() -> int:
275     args = parse_args()
276     if args.max_new_tokens < 0:
277         print("[lcqi-compare] --max-new-tokens cannot be negative", file=sys.↵
↵ stderr)

```

```

278         return 1
279     try:
280         run_gpt2_smoke.ensure_model_files(
281             args.cache_dir,
282             no_download=args.no_download,
283             timeout_seconds=args.timeout_seconds,
284         )
285         binary = run_gpt2_smoke.build_lcqi_gpt2(args.build_dir, args.jobs)
286         runs = [
287             run_lcqi(
288                 binary,
289                 args.cache_dir,
290                 engine="full",
291                 prompt=args.prompt,
292                 max_new_tokens=args.max_new_tokens,
293                 timeout_seconds=args.timeout_seconds,
294             ),
295             run_lcqi(
296                 binary,
297                 args.cache_dir,
298                 engine="cached",
299                 prompt=args.prompt,
300                 max_new_tokens=args.max_new_tokens,
301                 timeout_seconds=args.timeout_seconds,
302             ),
303         ]
304         if not args.skip_llama:
305             llama_run = run_llama_bench(
306                 args.smollm2_cache_dir.expanduser().resolve(),
307                 timeout_seconds=args.timeout_seconds,
308             )
309             if llama_run is not None:
310                 runs.append(llama_run)
311             write_report(args.report, runs, prompt=args.prompt, max_new_tokens=args↵
↵ .max_new_tokens)
312         except Exception as error:
313             print(f"[lcqi-compare] error: {error}", file=sys.stderr)
314             return 1
315
316         for run in runs:
317             print(f"{run.name}: returncode={run.returncode}")
318             for key in sorted(run.metrics):
319                 print(f"    {key}={run.metrics[key]}")
320             print(f"report={args.report}")
321         return 0 if all(run.returncode == 0 for run in runs) else 1
322
323
324 if __name__ == "__main__":
325     raise SystemExit(main())

```

`run_lcqi_benchmark.sh` 是本地 benchmark 收集脚本。它适合在固定机器上反复运行，把同一组 shape 和 backend 的结果落到报告文件。

Listing 16.10 tools/run\_lcqi\_benchmark.sh

```

1  #!/usr/bin/env bash
2  set -euo pipefail
3
4  SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
5  VOLUME_DIR="$(cd "${SCRIPT_DIR}/.." && pwd)"
6  REPO_DIR="$(cd "${VOLUME_DIR}/../.." && pwd)"
7  BUILD_DIR="${VOLUME_DIR}/build/lcqi-release"
8  RESULT_DIR="${VOLUME_DIR}/results"
9  REPEAT="${1:-200}"
10 STAMP="$(date -u +%Y%m%dT%H%M%S)"
11 CSV_PATH="${RESULT_DIR}/lcqi-bench-${STAMP}.csv"
12 META_PATH="${RESULT_DIR}/lcqi-bench-${STAMP}.meta.txt"
13 PERF_PATH="${RESULT_DIR}/lcqi-bench-${STAMP}.perf.txt"
14
15 mkdir -p "${RESULT_DIR}"
16
17 cmake -S "${VOLUME_DIR}" -B "${BUILD_DIR}" -DCMAKE_BUILD_TYPE=Release
18 cmake --build "${BUILD_DIR}"
19 ctest --test-dir "${BUILD_DIR}" --output-on-failure
20
21 {
22     echo "timestamp_utc=${STAMP}"
23     echo "repeat=${REPEAT}"
24     echo "commit=$(git -C "${REPO_DIR}" rev-parse --short HEAD)"
25     echo "uname=$(uname -a)"
26     echo "compiler=$((${CXX:-c++} --version | sed -n '1p')"
27     echo "cmake=$(cmake --version | sed -n '1p')"
28     echo "bench=${BUILD_DIR}/labs/linux_cpu_inference/lcqi_bench"
29 } > "${META_PATH}"
30
31 "${BUILD_DIR}/labs/linux_cpu_inference/lcqi_bench" "${REPEAT}" > "${CSV_PATH}"
32
33 if command -v perf >/dev/null 2>&1; then
34     perf stat \
35         -e cycles,instructions,cache-misses,branches,branch-misses \
36         "${BUILD_DIR}/labs/linux_cpu_inference/lcqi_bench" "${REPEAT}" \
37         > /dev/null 2> "${PERF_PATH}" || true
38 else
39     echo "perf_unavailable=1" > "${PERF_PATH}"
40 fi
41
42 echo "csv=${CSV_PATH}"
43 echo "meta=${META_PATH}"
44 echo "perf=${PERF_PATH}"

```

## 完整测试

`lcqi_tests.cpp` 把 safetensors、tokenizer、tiny MLP、int8 kernel、reference decoder、GPT-2 tiny path、trace、ledger 和 serving probe 串成验收网。新增优化或 shape 时，必须先扩展这里的 golden 和边界样例。

Listing 16.11 tests/lcqi\_tests.cpp

```

1  #include <lcqi/gpt2_reference.hpp>
2  #include <lcqi/inference.hpp>
3  #include <lcqi/int8_kernels.hpp>
4  #include <lcqi/model.hpp>
5  #include <lcqi/reference_decoder.hpp>
6  #include <lcqi/safetensors.hpp>
7  #include <lcqi/tokenizer.hpp>
8
9  #include <array>
10 #include <cmath>
11 #include <cstring>
12 #include <cstdlib>
13 #include <filesystem>
14 #include <fstream>
15 #include <iostream>
16 #include <span>
17 #include <sstream>
18 #include <stdexcept>
19 #include <string>
20 #include <utility>
21 #include <vector>
22
23 namespace {
24
25 constexpr float LCQI_TEST_EPSILON = 1.0e-5F;
26 constexpr std::int32_t LCQI_TEST_INPUT_SIZE = 4;
27 constexpr std::int32_t LCQI_TEST_HIDDEN_SIZE = 3;
28 constexpr std::int32_t LCQI_TEST_OUTPUT_SIZE = 2;
29 constexpr std::int32_t LCQI_TEST_PREDICTED_CLASS = 1;
30 constexpr float LCQI_TEST_LOGIT_0 = -0.48F;
31 constexpr float LCQI_TEST_LOGIT_1 = 1.06F;
32 constexpr float LCQI_TEST_HIDDEN_0 = 0.0F;
33 constexpr float LCQI_TEST_HIDDEN_1 = 0.4F;
34 constexpr float LCQI_TEST_HIDDEN_2 = 1.4F;
35 constexpr float LCQI_RMS_EXPECTED_0 = 0.848528F;
36 constexpr float LCQI_RMS_EXPECTED_1 = 0.565685F;
37 constexpr float LCQI_ATTENTION_EXPECTED_0 = 2.0F;
38 constexpr float LCQI_ATTENTION_EXPECTED_1 = 3.0F;
39 constexpr float LCQI_KERNEL_MAX_DIFF = 1.0e-5F;
40 constexpr std::int32_t LCQI_SIMD_TEST_BLOCK = 8;
41 constexpr std::int32_t LCQI_REFERENCE_DECODER_PREDICTED_TOKEN = 0;
42 constexpr std::size_t LCQI_REFERENCE_DECODER_VOCAB_SIZE = 5;
43 constexpr std::array<float, LCQI_REFERENCE_DECODER_VOCAB_SIZE> ←
    ↪ LCQI_REFERENCE_DECODER_LOGITS{
44     0.352516979F,
45     -0.126884326F,
46     0.0797837451F,
47     -0.29797405F,
48     0.127030417F,
49 };
50 constexpr std::int32_t LCQI_REFERENCE_TRACE_TOKEN_COUNT = 3;

```

```

51 constexpr std::int32_t LCQI_REFERENCE_TRACE_HIDDEN_SIZE = 4;
52 constexpr std::int32_t LCQI_REFERENCE_TRACE_FIRST_TOKEN = 0;
53 constexpr std::uint64_t LCQI_TOKENIZER_HASH_EXPECTED = 13452902845388333734ULL;
54 constexpr std::int32_t LCQI_SAFETENSORS_PREFIX_BYTES = 8;
55 constexpr std::int32_t LCQI_TEST_BYTE_BITS = 8;
56 constexpr std::uint64_t LCQI_SAFETENSORS_VALID_PAYLOAD_BYTES = 48;
57 constexpr std::uint64_t LCQI_SAFETENSORS_EMBED_BEGIN = 0;
58 constexpr std::uint64_t LCQI_SAFETENSORS_EMBED_END = 24;
59 constexpr std::uint64_t LCQI_SAFETENSORS_QPROJ_BEGIN = 24;
60 constexpr std::uint64_t LCQI_SAFETENSORS_QPROJ_END = 48;
61 constexpr unsigned int LCQI_BYTE_MASK = 0xFFU;
62 constexpr std::int32_t LCQI_GPT2_PREDICTED_TOKEN = 3;
63 constexpr std::size_t LCQI_GPT2_VOCAB_SIZE = 6;
64 constexpr std::array<float, LCQI_GPT2_VOCAB_SIZE> LCQI_GPT2_LOGITS{
65     0.407358F,
66     -0.504049F,
67     -0.107834F,
68     0.840337F,
69     -0.450123F,
70     -0.156763F,
71 };
72 constexpr std::int32_t LCQI_GPT2_GENERATED_LENGTH = 5;
73 constexpr std::int32_t LCQI_GPT2_TOKENIZER_HELLO_ID = 7;
74 constexpr std::uint16_t LCQI_TEST_F16_ONE = 0x3C00U;
75 constexpr std::uint16_t LCQI_TEST_F16_NEGATIVE_TWO = 0xC000U;
76 constexpr std::uint16_t LCQI_TEST_BF16_ONE_POINT_FIVE = 0x3FC0U;
77 constexpr std::uint16_t LCQI_TEST_BF16_NEGATIVE_HALF = 0xBF00U;
78
79 struct KernelShapeCase {
80     std::int32_t input_size = 0;
81     std::int32_t output_size = 0;
82     std::int32_t output_block_size = 0;
83 };
84
85 struct RawTensorFixture {
86     std::string name;
87     std::string dtype;
88     std::vector<std::int64_t> shape;
89     std::vector<std::uint8_t> bytes;
90 };
91
92 struct FloatTensorFixture {
93     std::string name;
94     std::vector<std::int64_t> shape;
95     std::vector<float> values;
96 };
97
98 constexpr std::array<KernelShapeCase, 4> LCQI_KERNEL_SHAPE_CASES{{
99     {17, 19, 8},
100    {33, 7, 8},
101    {64, 32, 16},
102    {65, 31, 16},

```

```

103 }};
104
105 void require(bool condition, const char* message) {
106     if (!condition) {
107         throw std::runtime_error(message);
108     }
109 }
110
111 void require_close(float actual, float expected, const char* message) {
112     if (std::fabs(actual - expected) > LCQI_TEST_EPSILON) {
113         throw std::runtime_error(message);
114     }
115 }
116
117 std::filesystem::path test_model_path() {
118     return std::filesystem::path(__FILE__).parent_path().parent_path() /
119         "models" / "tiny_mlp_i8.txt";
120 }
121
122 std::filesystem::path test_tokenizer_path() {
123     return std::filesystem::path(__FILE__).parent_path().parent_path() /
124         "models" / "tiny_tokenizer.txt";
125 }
126
127 std::filesystem::path temp_file_path(const std::string& name) {
128     return std::filesystem::temp_directory_path() / name;
129 }
130
131 void write_le_u64(std::ofstream& output, std::uint64_t value) {
132     for (std::int32_t i = 0; i < LCQI_SAFETENSORS_PREFIX_BYTES; ++i) {
133         const char byte = static_cast<char>(
134             (value >> (LCQI_TEST_BYTE_BITS * i)) & LCQI_BYTE_MASK);
135         output.write(&byte, 1);
136     }
137 }
138
139 void append_le_u16(std::vector<std::uint8_t>& output, std::uint16_t value) {
140     output.push_back(static_cast<std::uint8_t>(value & LCQI_BYTE_MASK));
141     output.push_back(static_cast<std::uint8_t>((value >> LCQI_TEST_BYTE_BITS) &
142         ↵ LCQI_BYTE_MASK));
143 }
144
145 void append_le_f32(std::vector<std::uint8_t>& output, float value) {
146     std::uint32_t bits = 0;
147     std::memcpy(&bits, &value, sizeof(bits));
148     for (std::int32_t i = 0; i < 4; ++i) {
149         output.push_back(static_cast<std::uint8_t>(
150             (bits >> (LCQI_TEST_BYTE_BITS * i)) & LCQI_BYTE_MASK));
151     }
152 }
153
154 void write_safetensors_fixture(const std::filesystem::path& path,

```



```

154         const std::string& header,
155         std::uint64_t payload_bytes) {
156     std::ofstream output(path, std::ios::binary);
157     if (!output) {
158         throw std::runtime_error("cannot create safetensors test fixture");
159     }
160     write_le_u64(output, static_cast<std::uint64_t>(header.size()));
161     output.write(header.data(), static_cast<std::streamsize>(header.size()));
162     for (std::uint64_t i = 0; i < payload_bytes; ++i) {
163         const char byte = static_cast<char>(i & 0xFFU);
164         output.write(&byte, 1);
165     }
166 }
167
168 std::string json_shape(std::span<const std::int64_t> shape) {
169     std::ostringstream stream;
170     stream << '[';
171     for (std::size_t index = 0; index < shape.size(); ++index) {
172         if (index != 0) {
173             stream << ',';
174         }
175         stream << shape[index];
176     }
177     stream << ']';
178     return stream.str();
179 }
180
181 void write_safetensors_raw_fixture(const std::filesystem::path& path,
182                                   std::span<const RawTensorFixture> tensors) {
183     std::ostringstream header;
184     header << "{\"__metadata__\":{\"format\":\"lcqi-test\"}";
185     std::uint64_t offset = 0;
186     for (const RawTensorFixture& tensor : tensors) {
187         header << ",\"" << tensor.name << "\":{\"dtype\":\"" << tensor.dtype
188             << "\",\"shape\":\"" << json_shape(tensor.shape)
189             << "\",\"data_offsets\":[\"" << offset << ','
190             << offset + tensor.bytes.size() << "\"]}\"";
191         offset += tensor.bytes.size();
192     }
193     header << '}'";
194
195     std::ofstream output(path, std::ios::binary);
196     if (!output) {
197         throw std::runtime_error("cannot create safetensors raw test fixture");
198     }
199     const std::string header_text = header.str();
200     write_le_u64(output, static_cast<std::uint64_t>(header_text.size()));
201     output.write(header_text.data(), static_cast<std::streamsize>(header_text.↵
↵ size()));
202     for (const RawTensorFixture& tensor : tensors) {
203         output.write(reinterpret_cast<const char*>(tensor.bytes.data()),
204                     static_cast<std::streamsize>(tensor.bytes.size()));

```

```

205     }
206 }
207
208 std::vector<std::uint8_t> f32_payload(std::span<const float> values) {
209     std::vector<std::uint8_t> bytes;
210     bytes.reserve(values.size() * 4);
211     for (const float value : values) {
212         append_le_f32(bytes, value);
213     }
214     return bytes;
215 }
216
217 std::vector<float> transpose_linear_to_hf_conv1d(const lcqi::Gpt2LinearF32& ←
    ↪ linear) {
218     std::vector<float> transposed(
219         static_cast<std::size_t>(linear.input_size) *
220         static_cast<std::size_t>(linear.output_size),
221         0.0F);
222     for (std::int32_t out = 0; out < linear.output_size; ++out) {
223         for (std::int32_t in = 0; in < linear.input_size; ++in) {
224             transposed[static_cast<std::size_t>(in) *
225                 static_cast<std::size_t>(linear.output_size) +
226                 static_cast<std::size_t>(out)] =
227                 linear.weights[static_cast<std::size_t>(out) *
228                     static_cast<std::size_t>(linear.input_size) ←
    ↪ +
229                     static_cast<std::size_t>(in)];
230         }
231     }
232     return transposed;
233 }
234
235 void add_f32_tensor(std::vector<RawTensorFixture>& tensors,
236                     std::string name,
237                     std::vector<std::int64_t> shape,
238                     std::span<const float> values) {
239     tensors.push_back(RawTensorFixture{
240         std::move(name),
241         "F32",
242         std::move(shape),
243         f32_payload(values),
244     });
245 }
246
247 void write_text_file(const std::filesystem::path& path, const std::string& text ←
    ↪ ) {
248     std::ofstream output(path);
249     if (!output) {
250         throw std::runtime_error("cannot create text fixture");
251     }
252     output << text;
253 }

```

```

254
255 void write_tiny_gpt2_safetensors(const std::filesystem::path& path,
256                                 const lcqi::Gpt2ReferenceModel& model) {
257     std::vector<RawTensorFixture> tensors;
258     add_f32_tensor(tensors,
259                   "transformer.wte.weight",
260                   {model.config.vocab_size, model.config.hidden_size},
261                   model.token_embedding);
262     add_f32_tensor(tensors,
263                   "transformer.wpe.weight",
264                   {model.config.max_positions, model.config.hidden_size},
265                   model.position_embedding);
266     for (std::int32_t layer_id = 0; layer_id < model.config.layer_count; ++
↵ layer_id) {
267         const lcqi::Gpt2LayerWeightsF32& layer =
268             model.layers[static_cast<std::size_t>(layer_id)];
269         const std::string prefix = "transformer.h." + std::to_string(layer_id)
↵ + ". ";
270         add_f32_tensor(tensors, prefix + "ln_1.weight", {model.config.
↵ hidden_size}, layer.ln_1_weight);
271         add_f32_tensor(tensors, prefix + "ln_1.bias", {model.config.hidden_size
↵ }, layer.ln_1_bias);
272         const std::vector<float> c_attn_hf = transpose_linear_to_hf_conv1d(
↵ layer.c_attn);
273         add_f32_tensor(tensors,
274                       prefix + "attn.c_attn.weight",
275                       {model.config.hidden_size,
276                       model.config.hidden_size * 3},
277                       c_attn_hf);
278         add_f32_tensor(tensors,
279                       prefix + "attn.c_attn.bias",
280                       {model.config.hidden_size * 3},
281                       layer.c_attn.bias);
282         const std::vector<float> c_proj_hf = transpose_linear_to_hf_conv1d(
↵ layer.c_proj);
283         add_f32_tensor(tensors,
284                       prefix + "attn.c_proj.weight",
285                       {model.config.hidden_size, model.config.hidden_size},
286                       c_proj_hf);
287         add_f32_tensor(tensors,
288                       prefix + "attn.c_proj.bias",
289                       {model.config.hidden_size},
290                       layer.c_proj.bias);
291         add_f32_tensor(tensors, prefix + "ln_2.weight", {model.config.
↵ hidden_size}, layer.ln_2_weight);
292         add_f32_tensor(tensors, prefix + "ln_2.bias", {model.config.hidden_size
↵ }, layer.ln_2_bias);
293         const std::vector<float> c_fc_hf = transpose_linear_to_hf_conv1d(layer.
↵ c_fc);
294         add_f32_tensor(tensors,
295                       prefix + "mlp.c_fc.weight",
296                       {model.config.hidden_size, model.config.

```

```

    ↪ intermediate_size},
297         c_fc_hf);
298     add_f32_tensor(tensors,
299         prefix + "mlp.c_fc.bias",
300         {model.config.intermediate_size},
301         layer.c_fc.bias);
302     const std::vector<float> mlp_proj_hf =
303         transpose_linear_to_hf_conv1d(layer.mlp_c_proj);
304     add_f32_tensor(tensors,
305         prefix + "mlp.c_proj.weight",
306         {model.config.intermediate_size, model.config.↪
    ↪ hidden_size},
307         mlp_proj_hf);
308     add_f32_tensor(tensors,
309         prefix + "mlp.c_proj.bias",
310         {model.config.hidden_size},
311         layer.mlp_c_proj.bias);
312 }
313 add_f32_tensor(tensors,
314     "transformer.ln_f.weight",
315     {model.config.hidden_size},
316     model.final_ln_weight);
317 add_f32_tensor(tensors,
318     "transformer.ln_f.bias",
319     {model.config.hidden_size},
320     model.final_ln_bias);
321 write_safetensors_raw_fixture(path, tensors);
322 }
323
324 void test_tiny_mlp() {
325     const lcqi::TinyMlpModel model = lcqi::load_model(test_model_path());
326     require(model.input_size == LCQI_TEST_INPUT_SIZE, "input size mismatch");
327     require(model.hidden_size == LCQI_TEST_HIDDEN_SIZE, "hidden size mismatch")↪
    ↪ ;
328     require(model.output_size == LCQI_TEST_OUTPUT_SIZE, "output size mismatch")↪
    ↪ ;
329
330     const std::vector<float> input{1.0F, 2.0F, -1.0F, 0.5F};
331     const lcqi::InferenceResult result = lcqi::run_inference(model, input);
332
333     require(result.predicted_class == LCQI_TEST_PREDICTED_CLASS, "unexpected ↪
    ↪ predicted class");
334     require(result.logits.size() == static_cast<std::size_t>(↪
    ↪ LCQI_TEST_OUTPUT_SIZE),
335         "unexpected logits size");
336     require_close(result.logits[0], LCQI_TEST_LOGIT_0, "logit 0 mismatch");
337     require_close(result.logits[1], LCQI_TEST_LOGIT_1, "logit 1 mismatch");
338
339     std::vector<float> hidden(static_cast<std::size_t>(LCQI_TEST_HIDDEN_SIZE), ↪
    ↪ 0.0F);
340     lcqi::linear_i8(model.hidden, input, hidden);
341     lcqi::relu_inplace(hidden);

```

```

342     require_close(hidden[0], LCQI_TEST_HIDDEN_0, "hidden 0 mismatch");
343     require_close(hidden[1], LCQI_TEST_HIDDEN_1, "hidden 1 mismatch");
344     require_close(hidden[2], LCQI_TEST_HIDDEN_2, "hidden 2 mismatch");
345 }
346
347 void test_tokenizer_contract() {
348     const lcqi::TokenizerModel tokenizer = lcqi::load_tokenizer(↵
↵ test_tokenizer_path());
349     require(tokenizer.bos_id == 0, "tokenizer BOS id mismatch");
350     require(tokenizer.eos_id == 4, "tokenizer EOS id mismatch");
351     require(tokenizer.unk_id == -1, "tokenizer UNK id mismatch");
352     require(tokenizer.chat_template_hash == "tiny-chat-template-v1",
353             "tokenizer template hash mismatch");
354
355     const std::vector<std::int32_t> token_ids =
356         lcqi::encode_prompt(tokenizer, "tok1 tok2 tok3");
357     const std::vector<std::int32_t> expected{0, 1, 2, 3, 4};
358     require(token_ids == expected, "tokenizer prompt ids mismatch");
359     require(lcqi::tokenizer_contract_hash(tokenizer) == ↵
↵ LCQI_TOKENIZER_HASH_EXPECTED,
360             "tokenizer contract hash mismatch");
361 }
362
363 void test_tokenizer_rejects_unknown_without_unk() {
364     const lcqi::TokenizerModel tokenizer = lcqi::load_tokenizer(↵
↵ test_tokenizer_path());
365     bool threw = false;
366     try {
367         static_cast<void>(lcqi::encode_prompt(tokenizer, "tok1 missing_token"))↵
↵ ;
368     } catch (const std::runtime_error&) {
369         threw = true;
370     }
371     require(threw, "tokenizer should reject unknown token when unk id is absent↵
↵ ");
372 }
373
374 void test_safetensors_manifest_contract() {
375     const std::filesystem::path path =
376         temp_file_path("lcqi_safetensors_manifest_contract.safetensors");
377     const std::string header =
378         R"({"__metadata__":{"format":"lcqi-fixture"},"model.embed_tokens.weight↵
↵ {"dtype":"F32","shape":[2,3],"data_offsets":[0,24]},"model.layers.0.↵
↵ attn.q_proj.weight":{"dtype":"F16","shape":[4,3],"data_offsets"↵
↵ :[24,48]}})";
379     write_safetensors_fixture(path, header, ↵
↵ LCQI_SAFETENSORS_VALID_PAYLOAD_BYTES);
380
381     const lcqi::SafeTensorManifest manifest = lcqi::load_safetensors_manifest(↵
↵ path);
382     std::filesystem::remove(path);
383

```

```

384     require(manifest.header_size == static_cast<std::uint64_t>(header.size()),
385             "safetensors header size mismatch");
386     require(manifest.data_start_offset ==
387             static_cast<std::uint64_t>(LCQI_SAFETENSORS_PREFIX_BYTES) +
388             static_cast<std::uint64_t>(header.size()),
389             "safetensors data start mismatch");
390     require(manifest.metadata.size() == 1, "safetensors metadata count mismatch"
391     ↪ );
392     require(manifest.metadata[0].first == "format" &&
393             manifest.metadata[0].second == "lcqi-fixture",
394             "safetensors metadata mismatch");
395     require(manifest.tensors.size() == 2, "safetensors tensor count mismatch");
396
397     const lcqi::SafeTensorEntry* embedding =
398         manifest.find_tensor("model.embed_tokens.weight");
399     require(embedding != nullptr, "safetensors embedding tensor missing");
400     require(embedding->dtype == "F32", "safetensors embedding dtype mismatch");
401     require(embedding->shape == std::vector<std::int64_t>({2, 3}),
402             "safetensors embedding shape mismatch");
403     require(embedding->data_begin == LCQI_SAFETENSORS_EMBED_BEGIN &&
404             embedding->data_end == LCQI_SAFETENSORS_EMBED_END,
405             "safetensors embedding offset mismatch");
406     require(embedding->byte_size() ==
407             LCQI_SAFETENSORS_EMBED_END - LCQI_SAFETENSORS_EMBED_BEGIN,
408             "safetensors embedding byte size mismatch");
409
410     const lcqi::SafeTensorEntry* q_proj =
411         manifest.find_tensor("model.layers.0.attn.q_proj.weight");
412     require(q_proj != nullptr, "safetensors q projection tensor missing");
413     require(q_proj->dtype == "F16", "safetensors q projection dtype mismatch");
414     require(q_proj->data_begin == LCQI_SAFETENSORS_QPROJ_BEGIN &&
415             q_proj->data_end == LCQI_SAFETENSORS_QPROJ_END,
416             "safetensors q projection offset mismatch");
417 }
418
419 void test_safetensors_rejects_bad_offsets() {
420     const std::filesystem::path path =
421         temp_file_path("lcqi_safetensors_bad_offsets.safetensors");
422     const std::string header =
423         R("{\"tensor\":{\"dtype\":\"F32\",\"shape\":[4],\"data_offsets\":[0,32]}}");
424     write_safetensors_fixture(path, header, 4);
425
426     bool threw = false;
427     try {
428         static_cast<void>(lcqi::load_safetensors_manifest(path));
429     } catch (const std::runtime_error&) {
430         threw = true;
431     }
432     std::filesystem::remove(path);
433     require(threw, "safetensors parser should reject offsets beyond payload");
434 }

```

```

435 void test_safetensors_rejects_bad_dtype_shape_size() {
436     const std::filesystem::path path =
437         temp_file_path("lcqi_safetensors_bad_dtype_shape_size.safetensors");
438     const std::string header =
439         R("{\"tensor\":{\"dtype\":\"F32\",\"shape\":[4],\"data_offsets\":[0,8]}}");
440     write_safetensors_fixture(path, header, 8);
441
442     bool threw = false;
443     try {
444         static_cast<void>(lcqi::load_safetensors_manifest(path));
445     } catch (const std::runtime_error&) {
446         threw = true;
447     }
448     std::filesystem::remove(path);
449     require(threw, "safetensors parser should reject dtype/shape byte mismatch" ←
    ↪ );
450 }
451
452 void test_safetensors_reads_float_tensors() {
453     const std::filesystem::path path =
454         temp_file_path("lcqi_safetensors_float_read.safetensors");
455     std::vector<std::uint8_t> f16_bytes;
456     append_le_u16(f16_bytes, LCQI_TEST_F16_ONE);
457     append_le_u16(f16_bytes, LCQI_TEST_F16_NEGATIVE_TWO);
458     std::vector<std::uint8_t> bf16_bytes;
459     append_le_u16(bf16_bytes, LCQI_TEST_BF16_ONE_POINT_FIVE);
460     append_le_u16(bf16_bytes, LCQI_TEST_BF16_NEGATIVE_HALF);
461
462     const std::array<RawTensorFixture, 3> tensors{{
463         {"f32", "F32", {2}, f32_payload(std::array<float, 2>{1.25F, -2.5F})},
464         {"f16", "F16", {2}, f16_bytes},
465         {"bf16", "BF16", {2}, bf16_bytes},
466     }};
467     write_safetensors_raw_fixture(path, tensors);
468
469     const lcqi::SafeTensorFile file = lcqi::load_safetensors_file(path);
470     std::filesystem::remove(path);
471     const std::vector<float> f32 = lcqi::read_safetensor_f32_tensor(file, "f32" ←
    ↪ );
472     const std::vector<float> f16 = lcqi::read_safetensor_f32_tensor(file, "f16" ←
    ↪ );
473     const std::vector<float> bf16 = lcqi::read_safetensor_f32_tensor(file, " ←
    ↪ bf16");
474
475     require_close(f32[0], 1.25F, "safetensors F32 value 0 mismatch");
476     require_close(f32[1], -2.5F, "safetensors F32 value 1 mismatch");
477     require_close(f16[0], 1.0F, "safetensors F16 value 0 mismatch");
478     require_close(f16[1], -2.0F, "safetensors F16 value 1 mismatch");
479     require_close(bf16[0], 1.5F, "safetensors BF16 value 0 mismatch");
480     require_close(bf16[1], -0.5F, "safetensors BF16 value 1 mismatch");
481 }
482

```



```

483 void test_gpt2_tiny_forward_and_generation() {
484     const lcqi::Gpt2ReferenceModel model = lcqi::make_tiny_gpt2_reference_model(
485         ↪ );
486     const std::vector<std::int32_t> tokens{1, 2, 3};
487     const lcqi::Gpt2ForwardResult result = lcqi::run_gpt2_forward(model, tokens(
488         ↪ );
489
490     require(result.logits.size() == LCQI_GPT2_VOCAB_SIZE, "GPT-2 logits size ↪
491         ↪ mismatch");
492     require(result.predicted_token == LCQI_GPT2_PREDICTED_TOKEN,
493         "GPT-2 predicted token mismatch");
494     for (std::size_t index = 0; index < LCQI_GPT2_LOGITS.size(); ++index) {
495         require_close(result.logits[index], LCQI_GPT2_LOGITS[index], "GPT-2 ↪
496         ↪ logit mismatch");
497     }
498
499     const std::vector<std::int32_t> generated =
500         lcqi::gpt2_generate_greedy(model, tokens, 2);
501     const std::vector<std::int32_t> expected{1, 2, 3, 3, 3};
502     require(generated == expected, "GPT-2 greedy generation mismatch");
503     require(static_cast<std::int32_t>(generated.size()) == ↪
504         ↪ LCQI_GPT2_GENERATED_LENGTH,
505         "GPT-2 generated length mismatch");
506
507     lcqi::Gpt2KvCache cache(model.config);
508     lcqi::Gpt2ForwardResult cached_result;
509     for (const std::int32_t token : tokens) {
510         cached_result = lcqi::run_gpt2_forward_cached(model, cache, token);
511     }
512     require(cache.filled_tokens() == static_cast<std::int32_t>(tokens.size()),
513         "GPT-2 KV cache filled token count mismatch");
514     require(cache.byte_size() > 0, "GPT-2 KV cache byte size should be non-zero(
515         ↪ );
516     require(cached_result.predicted_token == result.predicted_token,
517         "cached GPT-2 predicted token mismatch");
518     require(cached_result.logits.size() == result.logits.size(),
519         "cached GPT-2 logits size mismatch");
520     for (std::size_t index = 0; index < result.logits.size(); ++index) {
521         require_close(cached_result.logits[index],
522             result.logits[index],
523             "cached GPT-2 logit mismatch");
524     }
525
526     const std::vector<std::int32_t> cached_generated =
527         lcqi::gpt2_generate_greedy_cached(model, tokens, 2);
528     require(cached_generated == expected, "cached GPT-2 greedy generation ↪
529         ↪ mismatch");
530 }
531
532 void test_gpt2_byte_bpe_tokenizer() {
533     const std::filesystem::path vocab_path = temp_file_path("lcqi_gpt2_vocab.(
534         ↪ json");

```



```

527     const std::filesystem::path merges_path = temp_file_path("lcqi_gpt2_merges.↵
↵ txt");
528     const std::string space_marker = "\xC4\xA0";
529     const std::string vocab =
530         "{ \"Hello\":7,\",\":8,\"\" + space_marker + "world\":9,\"!\":10}";
531     const std::string merges =
532         "#version: 0.2\n"
533         "H e\n"
534         "He l\n"
535         "Hel l\n"
536         "Hell o\n" +
537         space_marker + " w\n" +
538         space_marker + "w o\n" +
539         space_marker + "wo r\n" +
540         space_marker + "wor l\n" +
541         space_marker + "worl d\n";
542     write_text_file(vocab_path, vocab);
543     write_text_file(merges_path, merges);
544
545     const lcqi::Gpt2Tokenizer tokenizer =
546         lcqi::load_gpt2_tokenizer(vocab_path, merges_path, -1, -1);
547     std::filesystem::remove(vocab_path);
548     std::filesystem::remove(merges_path);
549
550     const std::vector<std::int32_t> ids =
551         lcqi::gpt2_encode(tokenizer, "Hello, world!");
552     const std::vector<std::int32_t> expected{
553         LCQI_GPT2_TOKENIZER_HELLO_ID,
554         8,
555         9,
556         10,
557     };
558     require(ids == expected, "GPT-2 BPE ids mismatch");
559     require(lcqi::gpt2_decode(tokenizer, ids) == "Hello, world!",
560         "GPT-2 BPE decode mismatch");
561 }
562
563 void test_gpt2_loads_hf_style_tiny_checkpoint() {
564     const std::filesystem::path directory =
565         temp_file_path("lcqi_gpt2_tiny_checkpoint_dir");
566     std::filesystem::remove_all(directory);
567     std::filesystem::create_directories(directory);
568
569     const lcqi::Gpt2ReferenceModel original = lcqi::↵
↵ make_tiny_gpt2_reference_model();
570     write_text_file(
571         directory / "config.json",
572         R("{ \"model_type\": \"gpt2\", \"n_embd\": 4, \"n_head\": 2, \"n_layer\": 1, \"n_positions\"↵
↵ : 8, \"n_ctx\": 8, \"vocab_size\": 6, \"n_inner\": 8, \"layer_norm_epsilon\": 0.00001, \"↵
↵ activation_function\": \"gelu_new\", \"bos_token_id\": 0, \"eos_token_id\": 5}"));
573     write_tiny_gpt2_safetensors(directory / "model.safetensors", original);
574

```

```

575     const lcqi::Gpt2ReferenceModel loaded = lcqi::load_gpt2_from_directory(↵
↵ directory);
576     const std::vector<std::int32_t> tokens{1, 2, 3};
577     const lcqi::Gpt2ForwardResult result = lcqi::run_gpt2_forward(loaded, ↵
↵ tokens);
578     std::filesystem::remove_all(directory);
579
580     require(result.predicted_token == LCQI_GPT2_PREDICTED_TOKEN,
581            "loaded GPT-2 predicted token mismatch");
582     for (std::size_t index = 0; index < LCQI_GPT2_LOGITS.size(); ++index) {
583         require_close(result.logits[index],
584            LCQI_GPT2_LOGITS[index],
585            "loaded GPT-2 logit mismatch");
586     }
587 }
588
589 void test_gpt2_loader_rejects_bad_shape() {
590     const std::filesystem::path directory =
591         temp_file_path("lcqi_gpt2_bad_shape_dir");
592     std::filesystem::remove_all(directory);
593     std::filesystem::create_directories(directory);
594
595     const lcqi::Gpt2ReferenceModel original = lcqi::↵
↵ make_tiny_gpt2_reference_model();
596     write_text_file(
597         directory / "config.json",
598         R"({"model_type":"gpt2","n_embd":5,"n_head":1,"n_layer":1,"n_positions"↵
↵ :8,"n_ctx":8,"vocab_size":6,"n_inner":8,"layer_norm_epsilon":0.00001,"↵
↵ activation_function":"gelu_new","bos_token_id":0,"eos_token_id":5})");
599     write_tiny_gpt2_safetensors(directory / "model.safetensors", original);
600
601     bool threw = false;
602     try {
603         static_cast<void>(lcqi::load_gpt2_from_directory(directory));
604     } catch (const std::runtime_error&) {
605         threw = true;
606     }
607     std::filesystem::remove_all(directory);
608     require(threw, "GPT-2 loader should reject bad tensor shape");
609 }
610
611 void test_reference_rms_norm() {
612     const std::vector<float> input{3.0F, 4.0F};
613     const std::vector<float> weight{1.0F, 0.5F};
614     std::vector<float> output(2, 0.0F);
615     lcqi::rms_norm(input, weight, 0.0F, output);
616     require_close(output[0], LCQI_RMS_EXPECTED_0, "RMSNorm output 0 mismatch");
617     require_close(output[1], LCQI_RMS_EXPECTED_1, "RMSNorm output 1 mismatch");
618 }
619
620 void test_reference_rope() {
621     std::vector<float> heads{1.0F, 0.0F, 0.0F, 1.0F};

```

```

622     lcqi::apply_rope(heads, 1, 4, 1, 10000.0F);
623     require_close(heads[0], std::cos(1.0F), "RoPE pair 0 cosine mismatch");
624     require_close(heads[1], std::sin(1.0F), "RoPE pair 0 sine mismatch");
625     require_close(heads[2], -std::sin(0.01F), "RoPE pair 1 negative sine ←
↪ mismatch");
626     require_close(heads[3], std::cos(0.01F), "RoPE pair 1 cosine mismatch");
627 }
628
629 void test_reference_kv_attention() {
630     lcqi::DecoderConfig config;
631     config.hidden_size = 4;
632     config.query_heads = 2;
633     config.kv_heads = 1;
634     config.head_dim = 2;
635     config.intermediate_size = 4;
636     config.vocab_size = 4;
637     config.max_context = 4;
638
639     lcqi::ReferenceKVCache cache(config, 1);
640     const std::vector<float> key{1.0F, 0.0F};
641     const std::vector<float> value{LCQI_ATTENTION_EXPECTED_0, ←
↪ LCQI_ATTENTION_EXPECTED_1};
642     cache.append(0, 0, key, value);
643     require(cache.filled_tokens() == 1, "KV filled token count mismatch");
644     require_close(cache.key(0, 0, 0)[0], 1.0F, "KV key read mismatch");
645
646     const std::vector<float> query{1.0F, 0.0F, 0.0F, 1.0F};
647     std::vector<float> output(4, 0.0F);
648     lcqi::attention_decode(config, cache, 0, 0, query, output);
649     require_close(output[0], LCQI_ATTENTION_EXPECTED_0, "attention head 0 dim 0←
↪ mismatch");
650     require_close(output[1], LCQI_ATTENTION_EXPECTED_1, "attention head 0 dim 1←
↪ mismatch");
651     require_close(output[2], LCQI_ATTENTION_EXPECTED_0, "attention head 1 dim 0←
↪ mismatch");
652     require_close(output[3], LCQI_ATTENTION_EXPECTED_1, "attention head 1 dim 1←
↪ mismatch");
653 }
654
655 void test_reference_decoder_end_to_end() {
656     const lcqi::ReferenceDecoderModel model = lcqi::←
↪ make_tiny_reference_decoder_model();
657     const std::vector<std::int32_t> tokens{1, 2, 3};
658     const lcqi::ReferenceDecodeResult result = lcqi::run_reference_decode(model←
↪ , tokens);
659     require(result.logits.size() == LCQI_REFERENCE_DECODER_VOCAB_SIZE,
660             "reference decoder logits size mismatch");
661     require(result.predicted_token == LCQI_REFERENCE_DECODER_PREDICTED_TOKEN,
662             "reference decoder predicted token mismatch");
663     for (std::size_t index = 0; index < LCQI_REFERENCE_DECODER_LOGITS.size(); ←
↪ ++index) {
664         require_close(result.logits[index],

```

```

665         LCQI_REFERENCE_DECODER_LOGITS[index],
666         "reference decoder golden logit mismatch");
667     }
668 }
669
670 void test_reference_decoder_trace_contract() {
671     const lcqi::ReferenceDecoderModel model = lcqi::↵
↵ make_tiny_reference_decoder_model();
672     const std::vector<std::int32_t> tokens{1, 2, 3};
673     const lcqi::ReferenceDecodeTraceResult trace =
674         lcqi::run_reference_decode_with_trace(model, tokens);
675
676     require(trace.result.predicted_token == ↵
↵ LCQI_REFERENCE_DECODER_PREDICTED_TOKEN,
677         "reference trace predicted token mismatch");
678     require(!trace.checkpoints.empty(), "reference trace checkpoints missing");
679
680     const lcqi::ReferenceTensorCheckpoint& first = trace.checkpoints.front();
681     require(first.name == "embedding", "first checkpoint should be embedding");
682     require(first.token_position == LCQI_REFERENCE_TRACE_FIRST_TOKEN,
683         "first checkpoint token position mismatch");
684     require(first.dtype == "f32", "first checkpoint dtype mismatch");
685     require(first.layout == "row_major", "first checkpoint layout mismatch");
686     require(first.shape.size() == 1 &&
687         first.shape[0] == LCQI_REFERENCE_TRACE_HIDDEN_SIZE,
688         "first checkpoint shape mismatch");
689     require(first.stride.size() == 1 && first.stride[0] == 1,
690         "first checkpoint stride mismatch");
691
692     bool saw_logits = false;
693     bool saw_kv_slot = false;
694     for (const lcqi::ReferenceTensorCheckpoint& checkpoint : trace.checkpoints)↵
↵ {
695         if (checkpoint.name == "kv_cache_key_slot") {
696             saw_kv_slot = true;
697             require(checkpoint.shape.size() == 4, "KV checkpoint rank mismatch"↵
↵ );
698             require(checkpoint.layout == "kv_contiguous", "KV checkpoint layout↵
↵ mismatch");
699         }
700         if (checkpoint.name == "logits") {
701             saw_logits = true;
702             require(checkpoint.token_position == ↵
↵ LCQI_REFERENCE_TRACE_TOKEN_COUNT - 1,
703                 "logits checkpoint token position mismatch");
704             require(checkpoint.shape.size() == 1 &&
705                 checkpoint.shape[0] ==
706                 static_cast<std::int32_t>(↵
↵ LCQI_REFERENCE_DECODER_VOCAB_SIZE),
707                 "logits checkpoint shape mismatch");
708             require(checkpoint.values.size() == LCQI_REFERENCE_DECODER_LOGITS.↵
↵ size(),

```

```

709         "logits checkpoint value size mismatch");
710         for (std::size_t index = 0; index < LCQI_REFERENCE_DECODER_LOGITS.↵
↵ size(); ++index) {
711             require_close(checkpoint.values[index],
712                           LCQI_REFERENCE_DECODER_LOGITS[index],
713                           "trace logits checkpoint mismatch");
714         }
715     }
716 }
717 require(saw_kv_slot, "KV checkpoint missing");
718 require(saw_logits, "logits checkpoint missing");
719 }
720
721 void test_packed_i8_kernel_matches_scalar() {
722     for (const KernelShapeCase& shape_case : LCQI_KERNEL_SHAPE_CASES) {
723         const lcqi::QuantizedLinearLayer layer =
724             lcqi::make_deterministic_i8_layer(shape_case.input_size,
725                                               shape_case.output_size,
726                                               0.03125F);
727         const lcqi::PackedLinearI8 packed =
728             lcqi::pack_linear_i8(layer, shape_case.output_block_size);
729         const std::vector<float> input = lcqi::make_deterministic_input(layer.↵
↵ input_size);
730         std::vector<float> scalar_output(static_cast<std::size_t>(layer.↵
↵ output_size), 0.0F);
731         std::vector<float> packed_output(static_cast<std::size_t>(layer.↵
↵ output_size), 0.0F);
732         lcqi::linear_i8(layer, input, scalar_output);
733         lcqi::linear_i8_packed(packed, input, packed_output);
734         require(lcqi::max_abs_diff(scalar_output, packed_output) <= ↵
↵ LCQI_KERNEL_MAX_DIFF,
735               "packed int8 kernel output differs from scalar");
736
737         const lcqi::PackedLinearI8 simd_packed =
738             lcqi::pack_linear_i8(layer, LCQI_SIMD_TEST_BLOCK);
739         if (lcqi::linear_i8_packed_avx2_available()) {
740             std::vector<float> avx2_output(static_cast<std::size_t>(layer.↵
↵ output_size), 0.0F);
741             lcqi::linear_i8_packed_avx2(simd_packed, input, avx2_output);
742             require(lcqi::max_abs_diff(scalar_output, avx2_output) <= ↵
↵ LCQI_KERNEL_MAX_DIFF,
743                   "AVX2 int8 kernel output differs from scalar");
744         }
745     }
746 }
747
748 } // namespace
749
750 int main() {
751     try {
752         test_tiny_mlp();
753         test_tokenizer_contract();

```

```
754     test_tokenizer_rejects_unknown_without_unk();
755     test_safetensors_manifest_contract();
756     test_safetensors_rejects_bad_offsets();
757     test_safetensors_rejects_bad_dtype_shape_size();
758     test_safetensors_reads_float_tensors();
759     test_gpt2_tiny_forward_and_generation();
760     test_gpt2_byte_bpe_tokenizer();
761     test_gpt2_loads_hf_style_tiny_checkpoint();
762     test_gpt2_loader_rejects_bad_shape();
763     test_reference_rms_norm();
764     test_reference_rope();
765     test_reference_kv_attention();
766     test_reference_decoder_end_to_end();
767     test_reference_decoder_trace_contract();
768     test_packed_i8_kernel_matches_scalar();
769
770     std::cout << "lcqi tests passed\n";
771     return EXIT_SUCCESS;
772 } catch (const std::exception& error) {
773     std::cerr << "lcqi test failure: " << error.what() << "\n";
774     return EXIT_FAILURE;
775 }
776 }
```

# 实践与代码总索引

本卷已经把实践任务、代码走读路线和完整代码清单合并到同一套内容目录。目录中不再存在一个独立的“源码卷部分”：构建入口、调用链、公共合同、tiny MLP、int8 kernel、reference decoder、GPT-2、工具脚本和测试文件都插入到相应主题之后，共同解释同一份 `books/cpu-volume-3/labs/linux_cpu_inference` 源码。

## 源码阅读任务

读者不需要再在实践卷和源码卷之间来回切换，也不需要先读完 16 个实践章再去末尾翻代码。实践章节提出任务，紧邻的代码走读解释完整实现，测试和 benchmark 章节给出回归证据；三者已经合并到这一本实践与代码卷里，共同构成一条可复查的工程证据链。

## 一、实践主题到代码走读的合并位置

- 第 1 章之后 插入“代码走读：构建入口、项目边界与调用链总图”，先把 CMake、README、模型资产、target、CLI、benchmark、trace、ledger、serving probe 和测试入口放到一张可复现地图里。
- 第 4 章之后 插入“代码走读：公共合同、Tiny MLP 与 Int8 Kernel”，把模型、推理、tokenizer、safetensors、reference decoder、GPT-2 头文件，以及 tiny MLP、基础 CLI、scalar/packed/AVX2 kernel 连起来。
- 第 10 章之后 插入“代码走读：Reference Decoder、Tokenizer、Safetensors 与 GPT-2”，把真实模型加载、byte-level BPE、safetensors、GPT-2 forward、cached KV 和 greedy generation 放在模型导入主题内部讲清楚。
- 第 16 章之后 插入“代码走读：工具、Smoke、Benchmark 与回归测试”，把 benchmark、trace、ledger、serving probe、SmolLM2 smoke、GPT-2 smoke、benchmark compare 和 CTest 回归证据集中收束。

## 二、最小复现命令

```
1 cmake -S books/cpu-volume-3 -B books/cpu-volume-3/build/lcqi-debug -↵  
   ↵ DCMMAKE_BUILD_TYPE=Debug  
2 cmake --build books/cpu-volume-3/build/lcqi-debug  
3 ctest --test-dir books/cpu-volume-3/build/lcqi-debug --output-on-failure  
4 make cpu3p-pdf  
5 make cpu3p-epub
```

完成这组命令后，读者应同时检查相邻实践章节的作业结果、对应代码走读、测试断言和报告字段，确认代码、测试、报告字段和章节讲解是一条证据链。

## 术语表

| 术语                | 实践含义                                                                |
|-------------------|---------------------------------------------------------------------|
| manifest          | 模型资产清单，记录 tensor name、dtype、shape、offset、checksum 和 tokenizer/配置版本。 |
| shape verifier    | 在 kernel 运行前检查 shape、dtype、layout、量化 descriptor 和 state edge 是否匹配。  |
| reference decoder | 以清楚正确为目标的 decoder 实现，用来生成 trace 和 golden。                           |
| oracle            | 比较候选实现与 reference 的判定器，可使用精确相等或误差界。                                 |
| trace checkpoint  | 推理链路中的可复查节点，至少包含 shape、stride、dtype、layout、checksum 和摘要值。           |
| KV cache          | attention 历史 K/V 状态，跨 token 存活，不能被普通 activation planner 复用。         |
| packing           | 把权重改写成更适合热循环和 SIMD 的布局。                                             |
| backend           | 执行计划落到机器的实现，例如 scalar、AVX2、AVX-512、GPU 或库调用。                        |
| fallback          | 后端能力不满足时显式降级到可验证路径，并记录原因。                                           |
| SLO               | 服务目标，例如 TTFT、TPOT、queue wait、拒绝率和取消回收时间。                            |